



# AAAI-24 Main Track

Title	Authors
Attention Disturbance and Dual-Path Constraint Network for Occluded Person Re-identification	Jiaer Xia; Lei Tan; Pingyang Dai; Mingbo Zhao; Yongjian Wu; Liujuan Cao
Zero-Shot Task Adaptation with Relevant Feature Information	Atsutoshi Kumagai; Tomoharu Iwata; Yasuhiro Fujiwara
DTL: Disentangled Transfer Learning for Visual Recognition	Minghao Fu; Ke Zhu; Jianxin Wu
ODTrack: Online Dense Temporal Token Learning for Visual Tracking	Yaozong Zheng; Bineng Zhong; Qihua Liang; Zhiyi Mo; Shengping Zhang; Xianxian Li
MM-Point: Multi-View Information-Enhanced Multi-Modal Self-Supervised 3D Point Cloud Understanding	HaiTao Yu; Mofei Song
Defying Imbalanced Forgetting in Class Incremental Learning	Shixiong Xu; Gaofeng Meng; Xing Nie; Bolin Ni; Bin Fan; Shiming Xiang
Towards Efficient and Effective Text-to-Video Retrieval with Coarse-to-Fine Visual Representation Learning	Kaibin Tian; Yanhua Cheng; Yi Liu; Xinglin Hou; Quan Chen; Han Li
Generative Model-Based Feature Knowledge Distillation for Action Recognition	Guiqin Wang; Peng Zhao; Yanjiang Shi; Cong Zhao; Shusen Yang
Robust 3D Tracking with Quality-Aware Shape Completion	Jingwen Zhang; Zikun Zhou; Guangming Lu; Jiandong Tian; Wenjie Pei
CFR-ICL: Cascade-Forward Refinement with Iterative Click Loss for Interactive Image Segmentation	Shoukun Sun; Min Xian; Fei Xu; Luca Capriotti; Tiankai Yao
AE-NeRF: Audio Enhanced Neural Radiance Field for Few Shot Talking Head Synthesis	Dongze Li; Kang Zhao; Wei Wang; Bo Peng; Yingya Zhang; Jing Dong; Tieniu Tan
DI-V2X: Learning Domain-Invariant Representation for Vehicle-Infrastructure Collaborative 3D Object Detection	Li Xiang; Junbo Yin; Wei Li; Chengzhong Xu; Ruigang Yang; Jianbing Shen
Leveraging Imagery Data with Spatial Point Prior for Weakly Semi-supervised 3D Object Detection	Hongzhi Gao; Zheng Chen; Zehui Chen; Lin Chen; Jiaming Liu; Shanghang Zhang; Feng Zhao
XKD: Cross-Modal Knowledge Distillation with Domain Alignment for Video Representation Learning	Pritam Sarkar; Ali Etemad
DiffusionEdge: Diffusion Probabilistic Model for Crisp Edge Detection	Yunfan Ye; Yuhang Huang; Renjiao Yi; Zhiping Cai; Kai Xu
Kernelized Normalizing Constant Estimation: Bridging Bayesian Quadrature and Bayesian Optimization	Xu Cai; Jonathan Scarlett
Aleth-NeRF: Illumination Adaptive NeRF with Concealing Field Assumption	Ziteng Cui; Lin Gu; Xiao Sun; Xianzheng Ma; Yu Qiao; Tatsuya Harada
HDMixer: Hierarchical Dependency with Extendable Patch for Multivariate Time Series Forecasting	Qihe Huang; Lei Shen; Ruixin Zhang; Jiahuan Cheng; Shouhong Ding; Zhengyang Zhou; Yang Wang
Cross-Covariate Gait Recognition: A Benchmark	Shinan Zou; Chao Fan; Jianbo Xiong; Chuanfu Shen; Shiqi Yu; Jin Tang
Which Is More Effective in Label Noise Cleaning, Correction or Filtering?	Gaoxia Jiang; Jia Zhang; Xuefei Bai; Wenjian Wang; Deyu Meng
Let All Be Whitened: Multi-Teacher Distillation for Efficient Visual Retrieval	Zhe Ma; Jianfeng Dong; Shouling Ji; Zhenguang Liu; Xuhong Zhang; Zonghui Wang; Sifeng He; Feng Qian; Xiaobo Zhang; Lei Yang
Identifiability of Direct Effects from Summary Causal Graphs	Simon Ferreira; Charles Assaad
NuScenes-QA: A Multi-Modal Visual Question Answering Benchmark for Autonomous Driving Scenario	Tianwen Qian; Jingjing Chen; Linhai Zhuo; Yang Jiao; Yu-Gang Jiang
CARAT: Contrastive Feature Reconstruction and Aggregation for Multi-Modal Multi-Label Emotion Recognition	Cheng Peng; Ke Chen; Lidan Shou; Gang Chen



ViSTec: Video Modeling for Sports Technique Recognition and Tactical PromptMRG: Diagnosis-Driven Prompts for Medical Report Generation	Yuchen He; Zeqing Yuan; Yihong Wu; Liqi Cheng; Dazhen Deng; Yingcai Wu Haibo Jin; Haoxuan Che; Yi Lin; Hao Chen
EMGAN: Early-Mix-GAN on Extracting Server-Side Model in Split Federated Learning	Jingtao Li; Xing Chen; Li Yang; Adnan Siraj Rakin; Deliang Fan; Chaitali Chakrabarti
FedDAT: An Approach for Foundation Model Finetuning in Multi-Modal Heterogeneous Federated Learning	Haokun Chen; Yao Zhang; Denis Krompass; Jindong Gu; Volker Tresp
Auto-Prox: Training-Free Vision Transformer Architecture Search via Automatic Proxy Discovery	Zimian Wei; Peijie Dong; Zheng Hui; Anggeng Li; Lujun Li; Menglong Lu; Hengyue Pan; Dongsheng Li
Follow Your Pose: Pose-Guided Text-to-Video Generation Using Pose-Free Videos	Yue Ma; Yingqing HE; Xiaodong Cun; Xintao Wang; Siran Chen; Xiu Li; Qifeng Chen
Style2Talker: High-Resolution Talking Head Generation with Emotion Style and Art Style	Shuai Tan; Bin Ji; Ye Pan
Colorizing Monochromatic Radiance Fields	Yean Cheng; Renjie Wan; Shuchen Weng; Chengxuan Zhu; Yakun Chang; Boxin Shi
SAM-PARSER: Fine-Tuning SAM Efficiently by Parameter Space Reconstruction	Zelin Peng; Zhengqin Xu; Zhilin Zeng; Xiaokang Yang; Wei Shen
Mimic: Speaking Style Disentanglement for Speech-Driven 3D Facial Animation	Hui Fu; Zeqing Wang; Ke Gong; Keze Wang; Tianshui Chen; Haojie Li; Haifeng Zeng; Wenxiong Kang
Tuning-Free Inversion-Enhanced Control for Consistent Image Editing	Xiaoyue Duan; Shuhao Cui; Guoliang Kang; Baochang Zhang; Zhengcong Fei; Mingyuan Fan; Junshi Huang
Frequency-Adaptive Pan-Sharpening with Mixture of Experts	Xuanhua He; Keyu Yan; Rui Li; Chengjun Xie; Jie Zhang; Man Zhou
Transformer-Based Video-Structure Multi-Instance Learning for Whole Slide Image Classification	Yingfan Ma; Xiaoyuan Luo; Kexue Fu; Manning Wang
Attribute-Missing Graph Clustering Network	Wenxuan Tu; Renxiang Guan; Sihang Zhou; Chuan Ma; Xin Peng; Zhiping Cai; Zhe Liu; Jieren Cheng; Xinwang Liu
Explicit Visual Prompts for Visual Object Tracking	Liangtao Shi; Bineng Zhong; Qihua Liang; Ning Li; Shengping Zhang; Xianxian
Task-Agnostic Privacy-Preserving Representation Learning for Federated Learning against Attribute Inference Attacks	Caridad Arroyo Arevalo; Sayedeh Leila Noorbakhsh; Yun Dong; Yuan Hong; Binghui Wang
LERE: Learning-Based Low-Rank Matrix Recovery with Rank Estimation	Zhengqin Xu; Yulun Zhang; Chao Ma; Yichao Yan; Zelin Peng; Shoulie Xie; Shiqian Wu; Xiaokang Yang
Adv-Diffusion: Imperceptible Adversarial Face Identity Attack via Latent Diffusion Model	Decheng Liu; Xijun Wang; Chunlei Peng; Nannan Wang; Ruimin Hu; Xinbo Gao
In-Hand 3D Object Reconstruction from a Monocular RGB Video	Shijian Jiang; Qi Ye; Rengan Xie; Yuchi Huo; Xiang Li; Yang Zhou; Jiming
Preference Ranking Optimization for Human Alignment	Feifan Song; Bowen Yu; Minghao Li; Haiyang Yu; Fei Huang; Yongbin Li; Houfeng Wang
Stability of Multi-Agent Learning in Competitive Networks: Delaying the Onset of Chaos	Aamal Hussain; Francesco Belardinelli
Peer Neighborhood Mechanisms: A Framework for Mechanism Generalization	Adam Richardson; Boi Faltings
Differentiable Auxiliary Learning for Sketch Re-identification	Xingyu Liu; Xu Cheng; Haoyu Chen; Hao Yu; Guoying Zhao
Disentangled Diffusion-Based 3D Human Pose Estimation with Hierarchical Spatial and Temporal Denoiser	Qingyuan Cai; Xuecai Hu; Saihui Hou; Li Yao; Yongzhen Huang



Rethinking Multi-Scale Representations in Deep Deraining Transformer	Hongming Chen; Xiang Chen; Jiyang Lu; Yufeng Li
Hard Regularization to Prevent Deep Online Clustering Collapse without Data Augmentation	Louis Mahon; Thomas Lukasiewicz
Sunshine to Rainstorm: Cross-Weather Knowledge Distillation for Robust 3D Object Detection	Xun Huang; Hai Wu; Xin Li; Xiaoliang Fan; Chenglu Wen; Cheng Wang
UVAGaze: Unsupervised 1-to-2 Views Adaptation for Gaze Estimation	Ruicong Liu; Feng Lu
Aspect-Based Sentiment Analysis with Explicit Sentiment Augmentations	Jihong Ouyang; Zhiyao Yang; Silong Liang; Bing Wang; Yimeng Wang; Ximing
Label Attentive Distillation for GNN-Based Graph Classification	Xiaobin Hong; Wenzhong Li; Chaoqun Wang; Mingkai Lin; Sanglu Lu
Using Symmetries to Lift Satisfiability Checking	Pierre Carbonnelle; Gottfried Schenner; Maurice Bruynooghe; Bart Bogaerts; Marc Denecker
Panoptic Scene Graph Generation with Semantics-Prototype Learning	Li Li; Wei Ji; Yiming Wu; Mengze Li; You Qin; Lina Wei; Roger Zimmermann
CEDFlow: Latent Contour Enhancement for Dark Optical Flow Estimation	Fengyuan Zuo; Zhaolin Xiao; Haiyan Jin; Haonan Su
Latent Space Editing in Transformer-Based Flow Matching	Vincent Tao Hu; Wei Zhang; Meng Tang; Pascal Mettes; Deli Zhao; Cees Snoek
Harnessing Holistic Discourse Features and Triadic Interaction for Sentiment Quadruple Extraction in Dialogues	Bobo Li; Hao Fei; Lizi Liao; Yu Zhao; Fangfang Su; Fei Li; Donghong Ji
ChromaFusionNet (CFNet): Natural Fusion of Fine-Grained Color Editing	Yi Dong; Yuxi Wang; Ruoxi Fan; Wenqi Ouyang; Zhiqi Shen; Peiran Ren; Xuansong Xie
SPD-DDPM: Denoising Diffusion Probabilistic Models in the Symmetric Positive Definite Space	Yunchen Li; Zhou Yu; Gaoqi He; Yunhang Shen; Ke Li; Xing Sun; Shaohui Lin
Few-Shot Neural Radiance Fields under Unconstrained Illumination	SeokYeong Lee; JunYong Choi; Seungryong Kim; Ig-Jae Kim; Junghyun Cho
Exploring Gradient Explosion in Generative Adversarial Imitation Learning: A Probabilistic Perspective	Wanying Wang; Yichen Zhu; Yirui Zhou; Chaomin Shen; Jian Tang; Zhiyuan Xu; Yaxin Peng; Yangchun Zhang
V2A-Mapper: A Lightweight Solution for Vision-to-Audio Generation by Connecting Foundation Models	Heng Wang; Jianbo Ma; Santiago Pascual; Richard Cartwright; Weidong Cai
STDiff: Spatio-Temporal Diffusion for Continuous Stochastic Video Prediction	Xi Ye; Guillaume-Alexandre Bilodeau
Multi-Prompts Learning with Cross-Modal Alignment for Attribute-Based Person Re-identification	Yajing Zhai; Yawen Zeng; Zhiyong Huang; Zheng Qin; Xin Jin; Da Cao
CDPNet: Cross-Modal Dual Phases Network for Point Cloud Completion	Zhenjiang Du; Jiale Dou; Zhitao Liu; Jiwei Wei; Guan Wang; Ning Xie; Yang Yang
BVT-IMA: Binary Vision Transformer with Information-Modified Attention	Zhenyu Wang; Hao Luo; Xuemei Xie; Fan Wang; Guangming Shi
Learning in Online Principle-Agent Interactions: The Power of Menus	Minbiao Han; Michael Albert; Haifeng Xu
SRFormer: Text Detection Transformer with Incorporated Segmentation and Regression	Qingwen Bu; Sungrae Park; Minsoo Khang; Yichuan Cheng
SAT-Based Algorithms for Regular Graph Pattern Matching	Miguel Terra-Neves; José Amaral; Alexandre Lemos; Rui Quintino; Pedro Resende; Antonio Alegria
BAND: Biomedical Alert News Dataset	Zihao Fu; Meiru Zhang; Zaiqiao Meng; Yannan Shen; David Buckeridge; Nigel Collier
SD-MVS: Segmentation-Driven Deformation Multi-View Stereo with Spherical Refinement and EM Optimization	Zhenlong Yuan; Jiakai Cao; Zhaoxin Li; Hao Jiang; Zhaoqi Wang
Dynamic Weighted Combiner for Mixed-Modal Image Retrieval	Fuxiang Huang; Lei Zhang; Xiaowei Fu; Suqi Song



MCL-NER: Cross-Lingual Named Entity Recognition via Multi-View Contrastive Learning	Ying Mo; Jian Yang; Jiahao Liu; Qifan Wang; Ruoyu Chen; Jingang Wang; Zhoujun Li
Iterative Token Evaluation and Refinement for Real-World Super-resolution	Chaofeng Chen; Shangchen Zhou; Liang Liao; Haoning Wu; Wenxiu Sun; Qiong Yan; Weisi Lin
CrossBind: Collaborative Cross-Modal Identification of Protein Nucleic-Acid-Binding Residues	Linglin Jing; Sheng Xu; Yifan Wang; Yuzhe Zhou; Tao Shen; Zhigang Ji; Hui Fang; Zhen Li; Siqi Sun
Optical Flow for Spike Camera with Hierarchical Spatial-Temporal Spike	Rui Zhao; Ruiqin Xiong; Jian Zhang; Xinfeng Zhang; Zhaoifei Yu; Tiejun Huang
Reviewing the Forgotten Classes for Domain Adaptation of Black-Box Predictors	Shaojie Zhang; Chun Shen; Shuai Lü; Zeyu Zhang
MatchDet: A Collaborative Framework for Image Matching and Object Detection	Jinxiang Lai; Wenlong Wu; Bin-Bin Gao; Jun Liu; Jiawei Zhan; Congchong Nie; Yi Zeng; Chengjie Wang
DGA-GNN: Dynamic Grouping Aggregation GNN for Fraud Detection	Mingjiang Duan; Tongya Zheng; Yang Gao; Gang Wang; Zunlei Feng; Xinyu Wang
FedDiv: Collaborative Noise Filtering for Federated Learning with Noisy Labels	Jichang Li; Guanbin Li; Hui Cheng; Zicheng Liao; Yizhou Yu
patchDPCC: A Patchwise Deep Compression Framework for Dynamic Point Clouds	Zirui Pan; Mengbai Xiao; Xu Han; Dongxiao Yu; Guanghui Zhang; Yao Liu
Regroup Median Loss for Combating Label Noise	Fengpeng Li; Kemou Li; Jinyu Tian; Jiantao Zhou
Multimodal Event Causality Reasoning with Scene Graph Enhanced Interaction Network	Jintao Liu; Kaiwen Wei; Chenglong Liu
TNPAR: Topological Neural Poisson Auto-Regressive Model for Learning Granger Causal Structure from Event Sequences	Yuequn Liu; Ruichu Cai; Wei Chen; Jie Qiao; Yuguang Yan; Zijian Li; Keli Zhang; Zhifeng Hao
Latent Diffusion Transformer for Probabilistic Time Series Forecasting	Shibo Feng; Chunyan Miao; Zhong Zhang; Peilin Zhao
Mining Gaze for Contrastive Learning toward Computer-Assisted Diagnosis	Zihao Zhao; Sheng Wang; Qian Wang; Dinggang Shen
Parallel Vertex Diffusion for Unified Visual Grounding	Zesen Cheng; Kehan Li; Peng Jin; Siheng Li; Xiangyang Ji; Li Yuan; Chang Liu; Jie Chen
G <sup>2</sup> SAM: Graph-Based Global Semantic Awareness Method for Multimodal Sarcasm Detection	Yiwei Wei; Shaozu Yuan; Hengyang Zhou; Longbiao Wang; Zhiling Yan; Ruosong Yang; Meng Chen
Meta-Learning-Based Adaptive Stability Certificates for Dynamical Systems	Amit Jena; Dileep Kalathil; Le Xie
Unsupervised Domain Adaptative Temporal Sentence Localization with Mutual Information Maximization	Daizong Liu; Xiang Fang; Xiaoye Qu; Jianfeng Dong; He Yan; Yang Yang; Pan Zhou; Yu Cheng
Fine-Grained Knowledge Selection and Restoration for Non-exemplar Class Incremental Learning	Jiang-Tian Zhai; Xialei Liu; Lu Yu; Ming-Ming Cheng
Watch Your Head: Assembling Projection Heads to Save the Reliability of Federated Models	Jinqian Chen; Jihua Zhu; Qinghai Zheng; Zhongyu Li; Zhiqiang Tian
SCTNet: Single Branch CNN with Transformer Semantic Information for Real-Time Segmentation	Zhengze Xu; Dongyue Wu; Changqian Yu; Xiangxiang Chu; Nong Sang; Changxin Gao
Domain Generalizable Person Search Using Unreal Dataset	Minyoung Oh; Duhyun Kim; Jae-Young Sim
Dynamic Tangled Derivative Logic of Metric Spaces	David Fernández-Duque; Yoàv Montacute
MaskDiff: Modeling Mask Distribution with Diffusion Probabilistic Model for Few-Shot Instance Segmentation	Minh-Quan Le; Tam Nguyen; Trung-Nghia Le; Thanh-Toan Do; Minh Do; Minh-Triet Tran
Stability in Online Coalition Formation	Martin Bullinger; René Romen



Learning Broadcast Protocols	Dana Fisman; Noa Izsak; Swen Jacobs
A Hybrid Global-Local Perception Network for Lane Detection	Qing Chang; Yifei Tong
Modeling Adaptive Inter-Task Feature Interactions via Sentiment-Aware Contrastive Learning for Joint Aspect-Sentiment Prediction	Wei Chen; Yuxuan Liu; Zhao Zhang; Fuzhen Zhuang; Jiang Zhong
COMMA: Co-articulated Multi-Modal Learning	Lianyu Hu; Liqing Gao; Zekang Liu; Chi-Man Pun; Wei Feng
IML-ViT: Benchmarking Image Manipulation Localization by Vision	Xiaochen Ma; Bo Du; Zhuohang Jiang; Ahmed Alhammadi; Jizhe Zhou
UniADS: Universal Architecture-Distiller Search for Distillation Gap	Liming Lu; Zhenghan Chen; Xiaoyu Lu; Shuchao Pang; Lujun Li; Yihang Rao
Fair Lotteries for Participatory Budgeting	Haris Aziz; Xinhang Lu; Mashbat Suzuki; Jeremy Vollen; Toby Walsh
Learning Planning Domains from Non-redundant Fully-Observed Traces: Theoretical Foundations and Complexity Analysis	Pascal Bachor; Gregor Behnke
Sparse3D: Distilling Multiview-Consistent Diffusion for Object Reconstruction from Sparse Views	Zixin Zou; Weihao Cheng; Yan-Pei Cao; Shi-Sheng Huang; Ying Shan; Song-Hai Zhang
DVANet: Disentangling View and Action Features for Multi-View Action Recognition	Nyle Siddiqui; Praveen Tirupattur; Mubarak Shah
SphereDiffusion: Spherical Geometry-Aware Distortion Resilient Diffusion Model	Tao Wu; Xuewei Li; Zhongang Qi; Di Hu; Xintao Wang; Ying Shan; Xi Li
Towards Transferable Adversarial Attacks with Centralized Perturbation	Shangbo Wu; Yu-an Tan; Yajie Wang; Ruinan Ma; Wencong Ma; Yuanzhang
Neural Time-Reversed Generalized Riccati Equation	Alessandro Betti; Michele Casoni; Marco Gori; Simone Marullo; Stefano Melacci; Matteo Tiezzi
MCA: Moment Channel Attention Networks	Yangbo Jiang; ZhiWei Jiang; Le Han; Zenan Huang; Nenggan Zheng
Compositional Inversion for Stable Diffusion Models	Xulu Zhang; Xiao-Yong Wei; Jinlin Wu; Tianyi Zhang; Zhen Lei; Zhaoxiang Zhang; Qing Li
Self-Supervised Representation Learning with Meta Comprehensive Regularization	Huijie Guo; Ying Ba; Jie Hu; Lingyu Si; Wenwen Qiang; Lei Shi
Towards Understanding Future: Consistency Guided Probabilistic Modeling for Action Anticipation	Zhao Xie; Yadong Shi; Kewei Wu; Yaru Cheng; Dan Guo
Toward Open-Set Human-Object Interaction Detection	Mingrui Wu; Yuqi Liu; Jiayi Ji; Xiaoshuai Sun; Rongrong Ji
SayCanPay: Heuristic Planning with Large Language Models Using Learnable Domain Knowledge	Rishi Hazra; Pedro Zuidberg Dos Martires; Luc De Raedt
Meta-Inverse Reinforcement Learning for Mean Field Games via Probabilistic Context Variables	Yang Chen; Xiao Lin; Bo Yan; Libo Zhang; Jiamou Liu; Neset TAN; Michael Witbrock
Coupled Confusion Correction: Learning from Crowds with Sparse Annotations	Hansong Zhang; Shikun Li; Dan Zeng; Chenggang Yan; Shiming Ge
MDGNN: Multi-Relational Dynamic Graph Neural Network for Comprehensive and Dynamic Stock Investment Prediction	Hao Qian; Hongting Zhou; Qian Zhao; Hao Chen; Hongxiang Yao; Jingwei Wang; Ziqi Liu; Fei Yu; Zhiqiang Zhang; Jun Zhou
Far3D: Expanding the Horizon for Surround-View 3D Object Detection	Xiaohui Jiang; Shuailin Li; Yingfei Liu; Shihao Wang; Fan Jia; Tiancai Wang; Lijin Han; Xiangyu Zhang
PPO-Clip Attains Global Optimality: Towards Deeper Understandings of	Nai-Chieh Huang; Ping-Chun Hsieh; Kuo-Hao Ho; I-Chen Wu
Preparing Lessons for Progressive Training on Language Models	Yu Pan; Ye Yuan; Yichun Yin; Jiabin Shi; Zenglin Xu; Ming Zhang; Lifeng Shang; Xin Jiang; Qun Liu
Towards Evidential and Class Separable Open Set Object Detection	Ruofan Wang; Rui-Wei Zhao; Xiaobo Zhang; Rui Feng



AACP: Aesthetics Assessment of Children's Paintings Based on Self-Supervised Learning	Shiqi Jiang; Ning Li; Chen Shi; Liping Guo; Changbo Wang; Chenhui Li
Exploring Equation as a Better Intermediate Meaning Representation for Numerical Reasoning of Large Language Models	Dingzirui Wang; Longxu Dou; Wenbin Zhang; Junyu Zeng; Wanxiang Che
MIND: Multi-Task Incremental Network Distillation	Francesco Pelosin; Jacopo Bonato; Luigi Sabetta; Alessandro Nicolosi
Towards Robust Image Stitching: An Adaptive Resistance Learning against Compatible Attacks	Zhiying Jiang; Xingyuan Li; Jinyuan Liu; Xin Fan; Risheng Liu
Independence of Irrelevant Alternatives under the Lens of Pairwise Distortion	Théo Delemazure; Jérôme Lang; Grzegorz Pierczyński
HyperEditor: Achieving Both Authenticity and Cross-Domain Capability in Image Editing via Hypernetworks	Hai Zhang; Chunwei Wu; Guitao Cao; Hailing Wang; Wenming Cao
Enhancing Low-Resource Relation Representations through Multi-View Decoupling	Chenghao Fan; Wei Wei; Xiaoye Qu; Zhenyi Lu; Xie Wenfeng; Yu Cheng; Danyang Chen
Racing Control Variable Genetic Programming for Symbolic Regression	Nan Jiang; Yexiang Xue
Mask-Homo: Pseudo Plane Mask-Guided Unsupervised Multi-Homography Estimation	Yasi Wang; Hong Liu; Chao Zhang; Lu Xu; Qiang Wang
Unifying Visual and Vision-Language Tracking via Contrastive Learning	Yinchao Ma; Yuyang Tang; Wenfei Yang; Tianzhu Zhang; Jinpeng Zhang; Mengxue Kang
Test-Time Domain Adaptation by Learning Domain-Aware Batch Normalization	Yanan Wu; Zhixiang Chi; Yang Wang; Konstantinos Plataniotis; Songhe Feng
SAMFlow: Eliminating Any Fragmentation in Optical Flow with Segment Anything Model	Shili Zhou; Ruian He; Weimin Tan; Bo Yan
A Comprehensive Analysis of the Effectiveness of Large Language Models as Automatic Dialogue Evaluators	Chen Zhang; Luis Fernando D'Haro; Yiming Chen; Malu Zhang; Haizhou Li
Understanding the Role of the Projector in Knowledge Distillation	Roy Miles; Krystian Mikolajczyk
Semantic Lens: Instance-Centric Semantic Alignment for Video Super-	Qi Tang; Yao Zhao; Meiqin Liu; Jian Jin; Chao Yao
Predicting Real-World Penny Auction Durations by Integrating Game Theory and Machine Learning	Yujia Wang; Haoran Yu
Federated Modality-Specific Encoders and Multimodal Anchors for Personalized Brain Tumor Segmentation	Qian Dai; Dong Wei; Hong Liu; Jinghan Sun; Liansheng Wang; Yefeng Zheng
Self-Paced Unified Representation Learning for Hierarchical Multi-Label Classification	Zixuan Yuan; Hao Liu; Haoyi Zhou; Denghui Zhang; Xiao Zhang; Hao Wang; Hui Xiong
HiHPQ: Hierarchical Hyperbolic Product Quantization for Unsupervised Image Retrieval	Zexuan Qiu; Jiahong Liu; Yankai Chen; Irwin King
VIGC: Visual Instruction Generation and Correction	Bin Wang; Fan Wu; Xiao Han; Jiahui Peng; Huaping Zhong; Pan Zhang; Xiaoyi Dong; Weijia Li; Wei Li; Jiaqi Wang; Conghui He
Few-Shot Learning from Augmented Label-Uncertain Queries in Bongard-HOI	Qinqian Lei; Bo Wang; Robby T. Tan
PG-LBO: Enhancing High-Dimensional Bayesian Optimization with Pseudo-Label and Gaussian Process Guidance	Taikai Chen; Yue Duan; Dong Li; Lei Qi; Yinghuan Shi; Yang Gao
Relaxed Stationary Distribution Correction Estimation for Improved Offline Policy Optimization	Woosung Kim; Donghyeon Ki; Byung-Jun Lee
Optimal Transport with Cyclic Symmetry	Shoichiro Takeda; Yasunori Akagi; Naoki Marumo; Kenta Niwa





Instrumental Variable Estimation for Causal Inference in Longitudinal Data with Time-Dependent Latent Confounders	Debo Cheng; Ziqi Xu; Jiuyong Li; Lin Liu; Jixue Liu; Wentao Gao; Thuc Duy Le
TextGT: A Double-View Graph Transformer on Text for Aspect-Based Sentiment Analysis	Shuo Yin; Guoqiang Zhong
Color Event Enhanced Single-Exposure HDR Imaging	Mengyao Cui; Zhigang Wang; Dong Wang; Bin Zhao; Xuelong Li
Efficient Deweather Mixture-of-Experts with Uncertainty-Aware Feature-Wise Linear Modulation	Rongyu Zhang; Yulin Luo; Jiaming Liu; Huanrui Yang; Zhen Dong; Denis Gudovskiy; Tomoyuki Okuno; Yohei Nakata; Kurt Keutzer; Yuan Du; Shanghang Zhang
Adversarial Socialbots Modeling Based on Structural Information Principles	Xianghua Zeng; Hao Peng; Angsheng Li
Wavelet Dynamic Selection Network for Inertial Sensor Signal Enhancement	Yifeng Wang; Yi Zhao
Adaptive Prompt Routing for Arbitrary Text Style Transfer with Pre-trained Language Models	Qingyi Liu; Jinghui Qin; Wenxuan Ye; Hao Mou; Yuxuan He; Keze Wang
Solving Spectrum Unmixing as a Multi-Task Bayesian Inverse Problem with Latent Factors for Endmember Variability	Dong Wu; Mingmin Chi; Xuan Zang; Bo Peng
Long-Tailed Learning as Multi-Objective Optimization	Weiqi Li; Fan Lyu; Fanhua Shang; Liang Wan; Wei Feng
DeS3: Adaptive Attention-Driven Self and Soft Shadow Removal Using ViT Similarity	Yeying Jin; Wei Ye; Wenhan Yang; Yuan Yuan; Robby T. Tan
Deep Linear Array Pushbroom Image Restoration: A Degradation Pipeline and Jitter-Aware Restoration Network	Zida Chen; Ziran Zhang; Haoying Li; Menghao Li; Yueting Chen; Qi Li; Huajun Feng; Zhihai Xu; Shiqi Chen
SQLdepth: Generalizable Self-Supervised Fine-Structured Monocular Depth Estimation	Youhong Wang; Yunji Liang; Hao Xu; Jiao Shaohui; Hongkai Yu
Leveraging Partial Symmetry for Multi-Agent Reinforcement Learning	Xin Yu; Rongye Shi; Pu Feng; Yongkai Tian; Simin Li; Shuhao Liao; Wenjun
Spotting the Unseen: Reciprocal Consensus Network Guided by Visual Archetypes	Wenbo Hu; Hongjian Zhan; Xinchun Ma; Yue Lu; Ching Suen
Learning Hierarchical Prompt with Structured Linguistic Knowledge for Vision-Language Models	Yubin Wang; Xinyang Jiang; De Cheng; Dongsheng Li; Cairong Zhao
DGL: Dynamic Global-Local Prompt Tuning for Text-Video Retrieval	Xiangpeng Yang; Linchao Zhu; Xiaohan Wang; Yi Yang
Multiobjective Lipschitz Bandits under Lexicographic Ordering	Bo Xue; Ji Cheng; Fei Liu; Yimu Wang; Qingfu Zhang
LF-ViT: Reducing Spatial Redundancy in Vision Transformer for Efficient Image Recognition	Youbing Hu; Yun Cheng; Anqi Lu; Zhiqiang Cao; Dawei Wei; Jie Liu; Zhijun Li
Diversity-Authenticity Co-constrained Stylization for Federated Domain Generalization in Person Re-identification	Fengxiang Yang; Zhun Zhong; Zhiming Luo; Yifan He; Shaozi Li; Nicu Sebe
Exploring Self- and Cross-Triplet Correlations for Human-Object Interaction Detection	Weibo Jiang; Weihong Ren; Jiandong Tian; Liangqiong Qu; Zhiyong Wang; Honghai Liu
AnomalyGPT: Detecting Industrial Anomalies Using Large Vision-Language Models	Zhaopeng Gu; Bingke Zhu; Guibo Zhu; Yingying Chen; Ming Tang; Jinqiao Wang
SwitchTab: Switched Autoencoders Are Effective Tabular Learners	Jing Wu; Suyao Chen; Qi Zhao; Renat Sergazinov; Chen Li; Shengjie Liu; Chongchao Zhao; Tianpei Xie; Hanqing Guo; Cheng Ji; Daniel Cociorva; Hakan Brunzell
Combinatorial CNN-Transformer Learning with Manifold Constraints for Semi-supervised Medical Image Segmentation	Huimin Huang; Yawen Huang; Shiao Xie; Lanfen Lin; Ruofeng Tong; Yen-Wei Chen; Yuexiang Li; Yefeng Zheng



Interactive Hyperparameter Optimization in Multi-Objective Problems via Preference Learning	Joseph Giovanelli; Alexander Tornede; Tanja Tornede; Marius Lindauer
FusionFormer: A Concise Unified Feature Fusion Transformer for 3D Pose Estimation	Yanlu Cai; Weizhong Zhang; Yuan Wu; Cheng Jin
Temporal Graph Contrastive Learning for Sequential Recommendation	Shengzhe Zhang; Liyi Chen; Chao Wang; Shuangli Li; Hui Xiong
Benchmarking Large Language Models in Retrieval-Augmented Generation	Jiawei Chen; Hongyu Lin; Xianpei Han; Le Sun
Backdoor Adjustment via Group Adaptation for Debiased Coupon Recommendations	Junpeng Fang; Gongduo Zhang; Qing Cui; Caizhi Tang; Lihong Gu; Longfei Li; Jinjie Gu; Jun Zhou
UniAP: Towards Universal Animal Perception in Vision via Few-Shot Learning	Meiqi Sun; Zhonghan Zhao; Wenhao Chai; Hanjun Luo; Shidong Cao; Yanting Zhang; Jenq-Neng Hwang; Gaoang Wang
Wavelet-Driven Spatiotemporal Predictive Learning: Bridging Frequency and Time Variations	Xuesong Nie; Yunfeng Yan; Siyuan Li; Cheng Tan; Xi CHEN; Haoyuan Jin; Zhihang Zhu; Stan Z. Li; Donglian Qi
WikiSQE: A Large-Scale Dataset for Sentence Quality Estimation in Wikipedia	Kenichiro Ando; Satoshi Sekine; Mamoru Komachi
Towards Diverse Perspective Learning with Selection over Multiple Temporal Poolings	Jihyeon Seong; Jungmin Kim; Jaesik Choi
Learning Ultrametric Trees for Optimal Transport Regression	Samantha Chen; Yusu Wang; Puoya Tabaghi
Controllable 3D Face Generation with Conditional Style Code Diffusion	Xiaolong Shen; Jianxin Ma; Chang Zhou; Zongxin Yang
Make RepVGG Greater Again: A Quantization-Aware Approach	Xiangxiang Chu; Liang Li; Bo Zhang
Exposing the Deception: Uncovering More Forgery Clues for Deepfake Detection	Zhongjie Ba; Qingyu Liu; Zhenguang Liu; Shuang Wu; Feng Lin; Li Lu; Kui Ren
On the Unstable Convergence Regime of Gradient Descent	Shuo Chen; Jiaying Peng; Xiaolong Li; Yao Zhao
Deep Hierarchical Video Compression	Ming Lu; Zhihao Duan; Fengqing Zhu; Zhan Ma
SurgicalSAM: Efficient Class Promptable Surgical Instrument Segmentation	Wenxi Yue; Jing Zhang; Kun Hu; Yong Xia; Jiebo Luo; Zhiyong Wang
Low-Latency Space-Time Supersampling for Real-Time Rendering	Ruian He; Shili Zhou; Yuqi Sun; Ri Cheng; Weimin Tan; Bo Yan
Learning Generalized Segmentation for Foggy-Scenes by Bi-directional Wavelet Guidance	Qi Bi; Shaodi You; Theo Gevers
Block-Level Goal Recognition Design	Tsz-Chiu Au
FoSp: Focus and Separation Network for Early Smoke Segmentation	Lujian Yao; Haitao Zhao; Jingchao Peng; Zhongze Wang; Kaijie Zhao
AMD: Autoregressive Motion Diffusion	Bo Han; Hao Peng; Minjing Dong; Yi Ren; Yixuan Shen; Chang Xu
Direct Amortized Likelihood Ratio Estimation	Adam Cobb; Brian Matejek; Daniel Elenius; Anirban Roy; Susmit Jha
Refined Characterizations of Approval-Based Committee Scoring Rules	Chris Dong; Patrick Lederer
Enhancing Cognitive Diagnosis Using Un-interacted Exercises: A Collaboration-Aware Mixed Sampling Approach	Haiping Ma; Changqian Wang; Hengshu Zhu; Shangshang Yang; Xiaoming Zhang; Xingyi Zhang
Training-Free Quantum Architecture Search	Zhimin He; Maijie Deng; Shenggen Zheng; Lvzhou Li; Haozhen Situ
Unsupervised Gene-Cell Collective Representation Learning with Optimal Transport	Jixiang Yu; Nanjun Chen; Ming Gao; Xiangtao Li; Ka-Chun Wong
Text Image Inpainting via Global Structure-Guided Diffusion Models	Shipeng Zhu; Pengfei Fang; Chenjie Zhu; Zuoyan Zhao; Qiang Xu; Hui Xue
Norm Tweaking: High-Performance Low-Bit Quantization of Large Language Models	Liang Li; Qingyuan Li; Bo Zhang; Xiangxiang Chu
On the Robustness of Neural-Enhanced Video Streaming against Adversarial Attacks	Qihua Zhou; Jingcai Guo; Song Guo; Ruibin Li; Jie Zhang; Bingjie Wang; Zhenda Xu





Hawkes-Enhanced Spatial-Temporal Hypergraph Contrastive Learning Based on Criminal Correlations	Ke Liang; Sihang Zhou; Meng Liu; Yue Liu; Wenxuan Tu; Yi Zhang; Liming Fang; Zhe Liu; Xinwang Liu
Participation Incentives in Approval-Based Committee Elections	Martin Bullinger; Chris Dong; Patrick Lederer; Clara Mehler
MGNet: Learning Correspondences via Multiple Graphs	Dai Luanyuan; Xiaoyu Du; Hanwang Zhang; Jinhui Tang
Revealing the Proximate Long-Tail Distribution in Compositional Zero-Shot Learning	Chenyi Jiang; Haofeng Zhang
MobileInst: Video Instance Segmentation on the Mobile	Renhong Zhang; Tianheng Cheng; Shusheng Yang; Haoyi Jiang; Shuai Zhang; Jiancheng Lyu; Xin Li; Xiaowen Ying; Dashan Gao; Wenyu Liu; Kai Ye; Tiejun Chen; Hua Wei; Liang Zhan
Uncertainty Regularized Evidential Regression	Zheng Xing; Weibing Zhao
Unsupervised Action Segmentation via Fast Learning of Semantically Consistent Actoms	
Reinforced Adaptive Knowledge Learning for Multimodal Fake News Detection	Litian Zhang; Xiaoming Zhang; Chaozhao Li; Linfeng Han; Ziyi Zhou; Feiran Huang
Continuous Piecewise-Affine Based Motion Model for Image Animation	Hexiang Wang; Fengqi Liu; Qianyu Zhou; Ran Yi; Xin Tan; Lizhuang Ma
EulerMormer: Robust Eulerian Motion Magnification via Dynamic Filtering within Transformer	Fei Wang; Dan Guo; Kun Li; Meng Wang
Learning Diffusions under Uncertainty	Hao Huang; Qian Yan; Keqi Han; Ting Gan; Jiawei Jiang; Quanqing Xu; Chuanhui Yang
Robustness-Guided Image Synthesis for Data-Free Quantization	Jianhong Bai; Yuchen Yang; Huanpeng Chu; Hualiang Wang; Zuozhu Liu; Ruizhe Chen; Xiaoxuan He; Lianrui Mu; Chengfei Cai; Haoji Hu
RadOcc: Learning Cross-Modality Occupancy Knowledge through Rendering Assisted Distillation	Haiming Zhang; Xu Yan; Dongfeng Bai; Jiantao Gao; Pan Wang; Liu Bingbing; Shuguang Cui; Zhen Li
GroundVLP: Harnessing Zero-Shot Visual Grounding from Vision-Language Pre-training and Open-Vocabulary Object Detection	Haozhan Shen; Tiancheng Zhao; Mingwei Zhu; Jianwei Yin
Enhanced Fine-Grained Motion Diffusion for Text-Driven Human Motion Synthesis	Dong Wei; Xiaoning Sun; Huaijiang Sun; Shengxiang Hu; Bin Li; Weiqing Li; Jianfeng Lu
Relational Distant Supervision for Image Captioning without Image-Text Pairs	Yayun Qi; Wentian Zhao; Xinxiao Wu
TurboSVM-FL: Boosting Federated Learning through SVM Aggregation for Lazy Clients	Mengdi Wang; Anna Bodonhelyi; Efe Bozkir; Enkelejda Kasneci
Dual-Window Multiscale Transformer for Hyperspectral Snapshot Compressive Imaging	Fulin Luo; Xi Chen; Xiuwen Gong; Weiwen Wu; Tan Guo
How to Trade Off the Quantity and Capacity of Teacher Ensemble: Learning Categorical Distribution to Stochastically Employ a Teacher for Distillation	Zixiang Ding; Guoqing Jiang; Shuai Zhang; Lin Guo; Wei Lin
DifAttack: Query-Efficient Black-Box Adversarial Attack via Disentangled Feature Space	Jun Liu; Jiantao Zhou; Jiandian Zeng; Jinyu Tian
Near-Optimal Resilient Aggregation Rules for Distributed Learning Using 1-Center and 1-Mean Clustering with Outliers	Yuhao Yi; Ronghui You; Hong Liu; Changxin Liu; Yuan Wang; Jiancheng Lv
Customizing Language Model Responses with Contrastive In-Context Learning	Xiang Gao; Kamalika Das
Balancing Humans and Machines: A Study on Integration Scale and Its Impact on Collaborative Performance	Rui Zou; Sannyuya Liu; Yawei Luo; Yaqi Liu; Jintian Feng; MengQi Wei; Jianwen Sun



FlightBERT++: A Non-autoregressive Multi-Horizon Flight Trajectory Prediction Framework	Dongyue Guo; Zheng Zhang; Zhen Yan; Jianwei Zhang; Yi Lin
When to Grow? A Fitting Risk-Aware Policy for Layer Growing in Deep Neural Networks	Haihang Wu; Wei Wang; Tamasha Malepathirana; Damith Senanayake; Denny Oetomo; Saman Halgamuge
A Cross-View Hierarchical Graph Learning Hypernetwork for Skill Demand-Supply Joint Prediction	Wenshuo Chao; Zhaopeng Qiu; Likang Wu; Zhuoning Guo; Zhi Zheng; Hengshu Zhu; Hao Liu
Instance-Aware Multi-Camera 3D Object Detection with Structural Priors Mining and Self-Boosting Learning	Yang Jiao; Zequn Jie; Shaoxiang Chen; Lechao Cheng; Jingjing Chen; Lin Ma; Yu-Gang Jiang
Agile Multi-Source-Free Domain Adaptation	Xinyao Li; Jingjing Li; Fengling Li; Lei Zhu; Ke Lu
Non-exemplar Domain Incremental Object Detection via Learning Domain Bias	Xiang Song; Yuhang He; Songlin Dong; Yihong Gong
QuerySum: A Multi-Document Query-Focused Summarization Dataset Augmented with Similar Query Clusters	Yushan Liu; Zili Wang; Ruifeng Yuan
Noise-Free Optimization in Early Training Steps for Image Super-resolution	MinKyu Lee; Jae-Pil Heo
Backdoor Attacks via Machine Unlearning	Zihao Liu; Tianhao Wang; Mengdi Huai; Chenglin Miao
AvatarVerse: High-Quality & Stable 3D Avatar Creation from Text and Pose	Huichao Zhang; Bowen Chen; Hao Yang; Liao Qu; Xu Wang; Li Chen; Chao Long; Feida Zhu; Daniel Du; Min Zheng
One Step Closer to Unbiased Aleatoric Uncertainty Estimation	Wang Zhang; Ziwen Martin Ma; Subhro Das; Lily Weng; Alexandre Megretsky; Luca Daniel; Lam Nguyen
Make Prompts Adaptable: Bayesian Modeling for Vision-Language Prompt Learning with Data-Dependent Prior	Youngjae Cho; HeeSun Bae; Seungjae Shin; Yeo Dong Youn; Weonyoung Joo; Il-Chul Moon
Contributing Dimension Structure of Deep Feature for Coreset Selection	Zhijing Wan; Zhixiang Wang; Yuran Wang; Zheng Wang; Hongyuan Zhu; Shin'ichi Satoh
Joint Learning Neuronal Skeleton and Brain Circuit Topology with Permutation Invariant Encoders for Neuron Classification	Minghui Liao; Guojia Wan; Bo Du
A Label Disambiguation-Based Multimodal Massive Multiple Instance Learning Approach for Immune Repertoire Classification	Fan Xu; Yu Zhao; Bingzhe Wu; Yueshan Huang; Qin Ren; Yang Xiao; Bing He; Jie Zheng; Jianhua Yao
ColNeRF: Collaboration for Generalizable Sparse Input Neural Radiance Field	Zhangkai Ni; Peiqi Yang; Wenhan Yang; Hanli Wang; Lin Ma; Sam Kwong
Reverse Multi-Choice Dialogue Commonsense Inference with Graph-of-	Li Zheng; Hao Fei; Fei Li; Bobo Li; Lizi Liao; Donghong Ji; Chong Teng
Can Large Language Models Serve as Rational Players in Game Theory? A Systematic Analysis	Caoyun Fan; Jindou Chen; Yaohui Jin; Hao He
Say Anything with Any Style	Shuai Tan; Bin Ji; Yu Ding; Ye Pan
Inducing Clusters Deep Kernel Gaussian Process for Longitudinal Data	Junjie Liang; Weijieying Ren; Sahar Hanifi; Vasant Honavar
Learning Deformable Hypothesis Sampling for Accurate PatchMatch Multi-View Stereo	Hongjie Li; Yao Guo; Xianwei Zheng; Hanjiang Xiong
Semi-supervised Open-World Object Detection	Sahal Shaji Mullappilly; Abhishek Singh Gehlot; Rao Muhammad Anwer; Fahad Shahbaz Khan; Hisham Cholakkal
Decoupling Representation and Knowledge for Few-Shot Intent Classification and Slot Filling	Jie Han; Yixiong Zou; Haozhao Wang; Jun Wang; Wei Liu; Yao Wu; Tao Zhang; Ruixuan Li
QI-IRA: Quantum-Inspired Interactive Ranking Aggregation for Person Re-identification	Chunyu Hu; Hong Zhang; Chao Liang; Hao Huang



Boosting Adversarial Transferability across Model Genus by Deformation-Constrained Warping	Qinliang Lin; Cheng Luo; Zenghao Niu; Xilin He; Weicheng Xie; Yuanbo Hou; Linlin Shen; Siyang Song
What Makes Good Collaborative Views? Contrastive Mutual Information Maximization for Multi-Agent Perception	Wanfang Su; Lixing Chen; Yang Bai; Xi Lin; Gaolei Li; Zhe Qu; Pan Zhou
Towards Compact 3D Representations via Point Feature Enhancement Masked Autoencoders	Yaohua Zha; Huizhen Ji; Jinmin Li; Rongsheng Li; Tao Dai; Bin Chen; Zhi Wang; Shu-Tao Xia
UMIE: Unified Multimodal Information Extraction with Instruction Tuning	Lin Sun; Kai Zhang; Qingyuan Li; Renze Lou
Exploring Base-Class Suppression with Prior Guidance for Bias-Free One-Shot Object Detection	Wenwen Zhang; Yun Hu; Hangguan Shan; Eryun Liu
Generalizing across Temporal Domains with Koopman Operators	QIUHAO Zeng; Wei Wang; Fan Zhou; Gezheng Xu; Ruizhi Pu; Changjian Shui; Christian Gagné; Shichun Yang; Charles Ling; Boyu Wang
eTag: Class-Incremental Learning via Embedding Distillation and Task-Oriented Generation	Libo Huang; Yan Zeng; Chuanguang Yang; Zhulin An; Boyu Diao; Yongjun Xu
Progressively Knowledge Distillation via Re-parameterizing Diffusion Reverse Process	Xufeng Yao; Fanbin Lu; Yuechen Zhang; Xinyun Zhang; Wenqian Zhao; Bei Yu
A Generalized Shuffle Framework for Privacy Amplification: Strengthening Privacy Guarantees and Enhancing Utility	Chen E; Yang Cao; Ge Yifei
Point-to-Spike Residual Learning for Energy-Efficient 3D Point Cloud Classification	Qiaoyun Wu; Quanyao Zhang; Chunyu Tan; Yun Zhou; Changyin Sun
SALSA: Semantically-Aware Latent Space Autoencoder	Kathryn Kirchoff; Travis Maxfield; Alexander Tropsha; Shawn Gomez
SAT-Based Techniques for Lexicographically Smallest Finite Models	Mikoláš Janota; Chaiwah Chow; João Araújo; Michael Codish; Petr Vojtěchovský
Joint Demosaicing and Denoising for Spike Camera	Yanchen Dong; Ruiqin Xiong; Jing Zhao; Jian Zhang; Xiaopeng Fan; Shuyuan Zhu; Tiejun Huang
Does Few-Shot Learning Suffer from Backdoor Attacks?	Xinwei Liu; Xiaojun Jia; Jindong Gu; Yuan Xun; Siyuan Liang; Xiaochun Cao
O <sup>2</sup> -Recon: Completing 3D Reconstruction of Occluded Objects in the Scene with a Pre-trained 2D Diffusion Model	Yubin Hu; Sheng Ye; Wang Zhao; Matthieu Lin; Yuze He; Yu-Hui Wen; Ying He; Yong-Jin Liu
Learning from Polar Representation: An Extreme-Adaptive Model for Long-Term Time Series Forecasting	Yanhong Li; Jack Xu; David Anastasiu
TF-CLIP: Learning Text-Free CLIP for Video-Based Person Re-identification	Chenyang Yu; Xuehu Liu; Yingquan Wang; Pingping Zhang; Huchuan Lu
Suppressing Uncertainty in Gaze Estimation	Shijing Wang; Yaping Huang
Rethinking Peculiar Images by Diffusion Models: Revealing Local Minima's	Jinhyeok Jang; Chan-Hyun Youn; Minsu Jeon; Changha Lee
RoPDA: Robust Prompt-Based Data Augmentation for Low-Resource Named Entity Recognition	Sihan Song; Furao Shen; Jian Zhao
Composite Sketch+Text Queries for Retrieving Objects with Elusive Names and Complex Interactions	Prajwal Gatti; Kshitij Parikh; Dhriti Paul; Manish Gupta; Anand Mishra
How to Protect Copyright Data in Optimization of Large Language Models?	Timothy Chu; Zhao Song; Chiwon Yang
Dialogue for Prompting: A Policy-Gradient-Based Discrete Prompt Generation for Few-Shot Learning	Chengzhengxu Li; Xiaoming Liu; Yichen Wang; Duyi Li; Yu Lan; Chao Shen
Runtime vs. Extracted Proof Size: An Exponential Gap for CDCL on QBFs	Olaf Beyersdorff; Benjamin Böhm; Meena Mahajan
Transformer as Linear Expansion of LearnGene	Shiyu Xia; Miaosen Zhang; Xu Yang; Ruiming Chen; Haokun Chen; Xin Geng



Transitivity-Preserving Graph Representation Learning for Bridging Local Connectivity and Role-Based Similarity	Van Thuy Hoang; O-Joun Lee
A Diffusion-Based Framework for Multi-Class Anomaly Detection	Haoyang He; Jiangning Zhang; Hongxu Chen; Xuhai Chen; Zhishan Li; Xu Chen; Yabiao Wang; Chengjie Wang; Lei Xie
Cross-Sentence Gloss Consistency for Continuous Sign Language	Qi Rao; Ke Sun; Xiaohan Wang; Qi Wang; Bang Zhang
Online Restless Multi-Armed Bandits with Long-Term Fairness Constraints	Shufan Wang; Guojun Xiong; Jian Li
M2SD: Multiple Mixing Self-Distillation for Few-Shot Class-Incremental	Jinhao Lin; Ziheng Wu; Weifeng Lin; Jun Huang; RongHua Luo
S3A: Towards Realistic Zero-Shot Classification via Self Structural Semantic Alignment	Sheng Zhang; Muzammal Naseer; Guangyi Chen; Zhiqiang Shen; Salman Khan; Kun Zhang; Fahad Shahbaz Khan
A Provably Accurate Randomized Sampling Algorithm for Logistic Regression	Agniva Chowdhury; Pradeep Ramuhalli
Towards Squeezing-Averse Virtual Try-On via Sequential Deformation	Sang-Heon Shim; Jiwoo Chung; Jae-Pil Heo
BARET: Balanced Attention Based Real Image Editing Driven by Target-Text Inversion	Yuming Qiao; Fanyi Wang; Jingwen Su; Yanhao Zhang; Yunjie Yu; Siyu Wu; Guo-Jun Qi
HORIZON: High-Resolution Semantically Controlled Panorama Synthesis	Kun Yan; Lei Ji; Jian Liang; Chenfei Wu; Ming Zhou; Nan Duan; Shuai Ma
Double-Layer Hybrid-Label Identification Feature Selection for Multi-View Multi-Label Learning	Pingting Hao; Kunpeng Liu; Wanfu Gao
Open-Set Facial Expression Recognition	Yuhang Zhang; Yue Yao; Xuannan Liu; Lixiong Qin; Wenjing Wang; Weihong Deng
SimCS: Simulation for Domain Incremental Online Continual Segmentation	Motasem Alfarra; Zhipeng Cai; Adel Bibi; Bernard Ghanem; Matthias Müller
Variance-Insensitive and Target-Preserving Mask Refinement for Interactive Image Segmentation	Chaowei Fang; Ziyin Zhou; Junye Chen; Hanjing Su; Qingyao Wu; Guanbin Li
DCLP: Neural Architecture Predictor with Curriculum Contrastive Learning	Shenghe Zheng; Hongzhi Wang; Tianyu Mu
UniCATS: A Unified Context-Aware Text-to-Speech Framework with Contextual VQ-Diffusion and Vocoding	Chenpeng Du; Yiwei Guo; Feiyu Shen; Zhijun Liu; Zheng Liang; Xie Chen; Shuai Wang; Hui Zhang; Kai Yu
Homophily-Related: Adaptive Hybrid Graph Filter for Multi-View Graph Clustering	Zichen Wen; Yawen Ling; Yazhou Ren; TianYi Wu; Jianpeng Chen; Xiaorong Pu; Zhifeng Hao; Lifang He
NaRuto: Automatically Acquiring Planning Models from Narrative Texts	Ruiqi Li; Leyang Cui; Songtuan Lin; Patrik Haslum
GLDL: Graph Label Distribution Learning	Yufei Jin; Richard Gao; Yi He; Xingquan Zhu
CoSTA: End-to-End Comprehensive Space-Time Entanglement for Spatio-Temporal Video Grounding	Yaoyuan Liang; Xiao Liang; Yansong Tang; Zhao Yang; Ziran Li; Jingang Wang; Wenbo Ding; Shao-Lun Huang
FedTGP: Trainable Global Prototypes with Adaptive-Margin-Enhanced Contrastive Learning for Data and Model Heterogeneity in Federated Learning	Jianqing Zhang; Yang Liu; Yang Hua; Jian Cao
Generalized Planning in PDDL Domains with Pretrained Large Language Models	Tom Silver; Soham Dan; Kavitha Srinivas; Joshua Tenenbaum; Leslie Kaelbling; Michael Katz
Combining Multiple Supervision for Robust Zero-Shot Dense Retrieval	Yan Fang; Qingyao Ai; Jingtao Zhan; Yiqun LIU; Xiaolong Wu; Zhao Cao
DPA-P2PNet: Deformable Proposal-Aware P2PNet for Accurate Point-Based Cell Detection	Zhongyi Shui; Sunyi Zheng; Chenglu Zhu; Shichuan Zhang; Xiaoxuan YU; Honglin Li; Jingxiong Li; Pingyi Chen; Lin Yang
SEA-GWNN: Simple and Effective Adaptive Graph Wavelet Neural Network	Swakshar Deb; Sejuti Rahman; Shafin Rahman
Navigating Open Set Scenarios for Skeleton-Based Action Recognition	Kunyu Peng; Cheng Yin; Junwei Zheng; Ruiping Liu; David Schneider; Jiaming Zhang; Kailun Yang; M. Saquib Sarfraz; Rainer Stiefelhausen; Alina Roitberg



E2HQV: High-Quality Video Generation from Event Camera via Theory-Inspired Model-Aided Deep Learning	Qiang Qu; Yiran Shen; Xiaoming Chen; Yuk Ying Chung; Tongliang Liu
Parallel Empirical Evaluations: Resilience despite Concurrency	Johannes Fichte; Tobias Geibinger; Markus Hecher; Matthias Schlögel
ANEDL: Adaptive Negative Evidential Deep Learning for Open-Set Semi-supervised Learning	Yang YU; Danruo Deng; Furui Liu; Dou Qi; Yueming Jin; Guangyong Chen; Pheng-Ann Heng
Test-Time Personalization with Meta Prompt for Gaze Estimation	Huan Liu; Julia Qi; Zhenhao Li; Mohammad Hassanpour; Yang Wang; Konstantinos Plataniotis; Yuanhao Yu
IINet: Implicit Intra-inter Information Fusion for Real-Time Stereo Matching	Ximeng Li; Chen Zhang; Wanjuan Su; Wenbing Tao
ALISON: Fast and Effective Stylometric Authorship Obfuscation	Eric Xing; Saranya Venkatraman; Thai Le; Dongwon Lee
LION: Implicit Vision Prompt Tuning	Haixin Wang; Jianlong Chang; Yihang Zhai; Xiao Luo; Jinan Sun; Zhouchen Lin; Qi Tian
Diffusion Language-Shapelets for Semi-supervised Time-Series Classification	Zhen Liu; Wenbin Pei; Disen Lan; Qianli Ma
Deep Structural Knowledge Exploitation and Synergy for Estimating Node Importance Value on Heterogeneous Information Networks	Yankai Chen; Yixiang Fang; Qiongyan Wang; Xin Cao; Irwin King
Bi-ViT: Pushing the Limit of Vision Transformer Quantization	Yanjing Li; Sheng Xu; Mingbao Lin; Xianbin Cao; Chuanjian Liu; Xiao Sun; Baochang Zhang
Almost Envy-Free Allocations of Indivisible Goods or Chores with Entitlements	Max Springer; MohammadTaghi Hajiaghayi; Hadi Yami
Unraveling Batch Normalization for Realistic Test-Time Adaptation	Zixian Su; Jingwei Guo; Kai Yao; Xi Yang; Qiufeng Wang; Kaizhu Huang
Inspecting Prediction Confidence for Detecting Black-Box Backdoor Attacks	Tong Wang; Yuan Yao; Feng Xu; Miao Xu; Shengwei An; Ting Wang
Layer Compression of Deep Networks with Straight Flows	Chengyue Gong; Xiaocong Du; Bhargav Bhushanam; Lemeng Wu; Xingchao Liu; Dhruv Choudhary; Arun Kejariwal; Qiang Liu
Spherical Pseudo-Cylindrical Representation for Omnidirectional Image Super-resolution	Qing Cai; Mu Li; Dongwei Ren; Jun Lyu; Haiyong Zheng; Junyu Dong; Yee Hong Yang
Task-Free Dynamic Sparse Vision Transformer for Continual Learning	Fei Ye; Adrian Bors
TriSampler: A Better Negative Sampling Principle for Dense Retrieval	Zhen Yang; Zhou Shao; Yuxiao Dong; Jie Tang
SuperJunction: Learning-Based Junction Detection for Retinal Image Registration	Yu Wang; Xiaoye Wang; Zaiwang Gu; Weide Liu; Wee Siong Ng; Weimin Huang; Jun Cheng
Towards More Faithful Natural Language Explanation Using Multi-Level Contrastive Learning in VQA	Chengen Lai; Shengli Song; Shiqi Meng; Jingyang Li; Sitong Yan; Guangneng Hu
NeRF-LiDAR: Generating Realistic LiDAR Point Clouds with Neural Radiance Fields	Junge Zhang; Feihu Zhang; Shaochen Kuang; Li Zhang
A Pre-convolved Representation for Plug-and-Play Neural Illumination Fields	Yiyu Zhuang; Qi Zhang; Xuan Wang; Hao Zhu; Ying Feng; Xiaoyu Li; Ying Shan; Xun Cao
ArtBank: Artistic Style Transfer with Pre-trained Diffusion Model and Implicit Style Prompt Bank	Zhanjie Zhang; Quanwei Zhang; Wei Xing; Guangyuan Li; Lei Zhao; Jiakai Sun; Zehua Lan; Junsheng Luan; Yiling Huang; Huaizhong Lin
Complexity of Neural Network Training and ETR: Extensions with Effectively Continuous Functions	Teemu Hankala; Miika Hannula; Juha Kontinen; Jonni Virtama
Identification of Causal Structure with Latent Variables Based on Higher Order Cumulants	Wei Chen; Zhiyi Huang; Ruichu Cai; Zhifeng Hao; Kun Zhang
Approximation Algorithms for Preference Aggregation Using CP-Nets	Abu Mohammad Hammad Ali; Boting Yang; Sandra Zilles



Uncertainty Quantification for Forward and Inverse Problems of PDEs via Latent Global Evolution	Tailin Wu; Willie Neiswanger; Hongtao Zheng; Stefano Ermon; Jure Leskovec
Learning from History: Task-agnostic Model Contrastive Learning for Image Restoration	Gang Wu; Junjun Jiang; Kui Jiang; Xianming Liu
Engineering an Exact Pseudo-Boolean Model Counter	Suwei Yang; Kuldeep S Meel
HR-Pro: Point-Supervised Temporal Action Localization via Hierarchical Reliability Propagation	Huaxin Zhang; Xiang Wang; Xiaohao Xu; Zhiwu Qing; Changxin Gao; Nong Sang
X4D-SceneFormer: Enhanced Scene Understanding on 4D Point Cloud Videos through Cross-Modal Knowledge Transfer	Linglin Jing; Ying Xue; Xu Yan; Chaoda Zheng; Dong Wang; Ruimao Zhang; Zhigang Wang; Hui Fang; Bin Zhao; Zhen Li
When Do Program-of-Thought Works for Reasoning?	Zhen Bi; Ningyu Zhang; Yinuo Jiang; Shumin Deng; Guozhou Zheng; Huajun Chen
T-SciQ: Teaching Multimodal Chain-of-Thought Reasoning via Large Language Model Signals for Science Question Answering	Lei Wang; Yi Hu; Jiabang He; Xing Xu; Ning Liu; Hui Liu; Heng Tao Shen
Monocular 3D Hand Mesh Recovery via Dual Noise Estimation	Hanhui Li; Xiaojian Lin; Xuan Huang; Zejun Yang; Zhisheng Wang; Xiaodan Liang
Effective Comparative Prototype Hashing for Unsupervised Domain	Hui Cui; Lihai Zhao; Fengling Li; Lei Zhu; Xiaohui Han; Jingjing Li
Neuromorphic Event Signal-Driven Network for Video De-raining	Chengjie Ge; Xueyang Fu; Peng He; Kunyu Wang; Chengzhi Cao; Zheng-Jun Zha
Navigating Real-World Partial Label Learning: Unveiling Fine-Grained Images with Attributes	Haoran Jiang; Zhihao Sun; Yingjie Tian
Spatial-Related Sensors Matters: 3D Human Motion Reconstruction Assisted with Textual Semantics	Yang Xueyuan; Chao Yao; Xiaojuan Ban
G-NAS: Generalizable Neural Architecture Search for Single Domain Generalization Object Detection	Fan Wu; Jinling Gao; Lanqing Hong; Xinbing Wang; Chenghu Zhou; Nanyang Ye
Multi-Domain Deep Learning from a Multi-View Perspective for Cross-Border E-commerce Search	Yiqian Zhang; Yinfu Feng; Wen-Ji Zhou; Yunan Ye; Min Tan; Rong Xiao; Haihong Tang; Jiajun Ding; Jun Yu
Binding-Adaptive Diffusion Models for Structure-Based Drug Design	Zhilin Huang; Ling Yang; Zaixi Zhang; Xiangxin Zhou; Yu Bao; Xiawu Zheng; Yuwei Yang; Yu Wang; Wenming Yang
Gaussian Mixture Proposals with Pull-Push Learning Scheme to Capture Diverse Events for Weakly Supervised Temporal Video Grounding	Sunoh Kim; Jungchan Cho; Joonsang Yu; YoungJoon Yoo; Jin Young Choi
Knowledge Graph Prompting for Multi-Document Question Answering	Yu Wang; Nedim Lipka; Ryan Rossi; Alexa Siu; Ruiyi Zhang; Tyler Derr
Multiscale Low-Frequency Memory Network for Improved Feature Extraction in Convolutional Neural Networks	Fuzhi Wu; Jiasong Wu; Youyong Kong; Chunfeng Yang; Guanyu Yang; Huazhong Shu; Guy Carrault; Lotfi Senhadji
Minimum Coverage Sets for Training Robust Ad Hoc Teamwork Agents	Muhammad Rahman; Jiaxun Cui; Peter Stone
Learning Multi-Object Positional Relationships via Emergent Communication	Yicheng Feng; Boshi An; Zongqing Lu
Fusion-Vital: Video-RF Fusion Transformer for Advanced Remote Physiological Measurement	Jae-Ho Choi; Kibong Kang; Kyung-Tae Kim
Spiking NeRF: Representing the Real-World Geometry by a Discontinuous Representation	Zhanfeng Liao; Yan Liu; Qian Zheng; Gang Pan
Swift-Mapping: Online Neural Implicit Dense Mapping in Urban Scenes	Ke Wu; Kaizhao Zhang; Mingzhe Gao; Jieru Zhao; Zhongxue Gan; Wenchao Ding





Three Heads Are Better than One: Improving Cross-Domain NER with Progressive Decomposed Network	Xuming Hu; Zhaochen Hong; Yong Jiang; ZhiChao Lin; Xiaobin Wang; Pengjun Xie; Philip Yu
Probabilistic Neural Circuits	Pedro Zuidberg Dos Martires
ExpCLIP: Bridging Text and Facial Expressions via Semantic Alignment	Yicheng Zhong; Huawei Wei; Peiji Yang; Zhisheng Wang
Improving Diffusion-Based Image Restoration with Error Contraction and Error Correction	Qiqi Bao; Zheng Hui; Rui Zhu; Peiran Ren; Xuansong Xie; Wenming Yang
Geometry-Guided Domain Generalization for Monocular 3D Object Detection	Fan Yang; Hui Chen; Yuwei He; Sicheng Zhao; Chenghao Zhang; Kai Ni; Guiguang Ding
Learning the Topology and Behavior of Discrete Dynamical Systems	Zirou Qiu; Abhijin Adiga; Madhav Marathe; S Ravi; Daniel Rosenkrantz; Richard Stearns; Anil Vullikanti
Any-Stereo: Arbitrary Scale Disparity Estimation for Iterative Stereo Matching	Zhaohuai Liang; Changhe Li
Robust Policy Learning via Offline Skill Diffusion	Woo Kyung Kim; Minjong Yoo; Honguk Woo
EAN: An Efficient Attention Module Guided by Normalization for Deep Neural Networks	Jiafeng Li; Zelin Li; Ying Wen
Learning to Optimize Permutation Flow Shop Scheduling via Graph-Based Imitation Learning	Longkang Li; Siyuan Liang; Zihao Zhu; Chris Ding; Hongyuan Zha; Baoyuan Wu
FG-EmoTalk: Talking Head Video Generation with Fine-Grained Controllable Facial Expressions	Zhaoxu Sun; Yuze Xuan; Fang Liu; Yang Xiang
Spectral Prompt Tuning: Unveiling Unseen Classes for Zero-Shot Semantic Segmentation	Wenhao Xu; Rongtao Xu; Changwei Wang; Shibiao Xu; Li Guo; Man Zhang; Xiaopeng Zhang
Decoupled Contrastive Learning for Long-Tailed Recognition	Shiyu Xuan; Shiliang Zhang
The Expected Loss of Preconditioned Langevin Dynamics Reveals the Hessian Rank	Amitay Bar; Rotem Mulayoff; Tomer Michaeli; Ronen Talmon
Improving Robustness for Joint Optimization of Camera Pose and Decomposed Low-Rank Tensorial Radiance Fields	BOYU Chen; Wei-Chen Chiu; Yu-Lun Liu
Exact Algorithms and Lowerbounds for Multiagent Path Finding: Power of Treelike Topology	Foivos Fioravantes; Dušan Knop; Jan Křišťan; Nikolaos Melissinos; Michal Opler
SiMA-Hand: Boosting 3D Hand-Mesh Reconstruction by Single-to-Multi-View Adaptation	Yinqiao Wang; Hao Xu; Pheng-Ann Heng; Chi-Wing Fu
Context-I2W: Mapping Images to Context-Dependent Words for Accurate Zero-Shot Composited Image Retrieval	Yuanmin Tang; Jing Yu; Keke Gai; Zhuang JiaMin; Gang Xiong; Yue Hu; Qi Wu
Modeling Continuous Motion for 3D Point Cloud Object Tracking	Zhipeng Luo; Gongjie Zhang; Changqing Zhou; Zhonghua Wu; Qingyi Tao; Lewei Lu; Shijian Lu
Symbolic Regression Enhanced Decision Trees for Classification Tasks	Kei Sen Fong; Mehul Motani
Trash to Treasure: Low-Light Object Detection via Decomposition-and-Aggregation	Xiaohan Cui; Long Ma; Tengyu Ma; Jinyuan Liu; Xin Fan; Risheng Liu
DocMSU: A Comprehensive Benchmark for Document-Level Multimodal Sarcasm Understanding	Hang Du; Guoshun Nan; Sicheng Zhang; Binzhu Xie; Junrui Xu; Hehe Fan; Qimei Cui; Xiaofeng Tao; Xudong Jiang
VadCLIP: Adapting Vision-Language Models for Weakly Supervised Video Anomaly Detection	Peng Wu; Xuerong Zhou; Guansong Pang; Lingru Zhou; Qingsen Yan; Peng Wang; Yanning Zhang



Conformal Crystal Graph Transformer with Robust Encoding of Periodic Invariance	Yingheng Wang; Shufeng Kong; John Gregoire; Carla Gomes
CI-STHPAN: Pre-trained Attention Network for Stock Selection with Channel-Independent Spatio-Temporal Hypergraph	Hongjie Xia; Huijie Ao; Long Li; Yu Liu; Sen Liu; Guangnan Ye; Hongfeng Chai
Effective Causal Discovery under Identifiable Heteroscedastic Noise Model	Naiyu Yin; Tian Gao; Yue Yu; Qiang Ji
Harnessing Edge Information for Improved Robustness in Vision Transformers	Yanxi Li; Chengbin Du; Chang Xu
Learning from Failure: Improving Meeting Summarization without Good	Ke Wang; Xiutian Zhao; Wei Peng
Deep Active Learning with Noise Stability	Xingjian Li; Pengkun Yang; Yangcheng Gu; Xueying ZHAN; Tianyang Wang; Min Xu; Chengzhong Xu
Stable Unlearnable Example: Enhancing the Robustness of Unlearnable Examples via Stable Error-Minimizing Noise	Yixin Liu; Kaidi Xu; Xun Chen; Lichao Sun
Adaptive FSS: A Novel Few-Shot Segmentation Framework via Prototype Enhancement	Jing Wang; Jiangyun Li; Chen Chen; Yisi Zhang; Haoran Shen; Tianxiang zhang
Efficient Lightweight Image Denoising with Triple Attention Transformer	Yubo Zhou; Jin Lin; Fangchen Ye; Yanyun Qu; Yuan Xie
Attention-Induced Embedding Imputation for Incomplete Multi-View Partial Multi-Label Classification	Chengliang Liu; Jia Jinlong; Jie Wen; Yabo Liu; Xiaoling Luo; Chao Huang; Yong Xu
Sample-Level Cross-View Similarity Learning for Incomplete Multi-View Clustering	Suyuan Liu; Junpu Zhang; Yi Wen; Xihong Yang; Siwei Wang; Yi Zhang; En Zhu; Chang Tang; Long Zhao; Xinwang Liu
Gaze Target Detection by Merging Human Attention and Activity Cues	Yaokun Yang; Yihan Yin; Feng Lu
SHoP: A Deep Learning Framework for Solving High-Order Partial Differential Equations	Tingxiong Xiao; Runzhao Yang; Yuxiao Cheng; Jinli Suo
ResMatch: Residual Attention Learning for Feature Matching	Yuxin Deng; Kaining Zhang; Shihua Zhang; Yansheng Li; Jiayi Ma
ISP-Teacher: Image Signal Process with Disentanglement Regularization for Unsupervised Domain Adaptive Dark Object Detection	Yin Zhang; Yongqiang Zhang; Zian Zhang; Man Zhang; Rui Tian; Mingli Ding
FacetCRS: Multi-Faceted Preference Learning for Pricking Filter Bubbles in Conversational Recommender System	Yongsen Zheng; Ziliang Chen; Jinghui Qin; Liang Lin
Adaptive Discovering and Merging for Incremental Novel Class Discovery	Guangyao Chen; Peixi Peng; Yangru Huang; Mengyue Geng; Yonghong Tian
IGAMT: Privacy Preserved Electronic Health Record Synthetic Approach with Heterogeneity and Irregularity	Wenjie Wang; Pengfei Tang; Jian Lou; Yuanming Shao; Li Xiong; Lance Waller; Yi-An Ko
Sample Efficient Reinforcement Learning with Partial Dynamics Knowledge	Meshal Alharbi; Mardavij Roozbehani; Munther Dahleh
Chain of Generation: Multi-Modal Gesture Synthesis via Cascaded Conditional Control	Zunnan Xu; Yachao Zhang; Sicheng Yang; Ronghui Li; Xiu Li
Locality Preserving Refinement for Shape Matching with Functional Maps	Yifan Xia; Yifan Lu; Yuan Gao; Jiayi Ma
ViTree: Single-Path Neural Tree for Step-Wise Interpretable Fine-Grained Visual Categorization	Danning Lao; Qi Liu; Jiazi Bu; Junchi Yan; Wei Shen
TIKP: Text-to-Image Knowledge Preservation for Continual Semantic Segmentation	Zhidong Yu; Wei Yang; Xike Xie; Zhenbo Shi
Symmetric Self-Paced Learning for Domain Generalization	Di Zhao; Yun Sing Koh; Gillian Dobbie; Hongsheng Hu; Philippe Fournier-Viger
Comprehensive View Embedding Learning for Single-Cell Multimodal	Zhenchao Tang; Jiehui Huang; Guanxing Chen; Calvin Yu-Chian Chen
A Multimodal, Multi-Task Adapting Framework for Video Action Recognition	Mengmeng Wang; Jiazheng Xing; Boyuan Jiang; Jun Chen; Jianbiao Mei; Xingxing Zuo; Guang Dai; Jingdong Wang; Yong Liu



SwiftPillars: High-Efficiency Pillar Encoder for Lidar-Based 3D Detection	Xin Jin; Kai Liu; Cong Ma; Ruining Yang; Fei Hui; Wei Wu
Image Content Generation with Causal Reasoning	Xiaochuan Li; Baoyu Fan; Runze Zhang; Liang Jin; Di Wang; Zhenhua Guo; Yaqian Zhao; Rengang Li
Cross-Modal and Uni-Modal Soft-Label Alignment for Image-Text Retrieval	Hailang Huang; Zhijie Nie; Ziqiao Wang; Ziyu Shang
FocalDreamer: Text-Driven 3D Editing via Focal-Fusion Assembly	Yuhan Li; Yishun Dou; Yue Shi; Yu Lei; Xuanhong Chen; Yi Zhang; Peng Zhou; Bingbing Ni
Enhancing RAW-to-sRGB with Decoupled Style Structure in Fourier Domain	Xuanhua He; Tao Hu; Guoli Wang; Zejin Wang; Run Wang; Qian Zhang; Keyu Yan; Ziyi Chen; Rui Li; Chengjun Xie; Jie Zhang; Man Zhou
PM-INR: Prior-Rich Multi-Modal Implicit Large-Scale Scene Neural Representation	Yiying Yang; Fukun Yin; Wen Liu; Jiayuan Fan; Xin Chen; Gang Yu; Tao Chen
DreamStyler: Paint by Style Inversion with Text-to-Image Diffusion Models	Namhyuk Ahn; Junsoo Lee; Chunggi Lee; Kunhee Kim; Daesik Kim; Seung-Hun Nam; Kibeom Hong
Learning to Approximate Adaptive Kernel Convolution on Graphs	Jaeyoon Sim; Sooyeon Jeon; InJun Choi; Guorong Wu; Won Hwa Kim
Domain Invariant Learning for Gaussian Processes and Bayesian Exploration	Xilong Zhao; Siyuan Bian; Yaoyun Zhang; Yuliang Zhang; Qinying Gu; Xinbing Wang; Chenghu Zhou; Nanyang Ye
Effect Size Estimation for Duration Recommendation in Online Experiments: Leveraging Hierarchical Models and Objective Utility Approaches	Yu Liu; Runzhe Wan; James McQueen; Doug Hains; Jinxiang Gu; Rui Song
Teaching Large Language Models to Translate with Comparison	Jiali Zeng; Fandong Meng; Yongjing Yin; Jie Zhou
MmAP: Multi-Modal Alignment Prompt for Cross-Domain Multi-Task Learning	Yi Xin; Junlong Du; Qiang Wang; Ke Yan; Shouhong Ding
On the Affinity, Rationality, and Diversity of Hierarchical Topic Modeling	Xiaobao Wu; Fengjun Pan; Thong Nguyen; Yichao Feng; Chaoqun Liu; Cong-Duy Nguyen; Anh Tuan Luu
ZOOM: Learning Video Mirror Detection with Extremely-Weak Supervision	Ke Xu; Tsun Wai Siu; Rynson W.H. Lau
Double Buffers CEM-TD3: More Efficient Evolution and Richer Exploration	Sheng Zhu; Chun Shen; Shuai Lü; Junhong Wu; Daolong An
Distilling Reliable Knowledge for Instance-Dependent Partial Label Learning	Dong-Dong Wu; Deng-Bao Wang; Min-Ling Zhang
Settling Decentralized Multi-Agent Coordinated Exploration by Novelty Sharing	Haobin Jiang; Ziluo Ding; Zongqing Lu
Information Design for Congestion Games with Unknown Demand	Svenja M. Griesbach; Martin Hoefer; Max Klimm; Tim Koglin
Hypercorrelation Evolution for Video Class-Incremental Learning	Sen Liang; Kai Zhu; Wei Zhai; Zhiheng Liu; Yang Cao
Improving Transferability for Cross-Domain Trajectory Prediction via Neural Stochastic Differential Equation	Daehee Park; Jaewoo Jeong; Kuk-Jin Yoon
Towards Efficient Diffusion-Based Image Editing with Instant Attention Masks	Siyu Zou; Jiji Tang; Yiyi Zhou; Jing He; Chaoyi Zhao; Rongsheng Zhang; Zhipeng Hu; Xiaoshuai Sun
ModWaveMLP: MLP-Based Mode Decomposition and Wavelet Denoising	Ke Sun; Pei Liu; Pengfei Li; Zhifang Liao
Model to Defeat Complex Structures in Traffic Forecasting	
V2Meow: Meowing to the Visual Beat via Video-to-Music Generation	Kun Su; Judith Li; Qingqing Huang; Dima Kuzmin; Joonseok Lee; Chris Donahue; Fei Sha; Aren Jansen; Yu Wang; Mauro Verzetti; Timo Denk
Explicitly Perceiving and Preserving the Local Geometric Structures for 3D Point Cloud Attack	Daizong Liu; Wei Hu
Boosting Residual Networks with Group Knowledge	Shengji Tang; Peng Ye; Baopu Li; Weihao Lin; Tao Chen; Tong He; Chong Yu; Wanli Ouyang
Understanding the Generalization of Pretrained Diffusion Models on Out-of-Distribution Data	Sai Niranjana Ramachandran; Rudrabha Mukhopadhyay; Madhav Agarwal; C.V. Jawahar; Vinay Nambodiri



Semantic Segmentation in Multiple Adverse Weather Conditions with Domain Knowledge Retention	Xin Yang; Wending Yan; Yuan Yuan; Michael Bi Mi; Robby T. Tan
Learning MDL Logic Programs from Noisy Data	Céline Hocquette; Andreas Niskanen; Matti Järvisalo; Andrew Cropper
Market-GAN: Adding Control to Financial Market Data Generation with Semantic Context	Haochong Xia; Shuo Sun; Xinrun Wang; Bo An
Who Knows the Answer? Finding the Best Model and Prompt for Each Query Using Confidence-Based Search	Walter Gerych; Yara Rizk; Vatche Isahagian; Vinod Muthusamy; Evelyn Duesterwald; Praveen Venkateswaran
Exact Inference for Continuous-Time Gaussian Process Dynamics	Katharina Ensinger; Nicholas Tagliapietra; Sebastian Ziesche; Sebastian Trimpe
Theoretical and Empirical Analysis of Cost-Function Merging for Implicit Hitting Set WCSP Solving	Javier Larrosa; Conrado Martínez; Emma Rollon
Where and How to Attack? A Causality-Inspired Recipe for Generating Counterfactual Adversarial Examples	Ruichu Cai; Yuxuan Zhu; Jie Qiao; Zefeng Liang; Furui Liu; Zhifeng Hao
Earthfarsser: Versatile Spatio-Temporal Dynamical Systems Modeling in One Model	Hao Wu; Yuxuan Liang; Wei Xiong; Zhengyang Zhou; Wei Huang; Shilong Wang; Kun Wang
Learning Spatially Collaged Fourier Bases for Implicit Neural Representation	Jason Chun Lok Li; Chang Liu; Binxiao Huang; Ngai Wong
Graph Reasoning Transformers for Knowledge-Aware Question Answering	Ruilin Zhao; Feng Zhao; Liang Hu; Guandong Xu
An Empirical Study of CLIP for Text-Based Person Search	Cao Min; Yang Bai; ziyin Zeng; Mang Ye; Min Zhang
HARDVS: Revisiting Human Activity Recognition with Dynamic Vision Sensors	Xiao Wang; Zongzhen Wu; Bo Jiang; Zhimin Bao; Lin Zhu; Guoqi Li; Yaowei Wang; Yonghong Tian
Enhancing Bilingual Lexicon Induction via Bi-directional Translation Pair Retrieving	Ding Qiuyu; Hailong Cao; Tiejun Zhao
High-Dimensional Analysis for Generalized Nonlinear Regression: From Asymptotics to Algorithm	Jian Li; Yong Liu; Weiping Wang
Underspecification in Language Modeling Tasks: A Causality-Informed Study of Gendered Pronoun Resolution	Emily McMillin
Beyond Grounding: Extracting Fine-Grained Event Hierarchies across Modalities	Hammad Ayyubi; Christopher Thomas; Lovish Chum; Rahul Lokesh; Long Chen; Yulei Niu; Xudong Lin; Xuande Feng; Jaywon Koo; Sounak Ray; Shih-Fangyikang Wang; Huminhao Zhu; Chao Zhang; Hanbin Zhao; Hui Qian
GAD-PVI: A General Accelerated Dynamic-Weight Particle-Based Variational Inference Framework	Katherine Metcalf; Miguel Sarabia; Masha Fedzechkina; Barry-John Theobald
Can You Rely on Synthetic Labellers in Preference-Based Reinforcement Learning? It's Complicated	XiaoHui Zhang; Jiangyan Yi; Chenglong Wang; Chu Yuan Zhang; Siding Zeng; Jianhua Tao
What to Remember: Self-Adaptive Continual Learning for Audio Deepfake Detection	Di Wu; Yuling Jiao; Li Shen; Haizhao Yang; Xiliang Lu
Neural Network Approximation for Pessimistic Offline Reinforcement Learning	Haziq Razali; Yiannis Demiris
Forecasting Bimanual Object Manipulation Sequences from Unimanual Observations	Ge Meng; Jingjia Huang; Yingying Wang; Zhenqi Fu; Xinghao Ding; Yue Huang
Progressive High-Frequency Reconstruction for Pan-Sharpening with Implicit Neural Representation	Pingping Cai; Deja Scott; Xiaoguang Li; Song Wang
Orthogonal Dictionary Guided Shape Completion Network for Point Cloud	Rune Kjærsgaard; AHCÈNE Boubekki; Line Clemmensen
Pantypes: Diverse Representatives for Self-Explainable Models	



Pseudo-Label Calibration Semi-supervised Multi-Modal Entity Alignment	Luyao Wang; Pengnian Qi; Xigang Bao; Chunlai Zhou; Biao Qin
Set Prediction Guided by Semantic Concepts for Diverse Video Captioning	Yifan Lu; Ziqi Zhang; Chunfeng Yuan; Peng Li; Yan Wang; Bing Li; Weiming
SFC: Shared Feature Calibration in Weakly Supervised Semantic	Xinqiao Zhao; Feilong Tang; Xiaoyang Wang; Jimin Xiao
Cross-Modal Feature Distribution Calibration for Few-Shot Visual Question Answering	Jing Zhang; Xiaoqiang Liu; Mingzhe Chen; Zhe Wang
EG-NAS: Neural Architecture Search with Fast Evolutionary Exploration	Zicheng Cai; Lei Chen; Peng Liu; Tongtao Ling; Yutao Lai
Variable Importance in High-Dimensional Settings Requires Grouping	Ahmad CHAMMA; Bertrand Thirion; Denis-Alexander Engemann
Selective Deep Autoencoder for Unsupervised Feature Selection	Wael Hassanieh; Abdallah Chehade
Beyond the Label Itself: Latent Labels Enhance Semi-supervised Point Cloud Panoptic Segmentation	Yujun Chen; Xin Tan; Zhizhong Zhang; Yanyun Qu; Yuan Xie
Continuous Rotation Group Equivariant Network Inspired by Neural Population Coding	Zhiqiang Chen; Yang Chen; xiaolong Zou; Shan Yu
Agile-Quant: Activation-Guided Quantization for Faster Inference of LLMs on the Edge	Xuan Shen; Peiyan Dong; Lei Lu; Zhenglun Kong; Zhengang Li; Ming Lin; Chao Wu; Yanzhi Wang
Point Deformable Network with Enhanced Normal Embedding for Point Cloud Analysis	Xingyilang Yin; Xi Yang; Liangchen Liu; Nannan Wang; Xinbo Gao
Generating and Reweighting Dense Contrastive Patterns for Unsupervised Anomaly Detection	Songmin Dai; Yifan Wu; Xiaoqiang Li; Xiangyang Xue
Patched Line Segment Learning for Vector Road Mapping	Jiakun Xu; Bowen Xu; Gui-Song Xia; Liang Dong; Nan Xue
CLIP-Gaze: Towards General Gaze Estimation via Visual-Linguistic Model	Pengwei Yin; Guanzhong Zeng; Jingjing Wang; Di Xie
Explaining Reinforcement Learning Agents through Counterfactual Action Outcomes	Yotam Amitai; Yael Friedler; Ofra Amir
PMET: Precise Model Editing in a Transformer	Xiaopeng Li; Shasha Li; Shezheng Song; Yang Jing; Jun Ma; Jie Yu
Sampling for Beyond-Worst-Case Online Ranking	Qingyun Chen; Sungjin Im; Benjamin Moseley; Chenyang Xu; Ruilong Zhang
MedBench: A Large-Scale Chinese Benchmark for Evaluating Medical Large Language Models	Yan Cai; Linlin Wang; Ye Wang; Gerard de Melo; Ya Zhang; Yan-Feng Wang; Liang He
Graph-Based Prediction and Planning Policy Network (GP3Net) for Scalable Self-Driving in Dynamic Environments Using Deep Reinforcement Learning	Jayabrata Chowdhury; Venkataramanan Shivaraman; Suresh Sundaram; Sujit PB
Electron Microscopy Images as Set of Fragments for Mitochondrial Segmentation	Naisong Luo; Rui Sun; Yuwen Pan; Tianzhu Zhang; Feng Wu
Multi-Cross Sampling and Frequency-Division Reconstruction for Image Compressed Sensing	Heping Song; Jingyao Gong; Hongying Meng; Yuping Lai
Zhongjing: Enhancing the Chinese Medical Capabilities of Large Language Model through Expert Feedback and Real-World Multi-Turn Dialogue	Songhua Yang; Hanjie Zhao; Senbin Zhu; Guangyu Zhou; Hongfei Xu; Yuxiang Jia; Zan Hongying
ExpeL: LLM Agents Are Experiential Learners	Andrew Zhao; Daniel Huang; Quentin Xu; Matthieu Lin; Yong-Jin Liu; Gao
PRP Rebooted: Advancing the State of the Art in FOND Planning	Christian Muise; Sheila McIlraith; J. Christopher Beck
All Beings Are Equal in Open Set Recognition	Chaohua Li; Enhao Zhang; Chuanxing Geng; Songcan Chen
Improving Expressive Power of Spectral Graph Neural Networks with Eigenvalue Correction	Kangkang Lu; Yanhua Yu; Hao Fei; Xuan Li; Zixuan Yang; Zirui Guo; Meiyu Liang; Mengran Yin; Tat-Seng Chua
Decentralized Sum-of-Nonconvex Optimization	Zhuanghua Liu; Bryan Kian Hsiang Low



MeDM: Mediating Image Diffusion Models for Video-to-Video Translation with Temporal Correspondence Guidance	Ernie Chu; Tzuhsuan Huang; Shuo-Yen LIN; Jun-Cheng Chen
B-spine: Learning B-spline Curve Representation for Robust and Interpretable Spinal Curvature Estimation	Hao Wang; Qiang Song; Ruofeng Yin; Rui Ma
Collaborative Synthesis of Patient Records through Multi-Visit Health State Inference	Hongda Sun; Hongzhan Lin; Rui Yan
Primitive-Based 3D Human-Object Interaction Modelling and Programming	Siqi Liu; Yong-Lu Li; Zhou FANG; Xinpeng Liu; Yang You; Cewu Lu
Link Prediction in Multilayer Networks via Cross-Network Embedding	Guojing Ren; Xiao Ding; Xiao-Ke Xu; Hai-Feng Zhang
An Effective Augmented Lagrangian Method for Fine-Grained Multi-View Optimization	Yuze Tan; Hecheng Cai; Shudong Huang; Shuping Wei; Fan Yang; Jiancheng Lv
Multiscale Attention Wavelet Neural Operator for Capturing Steep Trajectories in Biochemical Systems	Jiayang Su; Junbo Ma; Songyang Tong; Enze Xu; Minghan Chen
TMFormer: Token Merging Transformer for Brain Tumor Segmentation with Missing Modalities	Zheyu Zhang; Gang Yang; Yueyi Zhang; Huanjing Yue; Aiping Liu; Yunwei Ou; Jian Gong; Xiaoyan Sun
Prompting Multi-Modal Image Segmentation with Semantic Grouping	Qibin He
Semi-supervised TEE Segmentation via Interacting with SAM Equipped with Noise-Resilient Prompting	Sen Deng; Yidan Feng; Haoneng Lin; Yiting Fan; Alex Pui-Wai Lee; Xiaowei Hu; Jing Qin
FedCD: Federated Semi-supervised Learning with Class Awareness Balance via Dual Teachers	Yuzhi Liu; Huisi Wu; Jing Qin
Transfer and Alignment Network for Generalized Category Discovery	Wenbin An; Feng Tian; Wenkai Shi; Yan Chen; Yaqiang Wu; Qianying Wang; Ping Chen
Distribution-Conditioned Adversarial Variational Autoencoder for Valid Instrumental Variable Generation	Xinshu Li; Lina Yao
On Unsupervised Domain Adaptation: Pseudo Label Guided Mixup for Adversarial Prompt Tuning	Fanshuang Kong; Richong Zhang; Ziqiao Wang; Yongyi Mao
Learning Coalition Structures with Games	Yixuan Even Xu; Chun Kai Ling; Fei Fang
Fusing Conditional Submodular GAN and Programmatic Weak Supervision	Kumar Shubham; Pranav Sastry; Prathosh A P
Enhancing Representation of Spiking Neural Networks via Similarity-Sensitive Contrastive Learning	Yuhan Zhang; Xiaode Liu; Yuanpei Chen; Weihang Peng; Yufei Guo; Xuhui Huang; Zhe Ma
Text2Analysis: A Benchmark of Table Question Answering with Advanced Data Analysis and Unclear Queries	Xinyi He; Mengyu Zhou; Xinrun Xu; Xiaojun Ma; Rui Ding; Lun Du; Yan Gao; Ran Jia; Xu Chen; Shi Han; Zejian Yuan; Dongmei Zhang
Are You Concerned about Limited Function Evaluations: Data-Augmented Pareto Set Learning for Expensive Multi-Objective Optimization	Yongfan Lu; Bingdong Li; Aimin Zhou
LRS: Enhancing Adversarial Transferability through Lipschitz Regularized Surrogate	Tao Wu; Tie Luo; Donald Wunsch
Towards Continual Learning Desiderata via HSIC-Bottleneck Orthogonalization and Equiangular Embedding	Depeng Li; Tianqi Wang; Junwei Chen; Qining Ren; Kenji Kawaguchi; Zhigang Zeng
Transformer-Based No-Reference Image Quality Assessment via Supervised Contrastive Learning	Jinsong Shi; Pan Gao; Jie Qin
Multi-View Randomized Kernel Classification via Nonconvex Optimization	Xiaojian Ding; Fan Yang





On the Structural Hardness of Answer Set Programming: Can Structure Efficiently Confine the Power of Disjunctions?	Markus Hecher; Rafael Kiesel
HEAP: Unsupervised Object Discovery and Localization with Contrastive Grouping	Xin Zhang; Jinheng Xie; Yuan Yuan; Michael Bi Mi; Robby T. Tan
Topic-VQ-VAE: Leveraging Latent Codebooks for Flexible Topic-Guided Document Generation	YoungJoon Yoo; Jongwon Choi
SasWOT: Real-Time Semantic Segmentation Architecture Search WithOut Training	Chendi Zhu; Lujun Li; Yuli Wu; Zhengxing Sun
Efficient Asynchronous Federated Learning with Prospective Momentum Aggregation and Fine-Grained Correction	Yu Zang; Zhe Xue; Shilong Ou; Lingyang Chu; Junping Du; Yunfei Long
PoetryDiffusion: Towards Joint Semantic and Metrical Manipulation in Poetry Generation	Zhiyuan Hu; Chumin Liu; Yue Feng; Anh Tuan Luu; Bryan Hooi
Video-Language Aligned Transformer for Video Question Answering	Linlin Zong; Jiahui Wan; Xianchao Zhang; Xinyue Liu; Wenxin Liang; Bo Xu
Spectrum Translation for Refinement of Image Generation (STIG) Based on Contrastive Learning and Spectral Filter Profile	Seokjun Lee; Seung-Won Jung; Hyunseok Seo
Gaze from Origin: Learning for Generalized Gaze Estimation by Embedding the Gaze Frontalization Process	Mingjie Xu; Feng Lu
Causal Walk: Debiasing Multi-Hop Fact Verification with Front-Door	Congzhi Zhang; Linhai Zhang; Deyu Zhou
Sketch and Refine: Towards Fast and Accurate Lane Detection	Chao Chen; Jie Liu; Chang Zhou; Jie Tang; Gangshan Wu
Causal Strategic Learning with Competitive Selection	Kiet Q. H. Vo; Muneeb Aadil; Siu Lun Chau; Krikamol Muandet
Scaling Few-Shot Learning for the Open World	Zhipeng Lin; Wenjing Yang; Haotian Wang; Haoang Chi; Long Lan; Ji Wang
Image Captioning with Multi-Context Synthetic Data	Feipeng Ma; Yizhou Zhou; Fengyun Rao; Yueyi Zhang; Xiaoyan Sun
Context Enhanced Transformer for Single Image Object Detection in Video Data	Seungjun An; Seonghoon Park; Gyeongnyeom Kim; Jeongyeol Baek; Byeongwon Lee; Seungryong Kim
Improving Panoptic Narrative Grounding by Harnessing Semantic Relationships and Visual Confirmation	Tianyu Guo; Haowei Wang; Yiwei Ma; Jiayi Ji; Xiaoshuai Sun
Semantic Complete Scene Forecasting from a 4D Dynamic Point Cloud Sequence	Zifan Wang; Zhuorui Ye; Haoran Wu; Junyu Chen; Li Yi
BAIT: Benchmarking (Embedding) Architectures for Interactive Theorem-A Score-Based Deterministic Diffusion Algorithm with Smooth Scores for General Distributions	Sean Lamont; Michael Norrish; Amir Dezfouli; Christian Walder; Paul Karthik Elamvazhuthi; Xuechen Zhang; Matthew Jacobs; Samet Oymak; Fabio Pasqualetti
On the Computational Complexity of Plan Verification, (Bounded) Plan-Optimality Verification, and Bounded Plan Existence	Songtuan Lin; Conny Olz; Malte Helmert; Pascal Bercher
OCEAN-MBRL: Offline Conservative Exploration for Model-Based Offline Reinforcement Learning	Fan Wu; Rui Zhang; Qi Yi; Yunkai Gao; Jiaming Guo; Shaohui Peng; Siming Lan; Husheng Han; Yansong Pan; Kaizhao Yuan; Pengwei Jin; Ruizhi Chen; Yunji Chen; Ling Li
NeuSurf: On-Surface Priors for Neural Surface Reconstruction from Sparse Input Views	Han Huang; Yulun Wu; Junsheng Zhou; Ge Gao; Ming Gu; Yu-Shen Liu
Boosting Multiple Instance Learning Models for Whole Slide Image Classification: A Model-Agnostic Framework Based on Counterfactual	Weiping Lin; Zhenfeng Zhuang; Lequan Yu; Liansheng Wang
DeblurSR: Event-Based Motion Deblurring under the Spiking Representation	Chen Song; Chandrajit Bajaj; Qixing Huang



Exploring Temporal Feature Correlation for Efficient and Stable Video Semantic Segmentation	Matthieu Lin; Jenny Sheng; Yubin Hu; Yangguang Li; Lu Qi; Andrew Zhao; Gao Huang; Yong-Jin Liu
Reducing Spatial Fitting Error in Distillation of Denoising Diffusion Models	Shengzhe Zhou; Zejian Li; Shengyuan Zhang; Lefan Hou; Changyuan Yang; Guang Yang; Yang Zhiyuan; Lingyun Sun
OntoFact: Unveiling Fantastic Fact-Skeleton of LLMs via Ontology-Driven Reinforcement Learning	Ziyu Shang; Ke Wenjun; Nana Xiu; Peng Wang; Jiajun Liu; Yanhui Li; Zhizhao Luo; Ke Ji
COMBHelper: A Neural Approach to Reduce Search Space for Graph Combinatorial Problems	Hao Tian; Sourav Medya; Wei Ye
Blind Face Restoration under Extreme Conditions: Leveraging 3D-2D Prior Fusion for Superior Structural and Texture Recovery	Zhengrui Chen; Liying Lu; Ziyang Yuan; Yiming Zhu; Yu Li; Chun Yuan; Weihong Deng
Learning Domain-Independent Heuristics for Grounded and Lifted Planning	Dillon Chen; Sylvie Thiebaux; Felipe Trevizan
TCNet: Continuous Sign Language Recognition from Trajectories and Correlated Regions	Hui Lu; Albert Ali Salah; Ronald Poppe
LAMPAT: Low-Rank Adaption for Multilingual Paraphrasing Using Adversarial Training	Khoi M. Le; Trinh Pham; Tho Quan; Anh Tuan Luu
Fine Structure-Aware Sampling: A New Sampling Training Scheme for Pixel-Aligned Implicit Models in Single-View Human Reconstruction	Kennard Yanting Chan; Fayao Liu; Guosheng Lin; Chuan Sheng Foo; Weisi Lin
Identification of Causal Structure in the Presence of Missing Data with Additive Noise Model	Jie Qiao; Zhengming Chen; Jianhua Yu; Ruichu Cai; Zhifeng Hao
Efficient Nonparametric Tensor Decomposition for Binary and Count Data	Zerui Tao; Toshihisa Tanaka; Qibin Zhao
IS-DARTS: Stabilizing DARTS through Precise Measurement on Candidate Importance	Hongyi He; Longjun Liu; Haonan Zhang; Nanning Zheng
Revisiting Open-Set Panoptic Segmentation	Yufei Yin; Hao Chen; Wengang Zhou; Jiajun Deng; Haiming Xu; Houqiang Li
StockMixer: A Simple Yet Strong MLP-Based Architecture for Stock Price Forecasting	Jinyong Fan; Yanyan Shen
Text-Based Occluded Person Re-identification via Multi-Granularity Contrastive Consistency Learning	Xinyi Wu; Wentao Ma; Dan Guo; Tongqing Zhou; Shan Zhao; Zhiping Cai
Block Image Compressive Sensing with Local and Global Information	Xiaoyu Kong; Yongyong Chen; Feng Zheng; Zhenyu He
SimDistill: Simulated Multi-Modal Distillation for BEV 3D Object Detection	Haimei Zhao; Qiming Zhang; Shanshan Zhao; Zhe Chen; Jing Zhang; Dacheng Tao
Inverse Weight-Balancing for Deep Long-Tailed Learning	Wenqi Dang; Zhou Yang; Weisheng Dong; Xin Li; Guangming Shi
Root Cause Analysis in Microservice Using Neural Granger Causal Discovery	Cheng-Ming Lin; Ching Chang; Wei-Yao Wang; Kuang-Da Wang; Wen-Chih Peng
Dynamic Reactive Spiking Graph Neural Network	Han Zhao; Xu Yang; Cheng Deng; Junchi Yan
AltNeRF: Learning Robust Neural Radiance Field via Alternating Depth-Pose Optimization	Kun Wang; Zhiqiang Yan; Huang Tian; Zhenyu Zhang; Xiang Li; Jun Li; Jian Yang
Audio Generation with Multiple Conditional Diffusion Model	Zhifang Guo; Jianguo Mao; Tao Rui; Long Yan; Kazushige Ouchi; Hong Liu; Xiangdong Wang
SNN-PDE: Learning Dynamic PDEs from Data with Simplicial Neural Networks	Jae Choi; Yuzhou Chen; Huikyo Lee; Hyun Kim; Yulia Gel
Variational Hybrid-Attention Framework for Multi-Label Few-Shot Aspect Category Detection	Cheng Peng; Ke Chen; Lidan Shou; Gang Chen



Focus Stacking with High Fidelity and Superior Visual Effects	Bo Liu; Bin Hu; Xiuli Bi; Weisheng Li; Bin Xiao
Watermarking Conditional Text Generation for AI Detection: Unveiling Challenges and a Semantic-Aware Watermark Remedy	Yu Fu; Deyi Xiong; Yue Dong
Multichannel AV-wav2vec2: A Framework for Learning Multichannel Multi-Modal Speech Representation	Qiushi Zhu; Jie Zhang; Yu Gu; Yuchen Hu; Lirong Dai
On Alternating-Time Temporal Logic, Hyperproperties, and Strategy Sharing	Raven Beutner; Bernd Finkbeiner
TACIT: A Target-Agnostic Feature Disentanglement Framework for Cross-Domain Text Classification	Rui Song; Fausto Giunchiglia; Yingji Li; Mingjie Tian; Hao Xu
Discrete Cycle-Consistency Based Unsupervised Deep Graph Matching	Siddharth Tourani; Muhammad Haris Khan; Carsten Rother; Bogdan Savchynskyy
Beyond Mimicking Under-Represented Emotions: Deep Data Augmentation with Emotional Subspace Constraints for EEG-Based Emotion Recognition	Zhi ZHANG; Sheng-hua Zhong; Yan Liu
Adaptive Feature Imputation with Latent Graph for Deep Incomplete Multi-View Clustering	Jingyu Pu; Chenhang Cui; Xinyue Chen; Yazhou Ren; Xiaorong Pu; Zhifeng Hao; Philip Yu; Lifang He
F3-Pruning: A Training-Free and Generalized Pruning Strategy towards Faster and Finer Text-to-Video Synthesis	Sitong Su; Jianzhi Liu; Lianli Gao; Jingkuan Song
Distilling Autoregressive Models to Obtain High-Performance Non-autoregressive Solvers for Vehicle Routing Problems with Faster Inference Speed	Yubin Xiao; Di Wang; Boyang Li; Mingzhao Wang; Xuan Wu; Changliang Zhou; You Zhou
DreamIdentity: Enhanced Editability for Efficient Face-Identity Preserved Image Generation	Zhuowei Chen; Shancheng Fang; Wei Liu; Qian He; Mengqi Huang; Zhendong Mao
Design Graph Guided Element Importance-Aware Layout Generation with Multi-Modality Cascade Transformer	Qiuyun Zhang; Bin Guo; Lina Yao; Hao Wang; Xiaotian Qiao; Ying Zhang; Zhiwen Yu
Encoding Constraints as Binary Constraint Networks Satisfying BTP	Ruiwei Wang
ProAgent: Building Proactive Cooperative Agents with Large Language Models	Ceyao Zhang; Kaijie Yang; Siyi Hu; Zihao Wang; Guanghe Li; Yihang Sun; Cheng Zhang; Zhaowei Zhang; Anji Liu; Song-Chun Zhu; Xiaojun Chang; Junge Zhang; Feng Yin; Yitao Liang; Yaodong Yang
Separate the Wheat from the Chaff: Model Deficiency Unlearning via Parameter-Efficient Module Operation	Xinshuo Hu; Dongfang Li; Zihao Zheng; Zhenyu Liu; Baotian Hu; Min Zhang
Unsupervised Continual Anomaly Detection with Contrastively-Learned Prompt	Jiaqi Liu; Kai Wu; Qiang Nie; Ying Chen; Bin-Bin Gao; Yong Liu; Jinbao Wang; Chengjie Wang; Feng Zheng
PPIDSG: A Privacy-Preserving Image Distribution Sharing Scheme with GAN in Federated Learning	Yuting Ma; Yuanzhi Yao; Xiaohua Xu
Understanding Distributed Representations of Concepts in Deep Neural Networks without Supervision	Wonjoon Chang; Dahee Kwon; Jaesik Choi
Continuous-Time Graph Representation with Sequential Survival Process	Abdulkadir Celikkanat; Nikolaos Nakis; Morten Mørup
Phoneme Hallucinator: One-Shot Voice Conversion via Set Expansion	Siyuan Shan; Yang Li; Amartya Banerjee; Junier Oliva
Robust Blind Text Image Deblurring via Maximum Consensus Framework	Zijian Min; Gundu Hassan; GeunSik Jo
Automatic Core-Guided Reformulation via Constraint Explanation and Condition Learning	Kevin Leo; Maria Garcia de la Banda; Mark Wallace; Grame Gange



Discriminatively Fuzzy Multi-View K-means Clustering with Local Structure Preserving	Jun Yin; Shiliang Sun; Lai Wei; Pei Wang
Linear-Time Algorithms for Front-Door Adjustment in Causal Graphs	Marcel Wienöbst; Benito van der Zander; Maciej Liskiewicz
Data-Driven Knowledge-Aware Inference of Private Information in Continuous Double Auctions	Lvye Cui; Haoran Yu
LGMRec: Local and Global Graph Learning for Multimodal Recommendation	Zhiqiang Guo; Jianjun Li; Guohui Li; Chaoyang Wang; Si Shi; Bin Ruan
A Comprehensive Augmentation Framework for Anomaly Detection	Lin Jiang; Yaping Yan
Robust Few-Shot Named Entity Recognition with Boundary Discrimination and Correlation Purification	Xiaojun Xue; Chunxia Zhang; Tianxiang Xu; Zhendong Niu
MULTISCRIP: Multimodal Script Learning for Supporting Open Domain Everyday Tasks	Jingyuan Qi; Minqian Liu; Ying Shen; Zhiyang Xu; Lifu Huang
Structural Information Enhanced Graph Representation for Link Prediction	Lei Shi; Bin Hu; Deng Zhao; Jianshan He; Zhiqiang Zhang; Jun Zhou
Voxel or Pillar: Exploring Efficient Point Cloud Representation for 3D Object Detection	Yuhao Huang; Sanping Zhou; Junjie Zhang; Jinpeng Dong; Nanning Zheng
FT-GAN: Fine-Grained Tune Modeling for Chinese Opera Synthesis	Meizhen Zheng; Peng Bai; Xiaodong Shi; Xun Zhou; Yiting Yan
Jointly Improving the Sample and Communication Complexities in Decentralized Stochastic Minimax Optimization	Xuan Zhang; Gabriel Mancino-Ball; Necdet Serhat Aybat; Yangyang Xu
H2GFormer: Horizontal-to-Global Voxel Transformer for 3D Semantic Scene Completion	Yu Wang; Chao Tong
Get an A in Math: Progressive Rectification Prompting	Zhenyu Wu; Meng Jiang; Chao Shen
Enhancing Semi-supervised Domain Adaptation via Effective Target Labeling	Jiujun He; Bin Liu; Guosheng Yin
Spatial-Semantic Collaborative Cropping for User Generated Content	Yukun Su; Yiwen Cao; Jingliang Deng; Fengyun Rao; Qingyao Wu
Confucius: Iterative Tool Learning from Introspection Feedback by Easy-to-Difficult Curriculum	Zhengliang Shi; Shen Gao; Minghang Zhu; Bowen Fang; Xin Xin; Pengjie Ren; Zhumin Chen; Jun Ma; Zhaochun Ren
Independency Adversarial Learning for Cross-Modal Sound Separation	Zhenkai Lin; Yanli Ji; Yang Yang
CPN: Complementary Proposal Network for Unconstrained Text Detection	Longhuang Wu; Shangxuan Tian; Youxin Wang; Pengfei Xiong
Visual Redundancy Removal for Composite Images: A Benchmark Dataset and a Multi-Visual-Effects Driven Incremental Method	Miaohui Wang; Rong Zhang; Lirong Huang; Yanshan Li
Reinforcement Learning as a Parsimonious Alternative to Prediction Cascades: A Case Study on Image Segmentation	Bharat Srikishan; Anika Tabassum; Srikanth Allu; Ramakrishnan Kannan; Nikhil Muralidhar
Expand-and-Quantize: Unsupervised Semantic Segmentation Using High-Dimensional Space and Product Quantization	Jiyoung Kim; Kyuhong Shim; Insu Lee; Byonghyo Shim
Learning Invariant Inter-pixel Correlations for Superpixel Generation	Sen Xu; Shikui Wei; Tao Ruan; Lixin Liao
Constrained Bayesian Optimization under Partial Observations: Balanced Improvements and Provable Convergence	Shengbo Wang; Ke Li
Augmented Commonsense Knowledge for Remote Object Grounding	Bahram Mohammadi; Yicong Hong; Yuankai Qi; Qi Wu; Shirui Pan; Javen Qinfeng Shi
An Efficient Knowledge Transfer Strategy for Spiking Neural Networks from Static to Event Domain	Xiang He; Dongcheng Zhao; Yang Li; Guobin Shen; Qingqun Kong; Yi Zeng
Sample-and-Bound for Non-convex Optimization	Yaoguang Zhai; Zhizhen Qin; Sicun Gao



FedGCR: Achieving Performance and Fairness for Federated Learning with Distinct Client Types via Group Customization and Reweighting	Shu-Ling Cheng; Chin-Yuan Yeh; Ting-An Chen; Eliana Pastor; Ming-Syan Chen
Safe Abductive Learning in the Presence of Inaccurate Rules	Xiao-Wen Yang; Jie-Jing Shao; Wei-Wei Tu; Yu-Feng Li; Wang-Zhou Dai; Zhi-Hua Zhou
Data Distribution Distilled Generative Model for Generalized Zero-Shot Recognition	Yijie Wang; Mingjian Hong; Luwen Huangfu; Sheng Huang
Online Markov Decision Processes Configuration with Continuous Decision Space	Davide Maran; Pierriccardo Olivieri; Francesco Emanuele Stradi; Giuseppe Urso; Nicola Gatti; Marcello Restelli
SAVSR: Arbitrary-Scale Video Super-resolution via a Learned Scale-Adaptive Network	Zekun Li; Hongying Liu; Fanhua Shang; Yuanyuan Liu; Liang Wan; Wei Feng
All Should Be Equal in the Eyes of LMs: Counterfactually Aware Fair Text Generation	Pragyan Banerjee; Abhinav Java; Surgan Jandial; Simra Shahid; Shaz Furniturewala; Balaji Krishnamurthy; Sumit Bhatia
Efficient Constraint Generation for Stochastic Shortest Path Problems	Johannes Schmalz; Felipe Trevizan
TC-LIF: A Two-Compartment Spiking Neuron Model for Long-Term Sequential Modelling	Shimin Zhang; Qu Yang; Chenxiang Ma; Jibin Wu; Haizhou Li; Kay Chen Tan
Towards Explainable Joint Models via Information Theory for Multiple Intent Detection and Slot Filling	Xianwei Zhuang; Xuxin Cheng; Yuexian Zou
Exploring One-Shot Semi-supervised Federated Learning with Pre-trained Diffusion Models	Mingzhao Yang; Shangchao Su; Bin Li; Xiangyang Xue
Critic-Guided Decision Transformer for Offline Reinforcement Learning	Yuanfu Wang; Chao Yang; Ying Wen; Yu Liu; Yu Qiao
EarthVQA: Towards Queryable Earth via Relational Reasoning-Based Remote Sensing Visual Question Answering	Junjue Wang; Zhuo Zheng; Zihang Chen; Ailong Ma; Yanfei Zhong
HONGAT: Graph Attention Networks in the Presence of High-Order Neighbors	Heng-Kai Zhang; Yi-Ge Zhang; Zhi Zhou; Yu-Feng Li
Exploiting Auxiliary Caption for Video Grounding	Hongxiang Li; Meng Cao; Xuxin Cheng; Yaowei Li; Zhihong Zhu; Yuexian Zou
The Logic of Doxastic Strategies	Junli Jiang; Pavel Naumov
Optimised Storage for Datalog Reasoning	Xinyue Zhang; Pan Hu; Yavor Nenov; Ian Horrocks
Shaping Up SHAP: Enhancing Stability through Layer-Wise Neighbor Selection	Gwladys Kelodjou; Laurence Rozé; Véronique Masson; Luis Galárraga; Romaric Gaudel; Maurice Tchuente; Alexandre Termier
ERL-TD: Evolutionary Reinforcement Learning Enhanced with Truncated Variance and Distillation Mutation	Qiuzhen Lin; Yangfan Chen; Lijia Ma; Wei-Neng Chen; Jianqiang Li
MetaMix: Meta-State Precision Searcher for Mixed-Precision Activation Quantization	Han-Byul Kim; Joo Hyung Lee; Sungjoo Yoo; Hong-Seok Kim
TCI-Former: Thermal Conduction-Inspired Transformer for Infrared Small Target Detection	Tianxiang Chen; Zhentao Tan; Qi Chu; Yue Wu; Bin Liu; Nenghai Yu
VELMA: Verbalization Embodiment of LLM Agents for Vision and Language Navigation in Street View	Raphael Schumann; Wanrong Zhu; Weixi Feng; Tsu-Jui Fu; Stefan Riezler; William Yang Wang
CAMEL: Capturing Metaphorical Alignment with Context Disentangling for Multimodal Emotion Recognition	Linhao Zhang; Li Jin; Guangluan Xu; Xiaoyu Li; Cai Xu; Kaiwen Wei; Nayu Liu; Haonan Liu
Heterogeneous Graph Reasoning for Fact Checking over Texts and Tables	Haisong Gong; Weizhi Xu; Shu Wu; Qiang Liu; Liang Wang
Aligner <sup>2</sup> : Enhancing Joint Multiple Intent Detection and Slot Filling via Adjustive and Forced Cross-Task Alignment	Zhihong Zhu; Xuxin Cheng; Yaowei Li; Hongxiang Li; Yuexian Zou



Rethinking Reverse Distillation for Multi-Modal Anomaly Detection	Zhihao Gu; Jiangning Zhang; Liang Liu; Xu Chen; Jinlong Peng; Zhenye Gan; Guannan Jiang; Annan Shu; Yabiao Wang; Lizhuang Ma
MELO: Enhancing Model Editing with Neuron-Indexed Dynamic LoRA	Lang Yu; Qin Chen; Jie Zhou; Liang He
Self-Supervised Multi-Modal Knowledge Graph Contrastive Hashing for Cross-Modal Search	Meiyu Liang; Junping Du; Zhengyang Liang; Yongwang Xing; Huang Wei; Zhe Xue
Towards Making Learnware Specification and Market Evolvable	Jian-Dong Liu; Zhi-Hao Tan; Zhi-Hua Zhou
Generalize for Future: Slow and Fast Trajectory Learning for CTR Prediction	Jian Zhu; Congcong Liu; Xue Jiang; Changping Peng; Zhangang Lin; Jingping Shao
Accelerating Text-to-Image Editing via Cache-Enabled Sparse Diffusion Inference	Zihao Yu; Haoyang Li; Fangcheng Fu; Xupeng Miao; Bin Cui
STAS: Spatial-Temporal Return Decomposition for Multi-Agent Reinforcement Learning	Sirui Chen; Zhaowei Zhang; Yaodong Yang; Yali Du
DexFuncGrasp: A Robotic Dexterous Functional Grasp Dataset Constructed from a Cost-Effective Real-Simulation Annotation System	Jinglue Hang; Xiangbo Lin; Tianqiang Zhu; Xuanheng Li; Rina Wu; Xiaohong Ma; Yi Sun
SUF: Stabilized Unconstrained Fine-Tuning for Offline-to-Online Reinforcement Learning	Jiaheng Feng; Mingxiao Feng; Haolin Song; Wengang Zhou; Houqiang Li
Enhancing Job Recommendation through LLM-Based Generative Adversarial Networks	Yingpeng Du; Di Luo; Rui Yan; Xiaopei Wang; Hongzhi Liu; Hengshu Zhu; Yang Song; Jie Zhang
Spear and Shield: Adversarial Attacks and Defense Methods for Model-Based Link Prediction on Continuous-Time Dynamic Graphs	Dongjin Lee; Juho Lee; Kijung Shin
Effectiveness of Constant Stepsize in Markovian LSA and Statistical Inference	Dongyan (Lucy) Huo; Yudong Chen; Qiaomin Xie
Fast Inter-frame Motion Prediction for Compressed Dynamic Point Cloud Attribute Enhancement	Wang Liu; Wei Gao; Xingming Mu
ASWT-SGNN: Adaptive Spectral Wavelet Transform-Based Self-Supervised Graph Neural Network	Ruyue Liu; Rong Yin; Yong Liu; Weiping Wang
DGCLUSTER: A Neural Framework for Attributed Graph Clustering via Modularity Maximization	Aritra Bhowmick; Mert Kosan; Zexi Huang; Ambuj Singh; Sourav Medya
Model Counting and Sampling via Semiring Extensions	Andreas Goral; Joachim Giesen; Mark Blacher; Christoph Staudt; Julien Klaus
Collaborative Weakly Supervised Video Correlation Learning for Procedure-Aware Instructional Video Analysis	Tianyao He; Huabin Liu; Yuxi Li; Xiao Ma; Cheng Zhong; Yang Zhang; Weiyao Lin
AdaFormer: Efficient Transformer with Adaptive Token Sparsification for Image Super-resolution	Xiaotong Luo; zekun Ai; Qiuyuan Liang; Ding Liu; Yuan Xie; Yanyun Qu; Yun Fu
The Causal Impact of Credit Lines on Spending Distributions	Yijun Li; Cheuk Hang Leung; Xiangqian Sun; Chaoqun Wang; Yiyang Huang; Xing Yan; Qi Wu; Dongdong Wang; Zhixiang Huang
Adversarial Robust Safeguard for Evading Deep Facial Manipulation	Jiazhi Guan; Yi Zhao; Zhuoer Xu; Changhua Meng; Ke Xu; Youjian Zhao
MetaRLEC: Meta-Reinforcement Learning for Discovery of Brain Effective Connectivity	Zuozhen Zhang; Junzhong Ji; Jinduo Liu
MetaCARD: Meta-Reinforcement Learning with Task Uncertainty Feedback via Decoupled Context-Aware Reward and Dynamics Components	Min Wang; Xin Li; Leiji Zhang; Mingzhong Wang
FMRNet: Image Deraining via Frequency Mutual Revision	Kui Jiang; Junjun Jiang; Xianming Liu; Xin Xu; Xianzheng Ma





Frozen CLIP Transformer Is an Efficient Point Cloud Encoder	Xiaoshui Huang; Zhou Huang; Sheng Li; Qu Wentao; Tong He; Yuenan Hou; Yifan Zuo; Wanli Ouyang
SkyScript: A Large and Semantically Diverse Vision-Language Dataset for Remote Sensing	Zhecheng Wang; Rajanie Prabha; Tianyuan Huang; Jiajun Wu; Ram Rajagopal
ACT: Empowering Decision Transformer with Dynamic Programming via Advantage Conditioning	Chen-Xiao Gao; Chenyang Wu; Mingjun Cao; Rui Kong; Zongzhang Zhang; Yang Yu
Data Shunt: Collaboration of Small and Large Models for Lower Costs and Better Performance	Dong Chen; Yueting Zhuang; Shuo Zhang; Jinfeng Liu; Su Dong; Siliang Tang
SC-NeuS: Consistent Neural Surface Reconstruction from Sparse and Noisy Views	Shi-Sheng Huang; Zixin Zou; Yichi Zhang; Yan-Pei Cao; Ying Shan
FedFixer: Mitigating Heterogeneous Label Noise in Federated Learning	Xinyuan Ji; Zhaowei Zhu; Wei Xi; Olga Gadyatskaya; Song Zilong; Yong Cai; Yang Liu
QAGait: Revisit Gait Recognition from a Quality Perspective	Zengbin Wang; Saihui Hou; Man Zhang; Xu Liu; Chunshui Cao; Yongzhen Huang; Peipei Li; Shibiao Xu
T2I-Adapter: Learning Adapters to Dig Out More Controllable Ability for Text-to-Image Diffusion Models	Chong Mou; Xintao Wang; Liangbin Xie; Yanze Wu; Jian Zhang; Zhongang Qi; Ying Shan
Improving PTM Site Prediction by Coupling of Multi-Granularity Structure and Multi-Scale Sequence Representation	Zhengyi Li; Menglu Li; Lida Zhu; Wen Zhang
Weighted Envy-Freeness for Submodular Valuations	Luisa Montanari; Ulrike Schmidt-Kraepelin; Warut Suksompong; Nicholas Teh
Exploring Large Language Model for Graph Data Understanding in Online Job Recommendations	Likang Wu; Zhaopeng Qiu; Zhi Zheng; Hengshu Zhu; Enhong Chen
Analyzing Generalization in Policy Networks: A Case Study with the Double-Integrator System	Ruining Zhang; Haoran Han; Maolong Lv; Qisong Yang; Jian Cheng
FFT-Based Dynamic Token Mixer for Vision	Yuki Tatsunami; Masato Taki
IPRemover: A Generative Model Inversion Attack against Deep Neural Network Fingerprinting and Watermarking	Wei Zong; Yang-Wai Chow; Willy Susilo; Joonsang Baek; Jongkil Kim; Seyit Camtepe
STEM: Unleashing the Power of Embeddings for Multi-Task Recommendation	Liangcai Su; Junwei Pan; Ximei Wang; Xi Xiao; Shijie Quan; Xihua Chen; Jie Jiang
Consistency-Guided Temperature Scaling Using Style and Content Information for Out-of-Domain Calibration	Wonjeong Choi; Jungwuk Park; Dong-Jun Han; YoungHyun Park; Jaekyun Moon
ConsistNER: Towards Instructive NER Demonstrations for LLMs with the Consistency of Ontology and Context	Chenxiao Wu; Ke Wenjun; Peng Wang; Zhizhao Luo; Guozheng Li; Wanyi Chen
U-mixer: An Unet-Mixer Architecture with Stationarity Correction for Time Series Forecasting	Xiang Ma; Xuemei Li; Lexin Fang; Tianlong Zhao; Caiming Zhang
Every Node Is Different: Dynamically Fusing Self-Supervised Tasks for Attributed Graph Clustering	Pengfei Zhu; Qian Wang; Yu Wang; Jialu Li; Qinghua Hu
Null Space Matters: Range-Null Decomposition for Consistent Multi-Contrast MRI Reconstruction	Jiacheng Chen; Jiawei Jiang; Fei Wu; Jianwei Zheng
Generalisation through Negation and Predicate Invention	David M. Cerna; Andrew Cropper
An Optimal Transport View for Subspace Clustering and Spectral Clustering	Yuguang Yan; Zhihao Xu; Canlin Yang; Jie ZHANG; Ruichu Cai; Michael Kwok-Po Ng



MGQFormer: Mask-Guided Query-Based Transformer for Image Manipulation Localization	Kunlun Zeng; Ri Cheng; Weimin Tan; Bo Yan
UniCell: Universal Cell Nucleus Classification via Prompt Learning	Junjia Huang; Haofeng Li; Xiang Wan; Guanbin Li
Fully-Connected Spatial-Temporal Graph for Multivariate Time-Series Data	Yucheng Wang; Yuecong Xu; Jianfei Yang; Min Wu; Xiaoli Li; Lihua Xie; Zhenghua Chen
SeGA: Preference-Aware Self-Contrastive Learning with Prompts for Anomalous User Detection on Twitter	Ying-Ying Chang; Wei-Yao Wang; Wen-Chih Peng
Root Cause Explanation of Outliers under Noisy Mechanisms	Phuoc Nguyen; Truyen Tran; Sunil Gupta; Thin Nguyen; Svetha Venkatesh
BLiRF: Bandlimited Radiance Fields for Dynamic Scene Modeling	Sameera Ramasinghe; Violetta Shevchenko; Gil Avraham; Anton Hengel
EVE: Efficient Vision-Language Pre-training with Masked Prediction and Modality-Aware MoE	Junyi Chen; Longteng Guo; Jia Sun; Shuai Shao; Zehuan Yuan; Liang Lin; Dongyu Zhang
De-biased Attention Supervision for Text Classification with Causality	Yiquan Wu; Yifei Liu; Ziyu Zhao; Weiming Lu; Yating Zhang; Changlong Sun; Fei Wu; Kun Kuang
Domain-Hallucinated Updating for Multi-Domain Face Anti-spoofing	Chengyang Hu; Ke-Yue Zhang; Taiping Yao; Shice Liu; Shouhong Ding; Xin Tan; Lizhuang Ma
N-gram Unsupervised Compoundation and Feature Injection for Better Symbolic Music Understanding	Jinhao Tian; Zuchao Li; Jiajia Li; Ping Wang
LRANet: Towards Accurate and Efficient Scene Text Detection with Low-Rank Approximation Network	Yuchen Su; Zhineng Chen; Zhiwen Shao; Yuning Du; Zhilong Ji; Jinfeng Bai; Yong Zhou; Yu-Gang Jiang
Quantifying and Analyzing Entity-Level Memorization in Large Language	Zhenhong Zhou; Jiuyang Xiang; Chaomeng Chen; Sen Su
Beyond Expected Return: Accounting for Policy Reproducibility When Evaluating Reinforcement Learning Algorithms	Manon Flageat; Bryan Lim; Antoine Cully
Self-Training Based Few-Shot Node Classification by Knowledge Distillation	Zongqian Wu; Yujie Mo; Peng Zhou; Shangbo Yuan; Xiaofeng Zhu
FairWASP: Fast and Optimal Fair Wasserstein Pre-processing	Zikai Xiong; Niccolo Dalmaso; Alan Mishler; Vamsi Potluru; Tucker Balch; Manuela Veloso
Learning Reduced Fluid Dynamics	Zherong Pan; Xifeng Gao; Kui Wu
MLNet: Mutual Learning Network with Neighborhood Invariance for Universal Domain Adaptation	Yanzuo Lu; Meng Shen; Andy J Ma; Xiaohua Xie; Jian-Huang Lai
CMG-Net: Robust Normal Estimation for Point Clouds via Chamfer Normal Distance and Multi-Scale Geometry	Yingrui Wu; Mingyang Zhao; Keqiang Li; Weize Quan; Tianqi Yu; Jianfeng Yang; Xiaohong Jia; Dong-Ming Yan
Elijah: Eliminating Backdoors Injected in Diffusion Models via Distribution Shift	Shengwei An; Sheng-Yen Chou; Kaiyuan Zhang; Qiuling Xu; Guanhong Tao; Guangyu Shen; Siyuan Cheng; Shiqing Ma; Pin-Yu Chen; Tsung-Yi Ho; Xiangyu Zhang
EDA: Evolving and Distinct Anchors for Multimodal Motion Prediction	Longzhong Lin; Xuewu Lin; Tianwei Lin; Lichao Huang; Rong Xiong; Yue
Exploiting Geometry for Treatment Effect Estimation via Optimal Transport	Yuguang Yan; Zeqin Yang; Weilin Chen; Ruichu Cai; Zhifeng Hao; Michael Kwok-Po Ng
Non-flat ABA Is an Instance of Bipolar Argumentation	Markus Ulbricht; Nico Potyka; Anna Rapberger; Francesca Toni
A New Mechanism for Eliminating Implicit Conflict in Graph Contrastive Learning	Dongxiao He; Jitao Zhao; Cuiying Huo; Huang Yongqi; Yuxiao Huang; Zhiyong Feng
Question Calibration and Multi-Hop Modeling for Temporal Question	Chao Xue; Di Liang; Pengfei Wang; Jing Zhang



Dynamic Feature Pruning and Consolidation for Occluded Person Re-identification	YuTeng Ye; Hang Zhou; Jiale Cai; Chenxing Gao; Youjia Zhang; Junle Wang; Qiang Hu; Junqing Yu; Wei Yang
Unveiling Details in the Dark: Simultaneous Brightening and Zooming for Low-Light Image Enhancement	Ziyu Yue; Jiaxin Gao; Zhixun Su
SPEAL: Skeletal Prior Embedded Attention Learning for Cross-Source Point Cloud Registration	Kezheng Xiong; Maoji Zheng; Qingshan Xu; Chenglu Wen; Siqi Shen; Cheng Wang
Attacks on Continual Semantic Segmentation by Perturbing Incremental	Zhidong Yu; Wei Yang; Xike Xie; Zhenbo Shi
A New Benchmark and Model for Challenging Image Manipulation Detection	Zhenfei Zhang; Mingyang Li; Ming-Ching Chang
Non-stationary Projection-Free Online Learning with Dynamic and Adaptive Regret Guarantees	Yibo Wang; Wenhao Yang; Wei Jiang; Shiyin Lu; Bing Wang; Haihong Tang; Yuanyu Wan; Lijun Zhang
Combinatorial Stochastic-Greedy Bandit	Fares Fourati; Christopher John Quinn; Mohamed-Slim Alouini; Vaneet
VMT-Adapter: Parameter-Efficient Transfer Learning for Multi-Task Dense Scene Understanding	Yi Xin; Junlong Du; Qiang Wang; Zhiwen Lin; Ke Yan
Structure-CLIP: Towards Scene Graph Knowledge to Enhance Multi-Modal Structured Representations	Yufeng Huang; Jiji Tang; Zhuo Chen; Rongsheng Zhang; Xinfeng Zhang; Weijie Chen; Zeng Zhao; Zhou Zhao; Tangjie Lv; Zhipeng Hu; Wen Zhang
Learning Diverse Risk Preferences in Population-Based Self-Play	Yuhua Jiang; Qihan Liu; Xiaoteng Ma; Chenghao Li; Yiqin Yang; Jun Yang; Bin Liang; Qianchuan Zhao
Zero-1-to-3: Domain-Level Zero-Shot Cognitive Diagnosis via One Batch of Early-Bird Students towards Three Diagnostic Objectives	Weibo Gao; Qi Liu; Hao Wang; Linan Yue; Haoyang Bi; Yin Gu; Fangzhou Yao; Zheng Zhang; Xin Li; Yuanjing He
A Compiler for Weak Decomposable Negation Normal Form	Petr Illner; Petr Kučera
Towards Balanced Alignment: Modal-Enhanced Semantic Modeling for Video Moment Retrieval	Zhihang Liu; Jun Li; Hongtao Xie; Pandeng Li; Jiannan Ge; Sun-Ao Liu; Guoqing Jin
Uncovering and Mitigating the Hidden Chasm: A Study on the Text-Text Domain Gap in Euphemism Identification	Yuxue Hu; Junsong Li; Mingmin Wu; Zhongqiang Huang; Gang Chen; Ying Sha
Perturbation-Invariant Adversarial Training for Neural Ranking Models: Improving the Effectiveness-Robustness Trade-Off	Yu-An Liu; Ruqing Zhang; Ming Kun Zhang; Wei Chen; Maarten de Rijke; Jiafeng Guo; Xueqi Cheng
A Non-parametric Graph Clustering Framework for Multi-View Data	Shengju Yu; Siwei Wang; Zhibin Dong; Wenxuan Tu; Suyuan Liu; Zhao Lv; Pan Li; Miao Wang; En Zhu
Recurrent Partial Kernel Network for Efficient Optical Flow Estimation	Henrique Morimitsu; Xiaobin Zhu; Xiangyang Ji; Xu-Cheng Yin
DGPO: Discovering Multiple Strategies with Diversity-Guided Policy Optimization	Wentse Chen; Shiyu Huang; Yuan Chiang; Tim Pearce; Wei-Wei Tu; Ting Chen; Jun Zhu
VITA: 'Carefully Chosen and Weighted Less' Is Better in Medication Recommendation	Taeri Kim; Jiho Heo; Hongil Kim; Kijung Shin; Sang-Wook Kim
Integrated Decision Gradients: Compute Your Attributions Where the Model Makes Its Decision	Chase Walker; Sumit Jha; Kenny Chen; Rickard Ewetz
Object-Aware Domain Generalization for Object Detection	WooJu Lee; Dasol Hong; Hyungtae Lim; Hyun Myung
Triple Feature Disentanglement for One-Stage Adaptive Object Detection	Haoan Wang; Shilong Jia; Tieyong Zeng; Guixu Zhang; Zhi Li
HybridGait: A Benchmark for Spatial-Temporal Cloth-Changing Gait Recognition with Hybrid Explorations	Yilan Dong; Chunlin Yu; Ruiyang Ha; Ye Shi; Yuexin Ma; Lan Xu; Yanwei Fu; Jingya Wang
OVD-Explorer: Optimism Should Not Be the Sole Pursuit of Exploration in Noisy Environments	Jinyi Liu; Zhi Wang; Yan Zheng; Jianye Hao; Chenjia Bai; Junjie Ye; Zhen Wang; Haiyin Piao; Yang Sun



Semi-supervised 3D Object Detection with PatchTeacher and PillarMix	Xiaopei Wu; Liang Peng; Liang Xie; Yuenan Hou; Binbin Lin; Xiaoshui Huang; Haifeng Liu; Deng Cai; Wanli Ouyang
Influential Exemplar Replay for Incremental Learning in Recommender	Xinni Zhang; Yankai Chen; Chenhao Ma; Yixiang Fang; Irwin King
Mining Fine-Grained Image-Text Alignment for Zero-Shot Captioning via Text-Only Training	Longtian Qiu; Shan Ning; Xuming He
ND-MRM: Neuronal Diversity Inspired Multisensory Recognition Model	Qixin Wang; Chaoqiong Fan; Tianyuan Jia; Yuyang Han; Xia Wu
Abstract and Explore: A Novel Behavioral Metric with Cyclic Dynamics in Reinforcement Learning	Anjie Zhu; Peng-Fei Zhang; Ruihong Qiu; Zetao Zheng; Zi Huang; Jie Shao
Combating Data Imbalances in Federated Semi-supervised Learning with Dual Regulators	Sikai Bai; Shuaicheng Li; Weiming Zhuang; Jie Zhang; Kunlin Yang; Jun Hou; Shuai Yi; Shuai Zhang; Junyu Gao
On Computing Makespan-Optimal Solutions for Generalized Sliding-Tile	Marcus Gozon; Jingjin Yu
Fast and Controllable Post-training Sparsity: Learning Optimal Sparsity Allocation with Global Constraint in Minutes	Ruihao Gong; Yang Yong; Zining Wang; Jinyang Guo; Xiuying Wei; Yuqing Ma; Xianglong Liu
Step Vulnerability Guided Mean Fluctuation Adversarial Attack against Conditional Diffusion Models	Hongwei Yu; Jiansheng Chen; Xinlong Ding; Yudong Zhang; Ting Tang; Huimin Ma
LDMVFI: Video Frame Interpolation with Latent Diffusion Models	Duolikun Danier; Fan Zhang; David Bull
Curved Representation Space of Vision Transformers	Juyeop Kim; Junha Park; Songkuk Kim; Jong-Seok Lee
SyFormer: Structure-Guided Synergism Transformer for Large-Portion Image Inpainting	Jie Wu; Yuchao Feng; Honghui Xu; Chuanmeng Zhu; Jianwei Zheng
Bias-Conflict Sample Synthesis and Adversarial Removal Debias Strategy for Temporal Sentence Grounding in Video	Zhaobo Qi; Yibo Yuan; Xiaowen Ruan; Shuhui Wang; Weigang Zhang; Qingming Huang
Data Disparity and Temporal Unavailability Aware Asynchronous Federated Learning for Predictive Maintenance on Transportation Fleets	Leonie von Wahl; Niklas Heidenreich; Prasenjit Mitra; Michael Nolting; Nicolas Tempelmeier
DLCA-Recon: Dynamic Loose Clothing Avatar Reconstruction from Monocular Videos	Chunjie Luo; Fei Luo; Yusen Wang; Enxu Zhao; Chunxia Xiao
Discretization-Induced Dirichlet Posterior for Robust Uncertainty Quantification on Regression	Xuanlong Yu; Gianni Franchi; Jindong Gu; Emanuel Aldea
Social Physics Informed Diffusion Model for Crowd Simulation	Hongyi Chen; Jingtao Ding; Yong Li; Yue Wang; Xiao-Ping Zhang
Sketched Newton Value Iteration for Large-Scale Markov Decision Processes	Jinsong Liu; Chenghan Xie; Qi Deng; DongDong Ge; Yinyu Ye
Piecewise Linear Transformation – Propagating Aleatoric Uncertainty in Neural Networks	Thomas Krapf; Michael Hagn; Paul Miethaner; Alexander Schiller; Lucas Luttner; Bernd Heinrich
Progressive Feature Self-Reinforcement for Weakly Supervised Semantic Segmentation	Jingxuan He; Lechao Cheng; Chaowei Fang; Zunlei Feng; Tingting Mu; Mingli Song
GOODAT: Towards Test-Time Graph Out-of-Distribution Detection	Luzhi Wang; Di Jin; He Zhang; Yixin Liu; Dongxiao He; Wenjie Wang; Shirui Pan; Tat-Seng Chua
Discriminative Forests Improve Generative Diversity for Generative Adversarial Networks	Junjie Chen; Jiahao Li; Chen Song; Bin Li; Qingcai Chen; Hongchang Gao; Wendy Hui Wang; Zenglin Xu; Xinghua Shi
Terrain Diffusion Network: Climatic-Aware Terrain Generation with Geological Sketch Guidance	Zexin Hu; Kun Hu; Clinton Mo; Lei Pan; Zhiyong Wang
CoPL: Contextual Prompt Learning for Vision-Language Understanding	Koustava Goswami; Srikrishna Karanam; Prateksha Udhayan; K J Joseph; Balaji Vasan Srinivasan



Fractional Deep Reinforcement Learning for Age-Minimal Mobile Edge Computing	Lyudong Jin; Ming Tang; Meng Zhang; Hao Wang
A Fast Exact Solver with Theoretical Analysis for the Maximum Edge-Weighted Clique Problem	Lu Liu; Mingyu Xiao; Yi Zhou
W2P: Switching from Weak Supervision to Partial Supervision for Semantic Segmentation	Fangyuan Zhang; Tianxiang Pan; Junhai Yong; Bin Wang
Imagine, Initialize, and Explore: An Effective Exploration Method in Multi-Agent Reinforcement Learning	Zeyang Liu; Lipeng Wan; Xinrui Yang; Zhuoran Chen; Xingyu Chen; Xuguang Lan
KD-Club: An Efficient Exact Algorithm with New Coloring-Based Upper Bound for the Maximum K-defective Clique Problem	Jiongzhi Zheng; Mingming Jin; Kun He
Directed Diffusion: Direct Control of Object Placement through Attention Guidance	Wan-Duo Ma; Avisek Lahiri; J. P. Lewis; W. Bastiaan Kleijn; Thomas Leung
Intra- and Inter-group Optimal Transport for User-Oriented Fairness in Recommender Systems	Zhongxuan Han; Chaochao Chen; Xiaolin Zheng; Meng Li; Weiming Liu; Binhui Yao; Yuyuan Li; Jianwei Yin
Visual Hallucination Elevates Speech Recognition	Fang Zhang; Yongxin Zhu; Xiangxiang Wang; Huang Chen; Xing Sun; Linli Xu
Causal Adversarial Perturbations for Individual Fairness and Robustness in Heterogeneous Data Spaces	Ahmad-Reza Ehyaei; Kiarash Mohammadi; Amir-Hossein Karimi; Samira Samadi; Golnoosh Farnadi
Comprehensive Visual Grounding for Video Description	Wenhui Jiang; Yibo Cheng; Liu Linxin; Yuming Fang; Yuxin Peng; Yang Liu
Multilevel Attention Network with Semi-supervised Domain Adaptation for Drug-Target Prediction	ZhouSan Xie; Shikui Tu; Lei Xu
The Complexity of Computing Robust Mediated Equilibria in Ordinal Games	Vincent Conitzer
STAIR: Spatial-Temporal Reasoning with Auditable Intermediate Results for Video Question Answering	Yueqian Wang; Yuxuan Wang; Kai Chen; Dongyan Zhao
CatFormer: Category-Level 6D Object Pose Estimation with Transformer	Sheng Yu; Di-Hua Zhai; Yuanqing Xia
Parallel Beam Search Algorithms for Domain-Independent Dynamic Programming	Ryo Kuroiwa; J. Christopher Beck
Deep Unfolded Network with Intrinsic Supervision for Pan-Sharpening	Hebaixu Wang; Meiqi Gong; Xiaoguang Mei; Hao Zhang; Jiayi Ma
Learning to Pivot as a Smart Expert	Tianhao Liu; Shanwen Pu; DongDong Ge; Yinyu Ye
Weakly Supervised Semantic Segmentation for Driving Scenes	Dongseob Kim; Seungho Lee; Junsuk Choe; Hyunjung Shim
Non-parametric Representation Learning with Kernels	Pascal Esser; Maximilian Fleissner; Debarghya Ghoshdastidar
Rethinking Robustness of Model Attributions	Sandesh Kamath; Sankalp Mittal; Amit Deshpande; Vineeth Balasubramanian
SURER: Structure-Adaptive Unified Graph Neural Network for Multi-View Clustering	Jing Wang; Songhe Feng; Gengyu Lyu; Jiazheng Yuan
Sparse Variational Student-t Processes	Jian Xu; Delu Zeng
GSDD: Generative Space Dataset Distillation for Image Super-resolution	Haiyu Zhang; Shaolin Su; Yu Zhu; Jinqiu Sun; Yanning Zhang
FedMut: Generalized Federated Learning via Stochastic Mutation	Ming Hu; Cao Yue; Anran Li; Zhiming Li; Chengwei Liu; Tianlin Li; Mingsong Chen; Yang Liu
DR-Label: Label Deconstruction and Reconstruction of GNN Models for Catalysis Systems	Bowen Wang; Chen Liang; Jiaze Wang; Jiezhong Qiu; Furui Liu; Shaogang Hao; Dong Li; Guangyong Chen; Xiaolong Zou; Pheng-Ann Heng
CLIP-Guided Federated Learning on Heterogeneity and Long-Tailed Data	Jiangming Shi; Shanshan Zheng; Xiangbo Yin; Yang Lu; Yuan Xie; Yanyun Qu



Point-PEFT: Parameter-Efficient Fine-Tuning for 3D Pre-trained Models	Yiwen Tang; Ray Zhang; Zoey Guo; Xianzheng Ma; Bin Zhao; Zhigang Wang; Dong Wang; Xuelong Li
Project-Fair and Truthful Mechanisms for Budget Aggregation	Rupert Freeman; Ulrike Schmidt-Kraepelin
Identification of Necessary Semantic Undertakers in the Causal View for Image-Text Matching	Huatian Zhang; Lei Zhang; Kun Zhang; Zhendong Mao
PointPatchMix: Point Cloud Mixing with Patch Scoring	Yi Wang; Jiaze Wang; Jinpeng Li; Zixu Zhao; Guangyong Chen; Anfeng Liu; Pheng-Ann Heng
Commonsense for Zero-Shot Natural Language Video Localization	Meghana Holla; Ismini Lourentzou
REPrune: Channel Pruning via Kernel Representative Selection	Mincheol Park; Dongjin Kim; Cheonjun Park; Yuna Park; Gyeong Eun Gong; Won Woo Ro; Suhyun Kim
Seeing Dark Videos via Self-Learned Bottleneck Neural Representation	Haofeng Huang; Wenhan Yang; Lingyu Duan; Jiaying Liu
NeBLa: Neural Beer-Lambert for 3D Reconstruction of Oral Structures from Panoramic Radiographs	Sihwa Park; Seongjun Kim; Doeyoung Kwon; Yohan Jang; In-Seok Song; Seung Baek
Towards Improved Proxy-Based Deep Metric Learning via Data-Augmented Domain Adaptation	Li Ren; Chen Chen; Liqiang Wang; Kien Hua
FD3D: Exploiting Foreground Depth Map for Feature-Supervised Monocular 3D Object Detection	Zizhang Wu; Yuanzhu Gan; Yunzhe Wu; Ruihao Wang; Xiaoquan Wang; Jian Pu
Implicit Modeling of Non-rigid Objects with Cross-Category Signals	Yuchun Liu; Benjamin Planche; Meng Zheng; Zhongpai Gao; Pierre Sibut-Bourde; Fan Yang; Terrence Chen; Ziyang Wu
Thompson Sampling for Real-Valued Combinatorial Pure Exploration of Multi-Armed Bandit	Shintaro Nakamura; Masashi Sugiyama
Improving Neural Network Generalization on Data-Limited Regression with Doubly-Robust Boosting	Hao Wang
Large-Scale Non-convex Stochastic Constrained Distributionally Robust Optimization	Qi Zhang; Yi Zhou; Ashley Prater-Bennette; Lixin Shen; Shaofeng Zou
QPEN: Quantum Projection and Quantum Entanglement Enhanced Network for Cross-Lingual Aspect-Based Sentiment Analysis	Xingqiang Zhao; Hai Wan; Kunxun Qi
Approximating the Shapley Value without Marginal Contributions	Patrick Kolpaczki; Viktor Bengs; Maximilian Muschalik; Eyke Hüllermeier
SHaRPOSE: Sparse High-Resolution Representation for Human Pose Estimation	Xiaoqi An; Lin Zhao; Chen Gong; Nannan Wang; Di Wang; Jian Yang
Anchoring Path for Inductive Relation Prediction in Knowledge Graphs	Zhixiang Su; Di Wang; Chunyan Miao; Lizhen Cui
Distribution Matching for Multi-Task Learning of Classification Tasks: A Large-Scale Study on Faces & Beyond	Dimitrios Kollias; Viktoriia Sharmanska; Stefanos Zafeiriou
FeatWalk: Enhancing Few-Shot Classification through Local View Leveraging	Dalong Chen; Jianjia Zhang; Wei-shi Zheng; Ruixuan Wang
End-to-End Learning of LTLf Formulae by Faithful LTLf Encoding	Hai Wan; Pingjia Liang; Jianfeng Du; Weilin Luo; Rongzhen Ye; Bo Peng
Pano-NeRF: Synthesizing High Dynamic Range Novel Views with Geometry from Sparse Low Dynamic Range Panoramic Images	Zhan Lu; Qian Zheng; Boxin Shi; Xudong Jiang
Backpropagation Through Agents	Zhiyuan Li; Wenshuai Zhao; Lijun Wu; Joni Pajarinen
Self-Prompt Mechanism for Few-Shot Image Recognition	Mingchen Song; Huiqiang Wang; Guoqiang Zhong
ECHO-GL: Earnings Calls-Driven Heterogeneous Graph Learning for Stock Movement Prediction	Mengpu Liu; Mengying Zhu; Xiuyuan Wang; Guofang Ma; Jianwei Yin; Xiaolin Zheng





An Embedding-Unleashing Video Polyp Segmentation Framework via Region Linking and Scale Alignment	Zhixue Fang; Xinrong Guo; Jingyin Lin; Huisi Wu; Jing Qin
A Goal Interaction Graph Planning Framework for Conversational Recommendation	Xiaotong Zhang; Xuefang Jia; Han Liu; Xinyue Liu; Xianchao Zhang
A Dynamic Learning Method towards Realistic Compositional Zero-Shot Learning	Xiaoming Hu; Zilei Wang
LAFA: Multimodal Knowledge Graph Completion with Link Aware Fusion and Aggregation	Bin Shang; Yinliang Zhao; Jun Liu; Di Wang
Learning Safe Action Models with Partial Observability	Hai Le; Brendan Juba; Roni Stern
DeepSaDe: Learning Neural Networks That Guarantee Domain Constraint Satisfaction	Kshitij Goyal; Sebastijan Dumancic; Hendrik Blockeel
Lifting by Image – Leveraging Image Cues for Accurate 3D Human Pose Estimation	Feng Zhou; Jianqin Yin; Pelyang Li
Fine-Tuning Large Language Model Based Explainable Recommendation with Explainable Quality Reward	Mengyuan Yang; Mengying Zhu; Yan Wang; Linxun Chen; Yilei Zhao; Xiuyuan Wang; Bing Han; Xiaolin Zheng; Jianwei Yin
Few Shot Part Segmentation Reveals Compositional Logic for Industrial Anomaly Detection	Soopil Kim; Sion An; Philip Chikontwe; Myeongkyun Kang; Ehsan Adeli; Kilian Pohl; Sanghyun Park
Temporal Adaptive RGBT Tracking with Modality Prompt	Hongyu Wang; Xiaotao Liu; Yifan Li; Meng Sun; Dian Yuan; Jing Liu
Improving Open-Domain Dialogue Response Generation with Multi-Source Multilingual Commonsense Knowledge	Sixing Wu; Jiong Yu; Jiahao Chen; Xiaofan Deng; Wei Zhou
Learning Image Demoireing from Unpaired Real Data	Yunshan Zhong; Zhou Yuyao; Yuxin Zhang; Fei Chao; Rongrong Ji
Task-Free Continual Generation and Representation Learning via Dynamic Expansionable Memory Cluster	Fei Ye; Adrian Bors
Greedy-Based Online Fair Allocation with Adversarial Input: Enabling Best-of-Many-Worlds Guarantees	Zongjun Yang; Luofeng Liao; Christian Kroer
Counterfactual-Enhanced Information Bottleneck for Aspect-Based Sentiment Analysis	Mingshan Chang; Min Yang; Qingshan Jiang; Ruifeng Xu
Hand-Centric Motion Refinement for 3D Hand-Object Interaction via Hierarchical Spatial-Temporal Modeling	Yuze Hao; Jianrong Zhang; Tao Zhuo; Fuan Wen; Hehe Fan
LogoStyleFool: Vitiating Video Recognition Systems via Logo Style Transfer	Yuxin Cao; Ziyu Zhao; Xi Xiao; Derui Wang; Minhui Xue; Jin Lu
Weakly-Supervised Mirror Detection via Scribble Annotations	Mingfeng Zha; Yunqiang Pei; Guoqing Wang; Tianyu Li; Yang Yang; Wenbin Qian; Heng Tao Shen
FRIH: Fine-Grained Region-Aware Image Harmonization	Jinlong Peng; Zekun Luo; Liang Liu; Boshen Zhang
Robustness Verification of Deep Reinforcement Learning Based Control Systems Using Reward Martingales	Dapeng Zhi; Peixin Wang; Cheng Chen; Min Zhang
Domain-Controlled Prompt Learning	Qinglong Cao; Zhengqin Xu; Yuntian Chen; Chao Ma; Xiaokang Yang
Prompt-Based Distribution Alignment for Unsupervised Domain Adaptation	Shuanghao Bai; Min Zhang; Wanqi Zhou; Siteng Huang; Zhirong Luan; Donglin Wang; Badong Chen
OSFFNet: Omni-Stage Feature Fusion Network for Lightweight Image Super-resolution	Yang Wang; Tao Zhang
Runtime Analysis of the SMS-EMOA for Many-Objective Optimization	Weijie Zheng; Benjamin Doerr



Towards the Robustness of Differentially Private Federated Learning	Tao Qi; Huili Wang; Yongfeng Huang
Envy-Free House Allocation under Uncertain Preferences	Haris Aziz; Isaiah Iliffe; Bo Li; Angus Ritossa; Ankang Sun; Mashbat Suzuki
Prompt to Transfer: Sim-to-Real Transfer for Traffic Signal Control with Prompt Learning	Longchao Da; Minquan Gao; Hua Wei; Hao Mei
Dynamic Sub-graph Distillation for Robust Semi-supervised Continual Knowledge-Aware Neuron Interpretation for Scene Classification	Yan Fan; Yu Wang; Pengfei Zhu; Qinghua Hu
Revisiting Gradient Pruning: A Dual Realization for Defending against Gradient Attacks	Yong Guan; Freddy Lecue; Jiaoyan Chen; Ru Li; Jeff Z. Pan
ZO-AdaMU Optimizer: Adapting Perturbation by the Momentum and Uncertainty in Zeroth-Order Optimization	Lulu Xue; Shengshan Hu; Ruizhi Zhao; Leo yu Zhang; Shengqing Hu; Lichao Sun; Dezhong Yao
G2P-DDM: Generating Sign Pose Sequence from Gloss Sequence with Discrete Diffusion Model	Shuoran Jiang; Qingcai Chen; Yang Xiang; 煜康 林; Youcheng Pan; Xiangping Wu; Chuanyi Liu; Xiaobao Song
Multi-View Dynamic Reflection Prior for Video Glass Surface Detection	Pan Xie; Qipeng Zhang; Peng Taiying; Hao Tang; Yao Du; Zexian Li
Recurrent Graph Neural Networks and Their Connections to Bisimulation and Logic	Fang Liu; Yuhao Liu; Jiaying Lin; Ke Xu; Rynson W.H. Lau
Tackling Vision Language Tasks through Learning Inner Monologues	Maximilian Pflüger; David Tena Cucala; Egor Kostylev
When CEGAR Meets Regression: A Love Story in Optimal Classical Planning	Diji Yang; Kezhen Chen; Jinmeng Rao; Xiaoyuan Guo; Yawen Zhang; Jie Yang; Yi Zhang
Weak Distribution Detectors Lead to Stronger Generalizability of Vision-Language Prompt Tuning	Martín Pozo; Alvaro Torralba; Carlos Linares Lopez
Learning by Erasing: Conditional Entropy Based Transferable Out-of-Distribution Detection	Kun Ding; Haojian Zhang; Qiang Yu; Ying Wang; Shiming Xiang; Chunhong Pan
Full-Body Motion Reconstruction with Sparse Sensing from Graph Perspective	Meng Xing; Zhiyong Feng; Yong Su; Changjae Oh
SENCR: A Span Enhanced Two-Stage Network with Counterfactual Rethinking for Chinese NER	Feiyu Yao; Zongkai Wu; Li Yi
U-trustworthy Models. Reliability, Competence, and Confidence in Decision-Making	Hang Zheng; Qingsong Li; Shen Chen; Yuxuan Liang; Li Liu
CR-SAM: Curvature Regularized Sharpness-Aware Minimization	Ritwik Vashistha; Arya Farahi
Semantic-Aware Transformation-Invariant RoI Align	Tao Wu; Tie Luo; Donald Wunsch
Mean Teacher DETR with Masked Feature Alignment: A Robust Domain Adaptive Detection Transformer Framework	Guo-Ye Yang; Kiyohiro Nakayama; Zi-Kai Xiao; Tai-Jiang Mu; Sharon Xiaolei Huang; Shi-Min Hu
Polyper: Boundary Sensitive Polyp Segmentation	Weixi Weng; Chun Yuan
An Implicit Trust Region Approach to Behavior Regularized Offline Reinforcement Learning	Hao Shao; Yang Zhang; Qibin Hou
Another Way to the Top: Exploit Contextual Clustering in Learned Image M-BEV: Masked BEV Perception for Robust Autonomous Driving	Zhe Zhang; Xiaoyang Tan
You Only Read Once: Constituency-Oriented Relational Graph Convolutional Network for Multi-Aspect Multi-Sentiment Classification	Yichi Zhang; Zhihao Duan; Ming Lu; Dandan Ding; Fengqing Zhu; Zhan Ma
LogFormer: A Pre-train and Tuning Pipeline for Log Anomaly Detection	Siran Chen; Yue Ma; Yu Qiao; Yali Wang
	Yongqiang Zheng; Xia Li
	hongcheng Guo; Jian Yang; Jiaheng Liu; Jiaqi Bai; Boyang Wang; Zhoujun Li; Tieqiao Zheng; Bo Zhang; Junran Peng; Qi Tian



Adaptive Meta-Learning Probabilistic Inference Framework for Long Sequence Prediction	Jianping Zhu; Xin Guo; Yang Chen; Yao Yang; Li Wenbo; Bo Jin; Fei Wu
CL2CM: Improving Cross-Lingual Cross-Modal Retrieval via Cross-Lingual Knowledge Transfer	Yabing Wang; Fan Wang; Jianfeng Dong; Hao Luo
DrFuse: Learning Disentangled Representation for Clinical Multi-Modal Fusion with Missing Modality and Modal Inconsistency	Wenfang Yao; Kejing Yin; William Cheung; Jia Liu; Jing Qin
Mutual-Modality Adversarial Attack with Semantic Perturbation	Jingwen Ye; Ruonan Yu; Songhua Liu; Xinchao Wang
Multi-Label Supervised Contrastive Learning	Pingyue Zhang; Mengyue Wu
Causal Discovery from Poisson Branching Structural Causal Model Using High-Order Cumulant with Path Analysis	Jie Qiao; Yu Xiang; Zhengming Chen; Ruichu Cai; Zhifeng Hao
Unsupervised Neighborhood Propagation Kernel Layers for Semi-supervised Node Classification	Sonny Achten; Francesco Tonin; Panagiotis Patrinos; Johan Suykens
Hyperbolic Graph Diffusion Model	Lingfeng Wen; Xuan Tang; Mingjie Ouyang; Xiangxiang Shen; Jian Yang; Daxin Zhu; Mingsong Chen; Xian Wei
Mitigating Label Noise through Data Ambiguation	Julian Lienen; Eyke Hüllermeier
Federated Graph Learning under Domain Shift with Generalizable Prototypes	Guancheng Wan; Wenke Huang; Mang Ye
Moral Uncertainty and the Problem of Fanaticism	Jazon Szabo; Natalia Criado; Jose Such; Sanjay Modgil
Sampling-Resilient Multi-Object Tracking	Zepeng Li; Dongxiang Zhang; Sai Wu; Mingli Song; Gang Chen
No More Shortcuts: Realizing the Potential of Temporal Self-Supervision	Ishan Rajendrakumar Dave; Simon Jenni; Mubarak Shah
TEILP: Time Prediction over Knowledge Graphs via Logical Reasoning	Siheng Xiong; Yuan Yang; Ali Payani; James C Kerce; Faramarz Fekri
VLN-Video: Utilizing Driving Videos for Outdoor Vision-and-Language Navigation	Jialu Li; Aishwarya Padmakumar; Gaurav Sukhatme; Mohit Bansal
AGS: Affordable and Generalizable Substitute Training for Transferable Adversarial Attack	Ruikui Wang; Yuanfang Guo; Yunhong Wang
AMSP-UOD: When Vortex Convolution and Stochastic Perturbation Meet Underwater Object Detection	Jingchun Zhou; Zongxin He; Kin-Man Lam; Yudong Wang; Weishi Zhang; Chunle Guo; Chongyi Li
Learning Continuous Implicit Field with Local Distance Indicator for Arbitrary-Scale Point Cloud Upsampling	Shujuan Li; Junsheng Zhou; Baorui Ma; Yu-Shen Liu; Zhizhong Han
Automatic Radiology Reports Generation via Memory Alignment Network	Hongyu Shen; Mingtao Pei; Juncai Liu; Zhaoxing Tian
RWMS: Reliable Weighted Multi-Phase for Semi-supervised Segmentation	Wensi Liu; Xiao-Yu Tang; Chong Yang; Chunjie Yang
MorphVAE: Advancing Morphological Design of Voxel-Based Soft Robots with Variational Autoencoders	Junru Song; Yang Yang; Wei Peng; Weien Zhou; Feifei Wang; Wen Yao
Improving Cross-Modal Alignment with Synthetic Pairs for Text-Only Image Captioning	Zhiyue Liu; Jinyuan Liu; Fanrong Ma
VPDETR: End-to-End Vanishing Point DETection TRansformers	Taiyan Chen; Xianghua Ying; Jinfa Yang; Ruibin Wang; Ruohao Guo; Bowei Xing; Ji Shi
Graph Learning in 4D: A Quaternion-Valued Laplacian to Enhance Spectral GCNs	Stefano Fiorini; Stefano Coniglio; Michele Ciavotta; Enza Messina
Limited Memory Online Gradient Descent for Kernelized Pairwise Learning with Dynamic Averaging	Hilal AlQuabeh; William de Vazelhes; Bin Gu



Privileged Prior Information Distillation for Image Matting	Cheng Lyu; Jiake Xie; Bo Xu; Cheng Lu; Han Huang; Xin Huang; Ming Wu; Chuang Zhang; Yong Tang
Low-Rank Kernel Tensor Learning for Incomplete Multi-View Clustering	Tingting Wu; Songhe Feng; Jiazheng Yuan
Successive POI Recommendation via Brain-Inspired Spatiotemporal Aware Representation	Gehua Ma; He Wang; Jingyuan Zhao; Rui Yan; Huajin Tang
Semantic-Aware Data Augmentation for Text-to-Image Synthesis	Zhaorui Tan; Xi Yang; Kaizhu Huang
Heterogeneous Test-Time Training for Multi-Modal Person Re-identification	Zi Wang; Huaibo Huang; Aihua Zheng; Ran He
Theoretical Aspects of Generating Instances with Unique Solutions: Pre-assignment Models for Unique Vertex Cover	Takashi Horiyama; Yasuaki Kobayashi; Hirotaka Ono; Kazuhisa Seto; Ryu Suzuki
A General Implicit Framework for Fast NeRF Composition and Rendering	Xinyu Gao; Ziyi Yang; Yunlu Zhao; Yuxiang Sun; Xiaogang Jin; Changqing Zou
SDGAN: Disentangling Semantic Manipulation for Facial Attribute Editing	Wenmin Huang; Weiqi Luo; Jiwu Huang; Xiaochun Cao
Intrinsic Phase-Preserving Networks for Depth Super Resolution	Xuanhong Chen; Hang Wang; Jinfan Liu; Kairui Feng; Jialiang Chen; Xiaohang Wang; Bingbing Ni; Weimin Zhang
Towards Automated Chinese Ancient Character Restoration: A Diffusion-Based Method with a New Dataset	Haolong Li; Chenghao Du; Ziheng Jiang; Yifan Zhang; Jiawei Ma; Chen Ye
FairSIN: Achieving Fairness in Graph Neural Networks through Sensitive Information Neutralization	Cheng Yang; Jixi Liu; Yunhe Yan; Chuan Shi
Zero-Sum Games between Mean-Field Teams: Reachability-Based Analysis under Mean-Field Sharing	Yue Guan; Mohammad Afshari; Panagiotis Tsiotras
Improving GNN Calibration with Discriminative Ability: Insights and Strategies	Yujie Fang; Xin Li; Qianyu Chen; Mingzhong Wang
Cross-Modal Match for Language Conditioned 3D Object Grounding	Yachao Zhang; Runze Hu; Ronghui Li; Yanyun Qu; Yuan Xie; Xiu Li
Semi-supervised Class-Agnostic Motion Prediction with Pseudo Label Regeneration and BEVMix	Kewei Wang; Yizheng Wu; Zhiyu Pan; Xingyi Li; Ke Xian; Zhe Wang; Zhiguo Cao; Guosheng Lin
Ced-NeRF: A Compact and Efficient Method for Dynamic Neural Radiance Fields	Youtian Lin
Learning Efficient and Robust Multi-Agent Communication via Graph Information Bottleneck	Shifei Ding; Wei Du; Ling Ding; Lili Guo; Jian Zhang
SemTra: A Semantic Skill Translator for Cross-Domain Zero-Shot Policy Adaptation	Sangwoo Shin; Minjong Yoo; Jeongwoo Lee; Honguk Woo
Bidirectional Temporal Plan Graph: Enabling Switchable Passing Orders for More Efficient Multi-Agent Path Finding Plan Execution	Yifan Su; Rishi Veerapaneni; Jiaoyang Li
Quality-Diversity Generative Sampling for Learning with Synthetic Data	Allen Chang; Matthew C. Fontaine; Serena Booth; Maja J. Mataric; Stefanos Nikolaidis
Operator-Learning-Inspired Modeling of Neural Ordinary Differential Equations	Woojin Cho; Seunghyeon Cho; Hyundong Jin; Jinsung Jeon; Kookjin Lee; Sanghyun Hong; Dongeun Lee; Jonghyun Choi; Noseong Park
Considering Nonstationary within Multivariate Time Series with Variational Hierarchical Transformer for Forecasting	Muyao Wang; Wenchao Chen; Bo Chen
Enhancing Multi-Scale Diffusion Prediction via Sequential Hypergraphs and Adversarial Learning	Pengfei Jiao; Hongqian Chen; Qing Bao; Wang Zhang; Huaming Wu
TaskLAMA: Probing the Complex Task Understanding of Language Models	Quan Yuan; Mehran Kazemi; Xin Xu; Isaac Noble; Vaiva Imbrasaite; Deepak Ramachandran



RPSC: Robust Pseudo-Labeling for Semantic Clustering	Sihang Liu; Wenming Cao; Ruigang Fu; Kaixiang Yang; Zhiwen Yu
VSFormer: Visual-Spatial Fusion Transformer for Correspondence Pruning	Tangfei Liao; Xiaoqin Zhang; Li Zhao; Tao Wang; Guobao Xiao
Novelty vs. Potential Heuristics: A Comparison of Hardness Measures for Satisficing Planning	Simon Dold; Malte Helmert
Exploiting Label Skews in Federated Learning with Model Concatenation	Yiqun Diao; Qinbin Li; Bingsheng He
High-Fidelity Gradient Inversion in Distributed Learning	Zipeng Ye; Wenjian Luo; Qi Zhou; Yubo Tang
WaveNet: Tackling Non-stationary Graph Signals via Graph Spectral Wavelets	Zhirui Yang; Yulan Hu; Sheng Ouyang; Jingyu Liu; Shuqiang Wang; Xibo Ma; Wenhan Wang; Hanjing Su; Yong Liu
Multitarget Device-Free Localization via Cross-Domain Wi-Fi RSS Training Data and Attentional Prior Fusion	Na FAN; Zeyue Tian; Amartansh DUBEY; Samruddhi Deshmukh; Ross Murch; Qifeng Chen
SelfPromer: Self-Prompt Dehazing Transformers with Depth-Consistency	Cong Wang; Jinshan Pan; Wanyu LIN; Jiangxin Dong; Wei Wang; Xiao-Ming Wu
Regret Analysis of Policy Gradient Algorithm for Infinite Horizon Average Reward Markov Decision Processes	Qinbo Bai; Washim Uddin Mondal; Vaneet Aggarwal
Exact, Fast and Expressive Poisson Point Processes via Squared Neural Families	Russell Tsuchida; Cheng Soon Ong; Dino Sejdinovic
Learning with Noisy Labels Using Hyperspherical Margin Weighting	Shuo Zhang; Yuwen Li; Zhongyu Wang; Jianqing Li; Chengyu Liu
Bootstrapping Cognitive Agents with a Large Language Model	Feiyu Zhu; Reid Simmons
Unveiling Implicit Deceptive Patterns in Multi-Modal Fake News via Neuro-Symbolic Reasoning	Yiqi Dong; Dongxiao He; Xiaobao Wang; Youzhu Jin; Meng Ge; Carl Yang; Di Jin
Learning Generalized Medical Image Segmentation from Decoupled Feature Queries	Qi Bi; Jingjun Yi; Hao Zheng; Wei Ji; Yawen Huang; Yuexiang Li; Yefeng Zheng
E2E-AT: A Unified Framework for Tackling Uncertainty in Task-Aware End-to-End Learning	Wangkun Xu; Jianhong Wang; Fei Teng
DDAE: Towards Deep Dynamic Vision BERT Pretraining	Honghao Chen; Xiangwen Kong; Xiangyu Zhang; Xin Zhao; Kaiqi Huang
WeakPCSOD: Overcoming the Bias of Box Annotations for Weakly Supervised Point Cloud Salient Object Detection	Jun Wei; S. Kevin Zhou; Shuguang Cui; Zhen Li
Towards Fair Graph Federated Learning via Incentive Mechanisms	Chenglu Pan; Jiarong Xu; Yue Yu; Ziqi Yang; Qingbiao Wu; Chunping Wang; Lei CHEN; Yang Yang
Multi-Energy Guided Image Translation with Stochastic Differential Equations for Near-Infrared Facial Expression Recognition	Bingjun Luo; Zewen Wang; Jinpeng Wang; Junjie Zhu; Xibin Zhao; Yue Gao
A Generalized Neural Diffusion Framework on Graphs	Yibo Li; Xiao Wang; Hongrui Liu; Chuan Shi
Occluded Person Re-identification via Saliency-Guided Patch Transfer	Lei Tan; Jiaer Xia; Wenfeng Liu; Pingyang Dai; Yongjian Wu; Liujuan Cao
Graph Contrastive Invariant Learning from the Causal Perspective	Yanhu Mo; Xiao Wang; Shaohua Fan; Chuan Shi
Robust Loss Functions for Training Decision Trees with Noisy Labels	Jonathan Wilton; Nan Ye
Local-Global Multi-Modal Distillation for Weakly-Supervised Temporal Video Grounding	Peijun Bao; Yong Xia; Wenhan Yang; Boon Poh Ng; Meng Hwa Er; Alex Kot
Geometric-Facilitated Denoising Diffusion Model for 3D Molecule Generation	Can Xu; Haosen Wang; Weigang Wang; Zheng Pengfei; Hongyang Chen
Online Sensitivity Optimization in Differentially Private Learning	Filippo Galli; Catuscia Palamidessi; Tommaso Cucinotta
Deciphering Compatibility Relationships with Textual Descriptions via Extraction and Explanation	Yu Wang; Zexue He; Zhankui He; Hao Xu; Julian McAuley



SEIT: Structural Enhancement for Unsupervised Image Translation in Frequency Domain	Zhifeng Zhu; Yaochen Li; Yifan Li; Jinhua Yang; Peijun Chen; Yuehu Liu
Structural Entropy Based Graph Structure Learning for Node Classification	Liang Duan; Xiang Chen; Wenjie Liu; Daliang Liu; Kun Yue; Angsheng Li
Accelerating Cutting-Plane Algorithms via Reinforcement Learning Surrogates	Kyle Mana; Fernando Acero; Stephen Mak; Parisa Zehtabi; Michael Cashmore; Daniele Magazzeni; Manuela Veloso
Bayesian Inference with Complex Knowledge Graph Evidence	Armin Toroghi; Scott Sanner
MagiCapture: High-Resolution Multi-Concept Portrait Customization	Junha Hyung; Jaeyo Shin; Jaegul Choo
Long-Tailed Partial Label Learning by Head Classifier and Tail Classifier Cooperation	Yuheng Jia; Xiaorui Peng; Ran Wang; Min-Ling Zhang
Towards More Likely Models for AI Planning	Turgay Caglar; Sirine Belhaj; Tathagata Chakraborty; Michael Katz; Sarath Sreedharan
Liberating Seen Classes: Boosting Few-Shot and Zero-Shot Text Classification via Anchor Generation and Classification Reframing	Han Liu; Siyang Zhao; Xiaotong Zhang; Feng Zhang; Wei Wang; Fenglong Ma; Hongyang Chen; Hong Yu; Xianchao Zhang
Designing Biological Sequences without Prior Knowledge Using Evolutionary Reinforcement Learning	Xi Zeng; Xiaotian Hao; Hongyao Tang; Zhentao Tang; Shaoqing Jiao; Dazhi Lu; Jiajie Peng
Fair Participation via Sequential Policies	Reilly Raab; Ross Boczar; Maryam Fazel; Yang Liu
Towards Automated RISC-V Microarchitecture Design with Reinforcement Learning	Chen BAI; Jianwang Zhai; Yuzhe Ma; Bei Yu; Martin Ding Fat Wong
A Twist for Graph Classification: Optimizing Causal Information Flow in Graph Neural Networks	Zhe Zhao; Pengkun Wang; HaiBin Wen; Yudong Zhang; Zhengyang Zhou; Yang Wang
HACDR-Net: Heterogeneous-Aware Convolutional Network for Diabetic Retinopathy Multi-Lesion Segmentation	QiHao Xu; Xiaoling Luo; Chao Huang; Chengliang Liu; Jie Wen; Jiale Wang; Yong Xu
DC-NAS: Divide-and-Conquer Neural Architecture Search for Multi-Modal Classification	Xinyan Liang; Pinhan Fu; Qian Guo; Keyin Zheng; Yuhua Qian
Cross-Domain Contrastive Learning for Time Series Clustering	Furong Peng; Jiachen Luo; Xuan Lu; Sheng Wang; Feijiang Li
Frugal LMs Trained to Invoke Symbolic Solvers Achieve Parameter-Efficient Arithmetic Reasoning	Subhabrata Dutta; Ishan Pandey; Joykirat Singh; Sunny Manchanda; Soumen Chakrabarti; Tanmoy Chakraborty
Ahpatron: A New Budgeted Online Kernel Learning Machine with Tighter Mistake Bound	Yun Liao; Junfan Li; Shizhong Liao; Qinghua Hu; Jianwu Dang
KG-TREAT: Pre-training for Treatment Effect Estimation by Synergizing Patient Data with Knowledge Graphs	Ruoqi Liu; Lingfei Wu; Ping Zhang
ResDiff: Combining CNN and Diffusion Model for Image Super-resolution	Shuyao Shang; Zhengyang Shan; Guangxing Liu; LunQian Wang; XingHua Wang; Zekai Zhang; Jinglin Zhang
Defeasible Normative Reasoning: A Proof-Theoretic Integration of Logical Argumentation	Ofer Arieli; Kees van Berkel; Christian Straßer
WaveFormer: Wavelet Transformer for Noise-Robust Video Inpainting	Zhiliang Wu; Changchang Sun; Hanyu Xuan; Gaowen Liu; Yan Yan
Learning Multi-Task Sparse Representation Based on Fisher Information	Yayu Zhang; Yuhua Qian; Guoshuai Ma; Keyin Zheng; Guoqing Liu; Qingfu Zhang
Adaptive Uncertainty-Based Learning for Text-Based Person Retrieval	Shenshen Li; Chen He; Xing Xu; Fumin Shen; Yang Yang; Heng Tao Shen
Learning Subject-Aware Cropping by Outpainting Professional Photos	James Hong; Lu Yuan; Michaël Gharbi; Matthew Fisher; Kayvon Fatahalian
GridFormer: Point-Grid Transformer for Surface Reconstruction	Shengtao Li; Ge Gao; Yudong Liu; Yu-Shen Liu; Ming Gu





Decoupling Meta-Reinforcement Learning with Gaussian Task Contexts and Skills	Hongcai He; Anjie Zhu; Shuang Liang; Feiyu Chen; Jie Shao
Multi-Dimensional Fair Federated Learning	Cong Su; Guoxian Yu; Jun Wang; Hui Li; Qingzhong Li; Han Yu
Nearly Equitable Allocations beyond Additivity and Monotonicity	Siddharth Barman; Umang Bhaskar; Yeshwant Pandit; Soumyajit Pyne
Mono3DVG: 3D Visual Grounding in Monocular Images	Yang Zhan; Yuan Yuan; Zhitong Xiong
Relightable and Animatable Neural Avatars from Videos	Wenbin Lin; Chengwei Zheng; Junhai Yong; Feng Xu
FedNS: A Fast Sketching Newton-Type Algorithm for Federated Learning	Jian Li; Yong Liu; Weiping Wang
GSN: Generalisable Segmentation in Neural Radiance Field	Vinayak Gupta; Rahul Goel; Sirikonda Dhawal; P. J. Narayanan
No Head Left Behind – Multi-Head Alignment Distillation for Transformers	Tianyang Zhao; Kunwar Singh; Srikar Appalaraju; Peng Tang; Vijay Mahadevan; R. Manmatha; Ying Nian Wu
CycleVTON: A Cycle Mapping Framework for Parser-Free Virtual Try-On	Chenghu Du; Junyin Wang; Yi Rong; Shuqing Liu; Kai Liu; Shengwu Xiong
Focus-Then-Decide: Segmentation-Assisted Reinforcement Learning	Chao Chen; Jiacheng Xu; Weijian Liao; Hao Ding; Zongzhang Zhang; Yang Yu; Rui Zhao
Multi-Level Cross-Modal Alignment for Image Clustering	Liping Qiu; Qin Zhang; Xiaojun Chen; Shaotian Cai
ParaGuide: Guided Diffusion Paraphrasers for Plug-and-Play Textual Style Transfer	Zachary Horvitz; Ajay Patel; Chris Callison-Burch; Zhou Yu; Kathleen McKeown
Transition-Informed Reinforcement Learning for Large-Scale Stackelberg Mean-Field Games	Pengdeng Li; Runsheng Yu; Xinrun Wang; Bo An
Reliable Conflictive Multi-View Learning	Cai Xu; Jiajun Si; Ziyu Guan; Wei Zhao; Yue Wu; Xiyue Gao
Painterly Image Harmonization by Learning from Painterly Objects	Li Niu; Junyan Cao; Yan Hong; Liqing Zhang
TDeLTA: A Light-Weight and Robust Table Detection Method Based on Learning Text Arrangement	Yang Fan; Xiangping Wu; Qingcai Chen; Heng Li; Yan Huang; Zhixiang Cai; Qitian Wu
Learning Cluster-Wise Anchors for Multi-View Clustering	Chao Zhang; Xiuyi Jia; Zechao Li; Chunlin Chen; Huaxiong Li
Modeling Stereo-Confidence out of the End-to-End Stereo-Matching Network via Disparity Plane Sweep	Jae Young Lee; Woonghyun Ka; Jaehyun Choi; Junmo Kim
S2WAT: Image Style Transfer via Hierarchical Vision Transformer Using Strips Window Attention	Chiyu Zhang; Xiaogang Xu; Lei Wang; Zaiyan Dai; Jun Yang
EcomGPT: Instruction-Tuning Large Language Models with Chain-of-Task Tasks for E-commerce	Li Yangning; Shirong Ma; Xiaobin Wang; Shen Huang; Chengyue Jiang; Hai-Tao Zheng; Pengjun Xie; Fei Huang; Yong Jiang
Open-Set Graph Domain Adaptation via Separate Domain Alignment	Yu Wang; Ronghang Zhu; Pengsheng Ji; Sheng Li
Probability-Polarized Optimal Transport for Unsupervised Domain Adaptation	Yan Wang; Chuan-Xian Ren; Yi-Ming Zhai; You-Wei Luo; Hong Yan
DS-AL: A Dual-Stream Analytic Learning for Exemplar-Free Class-Incremental Learning	Huiping Zhuang; Run He; Kai Tong; Ziqian Zeng; Cen Chen; Zhiping Lin
Temporally and Distributionally Robust Optimization for Cold-Start Recommendation	Xinyu Lin; Wenjie Wang; Jujia Zhao; Yongqi Li; Fuli Feng; Tat-Seng Chua
KGTS: Contrastive Trajectory Similarity Learning over Prompt Knowledge Graph Embedding	Zhen Chen; Dalin Zhang; Shanshan Feng; Kaixuan Chen; Lisi Chen; Peng Han; Shuo Shang
Supervision Interpolation via LossMix: Generalizing Mixup for Object Detection and Beyond	Thanh Vu; Baochen Sun; Bodi Yuan; Alex Ngai; Yueqi Li; Jan-Michael Frahm
Image Safeguarding: Reasoning with Conditional Vision Language Model and Obfuscating Unsafe Content Counterfactually	Mazal Bethany; Brandon Wherry; Nishant Vishwamitra; Peyman Najafirad



Stitching Segments and Sentences towards Generalization in Video-Text Pre-training	Fan Ma; Xiaojie Jin; Heng Wang; Jingjia Huang; Linchao Zhu; Yi Yang
BOK-VQA: Bilingual outside Knowledge-Based Visual Question Answering via Graph Representation Pretraining	Minjun Kim; SeungWoo Song; Youhan Lee; Haneol Jang; KyungTae Lim
On the Role of Server Momentum in Federated Learning	Jianhui Sun; Xidong Wu; Heng Huang; Aidong Zhang
What Does a Query Answer Tell You? Informativeness of Query Answers for Knowledge Bases	Luca Andolfi; Gianluca Cima; Marco Console; Maurizio Lenzerini
Multi-Modality Affinity Inference for Weakly Supervised 3D Semantic Segmentation	Xiawei Li; Qingyuan Xu; Jing Zhang; Tianyi Zhang; Qian Yu; Lu Sheng; Dong Xu
Minibatch Stochastic Three Points Method for Unconstrained Smooth Minimization	Soumia Boucherouite; Grigory Malinovsky; Peter Richtarik; El houcin Bergou
VLCounter: Text-Aware Visual Representation for Zero-Shot Object Counting	Seunggu Kang; WonJun Moon; Euiyeon Kim; Jae-Pil Heo
Monte Carlo Tree Search in the Presence of Transition Uncertainty	Farnaz Kohankhaki; Kiarash Aghakasiri; Hongming Zhang; Ting-Han Wei; Chao Gao; Martin Müller
Permutation-Based Hypothesis Testing for Neural Networks	Francesca Mandel; Ian Barnett
Out-of-Distribution Detection in Long-Tailed Recognition with Calibrated Outlier Class Learning	Wenjun Miao; Guansong Pang; Xiao Bai; Tianqi Li; Jin Zheng
SoftCLIP: Softer Cross-Modal Alignment Makes CLIP Stronger	Yuting Gao; Jinfeng Liu; Zihan Xu; Tong Wu; Enwei Zhang; Ke Li; Jie Yang; Wei Liu; Xing Sun
Learning Accurate and Bidirectional Transformation via Dynamic Embedding Transportation for Cross-Domain Recommendation	Weiming Liu; Chaochao Chen; Xinting Liao; Mengling Hu; Yanchao Tan; Fan Wang; Xiaolin Zheng; Yew Soon Ong
SAME: Sample Reconstruction against Model Extraction Attacks	Yi Xie; Jie Zhang; Shiqian Zhao; Tianwei Zhang; Xiaofeng Chen
StyleSinger: Style Transfer for Out-of-Domain Singing Voice Synthesis	Yu Zhang; Rongjie Huang; Ruiqi Li; Jinzheng He; Yan Xia; Chen Feiyang; Xinyu Duan; Baoxing Huai; Zhou Zhao
DeepAccident: A Motion and Accident Prediction Benchmark for V2X Autonomous Driving	Tianqi Wang; Sukmin Kim; Ji Wenxuan; Enze Xie; Chongjian Ge; Junsong Chen; Zhenguo Li; Ping Luo
Reinforcement Learning and Data-Generation for Syntax-Guided Synthesis	Julian Parsert; Elizabeth Polgreen
The Irrelevance of Influencers: Information Diffusion with Re-Activation and Immunity Lasts Exponentially Long on Social Network Models	Tobias Friedrich; Andreas Göbel; Nicolas Klodt; Martin S. Krejca; Marcus Pappik
Can Large Language Models Understand Real-World Complex Instructions?	Qianyu He; Jie Zeng; Wenhao Huang; Lina Chen; Xiao Jin; 仟僖 何; Xunzhe Zhou; Jiaqing Liang; Yanghua Xiao
LMD: Faster Image Reconstruction with Latent Masking Diffusion	Zhiyuan Ma; Zhihuan Yu; Jianjun Li; Bowen Zhou
Hierarchical and Incremental Structural Entropy Minimization for Unsupervised Social Event Detection	Yuwei Cao; Hao Peng; Zhengtao Yu; Philip Yu
On Optimal Tradeoffs between EFX and Nash Welfare	Michal Feldman; Simon Mauras; Tomasz Ponitka
MoDE: A Mixture-of-Experts Model with Mutual Distillation among the Experts	Zhitian Xie; Yinger Zhang; Chenyi Zhuang; Qitao Shi; Zhining Liu; Jinjie Gu; Guannan Zhang
UPDP: A Unified Progressive Depth Pruner for CNN and Vision Transformer	Ji Liu; Dehua Tang; Yuanxian Huang; Li Zhang; Zeng Xiaocheng; Dong Li; Mingjie Lu; Jinzhang Peng; Yu Wang; Fan Jiang; Lu Tian; Ashish Sirasao
Editing Language Model-Based Knowledge Graph Embeddings	Siyuan Cheng; Ningyu Zhang; Bozhong Tian; Xi Chen; Qingbin Liu; Huajun



Spanning the Spectrum of Hatred Detection: A Persian Multi-Label Hate Speech Dataset with Annotator Rationales	Zahra Delbari; Nafise Sadat Moosavi; Mohammad Taher Pilehvar
Intentional Evolutionary Learning for Untrimmed Videos with Long Tail Distribution	Yuxi Zhou; Xiujie Wang; Jianhua Zhang; 佳佳 王; Jie Yu; Hao Zhou; Yi Gao; Shengyong Chen
VQ-FONT: Few-Shot Font Generation with Structure-Aware Enhancement and Quantization	Mingshuai Yao; Yabo Zhang; Xianhui Lin; Xiaoming Li; Wangmeng Zuo
Image as a Language: Revisiting Scene Text Recognition via Balanced, Unified and Synchronized Vision-Language Reasoning Network	Jiajun Wei; Hongjian Zhan; Yue Lu; Xiao Tu; Bing Yin; Cong Liu; Umapada Pal
Learning Bayesian Network Classifiers to Minimize the Class Variable Parameters	Shouta Sugahara; Koya Kato; Maomi Ueno
Twice Class Bias Correction for Imbalanced Semi-supervised Learning	Lan Li; Bowen Tao; Lu Han; De-Chuan Zhan; Han-Jia Ye
End-to-End RGB-D Image Compression via Exploiting Channel-Modality Redundancy	Huiming Zheng; Wei Gao
Fair Allocation of Items in Multiple Regions	Houyu Zhou; Tianze Wei; Biaoshuai Tao; Minming Li
CoreRec: A Counterfactual Correlation Inference for Next Set Recommendation	KeXin Li; Chengjiang Long; Shengyu Zhang; Xudong Tang; Zhichao Zhai; Kun Kuang; Jun Xiao
AesFA: An Aesthetic Feature-Aware Arbitrary Neural Style Transfer	Joonwoo Kwon; Sooyoung Kim; Yuewei Lin; Shinjae Yoo; Jiook Cha
Hyp-OW: Exploiting Hierarchical Structure Learning with Hyperbolic Distance Enhances Open World Object Detection	Thang Doan; Xin Li; Sima Behpour; Wenbin He; Liang Gou; Liu Ren
Generalising Planning Environment Redesign	Alberto Pozanco; Ramon Fraga Pereira; Daniel Borrajo
MusER: Musical Element-Based Regularization for Generating Symbolic Music with Emotion	Shulei Ji; Xinyu Yang
Weakly Supervised Open-Vocabulary Object Detection	Jianghang Lin; Yunhang Shen; Bingquan Wang; Shaohui Lin; Ke Li; Liujuan
Towards Stability and Generalization Bounds in Decentralized Minibatch Stochastic Gradient Descent	Jiahuan Wang; Hong Chen
Trend-Aware Supervision: On Learning Invariance for Semi-supervised Facial Action Unit Intensity Estimation	Yingjie Chen; Jiarui Zhang; Tao Wang; Yun Liang
Generalizable Fourier Augmentation for Unsupervised Video Object Segmentation	Huihui Song; Tiankang Su; Yuhui Zheng; Kaihua Zhang; Bo Liu; Dong Liu
Zero-Shot Aerial Object Detection with Visual Description Regularization	Chenyu Lin; Zhengqing Zang; Chenwei Tang; Tao Wang; Jiancheng Lv
CSL: Class-Agnostic Structure-Constrained Learning for Segmentation including the Unseen	Hao Zhang; Fang Li; Lu Qi; Ming-Hsuan Yang; Narendra Ahuja
Point2Real: Bridging the Gap between Point Cloud and Realistic Image for Open-World 3D Recognition	Hanxuan Li; Bin Fu; Ruiping Wang; Xilin Chen
GSENet: Global Semantic Enhancement Network for Lane Detection	Junhao Su; Zhenghan Chen; Chenghao He; Dongzhi Guan; Changpeng Cai; Tongxi Zhou; Jiashen Wei; Wenhua Tian; Zhihui Xie
DocNLC: A Document Image Enhancement Framework with Normalized and Latent Contrastive Representation for Multiple Degradations	Ruilu Wang; Yang Xue; Lianwen Jin
PrefAce: Face-Centric Pretraining with Self-Structure Aware Distillation	Siyuan Hu; Zheng Wang; Peng Hu; Xi Peng; Jie Wu; Hongyuan Zhu; Yew Soon Ong



Impartial Adversarial Distillation: Addressing Biased Data-Free Knowledge Distillation via Adaptive Constrained Optimization	Dongping Liao; Xitong Gao; Chengzhong Xu
GIN-SD: Source Detection in Graphs with Incomplete Nodes via Positional Encoding and Attentive Fusion	Le Cheng; Peican Zhu; Keke Tang; Chao Gao; Zhen Wang
On the Expressivity of Recurrent Neural Cascades	Nadezda Knorozova; Alessandro Ronca
Gramformer: Learning Crowd Counting via Graph-Modulated Transformer	Hui LIN; Zhiheng Ma; Xiaopeng Hong; Qinnan Shangguan; Deyu Meng
NightRain: Nighttime Video Deraining via Adaptive-Rain-Removal and Adaptive-Correction	Beibei Lin; Yeying Jin; Wending Yan; Wei Ye; Yuan Yuan; Shunli Zhang; Robby T. Tan
Federated Label-Noise Learning with Local Diversity Product Regularization	Xiaochen Zhou; Xudong Wang
A User-Friendly Framework for Generating Model-Preferred Prompts in Text-to-Image	Nailei Hei; Qianyu Guo; Zihao Wang; Yan Wang; Haofen Wang; Wenqiang Zhang
Modeling Knowledge Graphs with Composite Reasoning	Wanyun Cui; Linqiu Zhang
Locally Rainbow Paths	Till Fluschnik; Leon Kellerhals; Malte Renken
PCE-Palm: Palm Crease Energy Based Two-Stage Realistic Pseudo-Palmprint Generation	Lei Shen; Jianlong Jin; Ruixin Zhang; Jingyun Zhang; Ge Jin; Shouhong Ding; Yang Zhao; Wei Jia; Chenglong Zhao
Pushing the Limit of Fine-Tuning for Few-Shot Learning: Where Feature Reusing Meets Cross-Scale Attention	Ying-Yu Chen; Jun-Wei Hsieh; Xin Li; Ming-Ching Chang
Parameterization of (Partial) Maximum Satisfiability above Matching in a Variable-Clause Graph	Vasily Alferov; Ivan Bliznets; Kirill Brilliantov
Multi-Domain Incremental Learning for Face Presentation Attack Detection	Keyao Wang; Guosheng Zhang; Haixiao Yue; Ajian Liu; Gang Zhang; Haocheng Feng; Junyu Han; Errui Ding; Jingdong Wang
PaintHuman: Towards High-Fidelity Text-to-3D Human Texturing via Denoised Score Distillation	Jianhui Yu; Hao Zhu; Liming Jiang; Chen Change Loy; Weidong Cai; Wayne Wu
Point Transformer with Federated Learning for Predicting Breast Cancer HER2 Status from Hematoxylin and Eosin-Stained Whole Slide Images	Bao Li; Zhenyu Liu; Lizhi Shao; Bensheng Qiu; Hong Bu; Jie Tian
Harnessing the Power of Beta Scoring in Deep Active Learning for Multi-Label Text Classification	Wei Tan; Ngoc Dang Nguyen; Lan Du; Wray Buntine
Decomposing Temporal Equilibrium Strategy for Coordinated Distributed Multi-Agent Reinforcement Learning	Chenyang Zhu; Wen Si; Jinyu Zhu; Zhihao Jiang
Rethinking Graph Masked Autoencoders through Alignment and Uniformity	Liang Wang; Xiang Tao; Qiang Liu; Shu Wu; Liang Wang
DiDA: Disambiguated Domain Alignment for Cross-Domain Retrieval with Partial Labels	Haoran Liu; Ying Ma; Ming Yan; Yingke Chen; Dezhong Peng; Xu Wang
Knowledge-Aware Parameter Coaching for Personalized Federated Learning	Mingjian Zhi; Yuanguo Bi; Wenchao Xu; Haozhao Wang; tianao Xiang
Cached Transformers: Improving Transformers with Differentiable Memory Cachde	Zhaoyang Zhang; Wenqi Shao; Yixiao Ge; Xiaogang Wang; Jinwei Gu; Ping Luo
Relative Policy-Transition Optimization for Fast Policy Transfer	Jiawei xu; Cheng Zhou; Yizheng Zhang; Baoxiang Wang; Lei Han
What Are the Rules? Discovering Constraints from Data	Boris Wiegand; Dietrich Klakow; Jilles Vreeken
Learning Task-Aware Language-Image Representation for Class-Incremental Object Detection	Hongquan Zhang; Bin-Bin Gao; Yi Zeng; Xudong Tian; Xin Tan; Zhizhong Zhang; Yanyun Qu; Jun Liu; Yuan Xie
When Are Two Lists Better than One?: Benefits and Harms in Joint Decision-Making	Kate Donahue; Sreenivas Gollapudi; Kostas Kollias



Referred by Multi-Modality: A Unified Temporal Transformer for Video Object Segmentation	Shilin Yan; Renrui Zhang; Ziyu Guo; Wenchao Chen; Wei Zhang; Hongyang Li; Yu Qiao; Hao Dong; Zhongjiang He; Peng Gao
Decentralized Gradient-Free Methods for Stochastic Non-smooth Non-convex Optimization	Zhenwei Lin; Jingfan Xia; Qi Deng; Luo Luo
Proxyformer: Nystrom-Based Linear Transformer with Trainable Proxy Tokens	SangHo Lee; Hayun Lee; Dongkun Shin
High-Fidelity 3D Head Avatars Reconstruction through Spatially-Varying Expression Conditioned Neural Radiance Field	Minghan Qin; Yifan Liu; Yuelang Xu; Xiaochen Zhao; Yebin Liu; Haoqian Wang
Learning to Stop Cut Generation for Efficient Mixed-Integer Linear	Haotian Ling; Zhihai Wang; Jie Wang
Noise-Aware Image Captioning with Progressively Exploring Mismatched	Zhongtian Fu; Kefei Song; Luping Zhou; Yang Yang
Neural Physical Simulation with Multi-Resolution Hash Grid Encoding	Haoxiang Wang; Tao Yu; Tianwei Yang; Hui Qiao; Dai Qionghai
SDAC: A Multimodal Synthetic Dataset for Anomaly and Corner Case Detection in Autonomous Driving	Lei Gong; Yu Zhang; Yingqing Xia; Yanyong Zhang; Jianmin Ji
R3CD: Scene Graph to Image Generation with Relation-Aware Compositional Contrastive Control Diffusion	Jinxu Liu; Qi Liu
Time-Aware Knowledge Representations of Dynamic Objects with Multidimensional Persistence	Baris Coskunuzer; Ignacio Segovia-Dominguez; Yuzhou Chen; Yulia Gel
Calibrated One Round Federated Learning with Bayesian Inference in the Predictive Space	Mohsin Hasan; Guojun Zhang; Kaiyang Guo; Xi Chen; Pascal Poupart
Dual-Channel Learning Framework for Drug-Drug Interaction Prediction via Relation-Aware Heterogeneous Graph Transformer	Xiaorui Su; Pengwei Hu; Zhu-Hong You; Philip Yu; Lun Hu
Class-Attribute Priors: Adapting Optimization to Heterogeneity and Fairness Objective	Xuechen Zhang; Mingchen Li; Jiasi Chen; Christos Thrampoulidis; Samet Oymak
Levenshtein Distance Embedding with Poisson Regression for DNA Storage	Xiang Wei; Alan J.X. Guo; Sihan Sun; Mengyi Wei; Wei Yu
From Retrieval to Generation: A Simple and Unified Generative Model for End-to-End Task-Oriented Dialogue	Zeyuan Ding; Zhihao Yang; Ling Luo; Yuanyuan Sun; Hongfei Lin
AVSegFormer: Audio-Visual Segmentation with Transformer	Shengyi Gao; Zhe Chen; Guo Chen; Wenhai Wang; Tong Lu
SCP: Spherical-Coordinate-Based Learned Point Cloud Compression	Ao Luo; Linxin Song; Keisuke Nonaka; Kyohei Unno; Heming Sun; Masayuki Goto; Jiro Katto
Transient Glimpses: Unveiling Occluded Backgrounds through the Spike	Jiyuan Zhang; Shiyan Chen; Yajing Zheng; Zhaofei Yu; Tiejun Huang
Pandora's Problem with Deadlines	Ben Berger; Tomer Ezra; Michal Feldman; Federico Fusco
Restoring Speaking Lips from Occlusion for Audio-Visual Speech Recognition	Jiadong Wang; Zexu Pan; Malu Zhang; Robby Tan; Haizhou Li
Trust Region Methods for Nonconvex Stochastic Optimization beyond Lipschitz Smoothness	Chenghan Xie; Chenxi Li; Chuwen Zhang; Qi Deng; DongDong Ge; Yinyu Ye
Hot or Cold? Adaptive Temperature Sampling for Code Generation with Large Language Models	Yuqi Zhu; Jia Li; Ge Li; YunFei Zhao; Jia Li; Zhi Jin; Hong Mei
Retrieval-Augmented Primitive Representations for Compositional Zero-Shot Learning	Chenchen Jing; Yukun Li; Hao Chen; Chunhua Shen
Cooper: Coordinating Specialized Agents towards a Complex Dialogue Goal	Yi Cheng; Wenge Liu; Jian Wang; Chak Tou Leong; Yi Ouyang; Wenjie Li; Xian Wu; Yefeng Zheng
Keypoint Fusion for RGB-D Based 3D Hand Pose Estimation	Xingyu Liu; Pengfei Ren; Yuanyuan Gao; Jingyu Wang; Haifeng Sun; Qi Qi; Zirui Zhuang; Jianxin Liao



Learning to Manipulate Artistic Images	Wei Guo; Yuqi Zhang; De Ma; Qian Zheng
NESTER: An Adaptive Neurosymbolic Method for Causal Effect Estimation	Abbavaram Gowtham Reddy; Vineeth N Balasubramanian
Decomposing Semantic Shifts for Composed Image Retrieval	Xingyu Yang; Daqing Liu; Heng Zhang; Yong Luo; Chaoyue Wang; Jing Zhang
Enhancing Multi-Label Classification via Dynamic Label-Order Learning	Jiangnan Li; Yice ZHANG; Shiwei Chen; Ruifeng Xu
Structured Probabilistic Coding	Dou Hu; Lingwei Wei; Yaxin Liu; Wei Zhou; Songlin Hu
Expressive Multi-Agent Communication via Identity-Aware Learning	Wei Du; Shifei Ding; Lili Guo; Jian Zhang; Ling Ding
AMD: Anatomical Motion Diffusion with Interpretable Motion Decomposition and Fusion	Beibei Jing; Youjia Zhang; Zikai Song; Junqing Yu; Wei Yang
PerFedRLNAS: One-for-All Personalized Federated Neural Architecture	Dixi Yao; Baochun Li
GigaHumanDet: Exploring Full-Body Detection on Gigapixel-Level Images	Chenglong Liu; Haoran Wei; Jinze Yang; Jintao Liu; Wenxi Li; Yuchen Guo; Lu Fang
Evidential Uncertainty-Guided Mitochondria Segmentation for 3D EM Images	Ruohua Shi; Lingyu Duan; Tiejun Huang; Tingting Jiang
Knowledge Guided Semi-supervised Learning for Quality Assessment of User Generated Videos	Shankhanil Mitra; Rajiv Soundararajan
Detect Any Keypoints: An Efficient Light-Weight Few-Shot Keypoint Detector	Changsheng Lu; Piotr Koniusz
CFEVER: A Chinese Fact Extraction and VERification Dataset	Ying-Jia Lin; ChunYi Lin; Chia-Jen Yeh; Yi-Ting Li; YunYu Hu; Chih-Hao Hsu; MeiFeng Lee; Hung-Yu Kao
DiffusionTrack: Diffusion Model for Multi-Object Tracking	Run Luo; Zikai Song; Lintao Ma; Jinlin Wei; Wei Yang; Min Yang
Beyond TreeSHAP: Efficient Computation of Any-Order Shapley Interactions for Tree Ensembles	Maximilian Muschalik; Fabian Fumagalli; Barbara Hammer; Eyke Hüllermeier
Stereo Vision Conversion from Planar Videos Based on Temporal Multiplane Images	Shanding Diao; Yuan Chen; Yang Zhao; Wei Jia; Zhao Zhang; Ronggang Wang
PTUS: Photo-Realistic Talking Upper-Body Synthesis via 3D-Aware Motion Decomposition Warping	Luoyang Lin; Zutao Jiang; Xiaodan Liang; Liqian Ma; Michael C. Kampffmeyer; Xiaochun Cao
Narrowing the Gap between Supervised and Unsupervised Sentence Representation Learning with Large Language Model	Mingxin Li; Richong Zhang; Zhijie Nie; Yongyi Mao
Sequential Fusion Based Multi-Granularity Consistency for Space-Time Transformer Tracking	Kun Hu; Wenjing Yang; Wanrong Huang; Xianchen Zhou; Mingyu Cao; Jing Ren; Huibin Tan
Large-Scale Multi-Robot Coverage Path Planning via Local Search	Jingtao Tang; Hang Ma
Approximate Distance Oracle for Fault-Tolerant Geometric Spanners	Kyungjin Cho; Jihun Shin; Eunjin Oh
PTMQ: Post-training Multi-Bit Quantization of Neural Networks	Ke Xu; Zhongcheng Li; Shanshan Wang; Xingyi Zhang
ContactGen: Contact-Guided Interactive 3D Human Generation for Partners	Dongjun Gu; Jaehyeok Shim; Jaehoon Jang; Changwoo Kang; Kyungdon Joo
Cheaper and Faster: Distributed Deep Reinforcement Learning with Serverless Computing	Hanfei Yu; Jian Li; Yang Hua; Xu Yuan; Hao Wang
Data-Free Generalized Zero-Shot Learning	Bowen Tang; Jing Zhang; Yan Long; Qian Yu; Lu Sheng; Dong Xu
PC-Conv: Unifying Homophily and Heterophily with Two-Fold Filtering	Bingheng Li; Erlin Pan; Zhao Kang
Auditable Algorithms for Approximate Model Counting	S Akshay; Supratik Chakraborty; Kuldeep S Meel
Labels Need Prompts Too: Mask Matching for Natural Language Understanding Tasks	Bo Li; Wei Ye; Quansen Wang; Wen Zhao; Shikun Zhang
Any-Way Meta Learning	JunHoo Lee; Yearim Kim; Hyunho Lee; Nojun Kwak
Cycle Self-Refinement for Multi-Source Domain Adaptation	Chaoyang Zhou; Zengmao Wang; Bo Du; Yong Luo





Federated Adaptive Prompt Tuning for Multi-Domain Collaborative Learning	Shangchao Su; Mingzhao Yang; Bin Li; Xiangyang Xue
Machine-Created Universal Language for Cross-Lingual Transfer	Yaobo Liang; Quanzhi Zhu; Junhe Zhao; Nan Duan
Rethinking Mesh Watermark: Towards Highly Robust and Adaptable Deep 3D Mesh Watermarking	Xingyu Zhu; Guanhui Ye; Xiapu Luo; Xuetao Wei
Minimal Macro-Based Rewritings of Formal Languages: Theory and Applications in Ontology Engineering (and Beyond)	Christian Kindermann; Anne-Marie George; Bijan Parsia; Uli Sattler
Journey to the Center of the Knowledge Neurons: Discoveries of Language-Independent Knowledge Neurons and Degenerate Knowledge Neurons	Yuheng Chen; Pengfei Cao; Yubo Chen; Kang Liu; Jun Zhao
TREE-G: Decision Trees Contesting Graph Neural Networks	Maya Bechler-Speicher; Amir Globerson; Ran Gilad-Bachrach
PSC-CPI: Multi-Scale Protein Sequence-Structure Contrasting for Efficient and Generalizable Compound-Protein Interaction Prediction	Lirong Wu; Yufei Huang; Cheng Tan; Zhangyang Gao; Bozhen Hu; Haitao Lin; Zicheng Liu; Stan Z. Li
Reproduce, Replicate, Reevaluate. The Long but Safe Way to Extend Machine Learning Methods	Luisa Werner; Nabil Layaïda; Pierre Genevès; Jérôme Euzenat; Damien Graux
Span Graph Transformer for Document-Level Named Entity Recognition	Hongli Mao; Xian-Ling Mao; Hanlin Tang; Yu-Ming Shang; Heyan Huang
MASTER: Market-Guided Stock Transformer for Stock Price Forecasting	Tong Li; Zhaoyang Liu; Yanyan Shen; Xue Wang; Haokun Chen; Sen Huang
Knowledge Graph Error Detection with Contrastive Confidence Adaption	Xiangyu Liu; Yang Liu; Wei Hu
Responding to the Call: Exploring Automatic Music Composition Using a Knowledge-Enhanced Model	Zhejing Hu; Yan Liu; Gong Chen; Xiao MA; Sheng-hua Zhong; Qianwen Luo
Spatial Transform Decoupling for Oriented Object Detection	Hongtian Yu; Yunjie Tian; Qixiang Ye; Yunfan Liu
Improved Graph Contrastive Learning for Short Text Classification	Yonghao Liu; Lan Huang; Fausto Giunchiglia; Xiaoyue Feng; Renchu Guan
Coverage-Guaranteed Prediction Sets for Out-of-Distribution Data	Xin Zou; Weiwei Liu
Music Style Transfer with Time-Varying Inversion of Diffusion Models	Sifei Li; Yuxin Zhang; Fan Tang; Chongyang Ma; Weiming Dong; Changsheng Xu
Semantic-Guided Generative Image Augmentation Method with Diffusion Models for Image Classification	Bohan Li; Xiao Xu; Xinghao Wang; Yutai Hou; Yunlong Feng; Feng Wang; Xuanliang Zhang; Qingfu Zhu; Wanxiang Che
Patch-Wise Graph Contrastive Learning for Image Translation	Chanyong Jung; Gihyun Kwon; Jong Chul Ye
Test-Time Adaptation via Style and Structure Guidance for Histological Image Registration	Shenglong Zhou; Zhiwei Xiong; Feng Wu
DocFormerv2: Local Features for Document Understanding	Srikar Appalaraju; Peng Tang; Qi Dong; Nishant Sankaran; Yichu Zhou; R. Manmatha
Diverse and Stable 2D Diffusion Guided Text to 3D Generation with Noise Recalibration	Xiaofeng Yang; Fayao Liu; Yi Xu; Hanjing Su; Qingyao Wu; Guosheng Lin
Interactive Visual Task Learning for Robots	Weiwei Gu; Anant Sah; Nakul Gopalan
Visual Chain-of-Thought Prompting for Knowledge-Based Visual Reasoning	Zhenfang Chen; Qinhong Zhou; Yikang Shen; Yining Hong; Zhiqing Sun; Dan Gutfreund; Chuang Gan
DMMR: Cross-Subject Domain Generalization for EEG-Based Emotion Recognition via Denoising Mixed Mutual Reconstruction	Yiming Wang; Bin Zhang; Yujiao Tang
RL-SeqISP: Reinforcement Learning-Based Sequential Optimization for Image Signal Processing	Xinyu Sun; Zhikun Zhao; Lili Wei; Congyan Lang; Mingxuan Cai; Longfei Han; Juan Wang; Bing Li; Yuxuan Guo
Multi-Modal Prompting for Open-Vocabulary Video Visual Relationship	Shuo Yang; Yongqi Wang; Xiaofeng Ji; Xinxiao Wu



Progressive Painterly Image Harmonization from Low-Level Styles to High-Level Styles	Li Niu; Yan Hong; Junyan Cao; Liqing Zhang
Spatial-Contextual Discrepancy Information Compensation for GAN Inversion	Ziqiang Zhang; Yan Yan; Jing-Hao Xue; Hanzi Wang
Weakly-Supervised Temporal Action Localization by Inferring Salient Snippet-Feature	Wulian Yun; Mengshi Qi; Chuanming Wang; Huadong Ma
PMAC: Personalized Multi-Agent Communication	Xiangrui Meng; Ying Tan
Scalable Enumeration of Trap Spaces in Boolean Networks via Answer Set Programming	Giang Trinh; Belaid Benhamou; Samuel Pastva; Sylvain Soliman
DiG-In-GNN: Discriminative Feature Guided GNN-based Fraud Detector against Inconsistencies in Multi-Relation Fraud Graph	Jinghui Zhang; Zhengjia Xu; Dingyang Lv; Zhan Shi; Dian Shen; Jiahui Jin; Fang Dong
Live and Learn: Continual Action Clustering with Incremental Views	Xiaoqiang Yan; Yingtao Gan; Yiqiao Mao; Yangdong Ye; Hui Yu
SimCalib: Graph Neural Network Calibration Based on Similarity between Nodes	Boshi Tang; Zhiyong Wu; Xixin Wu; Qiaochu Huang; Jun Chen; Shun Lei; Helen Meng
SA <sup>2</sup> VP: Spatially Aligned-and-Adapted Visual Prompt	Wenjie Pei; Tongqi Xia; Fanglin Chen; Jinsong Li; Jiandong Tian; Guangming
Hypergraph-Guided Disentangled Spectrum Transformer Networks for Near-Infrared Facial Expression Recognition	Bingjun Luo; Haowen Wang; Jinpeng Wang; Junjie Zhu; Xibin Zhao; Yue Gao
Exploring Sparse Visual Prompt for Domain Adaptive Dense Prediction	Senqiao Yang; Jiarui Wu; Jiaming Liu; Xiaoqi Li; Qizhe Zhang; Mingjie Pan; Yulu Gan; Zehui Chen; Shanghang Zhang
Scalable Geometric Fracture Assembly via Co-creation Space among Assemblers	Ruiyuan Zhang; Jiaxiang Liu; Zexi Li; Hao Dong; Jie Fu; Chao Wu
Spot the Error: Non-autoregressive Graphic Layout Generation with Wireframe Locator	Jieru Lin; Danqing Huang; Tiejun Zhao; Dechen Zhan; Chin-Yew Lin
High-Fidelity Diffusion-Based Image Editing	Chen Hou; Guoqiang Wei; Zhibo Chen
SIG: Speaker Identification in Literature via Prompt-Based Generation	Zhenlin Su; Liyan Xu; Jin Xu; Jiangnan Li; Mingdu Huangfu
RLfOLD: Reinforcement Learning from Online Demonstrations in Urban Autonomous Driving	Daniel Coelho; Miguel Oliveira; Vitor Santos
Reward Penalties on Augmented States for Solving Richly Constrained RL Effectively	Hao Jiang; Tien Mai; Pradeep Varakantham; Huy Hoang
Parsing All Adverse Scenes: Severity-Aware Semantic Segmentation with Mask-Enhanced Cross-Domain Consistency	Fuhao Li; Ziyang Gong; Yupeng Deng; Xianzheng Ma; Renrui Zhang; Zhenming Ji; Xiangwei Zhu; Hong Zhang
Context-Aware Iteration Policy Network for Efficient Optical Flow Estimation	Ri Cheng; Ruian He; Xuhao Jiang; Shili Zhou; Weimin Tan; Bo Yan
Social-Aware Group Display Configuration in VR Conference	Bay-Yuan Hsu; Chih-Ya Shen; Hao Shan Yuan; Wang-Chien Lee; De-Nian Yang
Regulating AI: Applying Insights from Behavioural Economics and Psychology to the Application of Article 5 of the EU AI Act	Huixin Zhong; Eamonn O'Neill; Janina Hoffmann
Online Boosting Adaptive Learning under Concept Drift for Multistream Classification	En Yu; Jie Lu; Bin Zhang; Guangquan Zhang
Region-Aware Exposure Consistency Network for Mixed Exposure Correction	Jin Liu; Huiyuan Fu; Chuanming Wang; Huadong Ma
Arbitrary-Scale Video Super-resolution Guided by Dynamic Context	Cong Huang; Jiahao Li; Lei Chu; Dong Liu; Yan Lu
Unified Framework for Diffusion Generative Models in SO(3): Applications in Computer Vision and Astrophysics	Yesukhei Jagvaral; Francois Lanasus; Rachel Mandelbaum



Selective and Orthogonal Feature Activation for Pedestrian Attribute Recognition	Junyi Wu; Yan Huang; Min Gao; Yuzhen Niu; Mingjing Yang; Zhipeng Gao; Jianqiang Zhao
Hardness of Random Reordered Encodings of Parity for Resolution and CDCL	Leroy Chew; Alexis de Colnet; Friedrich Slivovsky; Stefan Szeider
Enhancing Evolving Domain Generalization through Dynamic Latent Representations	Binghui Xie; Yongqiang Chen; Jiaqi Wang; Kaiwen Zhou; Bo Han; Wei Meng; James Cheng
Diverse and Aligned Audio-to-Video Generation via Text-to-Video Model Adaptation	Guy Yariv; Itai Gat; Sagie Benaim; Lior Wolf; Idan Schwartz; Yossi Adi
A Convolutional Neural Network Interpretable Framework for Human Ventral Visual Pathway Representation	Mufan Xue; Xinyu Wu; Jinlong Li; Xuesong Li; Guoyuan Yang
Runtime Analysis of the $(\mu + 1)$ GA: Provable Speed-Ups from Strong Drift towards Diverse Populations	Benjamin Doerr; Aymen Echarchaoui; Mohammed Jamal; Martin S. Krejca
BAT: Behavior-Aware Human-Like Trajectory Prediction for Autonomous Driving	Haicheng Liao; Zhenning Li; Huanming Shen; Wenxuan Zeng; Dongping Liao; Guofa Li; Chengzhong Xu
Learning from Ambiguous Demonstrations with Self-Explanation Guided Reinforcement Learning	Yantian Zha; Lin Guan; Subbarao Kambhampati
CUDC: A Curiosity-Driven Unsupervised Data Collection Method with Adaptive Temporal Distances for Offline Reinforcement Learning	Chenyu Sun; Hangwei Qian; Chunyan Miao
MDFL: Multi-Domain Diffusion-Driven Feature Learning	Daixun Li; Weiying Xie; Jiaqing Zhang; Yunsong Li
Sample-Constrained Black Box Optimization for Audio Personalization	Rajalaxmi Rajagopalan; Yu-Lin Wei; Romit Roy Choudhury
Maxileximin Envy Allocations and Connected Goods	Gianluigi Greco; Francesco Scarcello
Shuffled Deep Regression	Masahiro Kohjima
Data-Augmented Curriculum Graph Neural Architecture Search under Distribution Shifts	Yang Yao; Xin Wang; Yijian Qin; Ziwei Zhang; Wenwu Zhu; Hong Mei
Improving the Adversarial Transferability of Vision Transformers with Virtual Dense Connection	Jianping Zhang; Yizhan Huang; Zhuoer Xu; Weilbin Wu; Michael Lyu
DTF-AT: Decoupled Time-Frequency Audio Transformer for Event Classification	Tony Alex; Sara Ahmed; Armin Mustafa; Muhammad Awais; Philip JB Jackson
Efficient Algorithms for Non-gaussian Single Index Models with Generative Priors	Junren CHEN; Zhaoqiang Liu
Embracing Language Inclusivity and Diversity in CLIP through Continual Language Learning	Bang Yang; Yong Dai; Xuxin Cheng; Yaowei Li; Asif Raza; Yuexian Zou
Omnidirectional Image Super-resolution via Bi-projection Fusion	Jiangang Wang; Yuning Cui; Yawen Li; Wenqi Ren; Xiaochun Cao
Video Event Extraction with Multi-View Interaction Knowledge Distillation	Kaiwen Wei; Du Runyan; Li Jin; Jian Liu; Jianhua Yin; Linhao Zhang; Jintao Liu; Nayu Liu; Jingyuan Zhang; Zhi Guo
History Matters: Temporal Knowledge Editing in Large Language Model	Xunjian Yin; Jin Jiang; Liming Yang; Xiaojun Wan
Fluctuation-Based Adaptive Structured Pruning for Large Language Models	Yongqi An; Xu Zhao; Tao Yu; Ming Tang; Jinqiao Wang
Hybrid-Supervised Dual-Search: Leveraging Automatic Learning for Loss-Free Multi-Exposure Image Fusion	Guanyao Wu; Hongming Fu; Jinyuan Liu; Long Ma; Xin Fan; Risheng Liu
Multi-Prototype Space Learning for Commonsense-Based Scene Graph Generation	Lianggangxu Chen; Youqi Song; Yiqing Cai; Jiale Lu; Yang Li; Yuan Xie; Changbo Wang; Gaoqi He



FPRF: Feed-Forward Photorealistic Style Transfer of Large-Scale 3D Neural Radiance Fields	GeonU Kim; Kim Youwang; Tae-Hyun Oh
Divide and Conquer: Hybrid Pre-training for Person Search	Yanling Tian; Di Chen; Yunan Liu; Jian Yang; Shanshan Zhang
Optimize & Reduce: A Top-Down Approach for Image Vectorization	Or Hirschorn; Amir Jevnisek; Shai Avidan
Proportional Aggregation of Preferences for Sequential Decision Making	Nikhil Chandak; Shashwat Goel; Dominik Peters
Tail-STEAK: Improve Friend Recommendation for Tail Users via Self-Training Enhanced Knowledge Distillation	Yijun Ma; Chaozhuo Li; Xiao Zhou
MA-Net: Rethinking Neural Unit in the Light of Astrocytes	Mengqiao Han; Liyuan Pan; Xiabi Liu
Limited Query Graph Connectivity Test	Mingyu Guo; Jialiang Li; Aneta Neumann; Frank Neumann; Hung Nguyen
Gated Attention Coding for Training High-Performance and Efficient Spiking Neural Networks	Xuerui Qiu; Rui-Jie Zhu; Yuhong Chou; Zhaorui Wang; Liang-Jian Deng; Guoqi Li
FedASMU: Efficient Asynchronous Federated Learning with Dynamic Staleness-Aware Model Update	Ji Liu; Juncheng Jia; Tianshi Che; Chao Huo; Jiaxiang Ren; Yang Zhou; Huaiyu Dai; Dejing Dou
Dynamic Budget Throttling in Repeated Second-Price Auctions	Zhaohua Chen; Chang Wang; Qian Wang; Yuqi Pan; Zhuming Shi; Zheng Cai; Yukun Ren; Zhihua Zhu; Xiaotie Deng
Towards a Theoretical Understanding of Why Local Search Works for Clustering with Fair-Center Representation	Zhen Zhang; Junfeng Yang; Limei Liu; Xuesong Xu; Guozhen Rong; Qilong Feng
ProCC: Progressive Cross-Primitive Compatibility for Open-World Compositional Zero-Shot Learning	Fushuo Huo; Wenchao Xu; Song Guo; Jingcai Guo; Haozhao Wang; Ziming Liu; Xiaocheng Lu
A Local-Ascending-Global Learning Strategy for Brain-Computer Interface	Dongrui Gao; Haokai Zhang; Pengrui Li; Tian Tang; Shihong Liu; Zhihong Zhou; Shaofei Ying; Ye Zhu; Yongqing Zhang
Data Roaming and Quality Assessment for Composed Image Retrieval	Matan Levy; Rami Ben-Ari; Nir Darshan; Dani Lischinski
Frequency Spectrum Is More Effective for Multimodal Representation and Fusion: A Multimodal Spectrum Rumor Detector	An Lao; Qi Zhang; Chongyang Shi; Longbing Cao; Kun Yi; Liang Hu; Duoqian Miao
Task-Driven Causal Feature Distillation: Towards Trustworthy Risk Prediction	Zhixuan Chu; Mengxuan Hu; Qing Cui; Longfei Li; Sheng Li
ConsistentEE: A Consistent and Hardness-Guided Early Exiting Method for Accelerating Language Models Inference	Ziqian Zeng; Yihuai Hong; Hongliang Dai; Huiping Zhuang; Cen Chen
CONSIDER: Commonalities and Specialties Driven Multilingual Code Retrieval Framework	Rui Li; LiYang He; Qi Liu; Yuze Zhao; Zheng Zhang; Zhenya Huang; Yu Su; Shijin WANG
Turning Dust into Gold: Distilling Complex Reasoning Capabilities from LLMs by Leveraging Negative Data	Yiwei Li; Peiwen Yuan; Shaoxiong Feng; Boyuan Pan; Bin Sun; Xinglin Wang; Heda Wang; Kan Li
Beyond Entities: A Large-Scale Multi-Modal Knowledge Graph with Triplet Fact Grounding	Jingping Liu; Mingchuan Zhang; Weichen Li; Chao Wang; Shuang Li; Haiyun Jiang; Sihang Jiang; Yanghua Xiao; Yunwen Chen
Cycle-Consistency Learning for Captioning and Grounding	Ning Wang; Jiajun Deng; Mingbo Jia
Analytically Tractable Models for Decision Making under Present Bias	Yasunori Akagi; Naoki Marumo; Takeshi Kurashima
Mitigating Large Language Model Hallucinations via Autonomous Knowledge Graph-Based Retrofitting	Xinyan Guan; Yanjiang Liu; Hongyu Lin; Yaojie Lu; Ben He; Xianpei Han; Le Sun
Underwater Organism Color Fine-Tuning via Decomposition and Guidance	Xiaofeng Cong; Jie Gui; Junming Hou
Neural Network Approximators for Marginal MAP in Probabilistic Circuits	Shivvrat Arya; Tahrira Rahman; Vibhav Gogate
Wikiformer: Pre-training with Structured Information of Wikipedia for Ad-Hoc Retrieval	Weihang Su; Qingyao Ai; Xiangsheng Li; Jia Chen; Yiqun LIU; Xiaolong Wu; Shengluan Hou



Rethinking Dimensional Rationale in Graph Contrastive Learning from Causal Perspective	Qirui Ji; Jiangmeng Li; Jie Hu; Rui Wang; Changwen Zheng; Fanjiang Xu
Text-to-Image Generation for Abstract Concepts	Jiayi Liao; Xu Chen; Qiang Fu; Lun Du; Xiangnan He; Xiang Wang; Shi Han; Dongmei Zhang
Diagnosing and Rectifying Fake OOD Invariance: A Restructured Causal Approach	Ziliang Chen; Yongsun Zheng; Zhao-Rong Lai; Quanlong Guan; Liang Lin
A Joint Framework with Heterogeneous-Relation-Aware Graph and Multi-Channel Label Enhancing Strategy for Event Causality Extraction	Ruili Pu; Yang Li; Jun Zhao; Suge Wang; Deyu Li; Jian Liao; Jianxing Zheng
Arbitrary-Scale Point Cloud Upsampling by Voxel-Based Network with Latent Geometric-Consistent Learning	Hang Du; Xuejun Yan; Jingjing Wang; Di Xie; Shiliang Pu
Pay to (Not) Play: Monetizing Impatience in Mobile Games	Taylor Lundy; Narun Raman; Hu Fu; Kevin Leyton-Brown
A Diffusion-Based Pre-training Framework for Crystal Property Prediction	Zixing Song; Ziqiao Meng; Irwin King
Exact ASP Counting with Compact Encodings	Mohimenul Kabir; Supratik Chakraborty; Kuldeep S Meel
Large Occluded Human Image Completion via Image-Prior Cooperating	Hengrun Zhao; Yu Zeng; Huchuan Lu; Lijun Wang
GMP-AR: Granularity Message Passing and Adaptive Reconciliation for Temporal Hierarchy Forecasting	Fan Zhou; Chen Pan; Lintao Ma; Yu Liu; Siqiao Xue; James Zhang; Jun Zhou; Hongyuan Mei; Weitao Lin; Zi Zhuang; Wenxin Ning; Yunhua Hu
Expressive Forecasting of 3D Whole-Body Human Motions	Pengxiang Ding; Qiongjie Cui; Haofan Wang; Min Zhang; Mengyuan Liu; Donglin Wang
SeqGPT: An Out-of-the-Box Large Language Model for Open Domain Sequence Understanding	Tianyu Yu; Chengyue Jiang; Chao Lou; Shen Huang; Xiaobin Wang; Wei Liu; Jiong Cai; Li Yangning; Yinghui Li; Kewei Tu; Hai-Tao Zheng; Ningyu Zhang; Pengjun Xie; Fei Huang; Yong Jiang
Non-excludable Bilateral Trade between Groups	Yixuan Even Xu; Hanrui Zhang; Vincent Conitzer
Learning to Learn in Interactive Constraint Acquisition	Dimosthenis Tsouros; Senne Berden; Tias Guns
CamoDiffusion: Camouflaged Object Detection via Conditional Diffusion	Zhongxi Chen; Ke Sun; Xianming Lin
Density Matters: Improved Core-Set for Active Domain Adaptive Segmentation	Shizhan Liu; Zhengkai Jiang; Yuxi Li; Jinlong Peng; Yabiao Wang; Weiyao Lin
Mastering Context-to-Label Representation Transformation for Event Causality Identification with Diffusion Models	Hieu Man; Franck Deroncourt; Thien Huu Nguyen
QDETRv: Query-Guided DETR for One-Shot Object Localization in Videos	Yogesh Kumar; Saswat Mallick; Anand Mishra; Sowmya Rasipuram; Anutosh Maitra; Roshni Ramnani
Imitate the Good and Avoid the Bad: An Incremental Approach to Safe Reinforcement Learning	Huy Hoang; Tien Mai; Pradeep Varakantham
Multiple Hypothesis Dropout: Estimating the Parameters of Multi-Modal Output Distributions	David Nguyen; David Liebowitz; Salil Kanhere; Surya Nepal
Batch Normalization Is Blind to the First and Second Derivatives of the Loss	Zhanpeng Zhou; Wen Shen; Huixin Chen; Ling Tang; Yuefeng Chen; Quanshi Zhang
Solving Satisfiability Modulo Counting for Symbolic and Statistical AI Integration with Provable Guarantees	Jinzhao Li; Nan Jiang; Yexiang Xue
Towards Equipping Transformer with the Ability of Systematic Compositionality	Chen Huang; Peixin Qin; Wenqiang Lei; Jiancheng Lv
Peer Learning: Learning Complex Policies in Groups from Scratch via Action Recommendations	Cedric Derstroff; Jannis Brugger; Mattia Cerrato; Jan Peters; Stefan Kramer
DeepPointMap: Advancing LiDAR SLAM with Unified Neural Descriptors	Xiaze Zhang; Ziheng Ding; Qi Jing; Yuejie Zhang; Wenchao Ding; Rui Feng





Rethinking Two-Stage Referring Expression Comprehension: A Novel Grounding and Segmentation Method Modulated by Point	Peizhi Zhao; Shiyi Zheng; Wenye Zhao; Dongsheng Xu; Pijian Li; Yi Cai; Qingbao Huang
Noisy Correspondence Learning with Self-Reinforcing Errors Mitigation	Zhuohang Dang; Minnan Luo; Chengyou Jia; Guang Dai; Xiaojun Chang; Jingdong Wang
Cost Minimization for Equilibrium Transition	Haoqiang Huang; Zihe Wang; Zhide Wei; Jie Zhang
Lost Domain Generalization Is a Natural Consequence of Lack of Training Domains	Yimu Wang; Yihan Wu; Hongyang Zhang
Neural Amortized Inference for Nested Multi-Agent Reasoning	Kunal Jha; Tuan Anh Le; Chuanyang Jin; Yen-Ling Kuo; Joshua Tenenbaum; Tianmin Shu
MotionMix: Weakly-Supervised Diffusion for Controllable Motion Generation	Nhat Hoang; Kehong Gong; Chuan Guo; Michael Bi Mi
Relational Programming with Foundational Models	Ziyang Li; Jiani Huang; Jason Liu; Felix Zhu; Eric Zhao; William Dodds; Neelay Velingker; Rajeev Alur; Mayur Naik
SEER: Backdoor Detection for Vision-Language Models through Searching Target Text and Image Trigger Jointly	Liuwan Zhu; Rui Ning; Jiang Li; Chunsheng Xin; Hongyi Wu
PVALane: Prior-Guided 3D Lane Detection with View-Agnostic Feature Alignment	Zewen Zheng; Xuemin Zhang; Yongqiang Mou; Xiang Gao; Chengxin Li; Guoheng Huang; Chi-Man Pun; Xiaochen Yuan
Simplifying Complex Observation Models in Continuous POMDP Planning with Probabilistic Guarantees and Practice	Idan Lev-Yehudi; Moran Barenboim; Vadim Indelman
Responsibility in Extensive Form Games	Qi SHI
Personalized Reinforcement Learning with a Budget of Policies	Dmitry Ivanov; Omer Ben-Porat
Repurposing Ensemble of Black-Box Models to New Task Domains	Minh Hoang; Nghia Hoang
Abstraction of Situation Calculus Concurrent Game Structures	Yves Lesperance; Giuseppe De Giacomo; Maryam Rostamigiv; Shakil Khan
Bridging the Semantic Latent Space between Brain and Machine: Similarity Is All You Need	Jiaxuan Chen; Yu Qi; Yueming Wang; Gang Pan
ViLT-CLIP: Video and Language Tuning CLIP with Multimodal Prompt Learning and Scenario-guided Optimization	Hao Wang; Fang Liu; Licheng Jiao; Jiahao Wang; Zehua Hao; Shuo Li; Lingling Li; Puhua Chen; Xu Liu
An Autoregressive Text-to-Graph Framework for Joint Entity and Relation Extraction	Urchade Zaratiana; Nadi Tomeh; Pierre Holat; Thierry Charnois
Reachability of Fair Allocations via Sequential Exchanges	Ayumi Igarashi; Naoyuki Kamiyama; Warut Suksompong; Sheung Man Yuen
HGE: Embedding Temporal Knowledge Graphs in a Product Space of Heterogeneous Geometric Subspaces	Jiaxin Pan; Mojtaba Nayyeri; Yinan Li; Steffen Staab
Causal Representation Learning via Counterfactual Intervention	Xiutian Li; Siqi Sun; Rui Feng
Improving Factual Error Correction by Learning to Inject Factual Errors	Xingwei He; Qianru Zhang; A-Long Jin; Jun Ma; Yuan Yuan; S.M. Yiu
Multi-Objective Bayesian Optimization with Active Preference Learning	Ryota Ozaki; Kazuki Ishikawa; Youhei Kanzaki; Shion Takeno; Ichiro Takeuchi; Masayuki Karasuyama
A Class of Topological Pseudodistances for Fast Comparison of Persistence Diagrams	Rolando Kindelan Nuñez; Mircea Petrache; Mauricio Cerda; Nancy Hitschfeld
CAVEN: An Embodied Conversational Agent for Efficient Audio-Visual Navigation in Noisy Environments	Xiulong Liu; Sudipta Paul; Moitreya Chatterjee; Anoop Cherian
Lyapunov-Stable Deep Equilibrium Models	Haoyu Chu; Shikui Wei; Ting Liu; Yao Zhao; Yuto Miyatake
TD <sup>2</sup> -Net: Toward Denoising and Debiasing for Video Scene Graph Generation	Xin Lin; Chong Shi; Yibing Zhan; Zuopeng Yang; Yaqi Wu; Dacheng Tao





Improving Automatic VQA Evaluation Using Large Language Models	Oscar Mañas; Benno Krojer; Aishwarya Agrawal
Layer-Wise Representation Fusion for Compositional Generalization	Yafang Zheng; Lei Lin; Shuangtao Li; Yuxuan Yuan; Zhaohong Lai; Shan Liu; Biao Fu; Yidong Chen; Xiaodong Shi
FedA3I: Annotation Quality-Aware Aggregation for Federated Medical Image Segmentation against Heterogeneous Annotation Noise	Nannan Wu; Zhaobin Sun; Zengqiang Yan; Li Yu
Procedural Level Generation with Diffusion Models from a Single Example	Shiqi Dai; Xuanyu Zhu; Naiqi Li; Tao Dai; Zhi Wang
MM-TTS: Multi-Modal Prompt Based Style Transfer for Expressive Text-to-Speech Synthesis	Wenhao Guan; Yishuang Li; Tao Li; Hukai Huang; Feng Wang; Jiayan Lin; Lingyan Huang; Lin Li; Qingyang Hong
Contrastive Tuning: A Little Help to Make Masked Autoencoders Forget	Johannes Lehner; Benedikt Alkin; Andreas Fürst; Elisabeth Rumetshofer; Lukas Miklautz; Sepp Hochreiter
Practical Privacy-Preserving MLaaS: When Compressive Sensing Meets Generative Networks	Jia Wang; Wuqiang Su; Zushu Huang; Jie Chen; Chengwen Luo; Jianqiang Li
Towards Fairness in Online Service with K Servers and Its Applications in Fair Food Delivery	Daman Deep Singh; Amit Kumar; Abhijnan Chakraborty
Generating Images of Rare Concepts Using Pre-trained Diffusion Models	Dvir Samuel; Rami Ben-Ari; Simon Raviv; Nir Darshan; Gal Chechik
CTO-SLAM: Contour Tracking for Object-Level Robust 4D SLAM	Xiaohan Li; Dong Liu; Jun Wu
RR-PU: A Synergistic Two-Stage Positive and Unlabeled Learning Framework for Robust Tax Evasion Detection	Shuzhi Cao; Jianfei Ruan; Bo Dong; Bin Shi; Qinghua Zheng
The Moderating Effect of Instant Runoff Voting	Kiran Tomlinson; Johan Ugander; Jon Kleinberg
How to Make Knockout Tournaments More Popular?	Juhi Chaudhary; Hendrik Molter; Meirav Zehavi
Complementary Knowledge Distillation for Robust and Privacy-Preserving Model Serving in Vertical Federated Learning	Dashan Gao; Sheng Wan; Lixin Fan; Xin Yao; Qiang Yang
Dealing with Numeric and Metric Time Constraints in PDDL3 via Compilation to Numeric Planning	Luigi Bonassi; Alfonso Emilio Gerevini; Enrico Scala
On Inference Stability for Diffusion Models	Viet Nguyen; Giang Vu; Tung Nguyen Thanh; Khoat Than; Toan Tran
QLABGrad: A Hyperparameter-Free and Convergence-Guaranteed Scheme for Deep Learning	Fang-Xiang Wu; Minghan Fu
Active Learning Guided by Efficient Surrogate Learners	Yunpyo An; Suyeong Park; Kwang In Kim
Manipulation-Robust Selection of Citizens' Assemblies	Bailey Flanigan; Jennifer Liang; Ariel Procaccia; Sven Wang
Implications of Distance over Redistricting Maps: Central and Outlier Maps	Seyed Esmaeili; Darshan Chakrabarti; Hayley Grape; Brian Brubach
Leveraging Normalization Layer in Adapters with Progressive Learning and Adaptive Distillation for Cross-Domain Few-Shot Learning	YongJin Yang; Taehyeon Kim; Se-Young Yun
Structure-Aware Multimodal Sequential Learning for Visual Dialog	Youngjin Kim; Min-Jun Kim; Kyunghwan An; Jinwoo Ahn; Jaeseok Kim; Yu-Jung Heo; Du-Seong Chang; Eun-Sol Kim
Domain-Aware Fine-Tuning: Enhancing Neural Network Adaptability	Seokhyeon Ha; Sunbeom Jeong; Jungwoo Lee
Content Filtering with Inattentive Information Consumers	Justin Grana; Alex Slivkins; Brendan Lucier; Nicole Immorlica; Ian Ball; James Bono
VLM2Scene: Self-Supervised Image-Text-LiDAR Learning with Foundation Models for Autonomous Driving Scene Understanding	Guibiao Liao; Jiankun Li; Xiaoqing Ye
Generalized Bradley-Terry Models for Score Estimation from Paired Comparisons	Julien Fageot; Lê-Nguyên Hoang; Oscar Villemaud; Sadegh Farhadkhani



Enhancing the Efficiency of Altruism and Taxes in Affine Congestion Games through Signalling	Vittorio Bilò; Cosimo Vinci
Optimal Quasi-clique: Hardness, Equivalence with Densest- $k$ -Subgraph, and Quasi-partitioned Community Mining	Aritra Konar; Nicholas Sidiropoulos
Learning Real-World Image De-weathering with Imperfect Supervision	Xiaohui Liu; Zhilu Zhang; Xiaohe Wu; Chaoyu Feng; Wang Xiaotao; Lei Lei; Wangmeng Zuo
RetouchFormer: Semi-supervised High-Quality Face Retouching Transformer with Prior-Based Selective Self-Attention	Xue Wen; Lianxin Xie; Le Jiang; Tianyi Chen; Si Wu; Cheng Liu; Hau San Wong
Approximate Integer Solution Counts over Linear Arithmetic Constraints	Cunjing Ge
Faster Stochastic Variance Reduction Methods for Compositional MiniMax Optimization	Jin Liu; Xiaokang Pan; Junwen Duan; Hong-Dong Li; Youqi Li; Zhe Qu
Multi-Class Support Vector Machine with Maximizing Minimum Margin	Feiping Nie; Zhezheng Hao; Rong Wang
A Two-Stage Information Extraction Network for Incomplete Multi-View Multi-Label Classification	Xin Tan; Ce Zhao; Chengliang Liu; Jie Wen; Zhanyan Tang
DanceMVP: Self-Supervised Learning for Multi-Task Primitive-Based Dance Performance Assessment via Transformer Text Prompting	Yun Zhong; Yiannis Demiris
Multi-Modal Disordered Representation Learning Network for Description-Based Person Search	Fan Yang; Wei Li; Menglong Yang; Binbin Liang; Jianwei Zhang
CK12: A Rounded K12 Knowledge Graph Based Benchmark for Chinese Holistic Cognition Evaluation	Weihaio You; Pengcheng Wang; Changlong Li; Zhilong Ji; Jinfeng Bai
Chronic Poisoning: Backdoor Attack against Split Learning	Fangchao Yu; Bo Zeng; Kai Zhao; Zhi Pang; Lina Wang
GO-DICE: Goal-Conditioned Option-Aware Offline Imitation Learning via Stationary Distribution Correction Estimation	Abhinav Jain; Vaibhav Unhelkar
Deep Variational Incomplete Multi-View Clustering: Exploring Shared Clustering Structures	Gehui Xu; Jie Wen; Chengliang Liu; Hu Bing; Yicheng Liu; Lunke Fei; Wei Wang
Neural Embeddings for kNN Search in Biological Sequence	Zhihao Chang; Linzhu Yu; Yanchao Xu; Wentao Hu
Online Conversion Rate Prediction via Multi-Interval Screening and Synthesizing under Delayed Feedback	Qiming Liu; Xiang Ao; Yuyao Guo; Qing He
RRL: Recommendation Reverse Learning	Xiaoyu You; Jianwei Xu; Mi Zhang; Zechen Gao; Min Yang
Motion Deblurring via Spatial-Temporal Collaboration of Frames and Events	Wen Yang; Jinjian Wu; Jupao Ma; Leida Li; Guangming Shi
Linear-Time Verification of Data-Aware Processes Modulo Theories via Covers and Automata	Alessandro Gianola; Marco Montali; Sarah Winkler
Tree-of-Reasoning Question Decomposition for Complex Question Answering with Large Language Models	Kun Zhang; Jiali Zeng; Fandong Meng; Yuanzhuo Wang; Shiqi Sun; Long Bai; Huawei Shen; Jie Zhou
Motif-Aware Riemannian Graph Neural Network with Generative-Contrastive Learning	Li Sun; Zhenhao Huang; Zixi Wang; Feiyang Wang; Hao Peng; Philip Yu
Non-monotone Sequential Submodular Maximization	Shaojie Tang; Jing Yuan
Automated Defect Report Generation for Enhanced Industrial Quality Control	Jiayuan Xie; Zhiping Zhou; Zihan Wu; Xinting Zhang; Jiexin Wang; Yi Cai; Qing Li
A Computation-Aware Shape Loss Function for Point Cloud Completion	Shunran Zhang; Xiubo Zhang; Tsz Nam (Edison) Chan; Shenghui Zhang; Leong Hou U



Towards Molecular Structure Discovery from Cryo-ET Density Volumes via Modelling Auxiliary Semantic Prototypes	Ashwin Nair; Xingjian Li; Mostofa Rafid Uddin; Min Xu
Emotion Rendering for Conversational Speech Synthesis with Heterogeneous Graph-Based Context Modeling	Rui Liu; Yifan Hu; Yi Ren; Xiang Yin; Haizhou Li
Finding Interpretable Class-Specific Patterns through Efficient Neural Search	Nils Philipp Walter; Jonas Fischer; Jilles Vreeken
Domain Generalization with Vital Phase Augmentation	Ingyun Lee; WooJu Lee; Hyun Myung
Adversarial Purification with the Manifold Hypothesis	Zhaoyuan Yang; Zhiwei Xu; Jing Zhang; Richard Hartley; Peter Tu
Cross-Gate MLP with Protein Complex Invariant Embedding Is a One-Shot Antibody Designer	Cheng Tan; Zhangyang Gao; Lirong Wu; Jun Xia; Jiangbin Zheng; Xihong Yang; Yue Liu; Bozhen Hu; Stan Z. Li
DiffAIL: Diffusion Adversarial Imitation Learning	Bingzheng Wang; Guoqiang Wu; Teng Pang; Yan Zhang; Yilong Yin
Learning to Unlearn: Instance-Wise Unlearning for Pre-trained Classifiers	Sungmin Cha; Sungjun Cho; Dasol Hwang; Honglak Lee; Taesup Moon; Moontae Lee
Learning-Augmented Online Algorithm for Two-Level Ski-Rental Problem	Keyuan Zhang; Zhongdong Liu; Nakjung Choi; Bo Ji
Visual Instruction Tuning with Polite Flamingo	Delong Chen; Jianfeng Liu; Wenliang Dai; Baoyuan Wang
Rethinking Propagation for Unsupervised Graph Domain Adaptation	Meihan Liu; Zeyu Fang; Zhen Zhang; Ming Gu; Sheng Zhou; Xin Wang; Jiajun
A General Theoretical Framework for Learning Smallest Interpretable Models	Sebastian Ordyniak; Giacomo Paesani; Mateusz banany; Stefan Szeider
Arithmetic Feature Interaction Is Necessary for Deep Tabular Learning	Yi Cheng; Renjun Hu; Haochao Ying; Xing Shi; Jian Wu; Wei Lin
Structural Information Guided Multimodal Pre-training for Vehicle-Centric Perception	Xiao Wang; Wentao Wu; Chenglong Li; Zhicheng Zhao; Zhe Chen; Yukai Shi; Jin Tang
DiffRAW: Leveraging Diffusion Model to Generate DSLR-Comparable Perceptual Quality sRGB from Smartphone RAW Images	Mingxin Yi; Kai Zhang; Pei Liu; Tanli Zuo; Jingduo Tian
Optimistic Model Rollouts for Pessimistic Offline Policy Optimization	Yuanzhao Zhai; Yiyang Li; Zijian Gao; Xudong Gong; Kele Xu; Dawei Feng; Ding Bo; Huaimin Wang
CIDR: A Cooperative Integrated Dynamic Refining Method for Minimal Feature Removal Problem	Qian Chen; Taolin Zhang; Dongyang Li; Xiaofeng He
Learning Multimodal Volumetric Features for Large-Scale Neuron Tracing	Qihua Chen; Xuejin Chen; Chenxuan Wang; Yixiong Liu; Zhiwei Xiong; Feng
DART: Dual-Modal Adaptive Online Prompting and Knowledge Retention for Test-Time Adaptation	Zichen Liu; Hongbo Sun; Yuxin Peng; Jiahuan Zhou
One-Step Forward and Backtrack: Overcoming Zig-Zagging in Loss-Aware Quantization Training	Lianbo Ma; Yuee Zhou; Jianlun Ma; Guo Yu; Qing Li
Not All Tasks Are Equally Difficult: Multi-Task Deep Reinforcement Learning with Dynamic Depth Routing	Jinmin He; Kai Li; Yifan Zang; Haobo Fu; Qiang Fu; Junliang Xing; Jian Cheng
VQAttack: Transferable Adversarial Attacks on Visual Question Answering via Pre-trained Models	Ziyi Yin; Muchao Ye; Tianrong Zhang; Jiaqi Wang; Han Liu; Jinghui Chen; Ting Wang; Fenglong Ma
Behavioral Recognition of Skeletal Data Based on Targeted Dual Fusion Strategy	Xiao Yun; Chenglong Xu; Kevin Riou; Kaiwen Dong; Yanjing Sun; Song Li; Kevin Subrin; Patrick Le Callet
Contextual Pre-planning on Reward Machine Abstractions for Enhanced Transfer in Deep Reinforcement Learning	Guy Azran; Mohamad H Danesh; Stefano Albrecht; Sarah Keren
EnMatch: Matchmaking for Better Player Engagement via Neural Combinatorial Optimization	Kai Wang; Haoyu Liu; Zhipeng Hu; Xiaochuan Feng; Minghao Zhao; Shiwei Zhao; Runze Wu; Xudong Shen; Tangjie Lv; Changjie Fan



Graph of Thoughts: Solving Elaborate Problems with Large Language Models	Maciej Besta; Nils Blach; Ales Kubicek; Robert Gerstenberger; Michal Podstawski; Lukas Gianinazzi; Joanna Gajda; Tomasz Lehmann; Hubert Niewiadomski; Piotr Nyczyk; Torsten Hoefler
Incomplete Contrastive Multi-View Clustering with High-Confidence Guiding	Guoqing Chao; Yi Jiang; Dianhui Chu
Symbolic Numeric Planning with Patterns	Matteo Cardellini; Enrico Giunchiglia; Marco Maratea
PHFormer: Multi-Fragment Assembly Using Proxy-Level Hybrid Transformer	Wenting Cui; Runzhao Yao; Shaoyi Du
Uncertainty Quantification for Data-Driven Change-Point Learning via Cross-Validation	Hui Chen; Yinxu Jia; Guanghui Wang; Changliang Zou
End-to-End Verification for Subgraph Solving	Stephan Gocht; Ciaran McCreesh; Magnus O. Myreen; Jakob Nordström; Andy Oertel; Yong Kiam Tan
FoX: Formation-Aware Exploration in Multi-Agent Reinforcement Learning	Yonghyeon Jo; Sunwoo Lee; Junghyuk Yeom; Seungyu Han
When Model Meets New Normals: Test-Time Adaptation for Unsupervised Time-Series Anomaly Detection	Dongmin Kim; Sunghyun Park; Jaegul Choo
A Unified Environmental Network for Pedestrian Trajectory Prediction	Yuchao Su; Yuanman Li; Wei Wang; Jiantao Zhou; Xia Li
Delivering Inflated Explanations	Yacine Izza; Alexey Ignatiev; Peter Stuckey; Joao Marques-Silva
Empowering CAM-Based Methods with Capability to Generate Fine-Grained and High-Faithfulness Explanations	Changqing Qiu; Fusheng Jin; Yining Zhang
BLADE: Box-Level Supervised Amodal Segmentation through Directed Expansion	Zhaochen Liu; Zhixuan Li; Tingting Jiang
Unsupervised Cross-Domain Image Retrieval via Prototypical Optimal	Bin Li; Ye Shi; Qian Yu; Jingya Wang
Cooperative Knowledge Distillation: A Learner Agnostic Approach	Michael Livanos; Ian Davidson; Stephen Wong
Knowledge Enhanced Representation Learning for Drug Discovery	Thanh Lam Hoang; Marco Luca Sbodio; Marcos Martinez; Mykhaylo Zayats; Raul Fernandez-Diaz; Victor Valls; Gabriele Picco; Cesar Berrospi; Vanessa Lopez
RGMComm: Return Gap Minimization via Discrete Communications in Multi-Agent Reinforcement Learning	Jingdi Chen; Tian Lan; Carlee Joe-Wong
CREAD: A Classification-Restoration Framework with Error Adaptive Discretization for Watch Time Prediction in Video Recommender Systems	Jie Sun; Zhao Ying Ding; Xiaoshuang Chen; Qi Chen; Yincheng Wang; Kaiqiao Zhan; Ben Wang
Unveiling the Significance of Toddler-Inspired Reward Transition in Goal-Oriented Reinforcement Learning	Junseok Park; Yoonsung Kim; Hee bin Yoo; Min Whoo Lee; Kibeom Kim; Won-Seok Choi; Minsu Lee; Byoung-Tak Zhang
Layer Collaboration in the Forward-Forward Algorithm	Guy Lorberbom; Itai Gat; Yossi Adi; Alexander Schwing; Tamir Hazan
One Step Learning, One Step Review	Huang Xiaolong; Qiankun Li; Xueran Li; Xuesong Gao
Synergistic Anchored Contrastive Pre-training for Few-Shot Relation Extraction	Da Luo; Yanglei Gan; Rui Hou; Run Lin; Qiao Liu; Yuxiang Cai; Wannian Gao
HOP to the Next Tasks and Domains for Continual Learning in NLP	Umberto Michieli; Mete Ozay
A Fixed-Parameter Tractable Algorithm for Counting Markov Equivalence Classes with the Same Skeleton	Vidya Sagar Sharma
Token-Level Contrastive Learning with Modality-Aware Prompting for Multimodal Intent Recognition	Qianrui Zhou; Hua Xu; Hao Li; Hanlei Zhang; Xiaohan Zhang; Yifan Wang; Kai Gao
Situation-Dependent Causal Influence-Based Cooperative Multi-Agent Reinforcement Learning	Xiao Du; Yutong Ye; Pengyu Zhang; Yaning Yang; Mingsong Chen; Ting Wang
HyperFast: Instant Classification for Tabular Data	David Bonet; Daniel Mas Montserrat; Xavier Giró-i-Nieto; Alexander Ioannidis



Make Lossy Compression Meaningful for Low-Light Images	Shilv Cai; Liqun Chen; Sheng Zhong; Luxin Yan; Jiahuan Zhou; Xu Zou
Semi-supervised Active Learning for Video Action Detection	Ayush Singh; Aayush J Rana; Akash Kumar; Shruti Vyas; Yogesh Singh
DiffSED: Sound Event Detection with Denoising Diffusion	Swapnil Bhosale; Sauradip Nag; Diptesh Kanojia; Jiankang Deng; Xiatian Zhu
Deep Copula-Based Survival Analysis for Dependent Censoring with Identifiability Guarantees	Weijia Zhang; Chun Kai Ling; Xuanhui Zhang
Task Contamination: Language Models May Not Be Few-Shot Anymore	Changmao Li; Jeffrey Flanigan
United We Stand: Using Epoch-Wise Agreement of Ensembles to Combat Overfit	Uri Stern; Daniel Shwartz; Daphna Weinshall
ESRL: Efficient Sampling-Based Reinforcement Learning for Sequence Generation	Chenglong Wang; Hang Zhou; Yimin Hu; Yifu Huo; Bei Li; Tongran Liu; Tong Xiao; Jingbo Zhu
Sparse Bayesian Deep Learning for Cross Domain Medical Image Reconstruction	Jiaxin Huang; Qi Wu; Yazhou Ren; Fan Yang; Aodi Yang; Qianqian Yang; Xiaorong Pu
GraFITi: Graphs for Forecasting Irregularly Sampled Time Series	Vijaya Krishna Yalavarthi; Kiran Madhusudhanan; Randolph Scholz; Nourhan Ahmed; Johannes Burchert; Shayan Jawed; Stefan Born; Lars Schmidt-Tianyi Chu; Wei Xing; Jiafu Chen; Zhizhong Wang; Jiakai Sun; Lei Zhao; Haibo Chen; Huaizhong Lin
Attack Deterministic Conditional Image Generative Models for Diverse and Controllable Generation	Qiu Yu Duan; Zhongyun Hua; Qing Liao; Yushu Zhang; Leo Yu Zhang
Conditional Backdoor Attack via JPEG Compression	Jiaxuan Liang; Jun Wang; Guoxian Yu; Shuyin Xia; Guoyin Wang
Multi-Granularity Causal Structure Learning	Harrison Rosenberg; Shimaa Ahmed; Guruprasad Ramesh; Kassem Fawaz; Ramya Korlakai Vinayak
Limitations of Face Image Generation	Yunpeng Wang; Meng Pang; Shengbo Chen; Hong Rao
Consistency-GAN: Training GANs with Consistency Model	Linsheng Chen; Guangrun Wang; Liuchun Yuan; Keze Wang; Ken Deng; Philip Torr
NeRF-VPT: Learning Novel View Representations with Neural Radiance Fields via View Prompt Tuning	Junkai Xu; Liang Peng; Haoran Cheng; Linxuan Xia; Qi Zhou; Dan Deng; Wei Qian; Wenxiao Wang; Deng Cai
Regulating Intermediate 3D Features for Vision-Centric Autonomous Driving	Yanmei Zou; Hongshan Yu; Zhengeng Yang; Zechuan Li; Naveed Akhtar
Improved MLP Point Cloud Processing with High-Dimensional Positional Encoding	Junwei He; Qianqian Xu; Yangbangyan Jiang; Zitai Wang; Qingming Huang
ADA-GAD: Anomaly-Denoised Autoencoders for Graph Anomaly Detection	Shreyas Chaudhari; David Arbour; Georgios Theodorou; Nikos Vlassis
Distributional Off-Policy Evaluation for Slate Recommendations	Ehsan Amid; Frank Nielsen; Richard Nock; Manfred K. Warmuth
Optimal Transport with Tempered Exponential Measures	Taehoon Kim; JaeYoo Park; Bohyung Han
Cross-Class Feature Augmentation for Class Incremental Learning	Cornelius Brand; Robert Ganian; Subrahmanyam Kalyanasundaram; Fionn McInerney
The Complexity of Optimizing Atomic Congestion	Yufei Huang; Siyuan Li; Lirong Wu; Jin Su; Haitao Lin; Odin Zhang; Zihan Liu; Zhangyang Gao; Jiangbin Zheng; Stan Z. Li
Protein 3D Graph Structure Learning for Robust Structure-Based Protein Property Prediction	Jiaying Zhu; Dong Li; Xueyang Fu; Gang Yang; Jie Huang; Aiping Liu; Zheng-Jun Zha
Learning Discriminative Noise Guidance for Image Forgery Detection and Localization	Wenjie Li; Qifan Song; Jean Honorio; Guang Lin
Federated X-armed Bandit	Yufei Cai; Yuxiang Wei; Zhilong Ji; Jinfeng Bai; Hu Han; Wangmeng Zuo
Decoupled Textual Embeddings for Customized Image Generation	Jun Wang; Ying Cui; Dongyan Guo; Junxia Li; Qingshan Liu; Chunhua Shen
PointAttN: You Only Need Attention for Point Cloud Completion	





SSMG: Spatial-Semantic Map Guided Diffusion Model for Free-Form Layout-to-Image Generation	Chengyou Jia; Minnan Luo; Zhuohang Dang; Guang Dai; Xiaojun Chang; Mengmeng Wang; Jingdong Wang
Efficient Spiking Neural Networks with Sparse Selective Activation for Continual Learning	Jiangrong Shen; Wenyao Ni; Qi Xu; Huajin Tang
PNeRFLoc: Visual Localization with Point-Based Neural Radiance Fields	Boming Zhao; Luwei Yang; Mao Mao; Hujun Bao; Zhaopeng Cui
What Do Hebbian Learners Learn? Reduction Axioms for Unstable Hebbian Learning	Caleb Schultz Kisby; Saúl Blanco; Larry Moss
SAT-Based Tree Decomposition with Iterative Cascading Policy Selection	Hai Xia; Stefan Szeider
Measuring Self-Supervised Representation Quality for Downstream Classification Using Discriminative Features	Neha Kalibhat; Kanika Narang; Hamed Firooz; Maziar Sanjabi; Soheil Feizi
Once and for All: Universal Transferable Adversarial Perturbation against Deep Hashing-Based Facial Image Retrieval	Long Tang; Dengpan Ye; Yunna Lv; Chuanxi Chen; Yunming Zhang
Benchmarking Large Language Models on Controllable Generation under Diversified Instructions	Yihan Chen; Benfeng Xu; Quan Wang; Yi Liu; Zhendong Mao
Transformer-Based Selective Super-resolution for Efficient Image Refinement	Tianyi Zhang; Kishore Kasichainula; Yaixin Zhuo; Baoxin Li; Jae-Sun Seo; Yu Cao
Point Cloud Part Editing: Segmentation, Generation, Assembly, and Selection	Kaiyi Zhang; Yang Chen; Ximing Yang; Weizhong Zhang; Cheng Jin
Neural Oscillators for Generalization of Physics-Informed Machine Learning	Taniya Kapoor; Abhishek Chandra; Daniel M. Tartakovsky; Hongrui Wang; Alfredo Nunez; Rolf Dollevoet
Maximizing Nash Social Welfare under Two-Sided Preferences	Pallavi Jain; Rohit Vaish
Fast Machine Unlearning without Retraining through Selective Synaptic Dampening	Jack Foster; Stefan Schoepf; Alexandra Brintrup
NeSyFOLD: A Framework for Interpretable Image Classification	Parth Padalkar; Huaduo Wang; Gopal Gupta
MESED: A Multi-Modal Entity Set Expansion Dataset with Fine-Grained Semantic Classes and Hard Negative Entities	Li Yangning; Tingwei Lu; Hai-Tao Zheng; Yinghui Li; Shulin Huang; Tianyu Yu; Jun Yuan; Rui Zhang
Decoupling Degradations with Recurrent Network for Video Restoration in Under-Display Camera	Chengxu Liu; Xuan Wang; Yuanling Fan; Shuai Li; Xueming Qian
Learning Temporal Resolution in Spectrogram for Audio Classification	Haohe Liu; Xubo Liu; Qiuqiang Kong; Wenwu Wang; Mark Plumbley
FLAME: A Small Language Model for Spreadsheet Formulas	Harshit Joshi; José Cambronero Sanchez; Abishai Ebenezer; Sumit Gulwani; Aditya Kanade; Vu Le; Gust Verbruggen; Ivan Radiček
A Robust Mutual-Reinforcing Framework for 3D Multi-Modal Medical Image Fusion Based on Visual-Semantic Consistency	Hao Zhang; Xuhui Zuo; Huabing Zhou; Tao Lu; Jiayi Ma
Devignet: High-Resolution Vignetting Removal via a Dual Aggregated Fusion Transformer with Adaptive Channel Expansion	Shenghong Luo; Xuhang Chen; Weiwen Chen; Zinuo Li; Shuqiang Wang; Chi-Man Pun
KGDM: A Diffusion Model to Capture Multiple Relation Semantics for Knowledge Graph Embedding	Xiao Long; Liansheng Zhuang; Aodi Li; Jiuchang Wei; Houqiang Li; Shafei Wang
Catalyst for Clustering-Based Unsupervised Object Re-identification: Feature Calibration	Huafeng Li; Qingsong Hu; Zhanxuan Hu
Proportional Representation in Metric Spaces and Low-Distortion Committee Selection	Yusuf Kalayci; David Kempe; Vikram Kher
Finding Visual Saliency in Continuous Spike Stream	Lin Zhu; Xianzhang Chen; Xiao Wang; Hua Huang





Completing Priceable Committees: Utilitarian and Representation Guarantees for Proportional Multiwinner Voting	Markus Brill; Jannik Peters
MindMap: Constructing Evidence Chains for Multi-Step Reasoning in Large Language Models	Yangyu Wu; Xu Han; Wei Song; Miao Cheng; Fei Li
Simplicity Bias in Overparameterized Machine Learning	Yakir Berchenko
Plug-In Diffusion Model for Sequential Recommendation	Haokai Ma; Ruobing Xie; Lei Meng; Xin Chen; Xu Zhang; Leyu Lin; Zhanhui Kang
Curvature-Invariant Adversarial Attacks for 3D Point Clouds	Jianping Zhang; Wenwei Gu; Yizhan Huang; Zhihan Jiang; Weibin Wu; Michael Lyu
Contrastive Continual Learning with Importance Sampling and Prototype-Instance Relation Distillation	Jiyong Li; Dilshod Azizov; Yang Li; Shangsong Liang
Adaptive Integration of Partial Label Learning and Negative Learning for Enhanced Noisy Label Learning	Mengmeng Sheng; Zeren Sun; Zhenhuang Cai; Tao Chen; Yichao Zhou; Yazhou Yao
Robust Beamforming for Downlink Multi-Cell Systems: A Bilevel Optimization Perspective	Xingdi Chen; Yu Xiong; Kai YANG
CoLAL: Co-learning Active Learning for Text Classification	Linh Le; Genghong Zhao; Xia Zhang; Guido Zuccon; Gianluca Demartini
Targeted Activation Penalties Help CNNs Ignore Spurious Signals	Dekai Zhang; Matt Williams; Francesca Toni
DIUSum: Dynamic Image Utilization for Multimodal Summarization	Min Xiao; Junnan Zhu; Feifei Zhai; Yu Zhou; Chengqing Zong
FedLF: Layer-Wise Fair Federated Learning	Zibin Pan; Chi Li; Fangchen Yu; Shuyi Wang; Haijin Wang; Xiaoying Tang; Junhua Zhao
Paths, Proofs, and Perfection: Developing a Human-Interpretable Proof System for Constrained Shortest Paths	Konstantin Sidorov; Gonalo Homem de Almeida Correia; Mathijs de Weerd; Emir Demirović
Robustly Improving Bandit Algorithms with Confounded and Selection Biased Offline Data: A Causal Approach	Wen Huang; Xintao Wu
VQCNIR: Clearer Night Image Restoration with Vector-Quantized Codebook	Wenbin Zou; Hongxia Gao; Tian Ye; Liang Chen; Weipeng Yang; Shasha Huang; Hongsheng Chen; Sixiang Chen
NILUT: Conditional Neural Implicit 3D Lookup Tables for Image Enhancement	Marcos Conde; Javier Vazquez-Corral; Michael Brown; Radu Timofte
On Estimating the Gradient of the Expected Information Gain in Bayesian Experimental Design	Ziqiao Ao; Jinglai Li
Depth-Guided Robust and Fast Point Cloud Fusion NeRF for Sparse Input Views	Shuai Guo; Qiuwen Wang; Yijie Gao; Rong Xie; Li Song
Discerning Temporal Difference Learning	Jianfei Ma
PA2D-MORL: Pareto Ascent Directional Decomposition Based Multi-Objective Reinforcement Learning	Tianmeng Hu; Biao Luo
SafeAR: Safe Algorithmic Recourse by Risk-Aware Policies	Haochen Wu; Shubham Sharma; Sunandita Patra; Sriram Gopalakrishnan
Communication-Efficient Collaborative Regret Minimization in Multi-Armed Bandits	Nikolai Karpov; Qin Zhang
Multi-Scene Generalized Trajectory Global Graph Solver with Composite Nodes for Multiple Object Tracking	Yan Gao; Haojun Xu; Jie Li; Nannan Wang; Xinbo Gao
Multimodal Graph Neural Architecture Search under Distribution Shifts	Jie Cai; Xin Wang; Haoyang Li; Ziwei Zhang; Wenwu Zhu
Summarizing Stream Data for Memory-Constrained Online Continual Learning	Jianyang Gu; Kai Wang; Wei Jiang; Yang You



Progressive Poisoned Data Isolation for Training-Time Backdoor Defense	Yiming Chen; Haiwei Wu; Jiantao Zhou
A Dynamic GCN with Cross-Representation Distillation for Event-Based Learning	Yongjian Deng; Hao Chen; Youfu Li
Vision Transformer Off-the-Shelf: A Surprising Baseline for Few-Shot Class-Agnostic Counting	Zhicheng Wang; Liwen Xiao; Zhiguo Cao; Hao Lu
Fine-Grained Distillation for Long Document Retrieval	Yucheng Zhou; Tao Shen; Xiubo Geng; Chongyang Tao; Jianbing Shen; Guodong Long; Can Xu; Daxin Jiang
Simultaneous Optimization of Bid Shading and Internal Auction for Demand-Side Platforms	Yadong Xu; Bonan Ni; Weiran Shen; Xun Wang; Zichen Wang; Yinsong Xue; Pingzhong Tang
Let There Be Sound: Reconstructing High Quality Speech from Silent Videos	Ji-Hoon Kim; Jaehun Kim; Joon Son Chung
Explaining Generalization Power of a DNN Using Interactive Concepts	Huilin Zhou; Hao Zhang; Huiqi Deng; Dongrui Liu; Wen Shen; Shih-Han Chan; Quanshi Zhang
Text2City: One-Stage Text-Driven Urban Layout Regeneration	Yiming Qin; Nanxuan Zhao; Bin Sheng; Rynson W.H. Lau
ShapeBoost: Boosting Human Shape Estimation with Part-Based Parameterization and Clothing-Preserving Augmentation	Siyuan Bian; Jiefeng Li; Jiasheng Tang; Cewu Lu
Graph Invariant Learning with Subgraph Co-mixup for Out-of-Distribution Generalization	Tianrui Jia; Haoyang Li; Cheng Yang; Tao Tao; Chuan Shi
Unsupervised Object Interaction Learning with Counterfactual Dynamics	Jongwook Choi; Sungtae Lee; Xinyu Wang; Sungryull Sohn; Honglak Lee
Symbolic Cognitive Diagnosis via Hybrid Optimization for Intelligent Education Systems	Junhao Shen; Hong Qian; Wei Zhang; Aimin Zhou
Value Kaleidoscope: Engaging AI with Pluralistic Human Values, Rights, and Duties	Taylor Sorensen; Liwei Jiang; Jena Hwang; Sydney Levine; Valentina Pyatkin; Peter West; Nouha Dziri; Ximing Lu; Kavel Rao; Chandra Bhagavatula; Maarten Sap; John Tasioulas; Yejin Choi
Unsupervised Extractive Summarization with Learnable Length Control Strategies	Renlong Jie; Xiaojun Meng; Xin Jiang; Qun Liu
Neural Gaussian Similarity Modeling for Differential Graph Structure Learning	Xiaolong Fan; Maoguo Gong; Yue Wu; Zedong Tang; Jieyi Liu
Dual Self-Paced Cross-Modal Hashing	Yuan Sun; Jian Dai; Zhenwen Ren; Yingke Chen; Dezhong Peng; Peng Hu
Optimal Survival Trees: A Dynamic Programming Approach	Tim Huisman; Jacobus van der Linden; Emir Demirović
Expediting Contrastive Language-Image Pretraining via Self-Distilled Encoders	Bumsoo Kim; Jinhyung Kim; Yeonsik Jo; Seung Hwan Kim
Dynamic Semantic-Based Spatial Graph Convolution Network for Skeleton-Based Human Action Recognition	Jianyang Xie; Yanda Meng; Yitian Zhao; Anh Nguyen; Xiaoyun Yang; Yalin Zheng
LISR: Learning Linear 3D Implicit Surface Representation Using Compactly Supported Radial Basis Functions	Atharva Pandey; Vishal Yadav; Rajendra Nagar; Santanu Chaudhury
FM-OV3D: Foundation Model-Based Cross-Modal Knowledge Blending for Open-Vocabulary 3D Detection	Dongmei Zhang; Chang Li; Renrui Zhang; Shenghao Xie; Wei Xue; Xiaodong Xie; Shanghang Zhang
FedLPS: Heterogeneous Federated Learning for Multiple Tasks with Local Parameter Sharing	Yongzhe Jia; Xuyun Zhang; Amin Beheshti; Wanchun Dou
Jointly Modeling Spatio-Temporal Features of Tactile Signals for Action Classification	Jimmy Lin; Junkai Li; Jiasi Gao; Weizhi Ma; Yang Liu
An Interpretable Approach to the Solutions of High-Dimensional Partial Differential Equations	Lulu Cao; Yufei Liu; Zhenzhong Wang; Dejun Xu; Kai Ye; Kay Chen Tan; Min Jiang



Compressing Image-to-Image Translation GANs Using Local Density Structures on Their Learned Manifold	Alireza Ganjdanesh; Shangqian Gao; Hiran Alipanah; Heng Huang
Detecting and Preventing Hallucinations in Large Vision Language Models	Anisha Gunjal; Jihan Yin; Erhan Bas
Debiased Novel Category Discovering and Localization	Juexiao Feng; Yuhong Yang; Yanchun Xie; Yaqian Li; Yandong Guo; Yuchen Guo; Yuwei He; Liuyu Xiang; Guiguang Ding
SHAP@k: Efficient and Probably Approximately Correct (PAC) Identification of Top-K Features	Sanjay Kariyappa; Leonidas Tsepenekas; Freddy Lecue; Daniele Magazzeni
Correlation Matching Transformation Transformers for UHD Image Restoration	Cong Wang; Jinshan Pan; Wei Wang; Gang Fu; Siyuan Liang; Mengzhu Wang; Xiao-Ming Wu; Jun Liu
Wasserstein Differential Privacy	Chengyi Yang; Jiayin Qi; Aimin Zhou
Exploiting Discrepancy in Feature Statistic for Out-of-Distribution Detection	Xiaoyuan Guan; Jiankang Chen; Shenshen Bu; Yuren Zhou; Wei-shi Zheng; Ruixuan Wang
Bi-directional Adapter for Multimodal Tracking	Bing Cao; Junliang Guo; Pengfei Zhu; Qinghua Hu
Using Stratified Sampling to Improve LIME Image Explanations	Muhammad Rashid; Elvio G. Amparore; Enrico Ferrari; Damiano Verda
Multi-Constellation-Inspired Single-Shot Global LiDAR Localization	Tongzhou Zhang; Gang Wang; Yu Chen; Hai Zhang; Jue Hu
M3SOT: Multi-Frame, Multi-Field, Multi-Space 3D Single Object Tracking	Jiaming Liu; Yue Wu; Maoguo Gong; Qiguang Miao; Wenping Ma; Cai Xu; Can Qin
MSGNet: Learning Multi-Scale Inter-series Correlations for Multivariate Time Series Forecasting	Wanlin Cai; Yuxuan Liang; Xianggen Liu; Jianshuai Feng; Yuankai Wu
Value at Adversarial Risk: A Graph Defense Strategy against Cost-Aware Attacks	Junlong Liao; Wenda Fu; Cong Wang; Zhongyu Wei; Jiarong Xu
HuTuMotion: Human-Tuned Navigation of Latent Motion Diffusion Models with Minimal Feedback	Gaoge Han; Shaoli Huang; Mingming Gong; Jinglei Tang
NN-Steiner: A Mixed Neural-Algorithmic Approach for the Rectilinear Steiner Minimum Tree Problem	Andrew Kahng; Robert Nerem; Yusu Wang; Chien-Yi Yang
Asymmetric Mutual Alignment for Unsupervised Zero-Shot Sketch-Based Image Retrieval	Zhihui Yin; Jiexi Yan; Chenghao Xu; Cheng Deng
MotionGPT: Finetuned LLMs Are General-Purpose Motion Generators	Yaqi Zhang; Di Huang; Bin Liu; Shixiang Tang; Yan Lu; Lu Chen; Lei Bai; Qi Chu; Nenghai Yu; Wanli Ouyang
SpaceGTN: A Time-Agnostic Graph Transformer Network for Handwritten Diagram Recognition and Segmentation	Haoxiang Hu; Cangjun Gao; YaoKun Li; Xiaoming Deng; Yukun Lai; Cuixia Ma; Yong-Jin Liu; Hongan Wang
A Reinforcement-Learning-Based Multiple-Column Selection Strategy for Column Generation	Haofeng Yuan; Lichang Fang; Shiji Song
A Primal-Dual Algorithm for Hybrid Federated Learning	Tom Overman; Garrett Blum; Diego Klabjan
IRPruneDet: Efficient Infrared Small Target Detection via Wavelet Structure-Regularized Soft Channel Pruning	Mingjin Zhang; Handi Yang; Jie Guo; Yunsong Li; Xinbo Gao; Jing Zhang
Negative Pre-aware for Noisy Cross-Modal Matching	Xu Zhang; Hao Li; Mang Ye
DALDet: Depth-Aware Learning Based Object Detection for Autonomous Driving	Ke Hu; Tongbo Cao; Yuan Li; Song Chen; Yi Kang
AdaCCD: Adaptive Semantic Contrasts Discovery Based Cross Lingual Adaptation for Code Clone Detection	Yangkai Du; Tengfei Ma; Lingfei Wu; Xuhong Zhang; Shouling Ji



Stratified GNN Explanations through Sufficient Expansion	Yuwen Ji; Lei Shi; Zhimeng Liu; Ge Wang
Backward Responsibility in Transition Systems Using General Power Indices	Christel Baier; Roxane van den Bossche; Sascha Klüppelholz; Johannes Lehmann; Jakob Piribauer
Review-Enhanced Hierarchical Contrastive Learning for Recommendation	Ke Wang; Yanmin Zhu; Tianzi Zang; Chunyang Wang; Mengyuan Jing
FashionERN: Enhance-and-Refine Network for Composed Fashion Image Retrieval	Yanzhe Chen; Huasong Zhong; Xiangteng He; Yuxin Peng; Jiahuan Zhou; Lele Cheng
Less Is More: Label Recommendation for Weakly Supervised Point Cloud Semantic Segmentation	Zhiyi Pan; Nan Zhang; Wei Gao; Shan Liu; Ge Li
SAUI: Scale-Aware Unseen Imagineer for Zero-Shot Object Detection	Jiahao Wang; Caixia Yan; Weizhan Zhang; Huan Liu; Hao Sun; Qinghua
Deciphering Raw Data in Neuro-Symbolic Learning with Provable Guarantees	Lue Tao; Yu-Xuan Huang; Wang-Zhou Dai; Yuan Jiang
What Effects the Generalization in Visual Reinforcement Learning: Policy Consistency with Truncated Return Prediction	Shuo Wang; Zhihao Wu; Xiaobo Hu; Jinwen Wang; Youfang Lin; Kai Lv
Scaling and Masking: A New Paradigm of Data Sampling for Image and Video Quality Assessment	Yongxu Liu; Yinghui Quan; Guoyao Xiao; Aobo Li; Jinjian Wu
FedCompetitors: Harmonious Collaboration in Federated Learning with Competing Participants	Shanli Tan; Hao Cheng; Xiaohu Wu; Han Yu; Tiantian He; Yew Soon Ong; Chongjun Wang; Xiaofeng Tao
COMBAT: Alternated Training for Effective Clean-Label Backdoor Attacks	Tran Huynh; Dang Nguyen; Tung Pham; Anh Tran
Gradient-Guided Modality Decoupling for Missing-Modality Robustness	Hao Wang; Shengda Luo; Guosheng Hu; Jianguo Zhang
Optimistic Policy Gradient in Multi-Player Markov Games with a Single Controller: Convergence beyond the Minty Property	Ioannis Anagnostides; Ioannis Panageas; Gabriele Farina; Tuomas Sandholm
Open-Vocabulary Video Relation Extraction	Wentao Tian; Zheng Wang; Yuqian Fu; Jingjing Chen; Lechao Cheng
Propagation Tree Is Not Deep: Adaptive Graph Contrastive Learning Approach for Rumor Detection	Chaoqun Cui; Caiyan Jia
Learning Robust Rationales for Model Explainability: A Guidance-Based Approach	Shuaibo Hu; Kui Yu
FontDiffuser: One-Shot Font Generation via Denoising Diffusion with Multi-Scale Content Aggregation and Style Contrastive Learning	Zhenhua Yang; Dezhi Peng; Yuxin Kong; Yuyi Zhang; Cong Yao; Lianwen Jin
Leveraging Diffusion Perturbations for Measuring Fairness in Computer Vision	Nicholas Lui; Bryan Chia; William Berrios; Candace Ross; Douwe Kiela
Using Clustering to Strengthen Decision Diagram Bounds for Discrete Optimization	Mohsen Nafar; Michael Römer
Rethinking the Paradigm of Content Constraints in Unpaired Image-to-Image Translation	Xiuding Cai; Yaoyao Zhu; Dong Miao; Linjie Fu; Yu Yao
A Unified Masked Autoencoder with Patchified Skeletons for Motion Synthesis	Esteve Valls Mascaro; Hyemin Ahn; Dongheui Lee
LAMM: Label Alignment for Multi-Modal Prompt Learning	Jingsheng Gao; Jiacheng Ruan; Suncheng Xiang; Zefang Yu; Ke Ji; Mingye Xie; Ting Liu; Yuzhuo Fu
Knowledge-Enhanced Historical Document Segmentation and Recognition	En-Hao Gao; Yu-Xuan Huang; Wen-Chao Hu; Xin-Hao Zhu; Wang-Zhou Dai
Beyond OOD State Actions: Supported Cross-Domain Offline Reinforcement Learning	Jinxin Liu; Ziqi Zhang; Zhenyu Wei; Zifeng Zhuang; Yachen Kang; Sibao Gai; Donglin Wang
Scribble Hides Class: Promoting Scribble-Based Weakly-Supervised Semantic Segmentation with Its Class Label	Xinliang Zhang; Lei Zhu; Hangzhou He; Lujia Jin; Yanye Lu
EAT: Towards Long-Tailed Out-of-Distribution Detection	Tong Wei; Bo-Lin Wang; Min-Ling Zhang



Angle Robustness Unmanned Aerial Vehicle Navigation in GNSS-Denied Scenarios	Yuxin Wang; Zunlei Feng; Haofei Zhang; Yang Gao; Jie Lei; Li Sun; Mingli Song
DenoSent: A Denoising Objective for Self-Supervised Sentence Representation Learning	Xinghao Wang; Junliang He; Pengyu Wang; Yunhua Zhou; Tianxiang Sun; Xipeng Qiu
G2L-CariGAN: Caricature Generation from Global Structure to Local Features	Xin Huang; Yunfeng Bai; Dong Liang; Feng Tian; Jinyuan Jia
The Complexity of Fair Division of Indivisible Items with Externalities	Argyrios Deligkas; Eduard Eiben; Viktoriia Korchemna; Simon Schierreich
A Unified Knowledge Transfer Network for Generalized Category Discovery	Wenkai Shi; Wenbin An; Feng Tian; Yan Chen; Yaqiang Wu; Qianying Wang; Ping Chen
Multi-Step Denoising Scheduled Sampling: Towards Alleviating Exposure Bias for Diffusion Models	Zhiyao Ren; Yibing Zhan; Liang Ding; Gaoang Wang; Chaoyue Wang; Zhongyi Fan; Dacheng Tao
TagFog: Textual Anchor Guidance and Fake Outlier Generation for Visual Out-of-Distribution Detection	Jiankang Chen; Tong Zhang; Wei-shi Zheng; Ruixuan Wang
Strategyproof Mechanisms for Group-Fair Obnoxious Facility Location Problems	Jiaqian Li; Minming Li; Hau Chan
Vision-Language Pre-training with Object Contrastive Learning for 3D Scene Understanding	Taolin Zhang; Sunan He; Tao Dai; Zhi Wang; Bin Chen; Shu-Tao Xia
Manifold-Based Verbalizer Space Re-embedding for Tuning-Free Prompt-Based Classification	Haochun Wang; Sendong Zhao; Chi Liu; Nuwa Xi; MuZhen Cai; Bing Qin; Ting Liu
Cross-Layer and Cross-Sample Feature Optimization Network for Few-Shot Fine-Grained Image Classification	Zhen-Xiang Ma; Zhen-Duo Chen; Li-Jun Zhao; Zi-Chao Zhang; Xin Luo; Xin-Shun Xu
Semi-supervised Blind Image Quality Assessment through Knowledge Distillation and Incremental Learning	Wensheng Pan; Timin Gao; Yan Zhang; Xiawu Zheng; Yunhang Shen; Ke Li; Runze Hu; Yutao Liu; Pingyang Dai
Graph-Aware Contrasting for Multivariate Time-Series Classification	Yucheng Wang; Yuecong Xu; Jianfei Yang; Min Wu; Xiaoli Li; Lihua Xie; Zhenghua Chen
Learning to Reweight for Generalizable Graph Neural Network	Zhengyu Chen; Teng Xiao; Kun Kuang; Zheqi Lv; Min Zhang; Jinluan Yang; Chengqiang Lu; Hongxia Yang; Fei Wu
s-ID: Causal Effect Identification in a Sub-population	Amir Mohammad Abouei; Ehsan Mokhtarian; Negar Kiyavash
LDS2AE: Local Diffusion Shared-Specific Autoencoder for Multimodal Remote Sensing Image Classification with Arbitrary Missing Modalities	Jiahui Qu; Yuanbo Yang; Wenqian Dong; Yufei Yang
Hidden Follower Detection: How Is the Gaze-Spacing Pattern Embodied in Frequency Domain?	Shu Li; Ruimin Hu; Suhui Li; Liang Liao
WebVLN: Vision-and-Language Navigation on Websites	Qi Chen; Dileepa Pitawela; Chongyang Zhao; Gengze Zhou; Hsiang-Ting Chen; Qi Wu
Fine-Grained Prototypes Distillation for Few-Shot Object Detection	Zichen Wang; Bo Yang; Haonan Yue; Zhenghao Ma
Advancing Video Synchronization with Fractional Frame Analysis: Introducing a Novel Dataset and Model	Yuxuan Liu; Haizhou Ai; Junliang Xing; Xuri Li; Xiaoyi Wang; Pin Tao
Unlocking the Power of Open Set: A New Perspective for Open-Set Noisy Label Learning	Wenhai Wan; Shao-Yuan Li; Xinrui Wang; Ming-Kun Xie; Sheng-Jun Huang; Songcan Chen
Parameterized Projected Bellman Operator	Théo Vincent; Alberto Maria Metelli; Boris Belousov; Jan Peters; Marcello Restelli; Carlo D'Eramo
Multi-Domain Multi-Scale Diffusion Model for Low-Light Image Enhancement	Kai Shang; Mingwen Shao; Chao Wang; Yuanshuo Cheng; Shuigen Wang



Weakly Supervised Multimodal Affordance Grounding for Egocentric Images	Lingjing Xu; Yang Gao; Wenfeng Song; Aimin Hao
A Dual-Way Enhanced Framework from Text Matching Point of View for Multimodal Entity Linking	Shezheng Song; Shan Zhao; ChengYu Wang; Tianwei Yan; Shasha Li; Xiaoguang Mao; Meng Wang
Learn How to See: Collaborative Embodied Learning for Object Detection and Camera Adjusting	Lingdong Shen; Chunlei Huo; Nuo Xu; Chaowei Han; Zichen Wang
SCD-Net: Spatiotemporal Clues Disentanglement Network for Self-Supervised Skeleton-Based Action Recognition	Cong Wu; Xiao-Jun Wu; Josef Kittler; Tianyang Xu; Sara Ahmed; Muhammad Awais; Zhenhua Feng
When to Show a Suggestion? Integrating Human Feedback in AI-Assisted Programming	Hussein Mozannar; Gagan Bansal; Adam Fourney; Eric Horvitz
Adaptive Hardness Negative Sampling for Collaborative Filtering	Riwei Lai; Rui Chen; Qilong Han; Chi Zhang; Li Chen
SeTformer Is What You Need for Vision and Language	Pourya Shamsolmoali; Masoumeh Zareapoor; Eric Granger; Michael Felsberg
Learning Small Decision Trees with Few Outliers: A Parameterized	Harmender Gahlawat; Meirav Zehavi
Refining Latent Homophilic Structures over Heterophilic Graphs for Robust Graph Convolution Networks	Chenyang Qiu; Guoshun Nan; Tianyu Xiong; Wendi Deng; Di Wang; Zhiyang Teng; Lijuan Sun; Qimei Cui; Xiaofeng Tao
CatmullRom Splines-Based Regression for Image Forgery Localization	Li Zhang; Mingliang Xu; Dong Li; Jianming Du; Rujing Wang
BaCon: Boosting Imbalanced Semi-supervised Learning via Balanced Feature-Level Contrastive Learning	Qianhan Feng; Lujing Xie; Shijie Fang; Tong Lin
MuLTI: Efficient Video-and-Language Understanding with Text-Guided MultiWay-Sampler and Multiple Choice Modeling	Jiaqi Xu; Bo Liu; Yunkuo Chen; Mengli Cheng; Xing Shi
Generalizable Sleep Staging via Multi-Level Domain Alignment	Jiquan Wang; Sha Zhao; Haiteng Jiang; Shijian Li; Tao Li; Gang Pan
Improving Knowledge Extraction from LLMs for Task Learning through Agent Analysis	James Kirk; Robert Wray; Peter Lindes; John Laird
GxVAEs: Two Joint VAEs Generate Hit Molecules from Gene Expression Profiles	Chen Li; Yoshihiro Yamanishi
Delving into Multimodal Prompting for Fine-Grained Visual Classification	Xin Jiang; Hao Tang; Junyao Gao; Xiaoyu Du; Shengfeng He; Zechao Li
Molecular Optimization Model with Patentability Constraint	Sally Turutov; Kira Radinsky
TAPE: Leveraging Agent Topology for Cooperative Multi-Agent Policy Gradient	Xingzhou Lou; Junge Zhang; Timothy Norman; Kaiqi Huang; Yali Du
Component Fourier Neural Operator for Singularly Perturbed Differential Equations	Ye Li; Ting Du; Yiwen Pang; Zhongyi Huang
Federated Causality Learning with Explainable Adaptive Optimization	Dezhi Yang; Xintong He; Jun Wang; Guoxian Yu; Carlotta Domeniconi; Jinglin Zhang
Diverse Person: Customize Your Own Dataset for Text-Based Person Search	Zifan Song; Guosheng Hu; Cairong Zhao
Recognizing Ultra-High-Speed Moving Objects with Bio-Inspired Spike	Junwei Zhao; Shiliang Zhang; Zhaofei Yu; Tiejun Huang
T2MAC: Targeted and Trusted Multi-Agent Communication through Selective Engagement and Evidence-Driven Integration	Chuxiong Sun; Zehua Zang; Jiabao Li; Jiangmeng Li; Xiao Xu; Rui Wang; Changwen Zheng
Region-Disentangled Diffusion Model for High-Fidelity PPG-to-ECG	Debaditya Shome; Pritam Sarkar; Ali Etemad
MFTN: Multi-Level Feature Transfer Network Based on MRI-Transformer for MR Image Super-resolution	Shuying Huang; Ge Chen; Yong Yang; Xiaozheng Wang; Chenbin Liang
Pay Attention to Target: Relation-Aware Temporal Consistency for Domain Adaptive Video Semantic Segmentation	Huayu Mai; Rui Sun; Yuan Wang; Tianzhu Zhang; Feng Wu





LaneGraph2Seq: Lane Topology Extraction with Language Model via Vertex-Edge Encoding and Connectivity Enhancement	Renyuan Peng; Xinyue Cai; Hang Xu; Jiachen Lu; Feng Wen; Wei Zhang; Li Zhang
Goal Alignment: Re-analyzing Value Alignment Problems Using Human-Aware Building Variable-Sized Models via Learngene Pool	Malek Mechergui; Sarath Sreedharan
Poincaré Differential Privacy for Hierarchy-Aware Graph Embedding	Boyu Shi; Shiyu Xia; Xu Yang; Haokun Chen; Zhiqiang Kou; Xin Geng
Generalized Planning for the Abstraction and Reasoning Corpus	Yuecen Wei; Haonan Yuan; Xingcheng Fu; Qingyun Sun; Hao Peng; Xianxian Li; Chunming Hu
IT3D: Improved Text-to-3D Generation with Explicit View Synthesis	Chao Lei; Nir Lipovetzky; Krista A. Ehinger
Quantum-Inspired Neural Network with Runge-Kutta Method	Yiwen Chen; Chi Zhang; Xiaofeng Yang; Zhongang Cai; Gang Yu; Lei Yang; Guosheng Lin
Chinese Spelling Correction as Rephrasing Language Model	Zipeng Fan; Jing Zhang; Peng Zhang; Qianxi Lin; Hui Gao
Stable Model Semantics for Description Logic Terminologies	Linfeng Liu; Hongqiu Wu; Hai Zhao
Detection-Based Intermediate Supervision for Visual Question Answering	Federica Di Stefano; Mantas Simkus
FlexKBQA: A Flexible LLM-Powered Framework for Few-Shot Knowledge Base Question Answering	Yuhang Liu; Daowan Peng; Wei Wei; Yuanyuan Fu; Xie Wenfeng; Dangyang Chen
DME: Unveiling the Bias for Better Generalized Monocular Depth Estimation	Zhenyu Li; Sunqi Fan; Yu Gu; Zhichao Duan; Bowen Dong; Ning Liu; Xiuxing Li; Jianyong Wang
Forced Exploration in Bandit Problems	Songsong Yu; Yifan Wang; Yunzhi Zhuge; Lijun Wang; Huchuan Lu
Optimal Attack and Defense for Reinforcement Learning	Qi Han; Li Zhu; Fei Guo
Optimal Mechanism in a Dynamic Stochastic Knapsack Environment	Jeremy McMahan; Young Wu; Xiaojin Zhu; Qiaomin Xie
Cross-Constrained Progressive Inference for 3D Hand Pose Estimation	Jihyeok Jung; Chan-Oi Song; Deok-Joo Lee; Kiho Yoon
Intelligent Calibration for Bias Reduction in Sentiment Corpora Annotation Process	Zhehan Kan; Xueting Hu; Zihan Liao; Ke Yu; Zhihai He
Fairness under Covariate Shift: Improving Fairness-Accuracy Tradeoff with Few Unlabeled Test Samples	Idan Toker; David Sarne; Jonathan Schler
Decomposing Constraint Networks for Calculating c-Representations	Shreyas Havaladar; Jatin Chauhan; Karthikeyan Shanmugam; Jay Nandy; Aravindan Raghuvier
Embedded Feature Selection on Graph-Based Multi-View Clustering	Marco Wilhelm; Gabriele Kern-Isberner
Intrinsic Action Tendency Consistency for Cooperative Multi-Agent Reinforcement Learning	Guangfei Li; Haizhou Yang; Quanxue Gao; Qianqian Wang; Wenhui Zhao
A Learnable Discrete-Prior Fusion Autoencoder with Contrastive Learning for Tabular Data Synthesis	Junkai Zhang; Yifan Zhang; Xi Sheryl Zhang; Yifan Zang; Jian Cheng
On the Convergence of an Adaptive Momentum Method for Adversarial Unsupervised Group Re-identification via Adaptive Clustering-Driven Progressive Learning	Rongchao Zhang; Yu Huang; Yiwei Lou; Dexuan Xu; Yongzhi Cao; Hanpin Wang
Well, Now We Know! Unveiling Sarcasm: Initiating and Exploring Multimodal Conversations with Reasoning	Sheng Long; Wei Tao; Shuohao Li; Jun Lei; Jun Zhang
Decoding AI's Nudge: A Unified Framework to Predict Human Behavior in AI-Assisted Decision Making	Hongxu Chen; Quan Zhang; Jian-Huang Lai; Xiaohua Xie
StegFormer: Rebuilding the Glory of Autoencoder-Based Steganography	Gopendra Singh; Mauajama Firdaus; Dushyant Singh Chauhan; Asif Ekbal; Pushpak Bhattacharyya
Learning to Prompt Knowledge Transfer for Open-World Continual Learning	Zhuoyan Li; Zhuoran Lu; Ming Yin
Transferable Video Moment Localization by Moment-Guided Query Prompting	Xiao Ke; Huanqi Wu; Wenzhong Guo
	Yujie Li; Xin Yang; Hao Wang; Xiangkun Wang; Tianrui Li
	Hao Jiang; Yang Yizhang; Yadong Mu



Energy-Efficient Streaming Time Series Classification with Attentive Power Iteration	Hao Huang; Tapan Shah; Scott Evans; Shinjae Yoo
BEV-MAE: Bird's Eye View Masked Autoencoders for Point Cloud Pre-training in Autonomous Driving Scenarios	ZhiWei Lin; Yongtao Wang; Shengxiang Qi; Nan Dong; Ming-Hsuan Yang
PORTAL: Automatic Curricula Generation for Multiagent Reinforcement Learning	Jizhou Wu; Jianye Hao; Tianpei Yang; Xiaotian Hao; Yan Zheng; Weixun Wang; Matthew Taylor
Towards Continual Knowledge Graph Embedding via Incremental Distillation	Jiajun Liu; Ke Wenjun; Peng Wang; Ziyu Shang; Gao Jinhua; Guozheng Li; Ke Ji; Yanhe Liu
Compositional Generalization for Multi-Label Text Classification: A Data-Augmentation Approach	Yuyang Chai; Zhuang Li; Jiahui Liu; Lei Chen; Fei Li; Donghong Ji; Chong Teng
M3D: Dataset Condensation by Minimizing Maximum Mean Discrepancy	Hansong Zhang; Shikun Li; Pengju Wang; Dan Zeng; Shiming Ge
Principle Component Trees and Their Persistent Homology	Ben Kizaric; Daniel Pimentel-Alarcon
Knowledge-Aware Explainable Reciprocal Recommendation	Kai-Huang Lai; Zhe-Rui Yang; Pei-Yuan Lai; Chang-Dong Wang; Mohsen Guizani; Min Chen
Threshold-Based Responsive Simulated Annealing for Directed Feedback Vertex Set Problem	Qingyun Zhang; YuMing Du; Zhipeng Lü; Chu-Min Li; Zhouxing Su; Junzhou Xu; Zhihuai Chen
MFABA: A More Faithful and Accelerated Boundary-Based Attribution Method for Deep Neural Networks	Zhiyu Zhu; Huaming Chen; Jiayu Zhang; Xinyi Wang; Zhibo Jin; Minhui Xue; Dongxiao Zhu; Kim-Kwang Raymond Choo
Data Poisoning to Fake a Nash Equilibria for Markov Games	Young Wu; Jeremy McMahan; Xiaojin Zhu; Qiaomin Xie
How to Use the Metropolis Algorithm for Multi-Objective Optimization?	Weijie Zheng; Mingfeng Li; Renzhong Deng; Benjamin Doerr
Exploring Diverse Representations for Open Set Recognition	Yu Wang; Junxian Mu; Pengfei Zhu; Qinghua Hu
Improving the Robustness of Knowledge-Grounded Dialogue via Contrastive Learning	Jiaan Wang; Jianfeng Qu; Kexin Wang; Zhixu Li; Wen Hua; Ximing Li; An Liu
Hyperspectral Image Reconstruction via Combinatorial Embedding of Cross-Channel Spatio-Spectral Clues	Xingxing Yang; Jie Chen; Zaifeng Yang
Ghost Noise for Regularizing Deep Neural Networks	Atli Kosson; Dongyang Fan; Martin Jaggi
Eliciting Kemeny Rankings	Anne-Marie George; Christos Dimitrakakis
A Sequentially Fair Mechanism for Multiple Sensitive Attributes	Francois HU; Philipp Ratz; Arthur Charpentier
From GARCH to Neural Network for Volatility Forecast	Pengfei Zhao; Haoren ZHU; Wilfred Ng; Dik-Lun LEE
Robust Visual Recognition with Class-Imbalanced Open-World Noisy Data	Na Zhao; Gim Hee Lee
Coupling Graph Neural Networks with Fractional Order Continuous Dynamics: A Robustness Study	Qiyu Kang; Kai Zhao; Yang Song; Yihang Xie; Yanan Zhao; Sijie Wang; Rui She; Wee Peng Tay
Text-Guided Molecule Generation with Diffusion Language Model	Haisong Gong; Qiang Liu; Shu Wu; Liang Wang
Three Heads Are Better than One: Complementary Experts for Long-Tailed Semi-supervised Learning	Chengcheng Ma; Ismail Elezi; Jiankang Deng; Weiming Dong; Changsheng Xu
TETRIS: Towards Exploring the Robustness of Interactive Segmentation	Andrey Moskalenko; Vlad Shakhuro; Anna Vorontsova; Anton Konushin; Anton Antonov; Alexander Krapukhin; Denis Shepelev; Konstantin Soshin
DVSAI: Diverse View-Shared Anchors Based Incomplete Multi-View Clustering	Shengju Yu; Siwei Wang; Pei Zhang; Miao Wang; Ziming Wang; Zhe Liu; Liming Fang; En Zhu; Xinwang Liu
SoundCount: Sound Counting from Raw Audio with Dyadic Decomposition Neural Network	Yuhang He; Zhuangzhuang Dai; Niki Trigoni; Long Chen; Andrew Markham



TagCLIP: A Local-to-Global Framework to Enhance Open-Vocabulary Multi-Label Classification of CLIP without Training	Yuqi Lin; Minghao Chen; Kaipeng Zhang; Hengjia Li; Mingming Li; Zheng Yang; Dongqin Lv; Binbin Lin; Haifeng Liu; Deng Cai
Semantic-Aware Autoregressive Image Modeling for Visual Representation Learning	Kaiyou Song; Shan Zhang; Tong Wang
Strong Baselines for Parameter-Efficient Few-Shot Fine-Tuning	Samyadeep Basu; Shell Hu; Daniela Massiceti; Soheil Feizi
Testing Self-Reducible Samplers	Rishiraj Bhattacharyya; Sourav Chakraborty; Yash Pote; Uddalok Sarkar; Sayantan Sen
MCSSME: Multi-Task Contrastive Learning for Semi-supervised Singing Melody Extraction from Polyphonic Music	Shuai Yu
End-to-End Real-Time Vanishing Point Detection with Transformer	Xin Tong; Shi Peng; Yufei Guo; Xuhui Huang
Direction-Aware Video Demoiréing with Temporal-Guided Bilateral Learning	Shuning Xu; Binbin SONG; Xiangyu Chen; Jiantao Zhou
Towards Optimal Subsidy Bounds for Envy-Freeable Allocations	Yasushi Kawase; Kazuhisa Makino; Hanna Sumita; Akihisa Tamura; Makoto Yokoo
Learning Hybrid Dynamics Models with Simulator-Informed Latent States	Katharina Ensinger; Sebastian Ziesche; Sebastian Trimpe
Factorized Diffusion Autoencoder for Unsupervised Disentangled Representation Learning	Ancong Wu; Wei-shi Zheng
Percentile Risk-Constrained Budget Pacing for Guaranteed Display Advertising in Online Optimization	Liang Dai; Kejie Lyu; Chengcheng Zhang; Guangming Zhao; Zhonglin Zu; Liang Wang; Bo Zheng
MERGE: Fast Private Text Generation	Zi Liang; Pinghui Wang; Ruofei Zhang; Nuo Xu; Shuo Zhang; Lifeng Xing; Haitao Bai; Ziyang Zhou
Working Memory Capacity of ChatGPT: An Empirical Study	Dongyu Gong; Xingchen Wan; Dingmin Wang
Everything2Motion: Synchronizing Diverse Inputs via a Unified Framework for Human Motion Synthesis	Zhaoxin Fan; Longbin Ji; Pengxin Xu; Fan Shen; Kai Chen
Adaptive Anytime Multi-Agent Path Finding Using Bandit-Based Large Neighborhood Search	Thomy Phan; Taoan Huang; Bistra Dilkina; Sven Koenig
AnomalyDiffusion: Few-Shot Anomaly Image Generation with Diffusion Model	Teng Hu; Jiangning Zhang; Ran Yi; Yuzhen Du; Xu Chen; Liang Liu; Yabiao Wang; Chengjie Wang
Sterling: Synergistic Representation Learning on Bipartite Graphs	Baoyu Jing; Yuchen Yan; Kaize Ding; Chanyoung Park; Yada Zhu; Huan Liu; Hanghang Tong
From Coarse to Fine: A Distillation Method for Fine-Grained Emotion-Causal Span Pair Extraction in Conversation	Xinhao Chen; Chong Yang; Changzhi Sun; Man LAN; Aimin Zhou
Decouple Content and Motion for Conditional Image-to-Video Generation	Cuifeng Shen; Yulu Gan; Chen Chen; Xiongwei Zhu; Lele Cheng; Tingting Gao; Jinzhi Wang
Roll with the Punches: Expansion and Shrinkage of Soft Label Selection for Semi-supervised Fine-Grained Learning	Yue Duan; Zhen Zhao; Lei Qi; Luping Zhou; Lei Wang; Yinghuan Shi
Compact HD Map Construction via Douglas-Peucker Point Transformer	Ruixin Liu; Zejian Yuan
An Eager Satisfiability Modulo Theories Solver for Algebraic Datatypes	Amar Shah; Federico Mora; Sanjit A. Seshia
Bridging the Gap between 2D and 3D Visual Question Answering: A Fusion Approach for 3D VQA	Wentao Mo; Yang Liu
Competition among Pairwise Lottery Contests	Xiaotie Deng; Hangxin Gan; Ningyuan Li; Weian Li; Qi Qi



S2CycleDiff: Spatial-Spectral-Bilateral Cycle-Diffusion Framework for Hyperspectral Image Super-resolution	Jiahui Qu; Jie He; Wenqian Dong; Jingyu Zhao
MathAttack: Attacking Large Language Models towards Math Solving Ability	Zihao Zhou; Qiufeng Wang; Mingyu Jin; Jie Yao; Jianan Ye; Wei Liu; Wei Wang; Xiaowei Huang; Kaizhu Huang
Relax Image-Specific Prompt Requirement in SAM: A Single Generic Prompt for Segmenting Camouflaged Objects	Jian Hu; Jiayi Lin; Shaogang Gong; Weitong Cai
Full Bayesian Significance Testing via Neural Networks	Zehua Liu; Zimeng Li; Jingyuan Wang; Yue He
Spatial-Temporal Interplay in Human Mobility: A Hierarchical Reinforcement Learning Approach with Hypergraph Representation	Zhaofan Zhang; Yanan Xiao; Lu Jiang; Dingqi Yang; Minghao Yin; Pengyang Wang
ConditionVideo: Training-Free Condition-Guided Video Generation	Bo Peng; Xinyuan Chen; Yaohui Wang; Chaochao Lu; Yu Qiao
Federated Learning with Extremely Noisy Clients via Negative Distillation	Yang Lu; Lin Chen; Yonggang Zhang; Yiliang Zhang; Bo Han; Yiu-ming Cheung; Hanzi Wang
On Disentanglement of Asymmetrical Knowledge Transfer for Modality-Task Agnostic Federated Learning	Jiayi Chen; Aidong Zhang
Unify Named Entity Recognition Scenarios via Contrastive Real-Time Updating Prototype	Yanhe Liu; Peng Wang; Ke Wenjun; Guozheng Li; Xiye Chen; Jiteng Zhao; Ziyu Shang
Union Subgraph Neural Networks	Jiaxing Xu; Aihu Zhang; Qingtian Bian; Vijay Prakash Dwivedi; Yiping Ke
Reliable Data Generation and Selection for Low-Resource Relation Extraction	Junjie Yu; Xing Wang; Wenliang Chen
H-ensemble: An Information Theoretic Approach to Reliable Few-Shot Multi-Source-Free Transfer	Yanru Wu; Jianning Wang; Weida Wang; Yang Li
UNEX-RL: Reinforcing Long-Term Rewards in Multi-Stage Recommender Systems with UNidirectional EXecution	Gengrui Zhang; Xiaoshuang Chen; Yao WANG; HongYi Qian; Kaiqiao Zhan; Ben Wang
Novel Class Discovery in Chest X-rays via Paired Images and Text	Jiaying Zhou; Yang Liu; Qingchao Chen
BiPFT: Binary Pre-trained Foundation Transformer with Low-Rank Estimation of Binarization Residual Polynomials	Xingrun Xing; Li Du; Xinyuan Wang; Xianlin Zeng; Yequan Wang; Zheng Zhang; Jiajun Zhang
Harnessing the Power of SVD: An SVA Module for Enhanced Signal Classification	Lei Zhai; Shuyuan Yang; Yitong Li; Zhixi Feng; Zhihao Chang; Quanwei Gao
Multi-Modal Latent Space Learning for Chain-of-Thought Reasoning in Language Models	Liqi He; Zuchao Li; Xiantao Cai; Ping Wang
Learning Dense Correspondence for NeRF-Based Face Reenactment	Songlin Yang; Wei Wang; Yushi Lan; Xiangyu Fan; Bo Peng; Lei Yang; Jing Dong
Mixed Geometry Message and Trainable Convolutional Attention Network for Knowledge Graph Completion	Bin Shang; Yinliang Zhao; Jun Liu; Di Wang
Multi-Domain Recommendation to Attract Users via Domain Preference Modeling	Hyunjun Ju; SeongKu Kang; Dongha Lee; Junyoung Hwang; Sanghwan Jang; Hwanjo Yu
Scaling Up Semi-supervised Learning with Unconstrained Unlabelled Data	Shuvendu Roy; Ali Etemad
Partial Multi-View Clustering via Self-Supervised Network	Qianqian Wang; Guoshuai Sheng; Quanxue Gao; Wei Feng; Zhiqiang Tao; Bo Dong
LLMRG: Improving Recommendations through Large Language Model Reasoning Graphs	Yan Wang; Zhixuan Chu; Xin Ouyang; Simeng Wang; Hongyan Hao; Yue Shen; Jinjie Gu; Siqiao Xue; James Zhang; Qing Cui; Longfei Li; Jun Zhou;



Selective Focus: Investigating Semantics Sensitivity in Post-training Quantization for Lane Detection	Yunqian Fan; Xiuying Wei; Ruihao Gong; Yuqing Ma; Xiangguo Zhang; Qi Zhang; Xianglong Liu
INFORMEDQX: Informed Conflict Detection for Over-Constrained Problems	Viet-Man Le; Alexander Felfernig; Thi Ngoc Trang Tran; Mathias Uta
TR-DETR: Task-Reciprocal Transformer for Joint Moment Retrieval and Highlight Detection	Hao Sun; Mingyao Zhou; Wenjing Chen; Wei Xie
MAPTree: Beating “Optimal” Decision Trees with Bayesian Decision Trees	Colin Sullivan; Mo Tiwari; Sebastian Thrun
Efficient Target Propagation by Deriving Analytical Solution	Yanhao Bao; Tatsukichi Shibuya; Ikuro Sato; Rei Kawakami; Nakamasa Inoue
Scalable Motion Style Transfer with Constrained Diffusion Generation	Wenjie Yin; Yi Yu; Hang Yin; Danica Kragic; Mårten Björkman
Debiasing Multimodal Sarcasm Detection with Contrastive Learning	Mengzhao Jia; Can Xie; Liqiang Jing
Large Language Models Are Clinical Reasoners: Reasoning-Aware Diagnosis Framework with Prompt-Generated Rationales	Taeyoon Kwon; Kai Ong; Dongjin Kang; Seungjun Moon; Jeong Ryong Lee; Dosik Hwang; Beomseok Sohn; Yongsik Sim; Dongha Lee; Jinyoung Yeo
Seed-Guided Fine-Grained Entity Typing in Science and Engineering Domains	Yu Zhang; Yunyi Zhang; Yanzhen Shen; Yu Deng; Lucian Popa; Larisa Shwartz; ChengXiang Zhai; Jiawei Han
HGPrompt: Bridging Homogeneous and Heterogeneous Graphs for Few-Shot Prompt Learning	Xingtong Yu; Yuan Fang; Zemin Liu; Xinming Zhang
Collaborative Consortium of Foundation Models for Open-World Few-Shot Learning	Shuai Shao; Yu Bai; Yan Wang; Baodi Liu; Bin Liu
Divergence-Guided Simultaneous Speech Translation	Xinjie Chen; Kai Fan; Wei Luo; Linlin Zhang; Libo Zhao; Xingao Liu; Zhongqiang Huang
ConVQG: Contrastive Visual Question Generation with Multimodal Guidance	Li Mi; Syrielle Montariol; Javiera Castillo Navarro; Xianjie Dai; Antoine Bosselut; Devis Tuia
Confusing Pair Correction Based on Category Prototype for Domain Adaptation under Noisy Environments	Churan Zhi; Junbao Zhuo; Shuhui Wang
A Novel Energy Based Model Mechanism for Multi-Modal Aspect-Based Sentiment Analysis	Tianshuo Peng; Zuchao Li; Ping Wang; Lefei Zhang; Hai Zhao
BLIVA: A Simple Multimodal LLM for Better Handling of Text-Rich Visual Questions	Wenbo Hu; Yifan Xu; Yi Li; Weiyue Li; Zeyuan Chen; Zhuowen Tu
Towards Inductive Robustness: Distilling and Fostering Wave-Induced Resonance in Transductive GCNs against Graph Adversarial Attacks	Ao Liu; Wenshan Li; Tao Li; Beibei Li; Hanyuan Huang; Pan Zhou
Conditional Variational Autoencoder for Sign Language Translation with Cross-Modal Alignment	Rui Zhao; Liang Zhang; Biao Fu; Cong Hu; Jinsong Su; Yidong Chen
Evaluate Geometry of Radiance Fields with Low-Frequency Color Prior	Qihang Fang; Yafei Song; Keqiang Li; Li Shen; Huaiyu Wu; Gang Xiong; Liefeng Bo
Learning Time Slot Preferences via Mobility Tree for Next POI	Tianhao Huang; Xuan Pan; Xiangrui Cai; Ying Zhang; Xiaojie Yuan
Amodal Scene Analysis via Holistic Occlusion Relation Inference and Generative Mask Completion	Bowen Zhang; Qing Liu; Jianming Zhang; Yilin Wang; Liyang Liu; Zhe Lin; Yifan Liu
MetaDiff: Meta-Learning with Conditional Diffusion for Few-Shot Learning	Baoquan Zhang; Chuyao Luo; Demin Yu; Xutao Li; Huiwei Lin; Yunming Ye; Bowen Zhang
Using Artificial Populations to Study Psychological Phenomena in Neural	Jesse Roberts; Kyle Moore; Drew Wilenzick; Douglas Fisher
Optimizing ADMM and Over-Relaxed ADMM Parameters for Linear Quadratic Problems	Song Jintao; Wenqi Lu; Yunwen Lei; Yuchao Tang; Zhenkuan Pan; Jinming Duan



Resource Democratization: Is Compute the Binding Constraint on AI	Rebecca Gelles; Veronica Kinoshita; Micah Musser; James Dunham
Amalgamating Multi-Task Models with Heterogeneous Architectures	Jidapa Thadarassiri; Walter Gerych; Xiangnan Kong; Elke Rundensteiner
Imitation of Life: A Search Engine for Biologically Inspired Design	Hen Emuna; Nadav Borenstein; Xin Qian; Hyeonsu Kang; Joel Chan; Aniket Kittur; Dafna Shahaf
Generating Universal Adversarial Perturbations for Quantum Classifiers	Gautham Anil; Vishnu Vinod; Apurva Narayan
A Separation and Alignment Framework for Black-Box Domain Adaptation	Mingxuan Xia; Junbo Zhao; Gengyu Lyu; Zenan Huang; Tianlei Hu; Gang Chen; Haobo Wang
Double-Bounded Optimal Transport for Advanced Clustering and Classification	Liangliang Shi; Zhaoqi Shen; Junchi Yan
Memory Asymmetry Creates Heteroclinic Orbits to Nash Equilibrium in Learning in Zero-Sum Games	Yuma Fujimoto; Kaito Ariu; Kenshi Abe
Efficient Look-Up Table from Expanded Convolutional Network for Accelerating Image Super-resolution	Kai Yin; Jie Shen
Quantum Interference Model for Semantic Biases of Glosses in Word Sense Disambiguation	Junwei Zhang; Ruifang He; Fengyu Guo; Chang Liu
Instance-Conditional Timescales of Decay for Nonstationary Learning	Nishant Jain; Pradeep Shenoy
Low-Distortion Clustering with Ordinal and Limited Cardinal Information	Jakob Burkhardt; Ioannis Caragiannis; Karl Fehrs; Matteo Russo; Chris Schwiegelshohn; Sudarshan Shyam
KPA-Tracker: Towards Robust and Real-Time Category-Level Articulated Object 6D Pose Tracking	Liu Liu; Anran Huang; Qi Wu; Dan Guo; Xun Yang; Meng Wang
Integer Is Enough: When Vertical Federated Learning Meets Rounding	Pengyu Qiu; Yuwen Pu; Yongchao Liu; Wenyan Liu; Yun Yue; Xiaowei Zhu; Lichun Li; Jinbao Li; Shouling Ji
Towards the Disappearing Truth: Fine-Grained Joint Causal Influences Learning with Hidden Variable-Driven Causal Hypergraphs	Kun Zhu; Chunhui Zhao
PICNN: A Pathway towards Interpretable Convolutional Neural Networks	Wengang Guo; Jiayi Yang; HuiLin YIN; Qijun Chen; Wei Ye
Object Attribute Matters in Visual Question Answering	Peize Li; Qingyi Si; Peng Fu; Zheng Lin; Yan Wang
BadRL: Sparse Targeted Backdoor Attack against Reinforcement Learning	Jing Cui; Yufei Han; Yuzhe Ma; Jianbin Jiao; Junge Zhang
Task-Disruptive Background Suppression for Few-Shot Segmentation	Suho Park; SuBeen Lee; Sangeek Hyun; Hyun Seok Seong; Jae-Pil Heo
Your Career Path Matters in Person-Job Fit	Zhuocheng Gong; Yang Song; Tao Zhang; Ji-Rong Wen; Dongyan Zhao; Rui Yan
Graph Neural Networks with Soft Association between Topology and Attribute	Yachao Yang; Yanfeng Sun; Shaofan Wang; Jipeng Guo; Junbin Gao; Fujiao Ju; Baocai Yin
Omni-Kernel Network for Image Restoration	Yuning Cui; Wenqi Ren; Alois Knoll
Decoupled Contrastive Multi-View Clustering with High-Order Random Walks	Yiding Lu; Yijie Lin; Mouxing Yang; Dezhong Peng; Peng Hu; Xi Peng
Unravelling Expressive Delegations: Complexity and Normative Analysis	Giannis Tyrovolas; Andrei Constantinescu; Edith Elkind
AdapterGNN: Parameter-Efficient Fine-Tuning Improves Generalization in GNNs	Shengrui Li; Xueting Han; Jing Bai
ContraNovo: A Contrastive Learning Approach to Enhance De Novo Peptide Sequencing	Zhi Jin; Sheng Xu; Xiang Zhang; Tianze Ling; Nanqing Dong; Wanli Ouyang; ZhiQiang Gao; Cheng Chang; Siqi Sun
Deletion-Robust Submodular Maximization with Knapsack Constraints	Shuang Cui; Kai Han; He Huang
D3: A Methodological Exploration of Domain Division, Modeling, and Balance in Multi-Domain Recommendations	Pengyue Jia; Yichao Wang; Shanru LIN; Xiaopeng Li; Xiangyu Zhao; Huifeng Guo; Ruiming Tang





Big Learning Expectation Maximization	Yulai Cong; Sijia Li
Offline and Online Optical Flow Enhancement for Deep Video Compression	Chuanbo Tang; Xihua Sheng; Zhuoyuan Li; Haotian Zhang; Li Li; Dong Liu
Non-exemplar Online Class-Incremental Continual Learning via Dual-Prototype Self-Augment and Refinement	Fushuo Huo; Wenchao Xu; Jingcai Guo; Haozhao Wang; Yunfeng Fan
United We Stand: Accelerating Privacy-Preserving Neural Inference by Conjunctive Optimization with Interleaved Nexus	Qiao Zhang; Tao Xiang; Chunsheng Xin; Hongyi Wu
Risk-Conditioned Reinforcement Learning: A Generalized Approach for Adapting to Varying Risk Measures	Gwangpyo Yoo; Jinwoo Park; Honguk Woo
DINGO: Towards Diverse and Fine-Grained Instruction-Following Evaluation	Zihui Gu; Xingwu Sun; Fengzong Lian; Zhanhui Kang; Chengzhong Xu; Ju Fan
MID-FiLD: MIDI Dataset for Fine-Level Dynamics	Jesung Ryu; Seungyeon Rhyu; Hong-Gyu Yoon; Eunchong Kim; Ju Young Yang; Taehyun Kim
A Positive-Unlabeled Metric Learning Framework for Document-Level Relation Extraction with Incomplete Labeling	Ye Wang; Huazheng Pan; Tao Zhang; Wen Wu; Wenxin Hu
SkeletonGait: Gait Recognition Using Skeleton Maps	Chao Fan; Jingzhe Ma; Dongyang Jin; Chuanfu Shen; Shiqi Yu
Learning to Rank in Generative Retrieval	Yongqi Li; Nan Yang; Liang Wang; Furu Wei; Wenjie Li
Autoregressive Omni-Aware Outpainting for Open-Vocabulary 360-Degree Image Generation	Zhuqiang Lu; Kun Hu; Chaoyue Wang; Lei Bai; Zhiyong Wang
Mitigating Label Bias in Machine Learning: Fairness through Confident	Yixuan Zhang; Boyu Li; Zenan Ling; Feng Zhou
Video Frame Prediction from a Single Image and Events	Juanjuan Zhu; Zhexiong Wan; Yuchao Dai
Learning to Learn Better Visual Prompts	Fengxiang Wang; Wanrong Huang; Shaowu Yang; Qi Fan; Long Lan
CEGAR-Based Approach for Solving Combinatorial Optimization Modulo Quantified Linear Arithmetics Problems	Kerian Thuillier; Anne Siegel; Loïc Paulevé
Maximizing the Success Probability of Policy Allocations in Online Systems	Artem Betlei; Mariia Vladimirova; Mehdi Sebbar; Nicolas Urien; Thibaud Rahier; Benjamin Heymann
Investigating the Effectiveness of Task-Agnostic Prefix Prompt for Instruction Following	Seonghyeon Ye; Hyeonbin Hwang; Sohee Yang; Hyeongu Yun; Yireun Kim; Minjoon Seo
Revisiting Document-Level Relation Extraction with Context-Guided Link Prediction	Monika Jain; Raghava Mutharaju; Ramakanth Kavuluru; Kuldeep Singh
Learning Discrete-Time Major-Minor Mean Field Games	Kai Cui; Gökçe Dayanıklı; Mathieu Laurière; Matthieu Geist; Olivier Pietquin; Heinz Koepl
Omnipotent Distillation with LLMs for Weakly-Supervised Natural Language Video Localization: When Divergence Meets Consistency	Peijun Bao; Zihao Shao; Wenhan Yang; Boon Poh Ng; Meng Hwa Er; Alex Kot
Operationalizing Essential Characteristics of Creativity in a Computational System for Music Composition	Paul Bodily; Dan Ventura
Fine-Grained Multi-View Hand Reconstruction Using Inverse Rendering	Qijun Gan; Wentong Li; Jinwei Ren; Jianke Zhu
Unsupervised Pan-Sharpening via Mutually Guided Detail Restoration	Huangxing Lin; Yuhang Dong; Xinghao Ding; Tianpeng Liu; Yongxiang Liu
Rethinking Causal Relationships Learning in Graph Neural Networks	Hang Gao; Chengyu Yao; Jiangmeng Li; Lingyu Si; Yifan Jin; Fengge Wu; Changwen Zheng; Huaping Liu
Curriculum-Enhanced Residual Soft An-Isotropic Normalization for Over-Smoothness in Deep GNNs	Jin Li; Qirong Zhang; Shuling Xu; Xinlong Chen; Longkun Guo; Yang-Geng Fu
Friendly Attacks to Improve Channel Coding Reliability	Anastasiia Kurmukova; Deniz Gunduz



Learning Not to Regret	David Sychrovský; Michal Sustr; Elnaz Davoodi; Michael Bowling; Marc Lanctot; Martin Schmid
MultiSum: A Multi-Facet Approach for Extractive Social Summarization Utilizing Semantic and Sociological Relationships	Tanglong Zhao; Ruifang He; Jing Xu; Bo Wang
Decentralized Monte Carlo Tree Search for Partially Observable Multi-Agent Pathfinding	Alexey Skrynnik; Anton Andreychuk; Konstantin Yakovlev; Aleksandr Panov
Generalizable Task Representation Learning for Offline Meta-Reinforcement Learning with Data Limitations	Renzhe Zhou; Chen-Xiao Gao; Zongzhang Zhang; Yang Yu
Towards Learning and Explaining Indirect Causal Effects in Neural Networks	Abbavaram Gowtham Reddy; Saketh Bachu; Harsharaj Pathak; Benin Godfrey L; Varshaneya V; Vineeth N Balasubramanian; Satyanarayan Kar
Relevant Intrinsic Feature Enhancement Network for Few-Shot Semantic Segmentation	Xiaoyi Bao; Jie Qin; Siyang Sun; Xingang Wang; Yun Zheng
Generating Novel Leads for Drug Discovery Using LLMs with Logical Feedback	Shreyas Bhat Brahmavar; Ashwin Srinivasan; Tirtharaj Dash; Sowmya Ramaswamy Krishnan; Lovekesh Vig; Arijit Roy; Raviprasad Aduri
Exact Policy Recovery in Offline RL with Both Heavy-Tailed Rewards and Data Corruption	Yiding Chen; Xuezhou Zhang; Qiaomin Xie; Xiaojin Zhu
Self-Distillation Regularized Connectionist Temporal Classification Loss for Text Recognition: A Simple Yet Effective Approach	Ziyin Zhang; Ning Lu; Minghui Liao; Yongshuai Huang; Cheng Li; Min Wang; Wei Peng
Limited-Supervised Multi-Label Learning with Dependency Noise	Yejiang Wang; Yuhai Zhao; Zhengkui Wang; Wen Shan; Xingwei Wang
Altruism in Facility Location Problems	Hau Chan; Minming Li; Houyu Zhou
Deep Semantic Graph Transformer for Multi-View 3D Human Pose Estimation	Lijun Zhang; Kangkang Zhou; Feng Lu; Xiang-Dong Zhou; Yu Shi
Task-Adaptive Prompted Transformer for Cross-Domain Few-Shot Learning	Jiamin Wu; Xin Liu; Xiaotian Yin; Tianzhu Zhang; Yongdong Zhang
Is a Large Language Model a Good Annotator for Event Extraction?	Ruirui Chen; Chengwei Qin; Weifeng Jiang; Dongkyu Choi
Fed-QSSL: A Framework for Personalized Federated Learning under Bitwidth and Data Heterogeneity	Yiyue Chen; Haris Vikalo; Chianing Wang
Dual-View Whitening on Pre-trained Text Embeddings for Sequential Recommendation	Lingzi Zhang; Xin Zhou; Zhiwei Zeng; Zhiqi Shen
Feature Fusion from Head to Tail for Long-Tailed Visual Recognition	Mengke Li; Zhikai HU; Yang Lu; Weichao Lan; Yiu-ming Cheung; Hui Huang
Adaptive Shortcut Debiasing for Online Continual Learning	Doyoung Kim; Dongmin Park; Yooju Shin; Jihwan Bang; Hwanjun Song; Jae-Gil Lee
Clarifying the Behavior and the Difficulty of Adversarial Training	Xu Cheng; Hao Zhang; Yue Xin; Wen Shen; Quanshi Zhang
Complexity of Credulous and Skeptical Acceptance in Epistemic Argumentation Framework	Gianvincenzo Alfano; Sergio Greco; Francesco Parisi; Irina Trubitsyna
Improve Robustness of Reinforcement Learning against Observation Perturbations via $L_\infty$ Lipschitz Policy Networks	Buqing Nie; Jingtian Ji; Yangqing Fu; Yue Gao
PathAsst: A Generative Foundation AI Assistant towards Artificial General Intelligence of Pathology	Yuxuan Sun; Chenglu Zhu; Sunyi Zheng; Kai Zhang; Lin Sun; Zhongyi Shui; Yunlong Zhang; Honglin Li; Lin Yang
DeepBranchTracer: A Generally-Applicable Approach to Curvilinear Structure Reconstruction Using Multi-Feature Learning	Chao Liu; Ting Zhao; Nenggan Zheng
Brush Your Text: Synthesize Any Scene Text on Images via Diffusion Model	Lingjun Zhang; Xinyuan Chen; Yaohui Wang; Yue Lu; Yu Qiao



Rolling-Unet: Revitalizing MLP's Ability to Efficiently Extract Long-Distance Dependencies for Medical Image Segmentation	Yutong Liu; Haijiang Zhu; Mengting Liu; Huaiyuan Yu; Zihan Chen; Jie Gao
Delegation-Relegation for Boolean Matrix Factorization	Florent Avellaneda; Roger Villemaire
Decoupled Training: Return of Frustratingly Easy Multi-Domain Learning	Ximei Wang; Junwei Pan; Xingzhuo Guo; Dapeng Liu; Jie Jiang
Deep Contrastive Graph Learning with Clustering-Oriented Guidance	Mulin Chen; Bocheng Wang; Xuelong Li
DOCTR: Disentangled Object-Centric Transformer for Point Scene Understanding	Xiaoxuan Yu; Hao Wang; Weiming Li; Qiang Wang; SoonYong Cho; Younghun Sung
Robust Evaluation Measures for Evaluating Social Biases in Masked Language Models	Yang Liu
Weakly Supervised Few-Shot Object Detection with DETR	Chenbo Zhang; Yinglu Zhang; Lu Zhang; Jiajia Zhao; Jihong Guan; Shuigeng Zhou
TopoGCL: Topological Graph Contrastive Learning	Yuzhou Chen; Jose Frias; Yulia Gel
Inference and Learning in Dynamic Decision Networks Using Knowledge Compilation	Gabriele Venturato; Vincent Derkinderen; Pedro Zuidberg Dos Martires; Luc De Raedt
Learning Small Decision Trees for Data of Low Rank-Width	Konrad K. Dabrowski; Eduard Eiben; Sebastian Ordyniak; Giacomo Paesani; Stefan Szeider
Causality-Inspired Invariant Representation Learning for Text-Based Person Retrieval	Yu Liu; Guihe Qin; Haipeng Chen; Zhiyong Cheng; Xun Yang
RedCore: Relative Advantage Aware Cross-Modal Representation Learning for Missing Modalities with Imbalanced Missing Rates	Jun Sun; Xinxin Zhang; Shoukang Han; Yu-Ping Ruan; Taihao Li
Learning Optimal Advantage from Preferences and Mistaking It for Reward	W Bradley Knox; Stephane Hatgis-Kessell; Sigurdur Orn Adalgeirsson; Serena Booth; Anca Dragan; Peter Stone; Scott Niekum
Discovering Sequential Patterns with Predictable Inter-event Delays	Joscha Cüppers; Paul Krieger; Jilles Vreeken
Ada-Retrieval: An Adaptive Multi-Round Retrieval Paradigm for Sequential Recommendations	Lei Li; Jianxun Lian; Xiao Zhou; Xing Xie
Attention Guided CAM: Visual Explanations of Vision Transformer Guided by Self-Attention	Saebom Leem; Hyunseok Seo
Weisfeiler and Lehman Go Paths: Learning Topological Features via Path Complexes	Quang Truong; Peter Chin
ViTEraser: Harnessing the Power of Vision Transformers for Scene Text Removal with SegMIM Pretraining	Dezhi Peng; Chongyu Liu; Yuliang Liu; Lianwen Jin
DTMFormer: Dynamic Token Merging for Boosting Transformer-Based Medical Image Segmentation	Zhehao Wang; Xian Lin; Nannan Wu; Li Yu; Kwang-Ting Cheng; Zengqiang Yan
FaceRSA: RSA-Aware Facial Identity Cryptography Framework	Zhongyi Zhang; Tianyi Wei; Wenbo Zhou; Hanqing Zhao; Weiming Zhang; Nenghai Yu
Language-Guided Transformer for Federated Multi-Label Classification	I-Jieh Liu; Ci-Siang Lin; Fu-En Yang; Yu-Chiang Frank Wang
iDet3D: Towards Efficient Interactive Object Detection for LiDAR Point Clouds	Dongmin Choi; Wonwoo Cho; Kangyeol Kim; Jaegul Choo
Boosting Neural Cognitive Diagnosis with Student's Affective State Modeling	Shanshan Wang; Zhen Zeng; Xun Yang; Ke Xu; Xingyi Zhang
Stitching Sub-trajectories with Conditional Diffusion Model for Goal-Conditioned Offline RL	Sungyeon Kim; Yunseon Choi; Daiki Matsunaga; Kee-Eung Kim



NegVSR: Augmenting Negatives for Generalized Noise Modeling in Real-world Video Super-Resolution	Yexing Song; Meilin Wang; Zhijing Yang; Xiaoyu Xian; Yukai Shi
Submodel Enumeration for CTL Is Hard	Nicolas Fröhlich; Arne Meier
TMPNN: High-Order Polynomial Regression Based on Taylor Map	Andrei Ivanov; Stefan Ailuro
Opponent-Model Search in Games with Incomplete Information	Junkang Li; Bruno Zanuttini; Véronique Ventos
Memory-Efficient Reversible Spiking Neural Networks	Hong Zhang; Yu Zhang
MKG-FENN: A Multimodal Knowledge Graph Fused End-to-End Neural Network for Accurate Drug–Drug Interaction Prediction	Di Wu; Wu Sun; Yi He; Zhong Chen; Xin Luo
A Variational Autoencoder for Neural Temporal Point Processes with Dynamic Latent Graphs	Sikun Yang; Hongyuan Zha
Teacher as a Lenient Expert: Teacher-Agnostic Data-Free Knowledge Distillation	Hyunjune Shin; Dong-Wan Choi
DA-Net: A Disentangled and Adaptive Network for Multi-Source Cross-Lingual Transfer Learning	Ling Ge; Chunming Hu; Guanghui Ma; Jihong Liu; Hong Zhang
Dense Projection for Anomaly Detection	Dazhi Fu; Zhao Zhang; Jicong Fan
Tensorized Label Learning on Anchor Graph	Jing Li; Quanyue Gao; Qianqian Wang; Wei Xia
Neighborhood-Enhanced 3D Human Pose Estimation with Monocular LiDAR in Long-Range Outdoor Scenes	Jingyi Zhang; Qihong Mao; Guosheng Hu; Siqi Shen; Cheng Wang
Adaptive Reactive Synthesis for LTL and LTLf Modulo Theories	Andoni Rodríguez; César Sánchez
LaViP: Language-Grounded Visual Prompting	Nilakshan Kunanantaseelan; Jing Zhang; Mehrtash Harandi
Entropy Induced Pruning Framework for Convolutional Neural Networks	Yiheng Lu; Ziyu Guan; Yaming Yang; Wei Zhao; Maoguo Gong; Cai Xu
1/2-Approximate MMS Allocation for Separable Piecewise Linear Concave Valuations	Chandra Chekuri; Pooja Kulkarni; Rucha Kulkarni; Ruta Mehta
Self-Interpretable Graph Learning with Sufficient and Necessary Explanations	Jiale Deng; Yanyan Shen
Hypothesis, Verification, and Induction: Grounding Large Language Models with Self-Driven Skill Learning	Shaohui Peng; Xing Hu; Qi Yi; Rui Zhang; Jiaming Guo; Di Huang; Zikang Tian; Ruizhi Chen; Zidong Du; Qi Guo; Yunji Chen; Ling Li
Adaptive Graph Learning for Multimodal Conversational Emotion Detection	Geng Tu; Tian Xie; Bin Liang; Hongpeng Wang; Ruifeng Xu
NondBREM: Nondeterministic Offline Reinforcement Learning for Large-Scale Order Dispatching	Hongbo Zhang; Guang Wang; Xu Wang; Zhengyang Zhou; Chen Zhang; Zheng Dong; Yang Wang
Learn the Force We Can: Enabling Sparse Motion Control in Multi-Object Video Generation	Aram Davtyan; Paolo Favaro
ConcaveQ: Non-monotonic Value Function Factorization via Concave Representations in Deep Multi-Agent Reinforcement Learning	Huiqun Li; Hanhan Zhou; Yifei Zou; Dongxiao Yu; Tian Lan
p-Laplacian Adaptation for Generative Pre-trained Vision-Language Models	Haoyuan Wu; Xinyun Zhang; Peng Xu; Peiyu Liao; Xufeng Yao; Bei Yu
Leveraging Local Variance for Pseudo-Label Selection in Semi-supervised Learning	Zeping Min; Jinfeng Bai; Chengfei Li
TOP-ReID: Multi-Spectral Object Re-identification with Token Permutation	Yuhao Wang; Xuehu Liu; Pingping Zhang; Hu Lu; Zhengzheng Tu; Huchuan
SDGMNet: Statistic-Based Dynamic Gradient Modulation for Local Descriptor Learning	Yuxin Deng; Jiayi Ma
Gradual Residuals Alignment: A Dual-Stream Framework for GAN Inversion and Image Attribute Editing	Hao Li; Mengqi Huang; Lei Zhang; Bo Hu; Yi Liu; Zhendong Mao



Optimistic Value Instructors for Cooperative Multi-Agent Reinforcement Learning	Chao Li; Yupeng Zhang; Jianqi Wang; Yujing Hu; Shaokang Dong; Wenbin Li; Tangjie Lv; Changjie Fan; Yang Gao
Hierarchize Pareto Dominance in Multi-Objective Stochastic Linear Bandits	Ji Cheng; Bo Xue; Jiaxiang Yi; Qingfu Zhang
Data-Free Hard-Label Robustness Stealing Attack	Xiaojian Yuan; Kejiang Chen; Wen Huang; Jie Zhang; Weiming Zhang; Nenghai Yu
Inertial Algorithm with Dry Fraction and Convolutional Sparse Coding for 3D Localization with Light Field Microscopy	Xiaofan Wang; Zhiyuan Deng; Changle Wang; Jinjia Wang
Segment beyond View: Handling Partially Missing Modality for Audio-Visual Semantic Segmentation	Renjie Wu; Hu Wang; Feras Dayoub; Hsiang-Ting Chen
RG-GAN: Dynamic Regenerative Pruning for Data-Efficient Generative Adversarial Networks	Divya Saxena; Jiannong Cao; Jiahao Xu; Tarun Kulshrestha
Explore 3D Dance Generation via Reward Model from Automatically-Ranked Demonstrations	Zilin Wang; Haolin Zhuang; Lu Li; Yinmin Zhang; Junjie Zhong; Jun Chen; Yu Yang; Boshi Tang; Zhiyong Wu
On the Outcome Equivalence of Extensive-Form and Behavioral Correlated Equilibria	Brian Zhang; Tuomas Sandholm
AltDiffusion: A Multilingual Text-to-Image Diffusion Model	Fulong Ye; Guang Liu; Xinya Wu; Ledell Wu
Automated Design of Affine Maximizer Mechanisms in Dynamic Settings	Michael Curry; Vinzenz Thoma; Darshan Chakrabarti; Stephen McAleer; Christian Kroer; Tuomas Sandholm; Niao He; Sven Seuken
Computing Nash Equilibria in Potential Games with Private Uncoupled Constraints	Nikolas Patris; Stelios Stavroulakis; Fivos Kalogiannis; Rose Zhang; Ioannis Panageas
Dependency Structure-Enhanced Graph Attention Networks for Event SGFormer: Semantic Graph Transformer for Point Cloud-Based 3D Scene Graph Generation	Qizhi Wan; Changxuan wan; Keli Xiao; Kun Lu; Chenliang Li; Xiping Liu; Dexi Changsheng Lv; Mengshi Qi; Xia Li; Zhengyuan Yang; Huadong Ma
HISR: Hybrid Implicit Surface Representation for Photorealistic 3D Human Reconstruction	Angtian Wang; Yuanlu Xu; Nikolaos Sarafianos; Robert Maier; Edmond Boyer; Alan Yuille; Tony Tung
Unifying Decision and Function Queries in Stochastic Boolean Satisfiability	Yu-Wei Fan; Jie-Hong Jiang
Self-Supervised Bird's Eye View Motion Prediction with Cross-Modality Signals	Shaoheng Fang; Zuhong Liu; Mingyu Wang; Chenxin Xu; Yiqi Zhong; Siheng Chen
Count What You Want: Exemplar Identification and Few-Shot Counting of Human Actions in the Wild	Yifeng Huang; Duc Nguyen; Nguyễn Lâm; Cuong Pham; Minh Hoai
Symmetric Q-learning: Reducing Skewness of Bellman Error in Online Reinforcement Learning	Motoki Omura; Takayuki Osa; Yusuke Mukuta; Tatsuya Harada
Enhancing Training of Spiking Neural Network with Stochastic Latency	Srinivas Anumasa; Bhaskar Mukhoty; Velibor Bojkovic; Giulia De Masi; Huan Xiong; Bin Gu
Grab What You Need: Rethinking Complex Table Structure Recognition with Flexible Components Deliberation	Hao Liu; Xin Li; Mingming Gong; Bing Liu; Yunfei Wu; Deqiang Jiang; Yinsong Liu; Xing Sun
LLMEval: A Preliminary Study on How to Evaluate Large Language Models	Yue Zhang; Ming Zhang; HaiPeng Yuan; Shichun Liu; Yongyao Shi; Tao Gui; Qi Zhang; Xuanjing Huang
SECap: Speech Emotion Captioning with Large Language Model	Yaoxun Xu; Hangting Chen; Jianwei Yu; Qiaochu Huang; Zhiyong Wu; Shi-Xiong Zhang; Guangzhi Li; Yi Luo; Rongzhi Gu
SeqRank: Sequential Ranking of Salient Objects	Huankang Guan; Rynson W.H. Lau



Tree Search-Based Evolutionary Bandits for Protein Sequence Optimization	Jiahao Qiu; Hui Yuan; Jinghong Zhang; Wentao Chen; Huazheng Wang; Mengdi Wang
GCNext: Towards the Unity of Graph Convolutions for Human Motion	Xinshun Wang; Qiongjie Cui; Chen Chen; Mengyuan Liu
Bilateral Gradual Semantics for Weighted Argumentation	Zongshun Wang; Yuping Shen
Robust Test-Time Adaptation for Zero-Shot Prompt Tuning	Ding-Chu Zhang; Zhi Zhou; Yu-Feng Li
What Makes Quantization for Large Language Model Hard? An Empirical Study from the Lens of Perturbation	Zhuocheng Gong; Jiahao Liu; Jingang Wang; Xunliang Cai; Dongyan Zhao; Rui Yan
Uncertainty-Aware Yield Prediction with Multimodal Molecular Features	Jiayuan Chen; Kehan Guo; Zhen Liu; Olexandr Isayev; Xiangliang Zhang
Preference Aware Dual Contrastive Learning for Item Cold-Start Recommendation	Wenbo Wang; Bingquan Liu; Lili Shan; Chengjie Sun; Ben Chen; Jian Guan
Learning Performance Maximizing Ensembles with Explainability Guarantees	Vincent Pisztor; Jia Li
Communication Efficient Distributed Newton Method over Unreliable Networks	Ming Wen; Chengchang Liu; Yuedong Xu
How to Evaluate the Generalization of Detection? A Benchmark for Comprehensive Open-Vocabulary Detection	Yiyang Yao; Peng Liu; Tiancheng Zhao; Qianqian Zhang; Jiajia Liao; Chunxin Fang; Kyusong Lee; Qing Wang
Axiomatic Aggregations of Abductive Explanations	Gagan Biradar; Yacine Izza; Elita Lobo; Vignesh Viswanathan; Yair Zick
Continual Vision-Language Retrieval via Dynamic Knowledge Rectification	Zhenyu Cui; Yuxin Peng; Xun Wang; Manyu Zhu; Jiahuan Zhou
KAM-CoT: Knowledge Augmented Multimodal Chain-of-Thoughts Reasoning	Debjyoti Mondal; Suraj Modi; Subhadarshi Panda; Rituraj Singh; Godawari Sudhakar Rao
MuST: Robust Image Watermarking for Multi-Source Tracing	Guanjie Wang; Zehua Ma; Chang Liu; Xi Yang; Han Fang; Weiming Zhang; Nenghai Yu
Efficient Conditional Diffusion Model with Probability Flow Sampling for Image Super-resolution	Yutao Yuan; Chun Yuan
Inconsistency-Based Data-Centric Active Open-Set Annotation	Ruiyu Mao; Ouyang Xu; Yunhui Guo
Hierarchical Planning and Learning for Robots in Stochastic Settings Using Zero-Shot Option Invention	Naman Shah; Siddharth Srivastava
Optimizing the Optimization of Planning Domains by Automatic Action Schema Splitting	Mojtaba Elahi; Jussi Rintanen
Frequency-Aware Deepfake Detection: Improving Generalizability through Frequency Space Domain Learning	Chuangchuang Tan; Yao Zhao; Shikui Wei; Guanghua Gu; Ping Liu; Yunchao Wei
Learning Persistent Community Structures in Dynamic Networks via Topological Data Analysis	Dexu Kong; Anping Zhang; Yang Li
PMRC: Prompt-Based Machine Reading Comprehension for Few-Shot Named Entity Recognition	Jin Huang; Danfeng Yan; Yuanqiang Cai
Efficient Learning in Polyhedral Games via Best-Response Oracles	Darshan Chakrabarti; Gabriele Farina; Christian Kroer
SlowTrack: Increasing the Latency of Camera-Based Perception in Autonomous Driving Using Adversarial Examples	Chen Ma; Ningfei Wang; Qi Alfred Chen; Chao Shen
Composite Active Learning: Towards Multi-Domain Active Learning with Theoretical Guarantees	Guang-Yuan Hao; Hengguan Huang; Haotian Wang; Jie Gao; Hao Wang
Fully Data-Driven Pseudo Label Estimation for Pointly-Supervised Panoptic Segmentation	Jing Li; Junsong Fan; Yuran Yang; Shuqi Mei; Jun Xiao; Zhaoxiang Zhang





From Toxic to Trustworthy: Using Self-Distillation and Semi-supervised Methods to Refine Neural Networks	Xianda Zhang; Baolin Zheng; Jianbao Hu; Chengyang Li; Xiaoying Bai
MINES: Message Intercommunication for Inductive Relation Reasoning over Neighbor-Enhanced Subgraphs	Ke Liang; Lingyuan Meng; Sihang Zhou; Wenxuan Tu; Siwei Wang; Yue Liu; Meng Liu; Long Zhao; Xiangjun Dong; Xinwang Liu
Evolving Parameterized Prompt Memory for Continual Learning	Muhammad Rifki Kurniawan; Xiang Song; Zhiheng Ma; Yuhang He; Yihong Gong; Yang Qi; Xing Wei
ACAMDA: Improving Data Efficiency in Reinforcement Learning through Guided Counterfactual Data Augmentation	Yuewen Sun; Erli Wang; Biwei Huang; Chaochao Lu; Lu Feng; Changyin Sun; Kun Zhang
Towards Safe Policy Learning under Partial Identifiability: A Causal Approach	Shalmali Joshi; Junzhe Zhang; Elias Bareinboim
Composing Biases by Using CP to Decompose Minimal Functional Dependencies for Acquiring Complex Formulae	Ramiz Gindullin; Nicolas Beldiceanu; Jovial Cheukam Ngouonou; Remi Douence; Claude-Guy Quimper
SocialCVAE: Predicting Pedestrian Trajectory via Interaction Conditioned Latents	Wei Xiang; Haoteng YIN; He Wang; Xiaogang Jin
Spectral-Based Graph Neural Networks for Complementary Item Recommendation	Haitong Luo; Xuying Meng; Suhang Wang; Hanyun Cao; Weiyao Zhang; Yequan Wang; Yujun Zhang
Double Auction on Diffusion Network	Miao Li; Yuhang Cao; Dengji Zhao
Understanding and Leveraging the Learning Phases of Neural Networks	Johannes Schneider; Mohit Prabhushankar
Computing the Why-Provenance for Datalog Queries via SAT Solvers	Marco Calautti; Ester Livshits; Andreas Pieris; Markus Schneider
FairTrade: Achieving Pareto-Optimal Trade-Offs between Balanced Accuracy and Fairness in Federated Learning	Maryam Badar; Sandipan Sikdar; Wolfgang Nejdl; Marco Fisichella
Deep Incomplete Multi-View Learning Network with Insufficient Label	Zhangqi Jiang; Tingjin Luo; Xinyan Liang
WeditGAN: Few-Shot Image Generation via Latent Space Relocation	Yuxuan Duan; Li Niu; Yan Hong; Liqing Zhang
Fact-Driven Logical Reasoning for Machine Reading Comprehension	Siru Ouyang; Zhuosheng Zhang; Hai Zhao
PREFER: Prompt Ensemble Learning via Feedback-Reflect-Refine	Chenrui Zhang; Lin Liu; Chuyuan Wang; Xiao Sun; Hongyu Wang; Jinpeng Wang; Mingchen Cai
An Attentive Inductive Bias for Sequential Recommendation beyond the Self-Attention	Yehjin Shin; Jeongwhan Choi; Hyowon Wi; Noseong Park
Improving Audio-Visual Segmentation with Bidirectional Generation	Dawei Hao; Yuxin Mao; Bowen He; Xiaodong Han; Yuchao Dai; Yiran Zhong
EarnHFT: Efficient Hierarchical Reinforcement Learning for High Frequency Trading	Molei Qin; Shuo Sun; Wentao Zhang; Haochong Xia; Xinrun Wang; Bo An
I-rebalance: Personalized Vehicle Repositioning for Supply Demand Balance	Haoyang Chen; Peiyan Sun; Qiyuan Song; Wanyuan Wang; Weiwei Wu; Wencan Zhang; Guanyu Gao; Yan Lyu
FaceCoresetNet: Differentiable Coresets for Face Set Recognition	Gil Shapira; Yosi Keller
Towards Multi-Mode Outlier Robust Tensor Ring Decomposition	Yuning Qiu; Guoxu Zhou; Andong Wang; Zhenhao Huang; Qibin Zhao
SPGroup3D: Superpoint Grouping Network for Indoor 3D Object Detection	Yun Zhu; Le Hui; Yaqi Shen; Jin Xie
Foreseeing Reconstruction Quality of Gradient Inversion: An Optimization Perspective	Hyeong Gwon Hong; Yooshin Cho; Hanbyel Cho; Jaesung Ahn; Junmo Kim
PNeSM: Arbitrary 3D Scene Stylization via Prompt-Based Neural Style Mapping	Jiafu Chen; Wei Xing; Jiakai Sun; Tianyi Chu; Yiling Huang; Boyan Ji; Lei Zhao; Huaizhong Lin; Haibo Chen; Zhizhong Wang
A Brain-Inspired Way of Reducing the Network Complexity via Concept-Regularized Coding for Emotion Recognition	Han Lu; Xiahai Zhuang; Qiang Luo



Coreference Graph Guidance for Mind-Map Generation	Zhuowei Zhang; Mengting Hu; Yinhao Bai; Zhen Zhang
CMDA: Cross-Modal and Domain Adversarial Adaptation for LiDAR-Based 3D Object Detection	Gyusam Chang; Wonseok Roh; Sujin Jang; Dongwook Lee; Daehyun Ji; Gyeongrok Oh; Jinsun Park; Jinkyu Kim; Sangpil Kim
Uncertainty Quantification in Heterogeneous Treatment Effect Estimation with Gaussian-Process-Based Partially Linear Model	Shunsuke Horii; Yoichi Chikahara
Generalized Variational Inference via Optimal Transport	Jinjin Chi; Zhichao Zhang; Zhiyao Yang; Jihong Ouyang; Hongbin Pei
Cumulative Difference Learning VAE for Time-Series with Temporally Correlated Inflow-Outflow	Tianchun Li; Chengxiang Wu; Pengyi Shi; Xiaoqian Wang
BCLNet: Bilateral Consensus Learning for Two-View Correspondence Pruning	Xiangyang Miao; Guobao Xiao; Shiping Wang; Jun Yu
Factorized Explainer for Graph Neural Networks	Rundong Huang; Farhad Shirani; Dongsheng Luo
Frame Semantic Role Labeling Using Arbitrary-Order Conditional Random Fields	Chaoyi Ai; Kewei Tu
TransGOP: Transformer-Based Gaze Object Prediction	Binglu Wang; Chenxi Guo; Yang Jin; Haisheng Xia; Nian Liu
Spatio-Temporal Pivotal Graph Neural Networks for Traffic Flow Forecasting	Weiyang Kong; Ziyu Guo; Yubao Liu
SOGDet: Semantic-Occupancy Guided Multi-View 3D Object Detection	Qiu Zhou; Jinming Cao; Hanchao Leng; Yifang Yin; Yu Kun; Roger
Exploiting Polarized Material Cues for Robust Car Detection	Wen Dong; Haiyang Mei; Ziqi Wei; Ao Jin; Sen Qiu; Qiang Zhang; Xin Yang
An Efficient Subgraph-Infering Framework for Large-Scale Heterogeneous Graphs	Wei Zhou; Hong Huang; Ruize Shi; Kehan Yin; Hai Jin
Signed Graph Neural Ordinary Differential Equation for Modeling Continuous-Time Dynamics	Lanlan Chen; Kai Wu; Jian Lou; Jing Liu
BBScore: A Brownian Bridge Based Metric for Assessing Text Coherence	Zhecheng Sheng; Tianhao Zhang; Chen Jiang; Dongyeop Kang
Empowering Dual-Level Graph Self-Supervised Pretraining with Motif Discovery	Pengwei Yan; Kaisong Song; Zhuoren Jiang; Yangyang Kang; Tianqianjin Lin; Changlong Sun; Xiaozhong Liu
Rectangle Search: An Anytime Beam Search	Sofia Lemons; Wheeler Ruml; Rob Holte; Carlos Linares Lopez
Exploring Domain Incremental Video Highlights Detection with the LiveFood Benchmark	Sen Pei; Shixiong Xu; Xiaojie Jin
Fewer Steps, Better Performance: Efficient Cross-Modal Clip Trimming for Video Moment Retrieval Using Language	Xiang Fang; Daizong Liu; Wanlong Fang; Pan Zhou; Zichuan Xu; Wenzheng Xu; Junyang Chen; Renfu Li
FAVOR: Full-Body AR-Driven Virtual Object Rearrangement Guided by Instruction Text	Kailin Li; Lixin Yang; Zenan Lin; Jian Xu; Xinyu Zhan; Yifei Zhao; Pengxiang Zhu; Wenxiong Kang; Kejian Wu; Cewu Lu
SkipDiff: Adaptive Skip Diffusion Model for High-Fidelity Perceptual Image Super-resolution	Xiaotong Luo; Yuan Xie; Yanyun Qu; Yun Fu
Self-Supervised Disentangled Representation Learning for Robust Target Speech Extraction	Zhaoxi Mu; Xinyu Yang; Sining Sun; Qing Yang
MWSIS: Multimodal Weakly Supervised Instance Segmentation with 2D Box Annotations for Autonomous Driving	Guangfeng Jiang; Jun Liu; Yuzhi Wu; Wenlong Liao; Tao He; Pai Peng
Feature Transportation Improves Graph Neural Networks	Moshe Eliasof; Eldad Haber; Eran Treister
PoseGen: Learning to Generate 3D Human Pose Dataset with NeRF	Mohsen Gholami; Rabab Ward; Z. Jane Wang
Harnessing Manycore Processors with Distributed Memory for Accelerated Training of Sparse and Recurrent Models	Jan Finkbeiner; Thomas Gmeinder; Mark Pupilli; Alexander Titterton; Emre Neftci



UFDA: Universal Federated Domain Adaptation with Practical Assumptions	Xinhui Liu; Zhenghao Chen; Luping Zhou; Dong Xu; Wei Xi; Gairui Bai; Yihan Zhao; Jizhong Zhao
How to Overcome Curse-of-Dimensionality for Out-of-Distribution Detection?	Soumya Suvra Ghosal; Yiyu Sun; Yixuan Li
ShareBERT: Embeddings Are Capable of Learning Hidden Layers	Jia Cheng Hu; Roberto Cavicchioli; Giulia Berardinelli; Alessandro Capotondi
Spatio-Temporal Fusion for Human Action Recognition via Joint Trajectory Graph	Yaolin Zheng; Hongbo Huang; Xiuying Wang; Xiaoxu Yan; Longfei Xu
Sparse Enhanced Network: An Adversarial Generation Method for Robust Augmentation in Sequential Recommendation	Junyang Chen; Guoxuan Zou; Pan Zhou; Wu Yirui; Zhenghan Chen; Houcheng Su; Huan Wang; Zhiguo Gong
Efficient Axiomatization of OWL 2 EL Ontologies from Data by Means of Formal Concept Analysis	Francesco Kriegel
Robust Node Classification on Graph Data with Graph and Label Noise	Yonghua Zhu; Lei Feng; Zhenyun Deng; Yang Chen; Robert Amor; Michael Witbrock
Sharpness-Aware Model-Agnostic Long-Tailed Domain Generalization	Houcheng Su; Weihao Luo; Daixian Liu; Mengzhu Wang; Jing Tang; Junyang Chen; Cong Wang; Zhenghan Chen
SpikingBERT: Distilling BERT to Train Spiking Language Models Using Implicit Differentiation	Malyaban Bal; Abhronil Sengupta
Temporal Correlation Vision Transformer for Video Person Re-Identification	Pengfei Wu; Le Wang; Sanping Zhou; Gang Hua; Changyin Sun
Learning Content-Enhanced Mask Transformer for Domain Generalized Urban-Scene Segmentation	Qi Bi; Shaodi You; Theo Gevers
CRA-PCN: Point Cloud Completion with Intra- and Inter-level Cross-Resolution Transformers	Yi Rong; Haoran Zhou; Lixin Yuan; Cheng Mei; Jiahao Wang; Tong Lu
QCS-SGM+: Improved Quantized Compressed Sensing with Score-Based Generative Models	Xiangming Meng; Yoshiyuki Kabashima
Shrinking Your TimeStep: Towards Low-Latency Neuromorphic Object Recognition with Spiking Neural Networks	Yongqi Ding; Lin Zuo; Mengmeng Jing; Pei He; Yongjun Xiao
Probabilities of Causation with Nonbinary Treatment and Effect	Ang Li; Judea Pearl
Enhancing Hyperspectral Images via Diffusion Model and Group-Autoencoder Super-resolution Network	Zhaoyang Wang; Dongyang Li; Mingyang Zhang; Hao Luo; Maoguo Gong
Towards Modeling Uncertainties of Self-explaining Neural Networks via Conformal Prediction	Wei Qian; Chenxu Zhao; Yangyi Li; Fenglong Ma; Chao Zhang; Mengdi Huai
Concept-Guided Prompt Learning for Generalization in Vision-Language	Yi Zhang; Ce Zhang; Ke Yu; Yushun Tang; Zhihai He
FRED: Towards a Full Rotation-Equivariance in Aerial Image Object Detection	Chanho Lee; Jinsu Son; Hyounguk Shon; Yunho Jeon; Junmo Kim
AT4CTR: Auxiliary Match Tasks for Enhancing Click-Through Rate Prediction	Qi Liu; Xuyang Hou; Defu Lian; Zhe Wang; Hanran Jin; Jia Cheng; Jun Lei
StegaStyleGAN: Towards Generic and Practical Generative Image Steganography	Wenkang Su; Jiangqun Ni; Yiyu Sun
Emergent Communication for Numerical Concepts Generalization	Enshuai Zhou; Yifan Hao; Rui Zhang; Yuxuan Guo; Zidong Du; Xishan Zhang; Xinkai Song; Chao Wang; Xuehai Zhou; Jiaming Guo; Qi Yi; Shaohui Peng; Di Huang; Ruizhi Chen; Qi Guo; Yunji Chen
Aligning Geometric Spatial Layout in Cross-View Geo-Localization via Feature Recombination	Qingwang Zhang; Yingying Zhu
Simple Weak Coresets for Non-decomposable Classification Measures	Jayesh Malaviya; Anirban Dasgupta; Rachit Chhaya



Learning Multi-Scale Video-Text Correspondence for Weakly Supervised Temporal Article Grounding	Wenjia Geng; Yong Liu; Lei Chen; Sujia Wang; Jie Zhou; Yansong Tang
A Surprisingly Simple Continuous-Action POMDP Solver: Lazy Cross-Entropy Search over Policy Trees	Marcus Hoerger; Hanna Kurniawati; Dirk Kroese; Nan Ye
YTCommentQA: Video Question Answerability in Instructional Videos	Saelyne Yang; Sunghyun Park; Yunseok Jang; Moontae Lee
Explainable Origin-Destination Crowd Flow Interpolation via Variational Multi-Modal Recurrent Graph Auto-Encoder	Qiang Zhou; Xinjiang Lu; Jingjing Gu; Zhe Zheng; Bo Jin; Jingbo Zhou
Towards Multi-Intent Spoken Language Understanding via Hierarchical Attention and Optimal Transport	Xuxin Cheng; Zhihong Zhu; Hongxiang Li; Yaowei Li; Xianwei Zhuang; Yuexian Zou
Large Language Models Are Neurosymbolic Reasoners	Meng Fang; Shilong Deng; Yudi Zhang; Zijing Shi; Ling Chen; Mykola Pechenizkiy; Jun Wang
Low Category Uncertainty and High Training Potential Instance Learning for Unsupervised Domain Adaptation	Xinyu Zhang; Meng Kang; Shuai Lü
FedST: Federated Style Transfer Learning for Non-IID Image Segmentation	Boyuan Ma; Yin Xiang; Jing Tan; Yongfeng Chen; Haiyou Huang; Hao Wang; Weihua Xue; Xiaojuan Ban
Efficient Learning of PDEs via Taylor Expansion and Sparse Decomposition into Value and Fourier Domains	Md Nasim; Yexiang Xue
Discrepancy and Uncertainty Aware Denoising Knowledge Distillation for Zero-Shot Cross-Lingual Named Entity Recognition	Ling Ge; Chunming Hu; Guanghui Ma; Jihong Liu; Hong Zhang
Foundations of Reactive Synthesis for Declarative Process Specifications	Andrey Rivkin; Luca Geatti; Marco Montali
Compound Text-Guided Prompt Tuning via Image-Adaptive Cues	Hao Tan; Jun Li; Yizhuang Zhou; Jun Wan; Zhen Lei; Xiangyu Zhang
Exploiting the Social-Like Prior in Transformer for Visual Reasoning	Yudong Han; Yupeng Hu; Xuemeng Song; Haoyu Tang; Mingzhu Xu; Liqiang Zechen Li; Weiming Huang; Kai Zhao; Min Yang; Yongshun Gong; Meng Chen
Urban Region Embedding via Multi-View Contrastive Prediction	Joseph Clements; Yingjie Lao
Resource Efficient Deep Learning Hardware Watermarks with Signature Alignment	
Memory-Efficient Prompt Tuning for Incremental Histopathology Classification	Yu Zhu; Kang Li; Lequan Yu; Pheng-Ann Heng
STAR: Boosting Low-Resource Information Extraction by Structure-to-Text Data Generation with Large Language Models	Mingyu Derek Ma; Xiaoxuan Wang; Po-Nien Kung; P. Jeffrey Brantingham; Nanyun Peng; Wei Wang
IOFM: Using the Interpolation Technique on the Over-Fitted Models to Identify Clean-Annotated Samples	Dongha Kim; Yongchan Choi; Kunwoong Kim; Ilsang Ohn; Yongdai Kim
Improving Distinguishability of Class for Graph Neural Networks	Dongxiao He; Shuwei Liu; Zhizhi Yu; Guangquan Xu; Meng Ge; Zhiyong Feng
Disjoint Partial Enumeration without Blocking Clauses	Giuseppe Spallitta; Roberto Sebastiani; Armin Biere
Graph Context Transformation Learning for Progressive Correspondence Pruning	Junwen Guo; Guobao Xiao; Shiping Wang; Jun Yu
A Multi-Modal Contrastive Diffusion Model for Therapeutic Peptide Generation	Yongkang Wang; Xuan Liu; Feng Huang; Zhankun Xiong; Wen Zhang
Adversarially Balanced Representation for Continuous Treatment Effect Estimation	Amirreza Kazemi; Martin Ester
X-RefSeg3D: Enhancing Referring 3D Instance Segmentation via Structured Cross-Modal Graph Neural Networks	Zhipeng Qian; Yiwei Ma; Jiayi Ji; Xiaoshuai Sun
Iterative Regularization with K-support Norm: An Important Complement to Sparse Recovery	William de Vazelhes; Bhaskar Mukhoty; Xiaotong Yuan; Bin Gu



Learning GAI-Decomposable Utility Models for Multiattribute Decision Making	Margot Herin; Patrice Perny; Nataliya Sokolovska
Multi-Region Text-Driven Manipulation of Diffusion Imagery	Yiming Li; Peng Zhou; Jun Sun; Yi Xu
Towards Fine-Grained HBOE with Rendered Orientation Set and Laplace Smoothing	Ruisi Zhao; Mingming Li; Zheng Yang; Binbin Lin; Xiaohui Zhong; Xiaobo Ren; Deng Cai; Boxi Wu
'Why Didn't You Allocate This Task to Them?' Negotiation-Aware Task Allocation and Contrastive Explanation Generation	Zahra Zahedi; Sailik Sengupta; Subbarao Kambhampati
Causal-Driven Skill Prerequisite Structure Discovery	Shenbao Yu; Yifeng Zeng; Fan Yang; Yinghui Pan
Repeated Fair Allocation of Indivisible Items	Ayumi Igarashi; Martin Lackner; Oliviero Nardi; Arianna Novaro
Audio Scanning Network: Bridging Time and Frequency Domains for Audio Classification	Liangwei Chen; Xiren Zhou; Huanhuan Chen
GAMC: An Unsupervised Method for Fake News Detection Using Graph Autoencoder with Masking	Shu Yin; Peican Zhu; Lianwei Wu; Chao Gao; Zhen Wang
IVP-VAE: Modeling EHR Time Series with Initial Value Problem Solvers	Jingge Xiao; Leonie Basso; Wolfgang Nejdl; Niloy Ganguly; Sandipan Sikdar
A Bregman Proximal Stochastic Gradient Method with Extrapolation for Nonconvex Nonsmooth Problems	Qingsong Wang; Zehui Liu; Chunfeng Cui; Deren Han
Learning Visual Abstract Reasoning through Dual-Stream Networks	Kai Zhao; Chang Xu; Bailu Si
High-Order Structure Based Middle-Feature Learning for Visible-Infrared Person Re-identification	Liuxiang Qiu; Si Chen; Yan Yan; Jing-Hao Xue; Da-Han Wang; Shunzhi Zhu
Mitigating Idiom Inconsistency: A Multi-Semantic Contrastive Learning Method for Chinese Idiom Reading Comprehension	Mingmin Wu; Yuxue Hu; Yongcheng Zhang; Zeng Zhi; Guixin Su; Ying Sha
Online Submodular Maximization via Online Convex Optimization	Tareq Si Salem; Gözde Özcan; Iasonas Nikolaou; Evimaria Terzi; Stratis Ioannidis
Approval-Based Committee Voting in Practice: A Case Study of (over-)Representation in the Polkadot Blockchain	Niclas Boehmer; Markus Brill; Alfonso Cevallos; Jonas Gehrlein; Luis Sanchez-Fernandez; Ulrike Schmidt-Kraepelin
Adversarial Attacks on Federated-Learned Adaptive Bitrate Algorithms	Ruixiao Zhang; Tianchi Huang
GOALNET: Interleaving Neural Goal Predicate Inference with Classical Planning for Generalization in Robot Instruction Following	Jigyasa Gupta; Shreya Sharma; Shreshth Tuli; Rohan Paul; Mausam
Conformal Autoregressive Generation: Beam Search with Coverage Guarantees	Nicolas Deutschmann; Marvin Alberts; María Rodríguez Martínez
Code-Style In-Context Learning for Knowledge-Based Question Answering	Zhijie Nie; Richong Zhang; Zhongyuan Wang; Xudong Liu
Graph Disentangled Contrastive Learning with Personalized Transfer for Cross-Domain Recommendation	Jing Liu; Lele Sun; Weizhi Nie; Peiguang Jing; Yu-ting Su
Two-Stage Evolutionary Reinforcement Learning for Enhancing Exploration and Exploitation	Qingling Zhu; Xiaoqiang Wu; Qiuzhen Lin; Wei-Neng Chen
Natural Strategic Ability in Stochastic Multi-Agent Systems	Raphaël Berthon; Joost-Pieter Katoen; Munyque Mittelman; Aniello Murano
Turning Waste into Wealth: Leveraging Low-Quality Samples for Enhancing Continuous Conditional Generative Adversarial Networks	Xin Ding; Yongwei Wang; Zuheng Xu
Manifold Constraints for Imperceptible Adversarial Attacks on Point Clouds	Keke Tang; Xu He; Weilong Peng; Jianpeng Wu; Yawen Shi; Daizong Liu; Pan Zhou; Wenping Wang; Zhihong Tian
SpFormer: Spatio-Temporal Modeling for Scanpaths with Transformer	Wenqi Zhong; Linzhi Yu; Chen Xia; Junwei Han; Dingwen Zhang
Object-Aware Adaptive-Positivity Learning for Audio-Visual Question	Zhangbin Li; Jinxing Zhou; Dan Guo; Jing Zhang; Meng Wang





Unit Selection with Nonbinary Treatment and Effect	Ang Li; Judea Pearl
Traffic Flow Optimisation for Lifelong Multi-Agent Path Finding	Zhe Chen; Daniel Harabor; Jiaoyang Li; Peter Stuckey
Enhancing Neural Radiance Fields with Adaptive Multi-Exposure Fusion: A Bilevel Optimization Approach for Novel View Synthesis	Yang Zou; Xingyuan Li; Zhiying Jiang; Jinyuan Liu
MolTailor: Tailoring Chemical Molecular Representation to Specific Tasks via Text Prompts	Haoqiang Guo; Sendong Zhao; Haochun Wang; Yanrui Du; Bing Qin
Accelerating the Global Aggregation of Local Explanations	Alon Mor; Yonatan Belinkov; Benny Kimelfeld
CUTS+: High-Dimensional Causal Discovery from Irregular Time-Series	Yuxiao Cheng; Lianglong Li; Tingxiong Xiao; Zongren Li; Jinli Suo; Kunlun He; Qionghai Dai
Sync-NeRF: Generalizing Dynamic NeRFs to Unsynchronized Videos	Seoha Kim; Jeongmin Bae; Youngsik Yun; Hahyun Lee; Gun Bang; Youngjung Uh
CF-NeRF: Camera Parameter Free Neural Radiance Fields with Incremental Learning	Qingsong Yan; Qiang Wang; Kaiyong Zhao; Jie Chen; Bo Li; Xiaowen Chu; Fei Deng
An Exercise in Tournament Design: When Some Matches Must Be Scheduled	Sushmita Gupta; Ramanujan Sridharan; Peter Strulo
Multi-Source Collaborative Gradient Discrepancy Minimization for Federated Domain Generalization	Yikang Wei; Yahong Han
Controller-Guided Partial Label Consistency Regularization with Unlabeled Data	Qian-Wei Wang; Bowen Zhao; Mingyan Zhu; Tianxiang Li; Zimo Liu; Shu-Tao Xia
Federated Contextual Cascading Bandits with Asynchronous Communication and Heterogeneous Users	Hantao Yang; Xutong Liu; Zhiyong Wang; Hong Xie; John C. S. Lui; Defu Lian; Enhong Chen
msLPCC: A Multimodal-Driven Scalable Framework for Deep LiDAR Point Cloud Compression	Miaohui Wang; Runnan Huang; Hengjin Dong; Di Lin; Song Yun; Wuyuan Xie
Prompting Segmentation with Sound Is Generalizable Audio-Visual Source Localizer	Yaoting Wang; Liu Weisong; Guangyao Li; Jian Ding; Di Hu; Xi Li
Cautiously-Optimistic Knowledge Sharing for Cooperative Multi-Agent Reinforcement Learning	Yanwen Ba; Xuan Liu; Xinning Chen; Hao Wang; Yang Xu; Kenli Li; Shigeng Zhang
A Unified View on Forgetting and Strong Equivalence Notions in Answer Set Programming	Zeynep G. Saribatur; Stefan Woltran
Graph Neural Prompting with Large Language Models	Yijun Tian; Huan Song; Zichen Wang; Haozhu Wang; Ziqing Hu; Fang Wang; Nitesh V. Chawla; Panpan Xu
Chain-of-Thought Improves Text Generation with Citations in Large Language Models	Bin Ji; Huijun Liu; Mingzhe Du; See-Kiong Ng
The Choice of Noninformative Priors for Thompson Sampling in Multiparameter Bandit Models	Jongyeong Lee; Chao-Kai Chiang; Masashi Sugiyama
Worst-Case VCG Redistribution Mechanism Design Based on the Lottery Ticket Hypothesis	Mingyu Guo
Hypergraph Neural Architecture Search	Wei Lin; Xu Peng; Zhengtao Yu; Taisong Jin
Box2Poly: Memory-Efficient Polygon Prediction of Arbitrarily Shaped and Rotated Text	Xuyang Chen; Dong Wang; Konrad Schindler; Mingwei Sun; Yongliang Wang; Nicolo Savioli; Liqiu Meng
Machine Learning-Powered Combinatorial Clock Auction	Ermis Nikiforos Soumalias; Jakob Weisstener; Jakob Heiss; Sven Seuken
Boosting Few-Shot Learning via Attentive Feature Regularization	Xingyu Zhu; Shuo Wang; Jinda Lu; Yanbin Hao; Haifeng Liu; Xiangnan He





A Fixed-Point Approach to Unified Prompt-Based Counting	Wei Lin; Antoni Chan
Mixup-Induced Domain Extrapolation for Domain Generalization	Meng Cao; Songcan Chen
InstructDoc: A Dataset for Zero-Shot Generalization of Visual Document Understanding with Instructions	Ryota Tanaka; Taichi Iki; Kyosuke Nishida; Kuniko Saito; Jun Suzuki
Parameterized Approximation Algorithms for Sum of Radii Clustering and Variants	Xianrun Chen; Dachuan Xu; Yicheng Xu; Yong Zhang
LSTKC: Long Short-Term Knowledge Consolidation for Lifelong Person Re-identification	Kunlun Xu; Xu Zou; Jiahuan Zhou
Interpretable3D: An Ad-Hoc Interpretable Classifier for 3D Point Clouds	Tuo Feng; Ruijie Quan; Xiaohan Wang; Wenguan Wang; Yi Yang
HDformer: A Higher-Dimensional Transformer for Detecting Diabetes Utilizing Long-Range Vascular Signals	Ella Lan
3D-STMN: Dependency-Driven Superpoint-Text Matching Network for End-to-End 3D Referring Expression Segmentation	Changli Wu; Yiwei Ma; Qi Chen; Haowei Wang; Gen Luo; Jiayi Ji; Xiaoshuai Sun
Meta-Reinforcement Learning via Exploratory Task Clustering	Zhendong Chu; Renqin Cai; Hongning Wang
TraceEvader: Making DeepFakes More Untraceable via Evading the Forgery Model Attribution	Mengjie Wu; Jingui Ma; Run Wang; Sidan Zhang; Ziyu Liang; Boheng Li; Chenhao Lin; Liming Fang; Lina Wang
Improved Metric Distortion via Threshold Approvals	Elliot Anshelevich; Aris Filos-Ratsikas; Christopher Jerrett; Alexandros A. Voudouris
Direct May Not Be the Best: An Incremental Evolution View of Pose	Yuelong Li; Tengfei Xiao; Lei Geng; Jianming Wang
AQ-DETR: Low-Bit Quantized Detection Transformer with Auxiliary Queries	Runqi Wang; Huixin Sun; Linlin Yang; Shaohui Lin; Chuanjian Liu; Yan Gao; Yao Hu; Baochang Zhang
Learning Only When It Matters: Cost-Aware Long-Tailed Classification	Yu-Cheng He; Yao-Xiang Ding; Han-Jia Ye; Zhi-Hua Zhou
Abstract Action Scheduling for Optimal Temporal Planning via OMT	Stefan Panjkovic; Andrea Micheli
MICA: Towards Explainable Skin Lesion Diagnosis via Multi-Level Image-Concept Alignment	Yequan Bie; Luyang Luo; Hao Chen
Pairwise-Label-Based Deep Incremental Hashing with Simultaneous Code Expansion	Dayan Wu; Qinghang Su; Bo Li; Weiping Wang
Talk Funny! A Large-Scale Humor Response Dataset with Chain-of-Humor Interpretation	Yuyan Chen; Yichen Yuan; Panjun Liu; Dayiheng Liu; Qinghao Guan; Mengfei Guo; Haiming Peng; Bang Liu; Zhixu Li; Yanghua Xiao
NaMa: Neighbor-Aware Multi-Modal Adaptive Learning for Prostate Tumor Segmentation on Anisotropic MR Images	Runqi Meng; Xiao Zhang; Shijie Huang; Yuning Gu; Guiqin Liu; Guangyu Wu; Nizhuan Wang; Kaicong Sun; Dinggang Shen
Translate Meanings, Not Just Words: IdiomKB's Role in Optimizing Idiomatic Translation with Language Models	Shuang Li; Jiangjie Chen; Siyu Yuan; Xinyi Wu; Hao Yang; Shimin Tao; Yanghua Xiao
Transferable Adversarial Attacks for Object Detection Using Object-Aware Significant Feature Distortion	Xinlong Ding; Jiansheng Chen; Hongwei Yu; Yu Shang; Yining Qin; Huimin Ma
Taxonomy Driven Fast Adversarial Training	Kun Tong; Chengze Jiang; Jie Gui; Yuan Cao
Residual Hyperbolic Graph Convolution Networks	Yangkai Xue; Jindou Dai; Zhipeng Lu; Yuwei Wu; Yunde Jia
LatestEval: Addressing Data Contamination in Language Model Evaluation through Dynamic and Time-Sensitive Test Construction	Yucheng Li; Frank Guerin; Chenghua Lin
Kumaraswamy Wavelet for Heterophilic Scene Graph Generation	Lianggangxu Chen; Youqi Song; Shaohui Lin; Changbo Wang; Gaoqi He
Efficient Online Crowdsourcing with Complex Annotations	Reshef Meir; Viet-An Nguyen; Xu Chen; Jagdish Ramakrishnan; Udi



BDIQA: A New Dataset for Video Question Answering to Explore Cognitive Reasoning through Theory of Mind	Yuanyuan Mao; Xin Lin; Qin Ni; Liang He
Semantic-Guided Novel Category Discovery	Weishuai Wang; Ting Lei; Qingchao Chen; Yang Liu
MedSegDiff-V2: Diffusion-based Medical Image Segmentation with Rating-Based Reinforcement Learning	Junde Wu; Wei Ji; Huazhu Fu; Min Xu; Yueming Jin; Yanwu Xu Devin White; Mingkang Wu; Ellen Novoseller; Vernon Lawhern; Nicholas Waytowich; Yongcan Cao
Advancing Spatial Reasoning in Large Language Models: An In-Depth Evaluation and Enhancement Using the StepGame Benchmark	Fangjun Li; David C. Hogg; Anthony G. Cohn
UniGen: A Unified Generative Framework for Retrieval and Question Answering with Large Language Models	Xiaoxi Li; Yujia Zhou; Zhicheng Dou
Colour Passing Revisited: Lifted Model Construction with Commutative Episodic Return Decomposition by Difference of Implicitly Assigned Sub-trajectory Reward	Malte Luttermann; Tanya Braun; Ralf Möller; Marcel Gehrke Haixin Lin; Hongqiu Wu; Jiaji Zhang; Yihao Sun; Junyin Ye; Yang Yu
TA&AT: Enhancing Task-Oriented Dialog with Turn-Level Auxiliary Tasks and Action-Tree Based Scheduled Sampling	Longxiang Liu; Xiuxing Li; Yang Feng
Efficient Representation Learning of Satellite Image Time Series and Their Fusion for Spatiotemporal Applications	Poonam Goyal; Arshveer Kaur; Arvind Ram; Navneet Goyal
ViT-Calibrator: Decision Stream Calibration for Vision Transformer	Lin Chen; Zhijie Jia; Lechao Cheng; Yang Gao; Jie Lei; Yijun Bei; Zunlei Feng
Parallel Ranking of Ads and Creatives in Real-Time Advertising Systems	Zhiguang Yang; Liufang Sang; Haoran Wang; Wenlong Chen; Lu Wang; Jie He; Changping Peng; Zhangang Lin; Chun Gan; Jingping Shao
Hypergraph Joint Representation Learning for Hypervertices and Hyperedges via Cross Expansion	Yuguang Yan; Yuanlin Chen; Shibo Wang; Hanrui Wu; Ruichu Cai
Towards Running Time Analysis of Interactive Multi-Objective Evolutionary Algorithms	Tianhao Lu; Chao Bian; Chao Qian
PreRoutGNN for Timing Prediction with Order Preserving Partition: Global Circuit Pre-training, Local Delay Learning and Attentional Cell Modeling	Ruizhe Zhong; Junjie Ye; Zhentao Tang; Shixiong Kai; Mingxuan Yuan; Jianye Hao; Junchi Yan
Upper Bounding Barlow Twins: A Novel Filter for Multi-Relational Clustering	Xiaowei Qian; Bingheng Li; Zhao Kang
Measuring Task Similarity and Its Implication in Fine-Tuning Graph Neural Networks	Renhong Huang; Jiarong Xu; Xin Jiang; Chenglu Pan; Zhiming Yang; Chunping Wang; Yang Yang
Continual Relation Extraction via Sequential Multi-Task Learning	Thanh-Thien Le; Manh Nguyen; Tung Nguyen; Linh Ngo Van; Thien Huu Nguyen
Revisiting Graph-Based Fraud Detection in Sight of Heterophily and Spectrum	Fan Xu; Nan Wang; Hao Wu; Xuezhi Wen; Xibin Zhao; Hai Wan
Progressive Text-to-Image Diffusion with Soft Latent Direction	YuTeng Ye; Jiale Cai; Hang Zhou; Guanwen Li; Youjia Zhang; Zikai Song; Chenxing Gao; Junqing Yu; Wei Yang
Generalization Analysis of Machine Learning Algorithms via the Worst-Case Data-Generating Probability Measure	Xinying Zou; Samir Perlaza; Inaki Esnaola; Eitan Altman
Bootstrapping Large Language Models for Radiology Report Generation	Chang Liu; Yuanhe Tian; Weidong Chen; Yan Song; Yongdong Zhang
Frequency-Controlled Diffusion Model for Versatile Text-Guided Image-to-Image Translation	Xiang Gao; Zhengbo Xu; Junhan Zhao; Jiaying Liu
Dynamic Regret of Adversarial MDPs with Unknown Transition and Linear Function Approximation	Long-Fei Li; Peng Zhao; Zhi-Hua Zhou



SciEval: A Multi-Level Large Language Model Evaluation Benchmark for Scientific Research	Liangtai Sun; Yang Han; Zihan Zhao; Da Ma; Zhennan Shen; Baocai Chen; Lu Chen; Kai Yu
DOGE-Train: Discrete Optimization on GPU with End-to-End Training	Ahmed Abbas; Paul Swoboda
Learning the Causal Structure of Networked Dynamical Systems under Latent Nodes and Structured Noise	Augusto Santos; Diogo Rente; Rui Seabra; José M. F. Moura
Improving Open Set Recognition via Visual Prompts Distilled from Common-Sense Knowledge	Seongyeop Kim; Hyung-Il Kim; Yong Man Ro
CrystalBox: Future-Based Explanations for Input-Driven Deep RL Systems	Sagar Patel; Sangeetha Abdu Jyothi; Nina Narodytska
GradTree: Learning Axis-Aligned Decision Trees with Gradient Descent	Sascha Marton; Stefan Lüdtke; Christian Bartelt; Heiner Stuckenschmidt
Redefining ABA+ Semantics via Abstract Set-to-Set Attacks	Yannis Dimopoulos; Wolfgang Dvorak; Matthias König; Anna Rapberger; Markus Ulbricht; Stefan Woltran
TexFit: Text-Driven Fashion Image Editing with Diffusion Models	Tongxin Wang; Mang Ye
Data Adaptive Traceback for Vision-Language Foundation Models in Image Classification	Wenshuo Peng; Kaipeng Zhang; Yue Yang; Hao Zhang; Yu Qiao
Personalized LoRA for Human-Centered Text Understanding	You Zhang; Jin Wang; Liang-Chih Yu; Dan Xu; Xuejie Zhang
Uncertainty-Aware GAN for Single Image Super Resolution	Chenxi Ma
DeepCalliFont: Few-Shot Chinese Calligraphy Font Synthesis by Integrating Dual-Modality Generative Models	Yitian Liu; Zhouhui Lian
Mitigating the Impact of False Negative in Dense Retrieval with Contrastive Confidence Regularization	Shiqi Wang; Yeqin Zhang; Cam-Tu Nguyen
DiffBEV: Conditional Diffusion Model for Bird's Eye View Perception	Jiayu Zou; Kun Tian; Zheng Zhu; Yun Ye; Xingang Wang
CORECODE: A Common Sense Annotated Dialogue Dataset with Benchmark Tasks for Chinese Large Language Models	Dan Shi; Chaobin You; Jian-Tao Huang; Taihao Li; Deyi Xiong
Incremental Quasi-Newton Methods with Faster Superlinear Convergence Rates	Zhuanghua Liu; Luo Luo; Bryan Kian Hsiang Low
Unsupervised Training Sequence Design: Efficient and Generalizable Agent Training	Wenjun Li; Pradeep Varakantham
SGNet: Structure Guided Network via Gradient-Frequency Awareness for Depth Map Super-resolution	Zhengxue Wang; Zhiqiang Yan; Jian Yang
Cell Graph Transformer for Nuclei Classification	Wei Lou; Guanbin Li; Xiang Wan; Haofeng Li
Patch-Aware Sample Selection for Efficient Masked Image Modeling	Zhengyang Zhuge; Jiaying Wang; Yong Li; Yongjun Bao; Peisong Wang; Jian Cheng
Towards Effective and General Graph Unlearning via Mutual Evolution	Xunkai Li; Yulin Zhao; Zhengyu Wu; Wentao Zhang; Rong-Hua Li; Guoren
PPEA-Depth: Progressive Parameter-Efficient Adaptation for Self-Supervised Monocular Depth Estimation	Yue-Jiang Dong; Yuan-Chen Guo; Ying-Tian Liu; Fang-Lue Zhang; Song-Hai Zhang
Enhance Sketch Recognition's Explainability via Semantic Component-Level Parsing	Guangming Zhu; Siyuan Wang; Tianci Wu; Liang Zhang
Prior and Prediction Inverse Kernel Transformer for Single Image Defocus Deblurring	Peng TANG; Zhiqiang Xu; Chunlai Zhou; Pengfei Wei; Peng Han; Xin Cao; Tobias Lasser
Dual-Perspective Knowledge Enrichment for Semi-supervised 3D Object Detection	Yucheng Han; Na Zhao; Weiling Chen; Keng Teck Ma; Hanwang Zhang



Removing Interference and Recovering Content Imaginatively for Visible Watermark Removal	Yicheng Leng; Chaowei Fang; Gen Li; Yixiang Fang; Guanbin Li
Label-Efficient Few-Shot Semantic Segmentation with Unsupervised Meta-Training	Jianwu Li; Kaiyue Shi; Guo-Sen Xie; Xiaofeng Liu; Jian Zhang; Tianfei Zhou
GINN-LP: A Growing Interpretable Neural Network for Discovering Multivariate Laurent Polynomial Equations	Nisal Ranasinghe; Damith Senanayake; Sachith Seneviratne; Malin Premaratne; Saman Halgamuge
SMILEtrack: SiMilarity LEarning for Occlusion-Aware Multiple Object Tracking	Wang Yu Hsiang; Jun-Wei Hsieh; Ping-Yang Chen; Ming-Ching Chang; Hung Hin So; Xin Li
FedCSL: A Scalable and Accurate Approach to Federated Causal Structure Learning	Xianjie Guo; Kui Yu; Lin Liu; Jiuyong Li
Constraint Latent Space Matters: An Anti-anomalous Waveform Transformation Solution from Photoplethysmography to Arterial Blood	Cheng Bian; Xiaoyu Li; Zhu Guangpu; Jiegeng Lyu; Qi Bi; Weile Zhang; Yelei Li; Zijing Zeng
Barely Supervised Learning for Graph-Based Fraud Detection	Hang Yu; Zhengyang Liu; Xiangfeng Luo
UCMCTrack: Multi-Object Tracking with Uniform Camera Motion	Kefu Yi; Kai Luo; Xiaolei Luo; Jiangui Huang; Hao Wu; Rongdong Hu; Wei Hao
AdapEdit: Spatio-Temporal Guided Adaptive Editing Algorithm for Text-Based Continuity-Sensitive Image Editing	Zhiyuan Ma; Guoli Jia; Bowen Zhou
How to Evaluate Behavioral Models	Greg d'Eon; Sophie Greenwood; Kevin Leyton-Brown; James R. Wright
OctOcc: High-Resolution 3D Occupancy Prediction with Octree	Wenzhe Ouyang; Xiaolin Song; Bailan Feng; Zenglin Xu
ReGCL: Rethinking Message Passing in Graph Contrastive Learning	Cheng Ji; Zixuan Huang; Qingyun Sun; Hao Peng; Xingcheng Fu; Qian Li; Jianxin Li
Recasting Regional Lighting for Shadow Removal	Yuhao Liu; Zhanghan Ke; Ke Xu; Fang Liu; Zhenwei Wang; Rynson W.H. Lau
Multiagent Gumbel MuZero: Efficient Planning in Combinatorial Action Spaces	Xiaotian Hao; Jianye Hao; Chenjun Xiao; Kai Li; Dong Li; Yan Zheng
LINGO-Space: Language-Conditioned Incremental Grounding for Space	Dohyun Kim; Nayoung Oh; Deokmin Hwang; Daehyung Park
ROG_PL: Robust Open-Set Graph Learning via Region-Based Prototype Learning	Qin Zhang; Jiexin Lu; Xiaowei Li; Liping Qiu; Shirui Pan; Xiaojun Chen; Junyang Chen
RadarMOSEVE: A Spatial-Temporal Transformer Network for Radar-Only Moving Object Segmentation and Ego-Velocity Estimation	Changsong Pang; Xieyuanli Chen; Yimin Liu; Huimin Lu; Yuwei Cheng
A Dual Stealthy Backdoor: From Both Spatial and Frequency Perspectives	Yudong Gao; Honglong Chen; Peng Sun; Junjian Li; Anqing Zhang; Zhibo Wang; Weifeng Liu
Towards Model Extraction Attacks in GAN-Based Image Translation via Domain Shift Mitigation	Di Mi; Yanjun Zhang; Leo Yu Zhang; Shengshan Hu; Qi Zhong; Haizhuan Yuan; Shirui Pan
Any-Size-Diffusion: Toward Efficient Text-Driven Synthesis for Any-Size HD Images	Qingping Zheng; Yuanfan Guo; Jiankang Deng; Jianhua Han; Ying Li; Songcen Xu; Hang Xu
Dual-Level Curriculum Meta-Learning for Noisy Few-Shot Learning Tasks	Xiaofan Que; Qi Yu
BeliefFlow: A Framework for Logic-Based Belief Diffusion via Iterated Belief Change	Nicolas Schwind; Katsumi Inoue; Sébastien Konieczny; Pierre Marquis
Hierarchical Aligned Multimodal Learning for NER on Tweet Posts	Pelpei Liu; Hong Li; Yimo Ren; Jie Liu; Shuaizong Si; Hongsong Zhu; Limin
No Prejudice! Fair Federated Graph Neural Networks for Personalized Recommendation	Nimesh Agrawal; Anuj Sirohi; Sandeep Kumar; Jayadeva Jayadeva
TimesURL: Self-Supervised Contrastive Learning for Universal Time Series Representation Learning	jiexi Liu; Songcan Chen



A Theory of Non-acyclic Generative Flow Networks	Leo Brunswic; Yinchuan Li; Yushun Xu; Yijun Feng; Shangling Jui; Lizhuang
Shadow Generation with Decomposed Mask Prediction and Attentive Shadow Filling	Xinhao Tao; Junyan Cao; Yan Hong; Li Niu
CLIM: Contrastive Language-Image Mosaic for Region Representation	Size Wu; Wenwei Zhang; Lumin XU; Sheng Jin; Wentao Liu; Chen Change
Existence Is Chaos: Enhancing 3D Human Motion Prediction with Uncertainty Consideration	Zhihao Wang; Yulin Zhou; Ningyu Zhang; Xiaosong Yang; Jun Xiao; Zhao Wang
Identification for Tree-Shaped Structural Causal Models in Polynomial Time	Aaryan Gupta; Markus Bläser
Superposed Atomic Representation for Robust High-Dimensional Data	Yulong Wang
Recovery of Multiple Low-Dimensional Structures	
An Approximate Skolem Function Counter	Arijit Shaw; Brendan Juba; Kuldeep S Meel
Uni-MIS: United Multiple Intent Spoken Language Understanding via Multi-View Intent-Slot Interaction	Shangjian Yin; Peijie Huang; Yuhong Xu
DHGCN: Dynamic Hop Graph Convolution Network for Self-Supervised Point Cloud Learning	Jincen Jiang; Lizhi Zhao; Xuequan Lu; Wei Hu; Imran Razzak; Meili Wang
Gaussian Process Neural Additive Models	Wei Zhang; Brian Barr; John Paisley
Neural Causal Abstractions	Kevin Xia; Elias Bareinboim
VVS: Video-to-Video Retrieval with Irrelevant Frame Suppression	Won Jo; Geuntaek Lim; Gwangjin Lee; Hyunwoo Kim; Byungsoo Ko; Yukyung Choi
ICAR: Image-Based Complementary Auto Reasoning	Xijun Wang; Anqi Liang; Junbang Liang; Ming Lin; Yu Lou; Shan Yang
DiSCO: Diffusion Schrödinger Bridge for Molecular Conformer Optimization	Danyeong Lee; Dohoon Lee; Dongmin Bang; Kim Sun
Synergistic Multiscale Detail Refinement via Intrinsic Supervision for Underwater Image Enhancement	Dehuan Zhang; Jingchun Zhou; Chunle Guo; Weishi Zhang; Chongyi Li
Colored Noise in PPO: Improved Exploration and Performance through Correlated Action Sampling	Jakob Hollenstein; Georg Martius; Justus Piater
Data Augmented Graph Neural Networks for Personality Detection	Yangfu Zhu; Yue Xia; Meiling Li; Tingting Zhang; Bin Wu
Generative Calibration of Inaccurate Annotation for Label Distribution Learning	Liang He; Yunan Lu; Weiwei Li; Xiuyi Jia
Enhancing Ensemble Clustering with Adaptive High-Order Topological	Jiaxuan Xu; Taiyong Li; Lei Duan
Enhancing Zero-Shot Multi-Speaker TTS with Negated Speaker Representations	Yejin Jeon; Yunsu Kim; Gary Geunbae Lee
DistilVPR: Cross-Modal Knowledge Distillation for Visual Place Recognition	Sijie Wang; Rui She; Qiyu Kang; Xingchao Jian; Kai Zhao; Yang Song; Wee Peng Tay
Decoupled Optimisation for Long-Tailed Visual Recognition	Cong Cong; Shiyu Xuan; Sidong Liu; Shiliang Zhang; Maurice Pagnucco; Yang Song
NodeMixup: Tackling Under-Reaching for Graph Neural Networks	Weigang Lu; Ziyu Guan; Wei Zhao; Yaming Yang; Long Jin
A Transfer Approach Using Graph Neural Networks in Deep Reinforcement Learning	Tianpei Yang; Heng You; Jianye Hao; Yan Zheng; Matthew Taylor
Recall-Oriented Continual Learning with Generative Adversarial Meta-Model	Haneol Kang; Dong-Wan Choi
An Effective Polynomial Technique for Compiling Conditional Effects Away	Alfonso Emilio Gerevini; Francesco Percassi; Enrico Scala
Multi-View People Detection in Large Scenes via Supervised View-Wise Contribution Weighting	Qi ZHANG; Yunfei Gong; Daijie Chen; Antoni Chan; Hui Huang



Decoding Global Preferences: Temporal and Cooperative Dependency Modeling in Multi-Agent Preference-Based Reinforcement Learning	Tianchen Zhu; Yue Qiu; Haoyi Zhou; Jianxin Li
LLM vs Small Model? Large Language Model Based Text Augmentation	Linmei Hu; Hongyu He; Duokang Wang; Ziwang Zhao; Yingxia Shao; Liqiang Nie
Enhanced Personality Detection Model	
Hybrid-SORT: Weak Cues Matter for Online Multi-Object Tracking	Mingzhan Yang; Guangxin Han; Bin Yan; Wenhua Zhang; Jinqing Qi; Huchuan Lu; Dong Wang
SpectralNeRF: Physically Based Spectral Rendering with Neural Radiance	Ru Li; Jia Liu; Guanghui Liu; Shengping Zhang; Bing Zeng; Shuaicheng Liu
SimPSI: A Simple Strategy to Preserve Spectral Information in Time Series Data Augmentation	Hyun Ryu; Sunjae Yoon; Hee Suk Yoon; Eunseop Yoon; Chang D. Yoo
DAG-Aware Variational Autoencoder for Social Propagation Graph Generation	Dongpeng Hou; Chao Gao; Xuelong Li; Zhen Wang
Fairness without Demographics through Shared Latent Space-Based	Rashidul Islam; Huiyuan Chen; Yiwei Cai
Exploring Channel-Aware Typical Features for Out-of-Distribution Detection	Rundong He; Yue Yuan; Zhongyi Han; Fan Wang; Wan Su; Yilong Yin; Tongliang Liu; Yongshun Gong
Disentangled Partial Label Learning	Wei-Xuan Bao; Yong Rui; Min-Ling Zhang
Disguise without Disruption: Utility-Preserving Face De-identification	Zikui Cai; Zhongpai Gao; Benjamin Planche; Meng Zheng; Terrence Chen; M. Salman Asif; Ziyang Wu
InterpretARA: Enhancing Hybrid Automatic Readability Assessment with Linguistic Feature Interpreter and Contrastive Learning	Jinshan Zeng; Xianchao Tong; Xianglong Yu; Wenyan Xiao; Qing Huang
Keep the Faith: Faithful Explanations in Convolutional Neural Networks for Case-Based Reasoning	Tom Nuno Wolf; Fabian Bongratz; Anne-Marie Rickmann; Sebastian Pölsterl; Christian Wachinger
Learning Multi-Modal Cross-Scale Deformable Transformer Network for Unregistered Hyperspectral Image Super-resolution	Wenqian Dong; Yang Xu; Jiahui Qu; Shaoxiong Hou
Generative Multi-Modal Knowledge Retrieval with Large Language Models	Xinwei Long; Jiali Zeng; Fandong Meng; Zhiyuan Ma; Kaiyan Zhang; Bowen Zhou; Jie Zhou
ScanERU: Interactive 3D Visual Grounding Based on Embodied Reference Understanding	Ziyang Lu; Yunqiang Pei; Guoqing Wang; Peiwei Li; Yang Yang; Yinjie Lei; Heng Tao Shen
ProxyDet: Synthesizing Proxy Novel Classes via Classwise Mixup for Open Vocabulary Object Detection	Joonhyun Jeong; Geondo Park; Jayeon Yoo; Hyungsik Jung; Heesu Kim
Model-Driven Deep Neural Network for Enhanced AoA Estimation Using 5G gNB	Shengheng Liu; Xingkang Li; Zihuan Mao; Peng Liu; Yongming Huang
AUC Optimization from Multiple Unlabeled Datasets	Zheng Xie; Yu Liu; Ming Li
Attacking Transformers with Feature Diversity Adversarial Perturbation	Chenxing Gao; Hang Zhou; Junqing Yu; YuTeng Ye; Jiale Cai; Junle Wang; Wei Yang
High-Quality Real-Time Rendering Using Subpixel Sampling Reconstruction	Boyu Zhang; Hongliang Yuan
Real3D the Curious Case of Neural Scene Degeneration	Dengsheng Chen; Jie Hu; Xiaoming Wei; Enhua Wu
Quad Bayer Joint Demosaicing and Denoising Based on Dual Encoder Network with Joint Residual Learning	Bolun Zheng; Li Haoran; Quan Chen; Tingyu Wang; Xiaofei Zhou; Zhenghui HU; Chenggang Yan
Invariant Random Forest: Tree-Based Model Solution for OOD Generalization	Yufan LIAO; Qi Wu; Xing Yan
Exploring Transformer Extrapolation	Zhen Qin; Yiran Zhong; Hui Deng
TiMix: Text-Aware Image Mixing for Effective Vision-Language Pre-training	Chaoya Jiang; Wei Ye; Haiyang Xu; Qinghao Ye; Ming Yan; Ji Zhang; Shikun Zhang





FACL-Attack: Frequency-Aware Contrastive Learning for Transferable Adversarial Attacks	Hunmin Yang; Jongoh Jeong; Kuk-Jin Yoon
Entropic Open-Set Active Learning	Bardia Safaei; Vibashan VS; Celso de Melo; Vishal Patel
Response Enhanced Semi-supervised Dialogue Query Generation	Jianheng Huang; Ante Wang; Linfeng Gao; Linfeng Song; Jinsong Su
READ-PVLA: Recurrent Adapter with Partial Video-Language Alignment for Parameter-Efficient Transfer Learning in Low-Resource Video-Language Modeling	Thong Nguyen; Xiaobao Wu; Xinshuai Dong; Khoi M. Le; Zhiyuan Hu; Cong-Duy Nguyen; See-Kiong Ng; Anh Tuan Luu
Deep Homography Estimation for Visual Place Recognition	Feng Lu; Shuting Dong; Lijun Zhang; Bingxi Liu; Xiangyuan Lan; Dongmei Jiang; Chun Yuan
Stop! Planner Time: Metareasoning for Probabilistic Planning Using Learned Performance Profiles	Matthew Budd; Bruno Lacerda; Nick Hawes
Optimizing Local Satisfaction of Long-Run Average Objectives in Markov Decision Processes	David Klaska; Antonin Kucera; Vojtěch Kůr; Vit Musil; Vojtech Rehak
Winnie: Task-Oriented Dialog System with Structure-Aware Contrastive Learning and Enhanced Policy Planning	Kaizhi Gao; Tianyu Wang; Zhongjing Ma; Suli Zou
Unifying Multi-Modal Uncertainty Modeling and Semantic Alignment for Text-to-Image Person Re-identification	Zhiwei Zhao; Bin Liu; Yan Lu; Qi Chu; Nenghai Yu
Dual-Prior Augmented Decoding Network for Long Tail Distribution in HOI Detection	Jiayi Gao; Kongming Liang; Tao Wei; Wei Chen; Zhanyu Ma; Jun Guo
Small Model Can Self-Correct	Haixia Han; Jiaqing Liang; Jie Shi; Qianyu He; Yanghua Xiao
GSO-Net: Grid Surface Optimization via Learning Geometric Constraints	Chaoyun Wang; Jingmin Xin; Nanning Zheng; Caigui Jiang
No Internal Regret with Non-convex Loss Functions	Dravyansh Sharma
DanceAnyWay: Synthesizing Beat-Guided 3D Dances with Randomized Temporal Contrastive Learning	Aneesh Bhattacharya; Manas Paranjape; Uttaran Bhattacharya; Aniket Bera
Prot2Text: Multimodal Protein's Function Generation with GNNs and Transformers	Hadi Abdine; Michail Chatzianastasis; Costas Bouyioukos; Michalis Vazirgiannis
A Hierarchical Network for Multimodal Document-Level Relation Extraction	Lingxing Kong; Jiuliang Wang; Zheng Ma; Qifeng Zhou; Jianbing Zhang; Liang He; Jiajun Chen
M2Doc: A Multi-Modal Fusion Approach for Document Layout Analysis	Ning Zhang; Hiuyi Cheng; Jiayu Chen; Zongyuan Jiang; Jun Huang; Yang Xue; Lianwen Jin
Low-Light Face Super-resolution via Illumination, Structure, and Texture Associated Representation	Chenyang Wang; Junjun Jiang; Kui Jiang; Xianming Liu
EPSD: Early Pruning with Self-Distillation for Efficient Model Compression	Dong Chen; Ning Liu; Yichen Zhu; Zhengping Che; Rui Ma; Fachao Zhang; Mou Xiaofeng; Yi Chang; Jian Tang
Protect Your Score: Contact-Tracing with Differential Privacy Guarantees	Rob Romijnders; Christos Louizos; Yuki Asano; Max Welling
Improved Anonymous Multi Agent Path Finding Algorithm	Zain Alabedeen Ali; Konstantin Yakovlev
Optimal Makespan in a Minute Timespan! A Scalable Multi-Robot Goal Assignment Algorithm for Minimizing Mission Time	Aakash; Indranil Saha
Generative-Based Fusion Mechanism for Multi-Modal Tracking	Zhangyong Tang; Tianyang Xu; Xiaojun Wu; Xue-Feng Zhu; Josef Kittler
Towards Epistemic-Doxastic Planning with Observation and Revision	Thorsten Engesser; Andreas Herzig; Elise Perrotin
Frequency Shuffling and Enhancement for Open Set Recognition	Lijun Liu; Rui Wang; Yuan Wang; Lihua Jing; Chuan Wang



NestE: Modeling Nested Relational Structures for Knowledge Graph	Bo Xiong; Mojtaba Nayyeri; Linhao Luo; Zihao Wang; Shirui Pan; Steffen
GMMFormer: Gaussian-Mixture-Model Based Transformer for Efficient	Yuting Wang; Jinpeng Wang; Bin Chen; Ziyun Zeng; Shu-Tao Xia
Partially Relevant Video Retrieval	
One Self-Configurable Model to Solve Many Abstract Visual Reasoning	Mikołaj Mańkiński; Jacek Mańdziuk
Problems	
From Past to Future: Rethinking Eligibility Traces	Dhawal Gupta; Scott Jordan; Shreyas Chaudhari; Bo Liu; Philip Thomas;
	Bruno C. da Silva
A Diffusion Model with State Estimation for Degradation-Blind Inverse Imaging	Liya Ji; ZheFan Rao; Sinno Jialin Pan; Chenyang Lei; Qifeng Chen
ConSequence: Synthesizing Logically Constrained Sequences for Electronic	Brandon Theodorou; Shruti Jain; Cao Xiao; Jimeng Sun
Health Record Generation	
Temporal-Distributed Backdoor Attack against Video Based Action	Xi Li; Songhe Wang; Ruiquan Huang; Mahanth Gowda; George Kesidis
Self-Supervised 3D Human Mesh Recovery from a Single Image with	Guoli Yan; Zichun Zhong; Jing Hua
Uncertainty-Aware Learning	
Spatial Voting with Incomplete Voter Information	Aviram Imber; Jonas Israel; Markus Brill; Hadas Shachnai; Benny Kimelfeld
SEC: More Accurate Clustering Algorithm via Structural Entropy	Junyu Huang; Qilong Feng; Jiahui Wang; Ziyun Huang; Jinhui Xu; Jianxin
Descanning: From Scanned to the Original Images with a Color Correction	Junghun Cha; Ali Haider; Seoyun Yang; Hoeyeong Jin; Subin Yang; A F M
Diffusion Model	Shahab Uddin; Jaehyoung Kim; Soo Ye Kim; Sung-Ho Bae
SparseGNV: Generating Novel Views of Indoor Scenes with Sparse RGB-D	Weihaio Cheng; Yan-Pei Cao; Ying Shan
Images	
Scale Optimization Using Evolutionary Reinforcement Learning for Object	Jialu Zhang; Xiaoying Yang; Wentao He; Jianfeng Ren; Qian Zhang; Yitian
Detection on Drone Imagery	Zhao; Ruibin Bai; Xiangjian He; Jiang Liu
Distributed Manifold Hashing for Image Set Classification and Retrieval	Xiaobo Shen; Peizhuo Song; Yun-Hao Yuan; Yuhui Zheng
Collaborative Tooth Motion Diffusion Model in Digital Orthodontics	Yeying Fan; Guangshun Wei; Chen Wang; Shaojie Zhuang; Wenping Wang;
	Yuanfeng Zhou
An Information-Flow Perspective on Algorithmic Fairness	Samuel Teuber; Bernhard Beckert
KeDuSR: Real-World Dual-Lens Super-resolution via Kernel-Free Matching	Huanjing Yue; Zifan Cui; Kun Li; Jingyu Yang
Robustly Train Normalizing Flows via KL Divergence Regularization	Kun Song; Ruben Solozabal Ochoa de Retana; Hao Li; Martin Takac; Lu Ren;
	Fakhri Karray
CAR-Transformer: Cross-Attention Reinforcement Transformer for Cross-	Yuang Cai; Yuyu Yuan
Lingual Summarization	
Simple Image-Level Classification Improves Open-Vocabulary Object	Ruohuan Fang; Guansong Pang; Xiao Bai
Comparing the Robustness of Modern No-Reference Image- and Video-	Anastasia Antsiferova; Khaled Abud; Aleksandr Gushchin; Ekaterina
Quality Metrics to Adversarial Attacks	Shumitskaya; Sergey Lavrushkin; Dmitriy Vatolin
CoVR: Learning Composed Video Retrieval from Web Video Captions	Lucas Ventura; Antoine Yang; Cordelia Schmid; Gul Varol
Double-Descent Curves in Neural Networks: A New Perspective Using	Ouns El Harzli; Bernardo Cuenca Grau; Guillermo Valle Perez; Ard Louis
Gaussian Processes	
Unknown-Aware Graph Regularization for Robust Semi-supervised Learning	Heejo Kong; Suneung Kim; Ho-Joong Kim; Seong-Whan Lee
from Uncurated Data	
Building Minimal and Reusable Causal State Abstractions for Reinforcement	Zizhao Wang; Caroline Wang; Xuesu Xiao; Yuke Zhu; Peter Stone
Learning	



Learning Encodings for Constructive Neural Combinatorial Optimization Needs to Regret	Rui Sun; Zhi Zheng; Zhenkun Wang
DeRDaVa: Deletion-Robust Data Valuation for Machine Learning	Xiao Tian; Rachael Hwee Ling Sim; Jue Fan; Bryan Kian Hsiang Low
Federated Partial Label Learning with Local-Adaptive Augmentation and Regularization	Yan Yan; Yuhong Guo
CGMGM: A Cross-Gaussian Mixture Generative Model for Few-Shot Semantic Segmentation	Junao Shen; Kun Kuang; Jiaheng Wang; Xinyu Wang; Tian Feng; Wei Zhang
NavGPT: Explicit Reasoning in Vision-and-Language Navigation with Large Language Models	Gengze Zhou; Yicong Hong; Qi Wu
Efficient Constrained K-center Clustering with Background Knowledge	Longkun Guo; Chaoqi Jia; Kewen Liao; Zhigang Lu; Minhui Xue
Taming the Sigmoid Bottleneck: Provably Argmaxable Sparse Multi-Label Classification	Andreas Grivas; Antonio Vergari; Adam Lopez
Ternary Spike: Learning Ternary Spikes for Spiking Neural Networks	Yufei Guo; Yuanpei Chen; Xiaode Liu; Weihang Peng; Yuhan Zhang; Xuhui Huang; Zhe Ma
Privacy Amplification by Iteration for ADMM with (Strongly) Convex Objective Functions	T-H. Hubert Chan; Hao Xie; Mengshi ZHAO
CGS-Mask: Making Time Series Predictions Intuitive for All	Feng Lu; Wei Li; Yifei Sun; Cheng Song; Ren Yufei; Albert Zomaya
Text Diffusion with Reinforced Conditioning	Yuxuan Liu; Tianchi Yang; Shaohan Huang; Zihan Zhang; Haizhen Huang; Furu Wei; Weiwei Deng; Feng Sun; Qi Zhang
Doubly Perturbed Task Free Continual Learning	Byung Hyun Lee; Min-hwan Oh; Se Young Chun
DP-AdamBC: Your DP-Adam Is Actually DP-SGD (Unless You Apply Bias Correction)	Qiaoyue Tang; Frederick Shpilevskiy; Mathias Lécuyer
Fine-Tuning Graph Neural Networks by Preserving Graph Generative Patterns	Yifei Sun; Qi Zhu; Yang Yang; Chunping Wang; Tianyu Fan; Jiajun Zhu; Lei CHEN
PAC-Bayes Generalisation Bounds for Dynamical Systems including Stable RNNs	Deividas Eringis; John Leth; Zheng-Hua Tan; Rafael Wisniewski; Mihaly Petreczky
Detection and Defense of Unlearnable Examples	Yifan Zhu; Iijia Yu; Xiao-Shan Gao
MEPSI: An MDL-Based Ensemble Pruning Approach with Structural	Xiao-Dong Bi; Shao-Qun Zhang; Yuan Jiang
Decentralized Scheduling with QoS Constraints: Achieving $O(1)$ QoS Regret of Multi-Player Bandits	Qingsong Liu; Zhixuan Fang
Hierarchical Multi-Marginal Optimal Transport for Network Alignment	Zhichen Zeng; Boxin Du; Si Zhang; Yinglong Xia; Zhining Liu; Hanghang Tong
Unsupervised Layer-Wise Score Aggregation for Textual OOD Detection	Maxime Darrin; Guillaume Staerman; Eduardo Dadalto Camara Gomes; Jackie Chi Kit Cheung; Pablo Piantanida; Pierre Colombo
Semi-supervised Learning of Dynamical Systems with Neural Ordinary Differential Equations: A Teacher-Student Model Approach	Yu Wang; Yuxuan Yin; Karthik Somayaji NS; Jan Drgona; Malachi Schram; Mahantesh Halappanavar; Frank Liu; Peng Li
A Plug-and-Play Quaternion Message-Passing Module for Molecular Conformation Representation	Angxiao Yue; Dixin Luo; Hongteng Xu
3D Visibility-Aware Generalizable Neural Radiance Fields for Interacting	Xuan Huang; Hanhui Li; Zejun Yang; Zhisheng Wang; Xiaodan Liang
Exploring Post-training Quantization in LLMs from Comprehensive Study to Low Rank Compensation	Zhewei Yao; Xiaoxia Wu; Cheng Li; Stephen Youn; Yuxiong He



New Classes of the Greedy-Applicable Arm Feature Distributions in the Sparse Linear Bandit Problem	Koji Ichikawa; Shinji Ito; Daisuke Hatano; Hanna Sumita; Takuro Fukunaga; Naonori Kakimura; Ken-ichi Kawarabayashi
Convolutional Channel-Wise Competitive Learning for the Forward-Forward Algorithm	Andreas Papachristodoulou; Christos Kyrkou; Stelios Timotheou; Theodoris Theodoridis
CcDPM: A Continuous Conditional Diffusion Probabilistic Model for Inverse Design	Yanxuan Zhao; Peng Zhang; Guopeng Sun; Zhigong Yang; Jianqiang Chen; Yueqing Wang
Deep Quantum Error Correction	Yoni Choukroun; Lior Wolf
DSD <sup>2</sup> : Can We Dodge Sparse Double Descent and Compress the Neural Network Worry-Free?	Victor Quetu; Enzo Tartaglione
Compositional Text-to-Image Synthesis with Attention Map Control of Diffusion Models	Ruichen Wang; Zekang Chen; Chen Chen; Jian Ma; Haonan Lu; Xiaodong Lin
Universal Weak Coreset	Ragesh Jaiswal; Amit Kumar
Contrastive Balancing Representation Learning for Heterogeneous Dose-Response Curves Estimation	Minqin Zhu; Anpeng Wu; Haoxuan Li; Ruoxuan Xiong; Bo Li; Xiaoqing Yang; Xuan Qin; Peng Zhen; Jiecheng Guo; Fei Wu; Kun Kuang
DePRL: Achieving Linear Convergence Speedup in Personalized Decentralized Learning with Shared Representations	Guojun Xiong; Gang Yan; Shiqiang Wang; Jian Li
RetroOOD: Understanding Out-of-Distribution Generalization in Retrosynthesis Prediction	Yemin Yu; Luotian Yuan; Ying WEI; Hanyu Gao; Fei Wu; Zhihua Wang; Xinhai Ye
Robust Visual Imitation Learning with Inverse Dynamics Representations	Siyuan Li; Xun Wang; Rongchang Zuo; Kewu Sun; Lingfei Cui; Jishiyu Ding; Peng Liu; Zhe Ma
Offline Model-Based Optimization via Policy-Guided Gradient Search	Yassine Chemingui; Aryan Deshwal; Nghia Hoang; Janardhan Rao Doppa
Generator Assisted Mixture of Experts for Feature Acquisition in Batch	Vedang Asgaonkar; Aditya Jain; Abir De
Progressive Distillation Based on Masked Generation Feature Method for Knowledge Graph Completion	Cunhang Fan; Yujie Chen; Jun Xue; Yonghui Kong; Jianhua Tao; Zhao Lv
Provably Powerful Graph Neural Networks for Directed Multigraphs	Beni Egressy; Luc von Niederhäusern; Jovan Blanuša; Erik Altman; Roger Wattenhofer; Kubilay Atas
MemoryBank: Enhancing Large Language Models with Long-Term Memory	Wanjuan Zhong; Lianghong Guo; Qiqi Gao; He Ye; Yanlin Wang
Formal Logic Enabled Personalized Federated Learning through Property Inference	Ziyan An; Taylor Johnson; Meiyi Ma
REGLO: Provable Neural Network Repair for Global Robustness Properties	Feisi Fu; Zhilu Wang; Weichao Zhou; Yixuan Wang; Jiameng Fan; Chao Huang; Zhu Qi; Xin Chen; Wenchao Li
Secure Distributed Sparse Gaussian Process Models Using Multi-Key Homomorphic Encryption	Adil Nawaz; Guopeng Chen; Muhammad Umair Raza; Zahid Iqbal; Jianqiang Li; Victor C.M. Leung; Jie Chen
Learn to Follow: Decentralized Lifelong Multi-Agent Pathfinding via Planning and Learning	Alexey Skrynnik; Anton Andreychuk; Maria Nesterova; Konstantin Yakovlev; Aleksandr Panov
A General Search-Based Framework for Generating Textual Counterfactual Explanations	Daniel Gilo; Shaul Markovitch
Robust Communicative Multi-Agent Reinforcement Learning with Active RewriteLM: An Instruction-Tuned Large Language Model for Text Rewriting	Lebin Yu; Yunbo Qiu; Quanming Yao; Yuan Shen; Xudong Zhang; Jian Wang; Lei Shu; Liangchen Luo; Jayakumar Hoskere; Yun Zhu; Yinxiao Liu; Simon Tong; Jindong Chen; Lei Meng
Catch-Up Mix: Catch-Up Class for Struggling Filters in CNN	Minsoo Kang; Minkoo Kang; Suhyun Kim



CaMIL: Causal Multiple Instance Learning for Whole Slide Image Classification	Kaitao Chen; Shiliang Sun; Jing Zhao
DDDM-VC: Decoupled Denoising Diffusion Models with Disentangled Representation and Prior Mixup for Verified Robust Voice Conversion	Ha-Yeong Choi; Sang-Hoon Lee; Seong-Whan Lee
Performative Federated Learning: A Solution to Model-Dependent and Heterogeneous Distribution Shifts	Kun Jin; Tongxin Yin; Zhongzhu Chen; Zeyu Sun; Xueru Zhang; Yang Liu; Mingyan Liu
Multi-Architecture Multi-Expert Diffusion Models	Yunsung Lee; Jin-Young Kim; Hyojun Go; Myeongho Jeong; Shinhyeok Oh; Seungtaek Choi
Eliciting Honest Information from Authors Using Sequential Review	Yichi Zhang; Grant Schoenebeck; Weijie Su
Approximation Scheme for Weighted Metric Clustering via Sherali-Adams	Dmitrii Avdiukhin; Vaggos Chatziafratis; Konstantin Makarychev; Grigory Yaroslavl'tsev
G-Adapter: Towards Structure-Aware Parameter-Efficient Transfer Learning for Graph Transformer Networks	Anchun Gui; Jinqiang Ye; Han Xiao
ConceptBed: Evaluating Concept Learning Abilities of Text-to-Image Diffusion Models	Maitreya Patel; Tejas Gokhale; Chitta Baral; Yezhou Yang
Inducing Point Operator Transformer: A Flexible and Scalable Architecture for Solving PDEs	Seungjun Lee; Tael L. Oh
GLOP: Learning Global Partition and Local Construction for Solving Large-Scale Routing Problems in Real-Time	Haoran Ye; Jiarui Wang; Helan Liang; Zhiguang Cao; Yong Li; Fanzhang Li
Accelerate Multi-Agent Reinforcement Learning in Zero-Sum Games with Subgame Curriculum Learning	Jiayu Chen; Zelai Xu; Yunfei Li; Chao Yu; Jiaming Song; Huazhong Yang; Fei Fang; Yu Wang; Yi Wu
Stochastic Bayesian Optimization with Unknown Continuous Context Distribution via Kernel Density Estimation	Xiaobin Huang; Lei Song; Ke Xue; Chao Qian
Regret Analysis of Repeated Delegated Choice	Suho Shin; Keivan Rezaei; Mohammad Hajiaghayi; Mohammad Mahdavi
No Prior Mask: Eliminate Redundant Action for Deep Reinforcement Learning	Dianyu Zhong; Yiqin Yang; Qianchuan Zhao
\$z\$-SignFedAvg: A Unified Stochastic Sign-Based Compression for Federated Learning	Zhiwei Tang; Yanmeng Wang; Tsung-Hui Chang
Contextual Pandora's Box	Alexia Atsidakou; Constantine Caramanis; Evangelia Gergatsouli; Orestis Papadigenopoulos; Christos Tzamos
Robust Distributed Gradient Aggregation Using Projections onto Gradient Manifolds	Kwang In Kim
Generative Model Perception Rectification Algorithm for Trade-Off between Diversity and Quality	Guipeng Lan; Shuai Xiao; Jiachen Yang; Jiabao Wen
Taming Binarized Neural Networks and Mixed-Integer Programs	Johannes Aspman; Georgios Korpas; Jakub Marecek
Partial Label Learning with a Partner	Chongjie Si; Zekun Jiang; Xuehui Wang; Yan Wang; Xiaokang Yang; Wei
Faithful Model Explanations through Energy-Constrained Conformal Counterfactuals	Patrick Altmeyer; Mojtaba Farmanbar; Arie Van Deursen; Cynthia Liem
A Graph Dynamics Prior for Relational Inference	Liming Pan; Cheng Shi; Ivan Dokmanic
Enhancing the Robustness of Spiking Neural Networks with Stochastic Gating Mechanisms	Jianhao Ding; Zhao Fei Yu; Tiejun Huang; Jian Liu
Towards Dynamic Spatial-Temporal Graph Learning: A Decoupled Perspective	Binwu Wang; Pengkun Wang; Yudong Zhang; Xu Wang; Zhengyang Zhou; Lei Bai; Yang Wang



A Closer Look at Curriculum Adversarial Training: From an Online Perspective	Lianghe Shi; Weiwei Liu
Learning Explicit Contact for Implicit Reconstruction of Hand-Held Objects from Monocular Images	Junxing Hu; Hongwen Zhang; Zerui Chen; Mengcheng Li; Yunlong Wang; Yebin Liu; Zhenan Sun
Get a Head Start: On-Demand Pedagogical Policy Selection in Intelligent Tutoring	Ge Gao; Xi Yang; Min Chi
Statistical Spatially Inhomogeneous Diffusion Inference	Yinuo Ren; Yiping Lu; Lexing Ying; Grant M. Rotskoff
Input Margins Can Predict Generalization Too	Coenraad Mouton; Marthinus Wilhelmus Theunissen; Marelle H Davel
Provably Convergent Federated Trilevel Learning	Yang Jiao; Kai YANG; Tiancheng Wu; Chengtao Jian; Jianwei Huang
Exploiting Symmetric Temporally Sparse BPTT for Efficient RNN Training	Xi Chen; Chang Gao; Zuowen Wang; Longbiao Cheng; Sheng Zhou; Shih-Chii Liu; Tobi Delbruck
Equity-Transformer: Solving NP-Hard Min-Max Routing Problems as Sequential Generation with Equity Context	Jiwoo Son; Minsu Kim; Sanghyeok Choi; Hyeonah Kim; Jinkyoo Park
OWQ: Outlier-Aware Weight Quantization for Efficient Fine-Tuning and Inference of Large Language Models	Changhun Lee; Jungyu Jin; Taesu Kim; Hyungjun Kim; Eunhyeok Park
DeepSpeed Data Efficiency: Improving Deep Learning Model Quality and Training Efficiency via Efficient Data Sampling and Routing	Conglong Li; Zhewei Yao; Xiaoxia Wu; Minjia Zhang; Connor Holmes; Cheng Li; Yuxiong He
Leaving the Nest: Going beyond Local Loss Functions for Predict-Then-Optimize	Sanket Shah; Bryan Wilder; Andrew Perrault; Milind Tambe
Robust Nonparametric Regression under Poisoning Attack	Puning Zhao; Zhiguo Wan
Scores for Learning Discrete Causal Graphs with Unobserved Confounders	Alexis Bellot; Junzhe Zhang; Elias Bareinboim
Dynamic Knowledge Injection for AIXI Agents	Samuel Yang-Zhao; Kee Siong Ng; Marcus Hutter
Finite-Time Frequentist Regret Bounds of Multi-Agent Thompson Sampling on Sparse Hypergraphs	Tianyuan Jin; Hao-Lun Hsu; William Chang; Pan Xu
Understanding and Improving Optimization in Predictive Coding Networks	Nicholas Alonso; Jeffrey Krichmar; Emre Neftci
PARSAC: Accelerating Robust Multi-Model Fitting with Parallel Sample Consensus	Florian Kluger; Bodo Rosenhahn
Factored Online Planning in Many-Agent POMDPs	Maris Galesloot; Thiago Simão; Sebastian Junges; Nils Jansen
Neural Reasoning about Agents' Goals, Preferences, and Actions	Matteo Bortoletto; Lei Shi; Andreas Bulling
Pareto Front-Diverse Batch Multi-Objective Bayesian Optimization	Alaleh Ahmadianshalchi; Syrine Belakaria; Janardhan Rao Doppa
Principal-Agent Reward Shaping in MDPs	Omer Ben-Porat; Yishay Mansour; Michal Moshkovitz; Boaz Taitler
Learning Uncertainty-Aware Temporally-Extended Actions	Joongkyu Lee; Seung Joon Park; Yunhao Tang; Min-hwan Oh
Adversarial Attacks on the Interpretation of Neuron Activation Maximization	Géraldin Nanfack; Alexander Fulleringer; Jonathan Marty; Michael Eickenberg; Eugene Belilovsky
Task Planning for Object Rearrangement in Multi-Room Environments	Karan Mirakhor; Sourav Ghosh; Dipanjan Das; Brojeshwar Bhowmick
Continuous Treatment Effect Estimation Using Gradient Interpolation and Kernel Smoothing	Lokesh Nagalapatti; Akshay Iyer; Abir De; Sunita Sarawagi
Feature Distribution Matching by Optimal Transport for Effective and Robust Coreset Selection	Weiwei Xiao; Yongyong Chen; Qiben Shan; Yaowei Wang; Jingyong Su
Dialogues Are Not Just Text: Modeling Cognition for Dialogue Coherence Evaluation	Xue Li; Jia Su; Yang Yang; Zipeng Gao; Xinyu Duan; Yi Guan





A Unified Self-Distillation Framework for Multimodal Sentiment Analysis with Uncertain Missing Modalities	Mingcheng Li; Dingkan Yang; Yuxuan Lei; Shunli Wang; Shuaibing Wang; Liuzhen Su; Kun Yang; Yuzheng Wang; Mingyang Sun; Lihua Zhang
Guiding a Harsh-Environments Robust Detector via RAW Data Characteristic Mining	Hongyang Chen; Hung-Shuo Tai; Kaisheng Ma
One at a Time: Progressive Multi-Step Volumetric Probability Learning for Reliable 3D Scene Perception	Bohan Li; Yasheng Sun; Jingxin Dong; Zheng Zhu; Jinming Liu; Xin Jin; Wenjun Zeng
LimeAttack: Local Explainable Method for Textual Hard-Label Adversarial	Hai Zhu; Qingyang Zhao; Weiwei Shang; Yuren Wu; Kai Liu
A Novel Skip Orthogonal List for Dynamic Optimal Transport Problem	Xiaoyang Xu; Hu Ding
Mixed-Effects Contextual Bandits	Kyungbok Lee; Myunghee Cho Paik; Min-hwan Oh; Gi-Soo Kim
On Partial Optimal Transport: Revising the Infeasibility of Sinkhorn and Efficient Gradient Methods	Anh Duc Nguyen; Tuan Dung Nguyen; Quang Nguyen; Hoang Nguyen; Lam Nguyen; Kim-Chuan Toh
Beyond Attention: Breaking the Limits of Transformer Context Length with Recurrent Memory	Aydar Bulatov; Yuri Kuratov; Yermek Kapushev; Mikhail Burtsev
Resisting Backdoor Attacks in Federated Learning via Bidirectional Elections and Individual Perspective	Zhen Qin; Feiyi Chen; Chen Zhi; Xueqiang Yan; Shuiguang Deng
Transportable Representations for Domain Generalization	Kasra Jalaldoust; Elias Bareinboim
Exponential Hardness of Optimization from the Locality in Quantum Neural Networks	Hao-Kai Zhang; Chengkai Zhu; Geng Liu; Xin Wang
VIXEN: Visual Text Comparison Network for Image Difference Captioning	Alexander Black; Jing Shi; Yifei Fan; Tu Bui; John Collomosse
MFOS: Model-Free & One-Shot Object Pose Estimation	JongMin Lee; Yohann Cabon; Romain Brégier; Sungjoo Yoo; Jerome Revaud
Hierarchical Topology Isomorphism Expertise Embedded Graph Contrastive Learning	Jiangmeng Li; Yifan Jin; Hang Gao; Wenwen Qiang; Changwen Zheng; Fuchun Sun
Cumulative Regret Analysis of the Piyavskii–Shubert Algorithm and Its Variants for Global Optimization	Kaan Gokcesu; Hakan Gökcesu
PDE+: Enhancing Generalization via PDE with Adaptive Distributional	Yige Yuan; Bingbing Xu; Bo Lin; Liang Hou; Fei Sun; Huawei Shen; Xueqi
Improved Bandits in Many-to-One Matching Markets with Incentive	Fang Kong; Shuai Li
Learning Representations on the Unit Sphere: Investigating Angular Gaussian and Von Mises-Fisher Distributions for Online Continual Learning	Nicolas Michel; Giovanni Chierchia; Romain Negrel; Jean-François BERCHER
Towards Real-World Test-Time Adaptation: Tri-net Self-Training with Balanced Normalization	Yongyi Su; Xun Xu; Kui Jia
Complete Neural Networks for Complete Euclidean Graphs	Snir Hordan; Tal Amir; Nadav Dym; Steven Gortler
Probabilistic Offline Policy Ranking with Approximate Bayesian Computation	Longchao Da; Porter Jenkins; Trevor Schwantes; Jeffrey Dotson; Hua Wei
DRF: Improving Certified Robustness via Distributional Robustness	Zekai Wang; Zhengyu Zhou; Weiwei Liu
Revisiting Disentanglement in Downstream Tasks: A Study on Its Necessity for Abstract Visual Reasoning	Ruiqian Nai; Zixin Wen; Ji Li; Yuanzhi Li; Yang Gao
Stealthy Adversarial Attacks on Stochastic Multi-Armed Bandits	Zhiwei Wang; Hongning Wang; Huazheng Wang
CutFreq: Cut-and-Swap Frequency Components for Low-Level Vision Augmentation	Hongyang Chen; Kaisheng Ma
PosDiffNet: Positional Neural Diffusion for Point Cloud Registration in a Large Field of View with Perturbations	Rui She; Sijie Wang; Qiyu Kang; Kai Zhao; Yang Song; Wee Peng Tay; Tianyu Geng; Xingchao Jian



Beyond Prototypes: Semantic Anchor Regularization for Better Representation Learning	Yanqi Ge; Qiang Nie; Ye Huang; Yong Liu; Chengjie Wang; Feng Zheng; Wen Li; Lixin Duan
Dynamic Spiking Graph Neural Networks	Yin Nan; Mengzhu Wang; Zhenghan Chen; Giulia De Masi; Huan Xiong; Bin Yiming Huang; Yujie Zeng; Qiang Wu; Linyuan Lü
Higher-Order Graph Convolutional Network with Flower-Petals Laplacians on Simplicial Complexes	
Xiezhi: An Ever-Updating Benchmark for Holistic Domain Knowledge Evaluation	Zhouhong Gu; Xiaoxuan Zhu; Haoning Ye; Lin Zhang; Jianchen Wang; Yixin Zhu; Sihang Jiang; Zhuozhi Xiong; Zihan Li; Weijie Wu; Qianyu He; Rui Xu; Wenhao Huang; Jingping Liu; Zili Wang; Shusen Wang; Weiguo Zheng; Hongwei Feng; Yanghua Xiao
A Perspective of Q-value Estimation on Offline-to-Online Reinforcement Learning	Yinmin Zhang; Jie Liu; Chuming Li; Yazhe Niu; Yaodong Yang; Yu Liu; Wanli Ouyang
Controllable Mind Visual Diffusion Model	Bohan Zeng; Shanglin Li; Xuhui Liu; Sicheng Gao; Xiaolong Jiang; Xu Tang; Yao Hu; Jianzhuang Liu; Baochang Zhang
Reconciling Predictive and Statistical Parity: A Causal Approach	Drago Plecko; Elias Bareinboim
Out of Thin Air: Exploring Data-Free Adversarial Robustness Distillation	Yuzheng Wang; Zhaoyu Chen; Dingkan Yang; Pinxue Guo; Kaixun Jiang; Wenqiang Zhang; Lizhe Qi
Towards Detailed Text-to-Motion Synthesis via Basic-to-Advanced Hierarchical Diffusion Model	Zhenyu Xie; Yang Wu; Xuehao Gao; Zhongqian Sun; Wei Yang; Xiaodan Liang
Better than Random: Reliable NLG Human Evaluation with Constrained Active Sampling	Jie Ruan; Xiao Pu; Mingqi Gao; Xiaojun Wan; Yuesheng Zhu
DUEL: Duplicate Elimination on Active Memory for Self-Supervised Class-Imbalanced Learning	Won-Seok Choi; Hyundo Lee; Dong-Sig Han; Junseok Park; Heeyeon Koo; Byoung-Tak Zhang
Dirichlet-Based Prediction Calibration for Learning with Noisy Labels	Chen-Chen Zong; Ye-Wen Wang; Ming-Kun Xie; Sheng-Jun Huang
HAGO-Net: Hierarchical Geometric Message Passing for Molecular Representation Learning	Hongbin Pei; Taile Chen; Chen A; Huiqi Deng; Jing Tao; Pinghui Wang; Xiaohong Guan